

Processes, Threads, Concurrency, and Synchronization

Lecture Overview

- **Introduction: Processes vs. Threads**
 - Design Aspects
 - Concurrency
- **Threads: Implementation Aspects**
 - User-level vs. Kernel-level
 - Multithreading patterns
- **Synchronization**
 - Mutual Exclusion
 - Hardware-based solutions
 - Software-based solutions
- **Threads in the real world**
 - Synchronization and the OS
 - Software solutions based on Threads

Processes vs. Threads

Process: is an instance of a program running on a computing system.

Process = Code + Data + Memory + Process Control Block

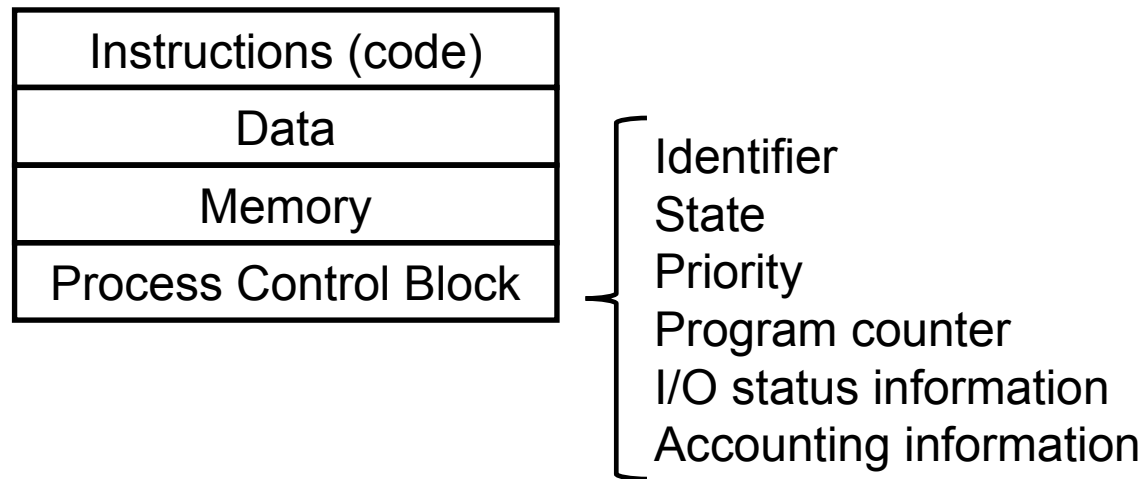



Fig 1. Simplified Process Image


Processes vs. Threads

Fork(): creating a process dynamically.

```
int main ()
{
    pid_t pid;
    pid = fork();
    if (pid == 0) // Child Process Code
        printf("Child process\n");
    else
        printf("Parent process\n");
    return (0);
}
```



Parent Process



Child Process

Processes do not share their memory address space, but they share the opened files.

Processes vs. Threads

Threads:

- The unit of dispatching is referred to as a thread or lightweight process.
- The unit of resource ownership is referred to as a process or task.

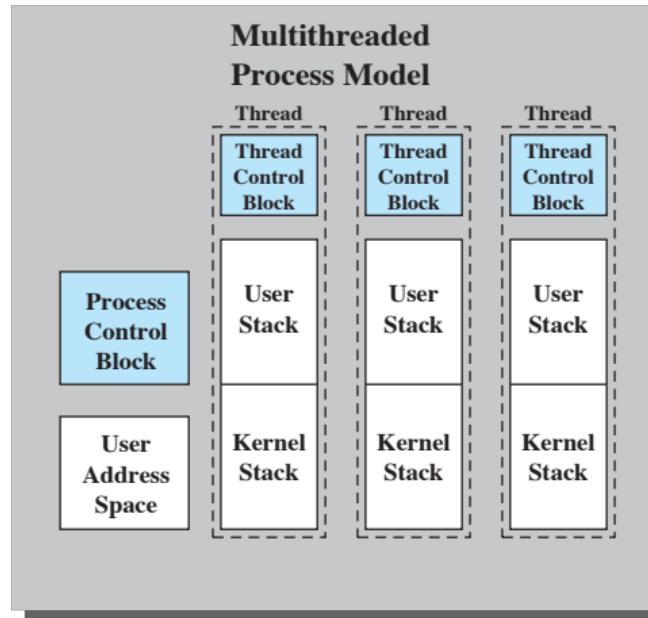


Fig 2. Simplified Thread Image (Stallings, OS8e)

Processes vs. Threads

Threads Challenges:

- No memory protection inside an address space.
- Lightweight processes can now interfere with each other.

```
void * child ()  
{  
    printf("Child process\n");  
}
```

```
int main ()  
{  
    thread_t tid;  
    pthread_create(&tid, NULL, child, NULL);  
    printf("Parent process\n");  
    return (0);  
}
```

Stack Address (NULL = anywhere)

Parameters (void pointer)

POSIX Thread Example

Processes vs. Threads

Threads Example:

```
#define COUNT 5
```

```
void *inc_x(void *x_void_ptr)
{
    int *x_ptr = (int *)x_void_ptr;
    for (int i=0; i< COUNT; i++)
        (*x_ptr)++;
    return NULL;
}
```

$(*x_ptr)++$ {
load reg, M
inc reg
store reg, M

Thread 1 loads 0

Threads 3, 4, 5 run to completion

Thread 2 runs through 4 of its 5 iterations.

Thread 1 increments 0 to 1 and stores the result (1).

Thread 2 loads 1

Thread 1 runs to completion

Thread 2 increments 1 to 2 and stores 2

POSIX Thread Example

Threads: Implementation

User-Level Threads (ULTs):

- User-level threads are managed by procedures within the task address space by the thread library.
- The kernel is not aware of the existence of threads.
- Example: POSIX Threads (POSIX = Portable OS interface)




Fig 3. User-level Threads (Stalling OS8e)

Threads: Implementation


User-Level Threads (ULTs):

- **Advantages:**

- Thread switching does not require kernel mode privileges (no performance penalty).
- Scheduling can be application specific.
- Can be implemented on any OS.

- **Disadvantages:**

- A blocking call from one thread will block all the threads from the same process.
Solution: Using non-blocking calls (not easy).
- Multithreaded programs cannot take advantage of multiprocessing



Threads: Implementation

Kernel-Level Threads (KLTs):

- Thread management is done by the kernel (one process table entry per thread).
- Best solution for multiprocessor architectures (kernel can allocate several processors to a single multithreaded task).





Fig 4. Kernel-level Threads (Stalling OS8e)

Threads: Implementation

Kernel-Level Threads (KLTs):

- **Advantages:**
 - The kernel can simultaneously schedule multiple threads from the same process on multiple processors.
 - If one thread in a process is blocked, the kernel can schedule another thread of the same process.
- **Disadvantages:**
 - Switching from one thread to another within the same process requires two context switches (overhead).




Threads: Implementation

Multithreading patterns:

Dispatcher-Worker pattern:

- Dispatcher Thread: reads the requests from the mailbox.
- Worker Thread: performs the task.
- **Challenges:**
 - Dispatcher can be a bottleneck.
 - Static / Dynamic Worker generation.




Threads: Implementation

Multithreading patterns:

Pipeline pattern:

- Threads assigned one subtask in the system.
- Entire task = Pipeline of threads.
- Multiple tasks concurrently run in the system, in different pipeline stages
- Shared buffer based communication between stages.
- **Challenges:**
 - Performance depends on weakest link.
 - Dividing requests in subtasks.



Threads: Implementation

Multithreading patterns:

Example:

For 5-step toy order application, we have the solutions using:

- Dispatcher-Worker pattern.
- Pipeline pattern.

Constraints:

- For both patterns, we have 5 threads.
- For the dispatcher-workers solution, a worker produces a toy order in 100 ms.
- For the pipeline solution, each one of the five stages takes 20 ms.
- Time to read and process the request is negligible.

How long will it take for these solutions to complete an 8-toy order and a 9-toy order?

Threads: Implementation

Multithreading patterns:

Dispatcher-Worker pattern

8-toy order:


0 ms (reading + processing) +
100 ms (processing 4 toy orders
concurrently) +
100 ms (processing 4 toy orders
concurrently) =

Total time = 200 ms.

9-toy order:

0 ms (reading + processing) +
100 ms (processing 4 toy orders
concurrently) +
100 ms (processing 4 toy orders
concurrently) +
100 ms (processing 1 toy order)

Total time = 300 ms



Threads: Implementation

Multithreading patterns:

Pipeline pattern

8-toy orders:


0 ms (reading + processing) +
100 ms (time to finish first toy order) +
20 ms * 7 (processing 7 toy orders) =

Total time = 240 ms.

9-toy orders:

0 ms (reading + processing) +
100 ms (time to finish first toy order) +
20 ms * 8 (processing 8 toy orders) =

Total time = 260 ms.



Threads: Synchronization

Mutual Exclusion:

Key Terms:

- **Atomic operation:** a function or action implemented as a sequence of one or more instructions that cannot be interrupted (indivisible).
- **Critical Section:** section of code within a process that requires access to shared resources that has to be executed as an atomic operation.
- **Mutual Exclusion:** property that guarantees that only one thread of execution will have access to a critical section for a particular shared resource at a particular time.
- **Racing Condition:** a situation when multiple threads or processes try to gain access to a particular shared resource at the same time.

Threads: Synchronization

Mutual Exclusion:

Challenges:

- **Concurrency**

- Detecting programming errors.
- Difficult for the OS to manage the allocation of resources optimally.

- **Busy waits**

Example: `while(TRUE);` Wasting CPU cycles.

- **Deadlocks**

Permanent blocking of a set of processes that compete for the same shared resource.

- **Starvation**

Postponing indefinitely the execution of a process based on a condition.

Threads: Synchronization

Mutual Exclusion:

Requirements:

- A process that halts must do so without interfering with other processes.
- No deadlock or starvation.
- A process must not be denied access to a critical section when there is no other process using it.
- No assumptions are made about relative process speeds or number of processes.
- A process remains inside its critical section for a finite period of time.

Threads: Synchronization

Mutual Exclusion: Hardware Solutions

Interrupt Disabling

- Disabling interrupts guarantees mutual exclusion

```
int main()
{
    disableInterrupt();
    // Critical section goes here
    enableInterrupt();
}
```

- Disadvantages:

- The performance of the system could be noticeably degraded.
- Not suitable for Multiprocessor platforms.

Threads: Synchronization

Mutual Exclusion: Hardware Solutions

Special Machine Instructions

- Test and Set
- Exchange

```
/* program mutualexclusion */
const int n = /* number of processes */;
int bolt;

void main()
{
    bolt = 0;
    parbegin (P(1), P(2), . . . ,P(n));
}
```

```
void P(int i)
{
    while (true)
    {
        while (!test_and_set(&bolt))
            /* do nothing */;
        /* critical section */;
        bolt = 0;
        /* non-critical section */;
    }
}
```

- Carried out atomically.

Threads: Synchronization

Mutual Exclusion: Hardware Solutions

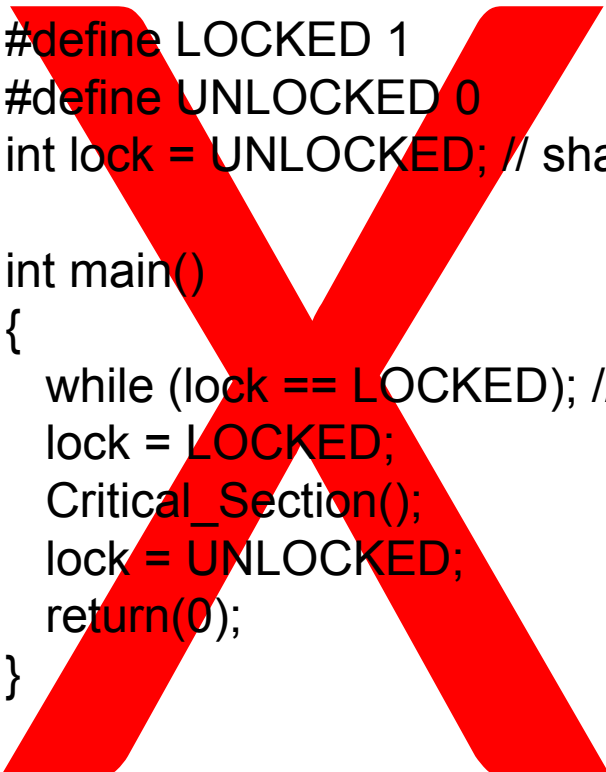
Special Machine Instructions

- Advantages:
 - Applicable to uniprocessors and multiprocessors.
 - Simple and easy to verify.
 - Supports multiple critical sections.
- Disadvantages:
 - Busy-wait is employed.
 - We have to avoid busy waits on single-core architectures.
 - We can use them only for short waits on multicore architectures.
 - Starvation is possible.

Threads: Synchronization

Mutual Exclusion: Software Solutions

Peterson's Algorithm



```
#define LOCKED 1
#define UNLOCKED 0
int lock = UNLOCKED; // shared

int main()
{
    while (lock == LOCKED); // busy wait
    lock = LOCKED;
    Critical_Section();
    lock = UNLOCKED;
    return(0);
}
```

Solution:

Move Lock=LOCKED before
while (lock == LOCKED); // busy wait

Result:

Deadlock

Threads: Synchronization

Mutual Exclusion: Software Solutions

Peterson's Algorithm

```
// shared variables
int reserved[2] = {FALSE, FALSE};
int mustwait; // tiebreaker
int main()
{
    int other= 1 - pid; //other process
    reserved[pid] = TRUE;
    must_wait = pid; // set tiebraker
    while (reserved[other] && must_wait==pid);
    // Critical Section
    reserved[pid] = FALSE;
    return(0);
}
```

Solution:

When two processes arrive in lockstep, last one must wait.

Result:

Mutual Exclusion

Threads: Synchronization

Mutual Exclusion: Software Solutions

Semaphores

- Special integer variables.
- May be initialized to a nonnegative integer value.
- Two atomic primitives
 - SemWait operation decrements the value (P() or Down)
 - The semSignal operation increments the value (V() or Up).

Threads: Synchronization

Mutual Exclusion: Software Solutions

Semaphores

- Normal implementation:
 - Processes waiting for a semaphore whose value is zero are put in the blocked state.
 - Busy waits are eliminated.
- Strong Semaphores: FIFO order
- Weak Semaphores: No specific order.

Threads: Synchronization

Mutual Exclusion: Software Solutions

Semaphores

- Producer – Consumer Problem: only one producer or consumer may access the buffer at any one time.

```
int n;  
binary_semaphore s = 1;  
binary_semaphore delay = 0;
```

```
void producer()  
{  
    while (true)  
    {  
        produce();  
        semWaitB(s);  
        append();  
        n++;  
        if (n==1)  
            semSignalB(delay);  
        semSignalB(s);  
    }  
}
```

```
void consumer()  
{  
    semWaitB(delay);  
    while (true)  
    {  
        semWaitB(s);  
        take();  
        n--;  
        semSignalB(s);  
        consume();  
        if (n==0)  
            semWaitB(delay);  
    }  
}
```

Threads: Synchronization

Semaphores: Producer - Consumer

	Producer	Consumer	s	n	Delay
1			1	0	0
2	semWaitB(s)		0	0	0
3	n++		0	1	0
4	if (n==1) (semSignalB(delay))		0	1	1
5	semSignalB(s)		1	1	1
6		semWaitB(delay)	1	1	0
7		semWaitB(s)	0	1	0
8		n--	0	0	0
9		semSignalB(s)	1	0	0
10	semWaitB(s)		0	0	0
11	n++		0	1	0
12	if (n==1) (semSignalB(delay))		0	1	1
13	semSignalB(s)		1	1	1
14		if (n==0) (semWaitB(delay))	1	1	1
15		semWaitB(s)	0	1	1
16		n--	0	0	1
17		semSignalB(s)	1	0	1
18		if (n==0) (semWaitB(delay))	1	0	0
19		semWaitB(s)	0	0	0
20		n--	0	-1	0
21		semSignalB(s)	1	-1	0

Fig 5. Producer - Consumer (Stalling OS8e)

Threads: Synchronization

Semaphores: Producer - Consumer

Solution 1

```
void consumer()
{
    semWaitB(delay);
    while (true)
    {
        semWaitB(s);
        take();
        n--;
        m=n;
        semSignalB(s);
        consume();
        if (m==0)
            semWaitB(delay);
    }
}
```

Solution 2

```
semaphore n = 0;
semaphore s = 1;
```

```
void producer()
{
    while (true)
    {
        produce();
        semWait(s);
        append();
        semSignal(s);
        semSignal(n);
    }
}
```

```
void consumer()
{
    while (true)
    {
        semWait(n);
        semWait(s);
        take();
        semSignal(s);
        consume();
    }
}
```

Threads: Synchronization

Mutual Exclusion: Software Solutions

Monitors

- Software module consisting of one or more procedures, an initialization sequence, and local data.
- Local data variables are accessible only by the monitor's procedures and not by any external procedure.
- Process enters monitor by invoking one of its procedures.
- Only one process may be executing in the monitor at a time.

Threads: Synchronization

Mutual Exclusion: Software Solutions

Monitors

- Implemented in a number of programming languages including Concurrent Pascal, Pascal-Plus, Modula-2, Modula-3, and Java
- Achieved by the use of condition variables using the primitives:
 - **cwait(c)**: suspend execution of the calling process on condition c.
 - **csignal(c)**: resume execution of some process blocked after a cwait(c). (Hoare semantics)
 - **cnotify(c)**: resume execution of all processes waiting for condition c. (Mesa semantics).

Threads: Synchronization

Mutual Exclusion: Software Solutions

Monitors

Producer – Consumer problem:

```
void append (char x)
{
    while(count == N) cwait(notfull);    /* buffer is full; avoid overflow */
    buffer[nextin] = x;
    nextin = (nextin + 1) % N;
    count++;                             /* one more item in buffer */
    cnotify(notempty);                   /* notify any waiting consumer */
}

void take (char x)
{
    while(count == 0) cwait(notempty); /* buffer is empty; avoid underflow */
    x = buffer[nextout];
    nextout = (nextout + 1) % N;
    count--;                             /* one fewer item in buffer */
    cnotify(notfull);                    /* notify any waiting producer */
}
```

Fig 6. Producer – Consumer with Mesa semantics (Stalling OS8e)

Threads in the real world

Synchronization and the OS

Windows

- Uses interrupt masks to protect access to global resources on uniprocessor systems.
- Uses spinlocks on multiprocessor systems (Spinlocking-thread will never be preempted).
- Also provides dispatcher objects (user-space) which may act as mutexes, semaphores, events, and timers (an event acts much like a condition variable).

Threads in the real world

Synchronization and the OS

Linux:

- Prior to kernel Version 2.6, disables interrupts to implement short critical sections.
- Version 2.6 and later, fully preemptive
- Linux provides:
 - Semaphores
 - Spinlocks
 - Reader-writer locks
- On single-cpu system, spinlocks replaced by enabling and disabling kernel preemption.

Threads in the real world

Synchronization and the OS

PThreads:

- Pthreads API is OS-independent
- It provides:
 - Mutex locks
 - Condition variables
- Non-portable extensions include:
 - Read-write locks
 - Spinlocks

Threads in the real world

Alternative Approaches

- **Transactional Memory**

A memory transaction is a sequence of read-write operations to memory that are performed atomically.

- **OpenMP**

- API that support parallel programming.
- `#pragma omp critical` directive is treated as a critical section and performed atomically.

- **Functional Programming Languages**

- Functional programming languages do not maintain state.
- Variables are treated as immutable and cannot change state once they have been assigned a value.
- Examples: Erlang and Scala

Threads in the real world

Applications

Core: basic server skeleton

Modules: per functionality

Flow of Control: Similar to Event Driven model

A combination of MP + MT
Each process =
dispatcher/worker with
dynamic thread pool
Number of processes can
also be dynamically
adjusted.




Fig 6. Producer – Apache Web Server (1)

1. <https://applied-programming.github.io/Operating-Systems-Notes/3-Threads-and-Concurrency/#event-driven-model>

Lecture Summary

- **Introduction: Processes vs. Threads**
 - Design Aspects
 - Concurrency
- **Threads: Implementation Aspects**
 - User-level vs. Kernel-level
 - Multithreading patterns
- **Synchronization**
 - Mutual Exclusion
 - Hardware-based solutions
 - Software-based solutions
- **Threads in the real world**
 - Synchronization and the OS
 - Software solutions based on Threads