

**COSC 3380 Spring 2024**

**Database Systems**

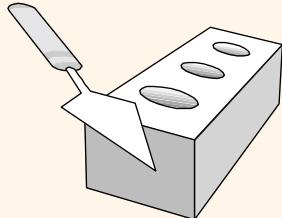
**M & W 4:00 to 5:30 PM**

Prof. **Victoria Hilford**

**PLEASE TURN your webcam ON (must have)**

**NO CHATTING during LECTURE**

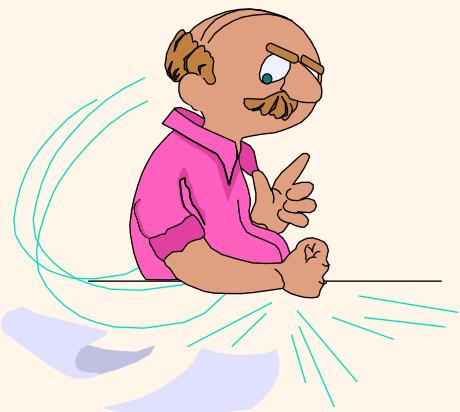
**VH, unhide Section 14: SET 3 STORAGE II**



# COSC 3380

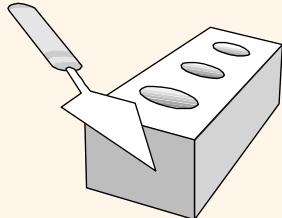
## 4 to 5:30

**PLEASE  
LOG IN  
CANVAS**



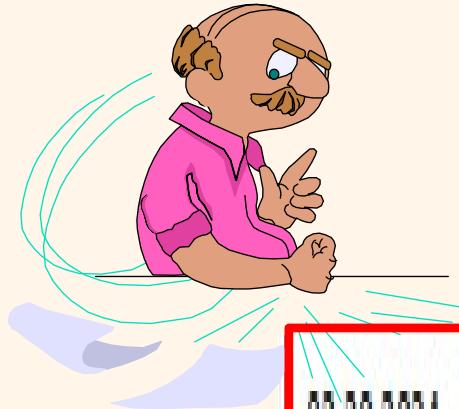
Please close all other windows.

03.20.2024 (17 - We)	ZyBook SET 3 - 2	Set 3 LECTURE 13 STORAGE II
03.25.2024 (18 - Mo)	ZyBook SET 3- 3	Set 3 LECTURE 14 STORAGE III B TREE INDEX
03.27.2024 (19 - We)	ZyBook SET 3- 4	Set 3 LECTURE 15 STORAGE IV HASH INDEX
04.01.2024 (20 - Mo)		EXAM 3 Practice (PART of 20 points)
04.03.2024 (21 - We)	IA Download  ZyBook SET 3 Sections  (4 PM)  (PART of 30 points)	EXAM 3 Review  (PART of 20 points)
04.08.2024 (22 - Mo)		EXAM 3  (PART of 50 points)



# COSC 3380

## Class 17



03.20.2024

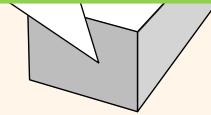
Book SET 3-2

Set 3

17 - Web

LECTURE 13 STORAGE II

**From 4:00 to 4:07 PM – 5 minutes.**



03.20.2024

ZyBook SET 3 · 2

(17 - We)

Set 3

LECTURE 13 STORAGE II

CLASS PARTICIPATION 20 points

20% of Total + :

## STORAGE II

Class 17 BEGIN PARTICIPATION

Not available until Mar 20 at 4:00pm | Due Mar 20 at 4:07pm | 100 pts

VH, publish

This is a synchronous online class.

Attendance is required.

Recording or distribution of class materials is prohibited.

1. At the beginning of selected classes there is an assessment in the first 10 minutes. (beige BOX in the Detailed Syllabus)

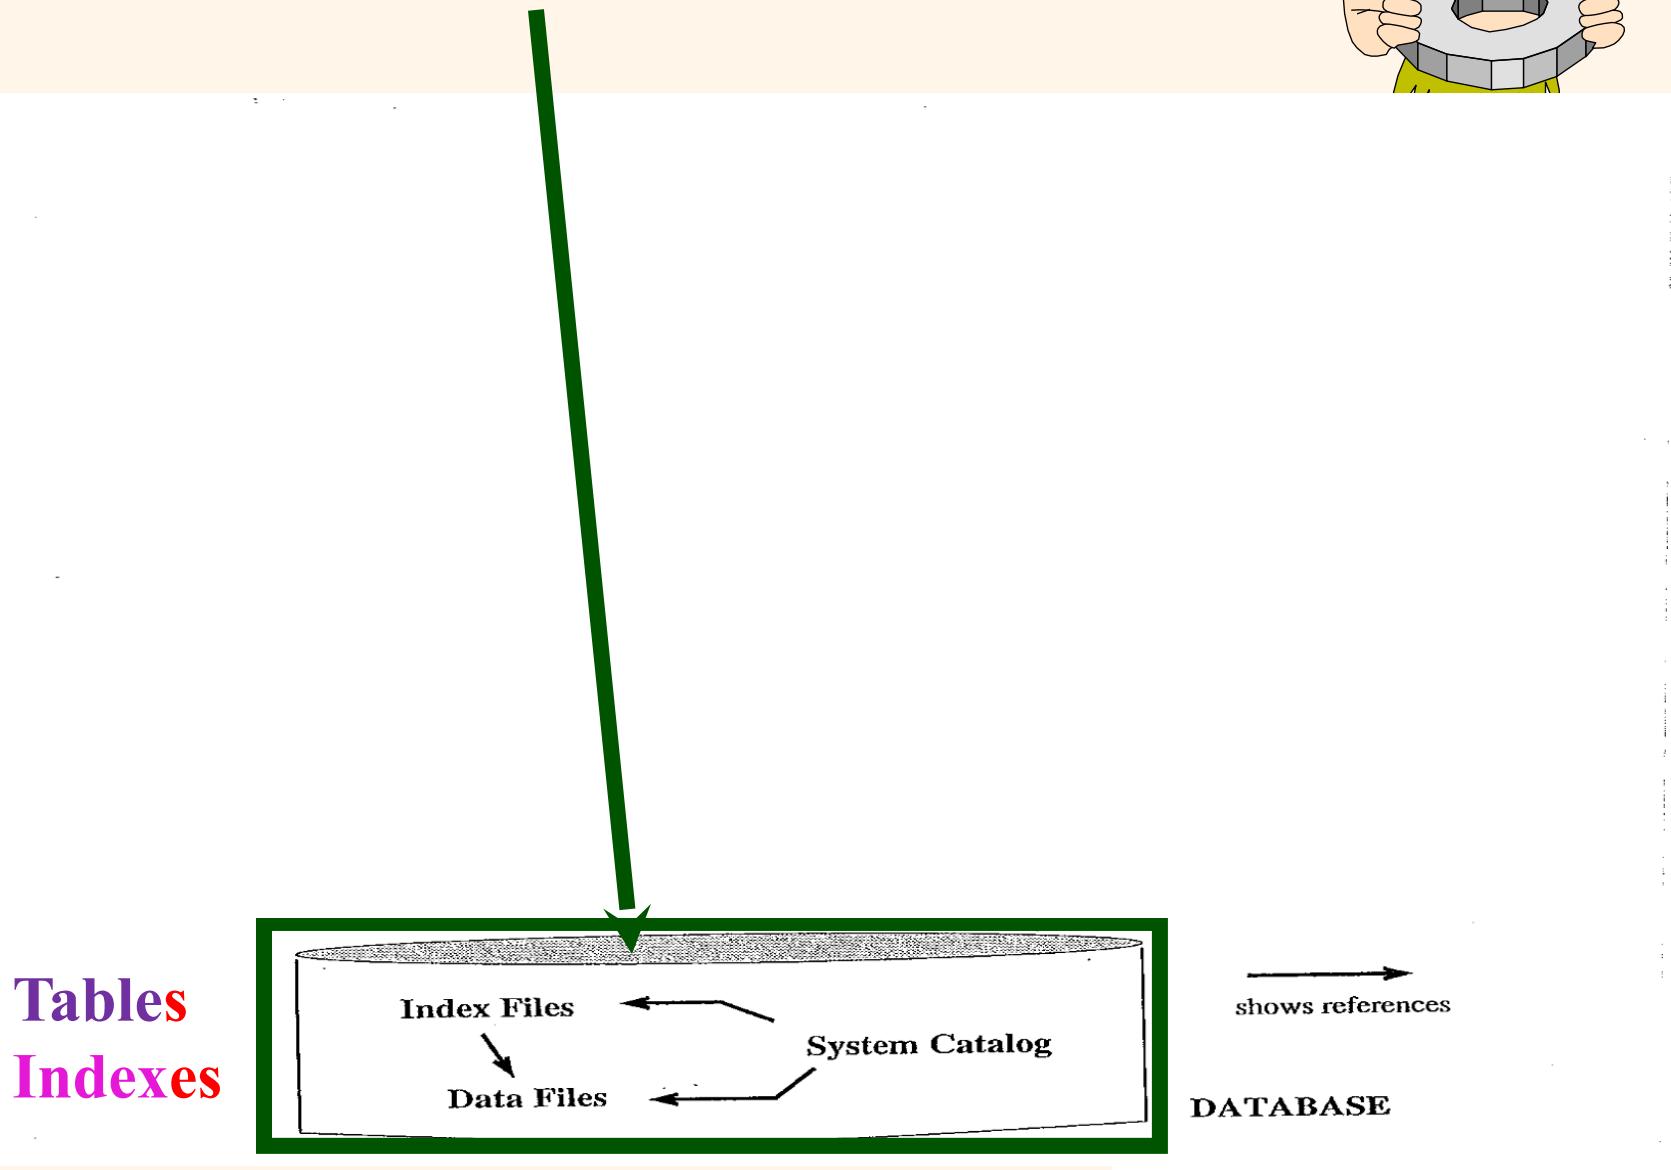
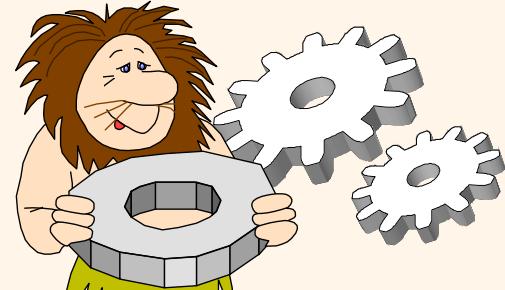
2. At the end of selected classes there is an assessment in the last 10 minutes. (blue BOX in the Detailed Syllabus)

3. ZyBook sections will be downloaded and used for 30% of Total Score on the dates specified in the Detailed Syllabus.

4. EXAMS are in CANVAS. No late EXAMS.

5. I have to be present in TEAMS in order to take any graded assignment assigned during that class.

# A *DBMS*



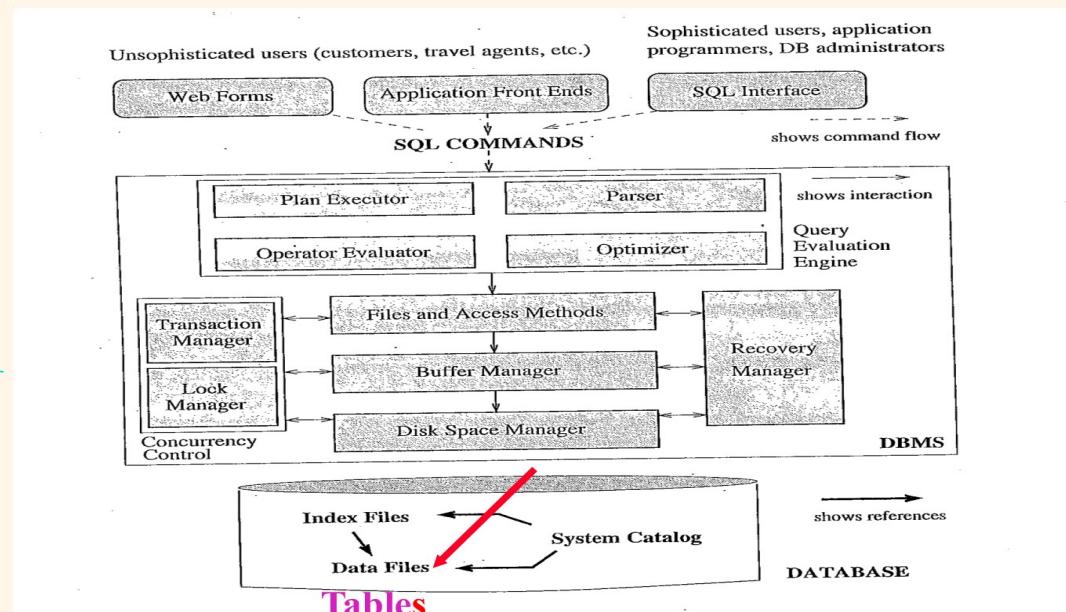
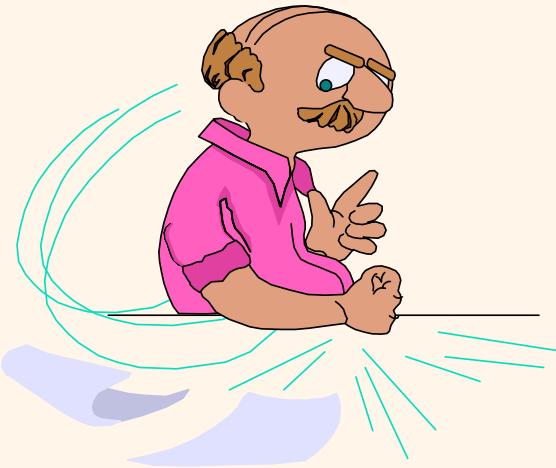
From 4:07 to 5:00 PM – 53 minutes.



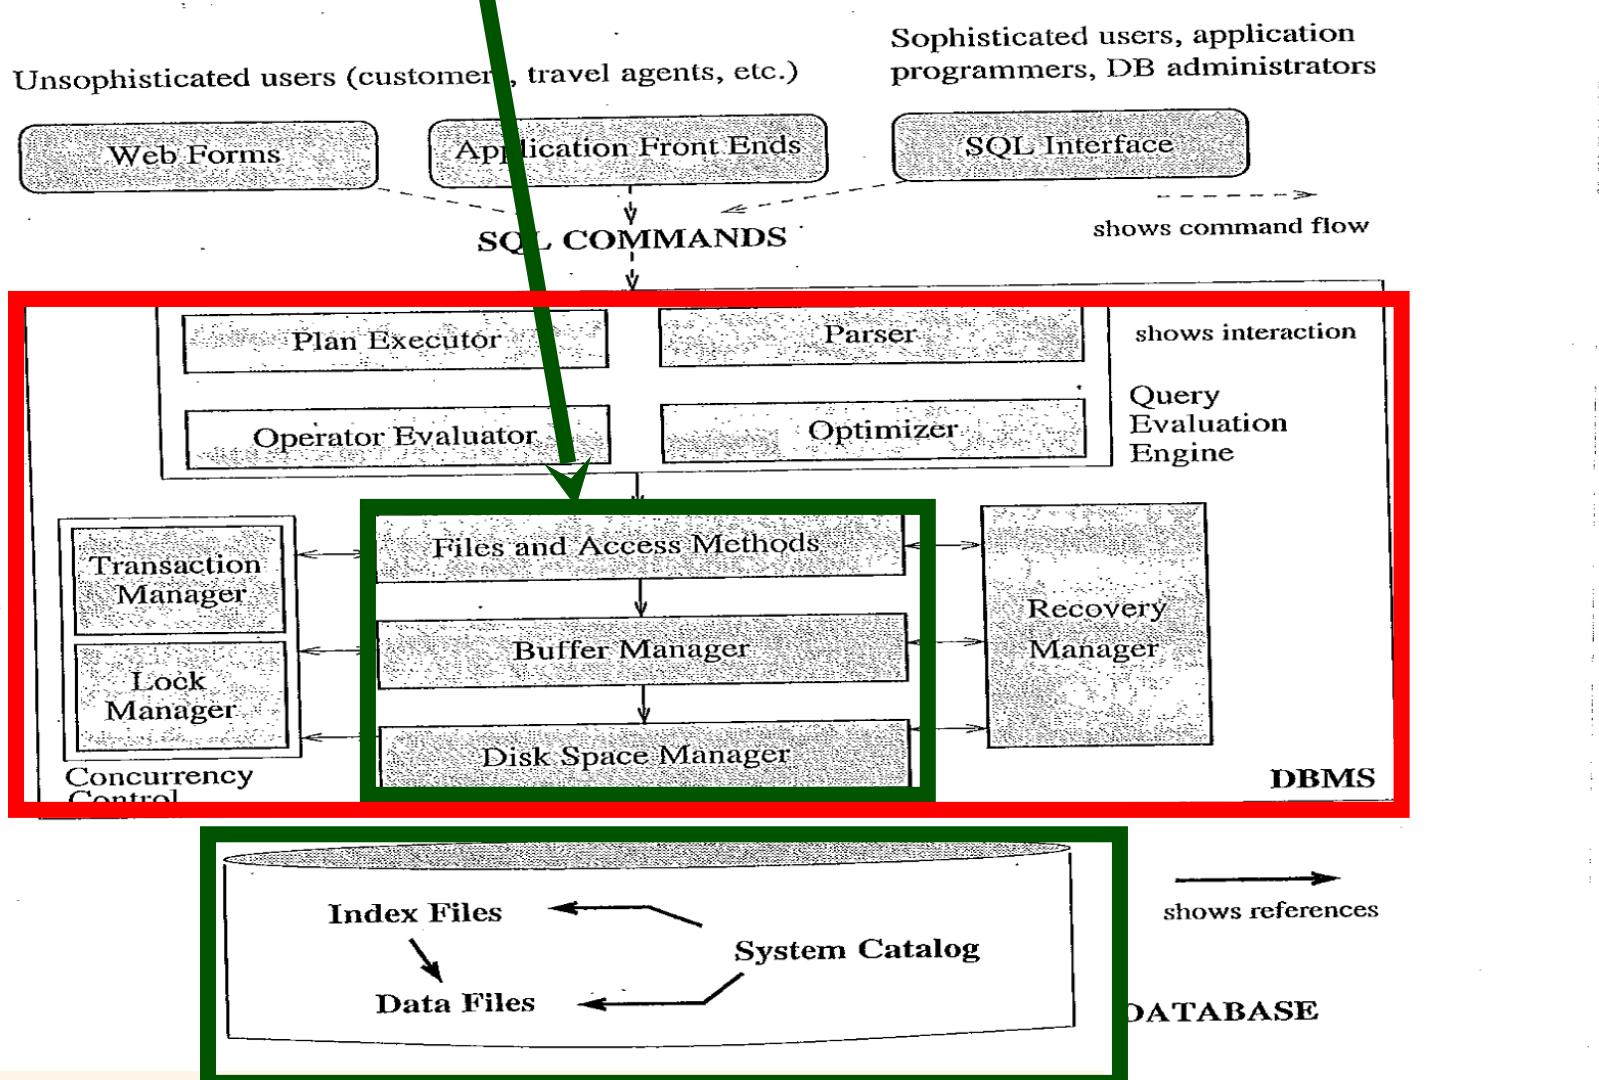
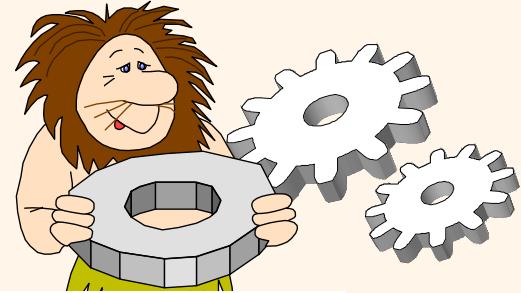
# COSC 3380

## Lecture 13

# *Storing Data: Disks and Files*



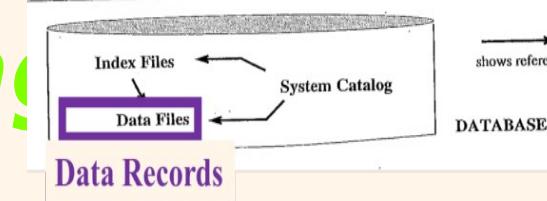
# A DBMS



# Data on External Storage

## Architecture:

- Files and Access Methods processes Pages obtained by asking the Buffer Manager to fetch the Page *rid*.
- Buffer Manager reads Pages from Disk into Memory for processing & writes Pages from Memory to Disk for persistent storage.
- Disk Space Manager allocates/deallocates Disk Pages when the Files and Access Methods needs additional space to hold new Data Records in a File/no longer needs one of its Disk Pages.



# Disk Space Management



Lowest layer of **DBMS** software manages space on Disk.

Higher levels call upon this layer to:

- allocate/de-allocate a **Page**
- read/write a **Page**

Request for a **sequence of Pages** must be satisfied by allocating the **Pages** sequentially on **Disk!**

Higher levels **don't need to know how this is done**, or **how free space is managed**.

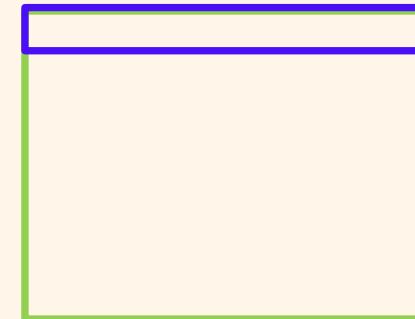


# *Data Record (Tuple/Row) Formats?*

Students(sid:string,name:string,login:string,age:integer,gpa:real)

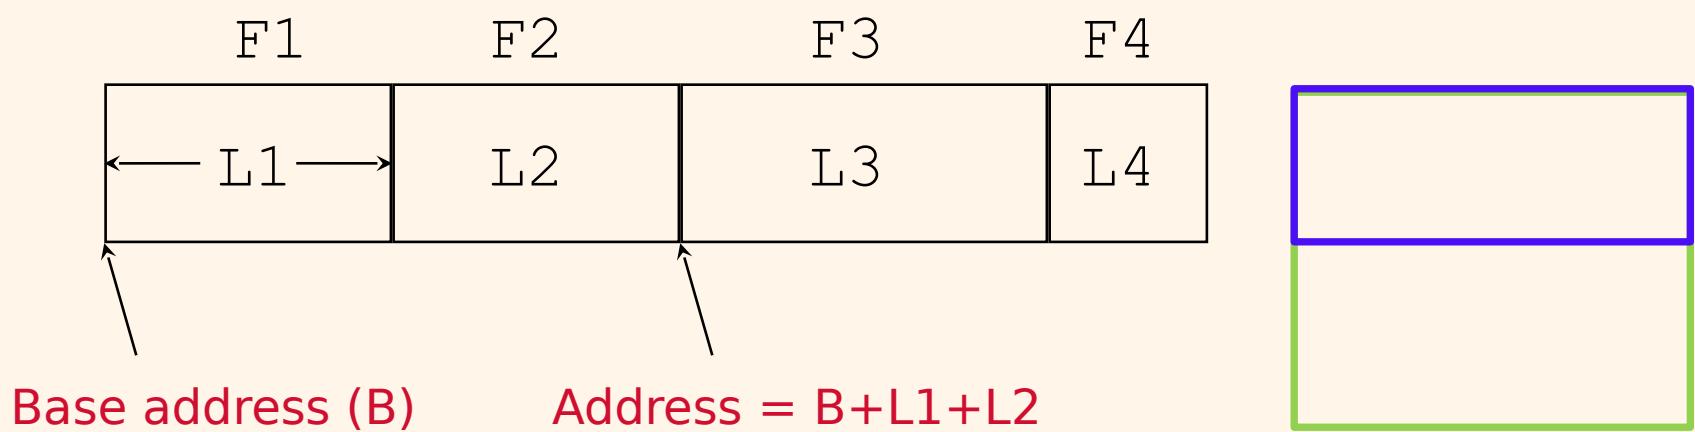
CREATE TABLE Students

(sid: CHAR(20),  
name: CHAR(20),  
login: CHAR(10),  
age: INTEGER,  
gpa: REAL)

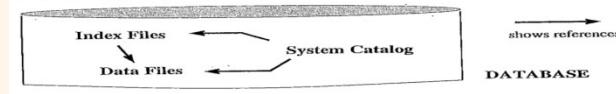


F1	F2	F3	F4	F5
sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eeecs	18	3.2
53650	Smith	smith@math	19	3.8

# Record (tuple) Formats: Fixed Length



- ❖ Information about field types same for **all Data Records** in a **File**; stored in *system Catalogs*.

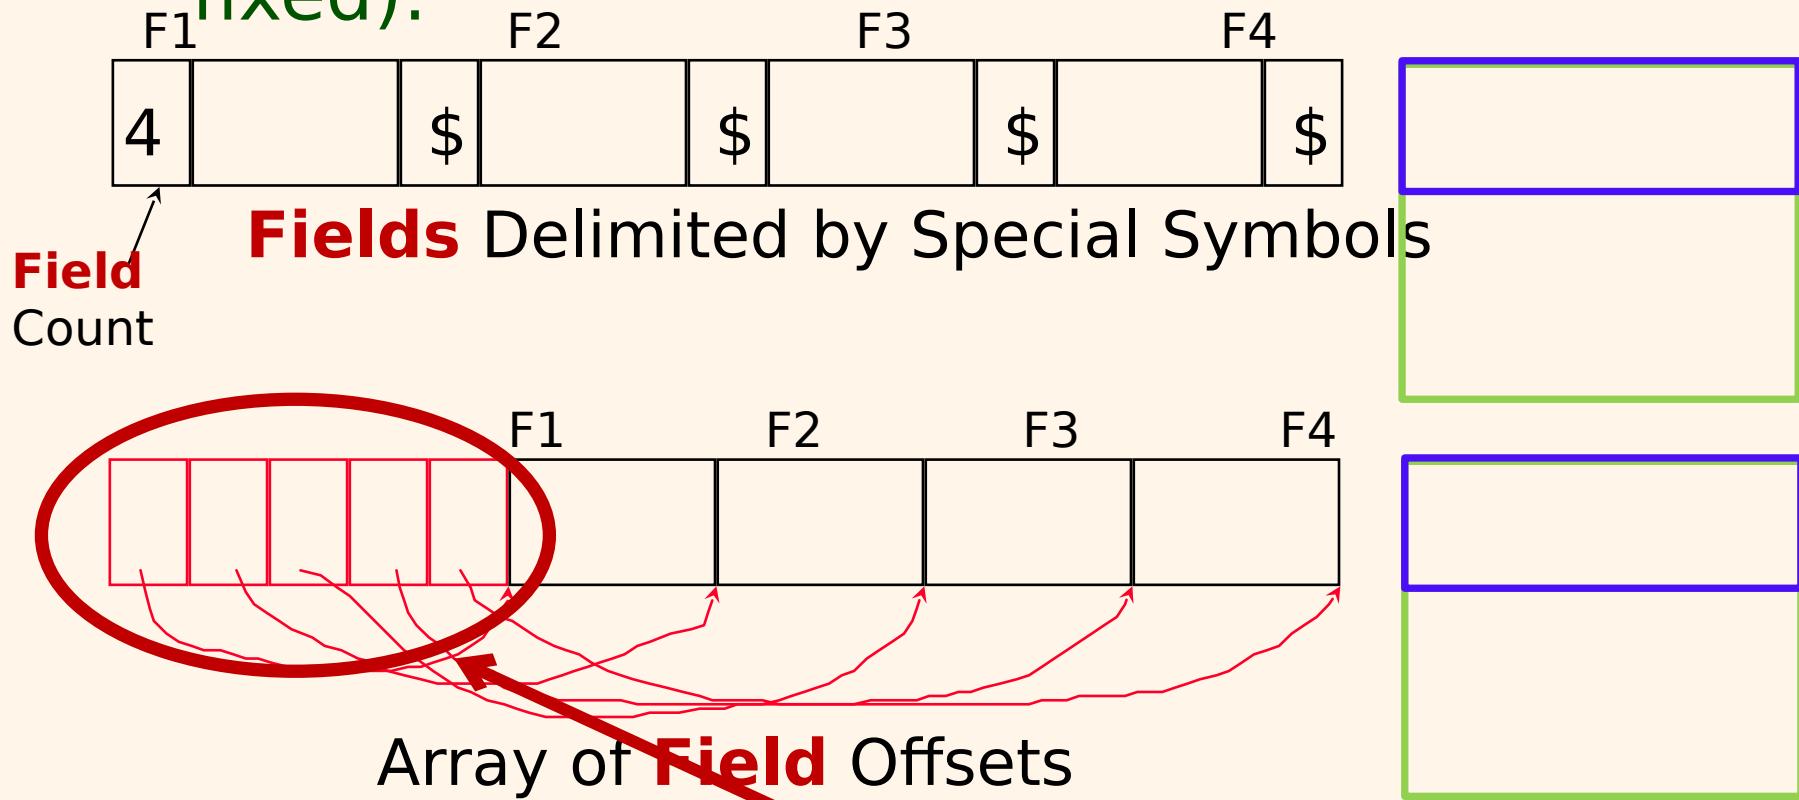


- ❖ Finding ***i'th field*** requires **scan of Data Record**.

# Record (tuple) Formats. **variable Length**



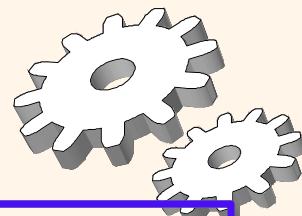
- ❖ Two alternative formats (**# fields** is fixed):



- Second offers direct access to **i'th field**, efficient storage

of **nulls** (special don't know value): **small Directory**

# Transcript File Stored as a Heap File



<i>StudId</i>	<i>CrsCode</i>	<i>Sem</i>	<i>Grade</i>
---------------	----------------	------------	--------------

666666	MGT123	F1994	4.0
123456	CS305	S1996	4.0
987654	CS305	F1995	2.0

B = 3  
R = 4

Page 0

R

717171	CS315	S1997	4.0
666666	EE101	S1998	3.0
765432	MAT123	S1996	2.0
515151	EE101	F1995	3.0

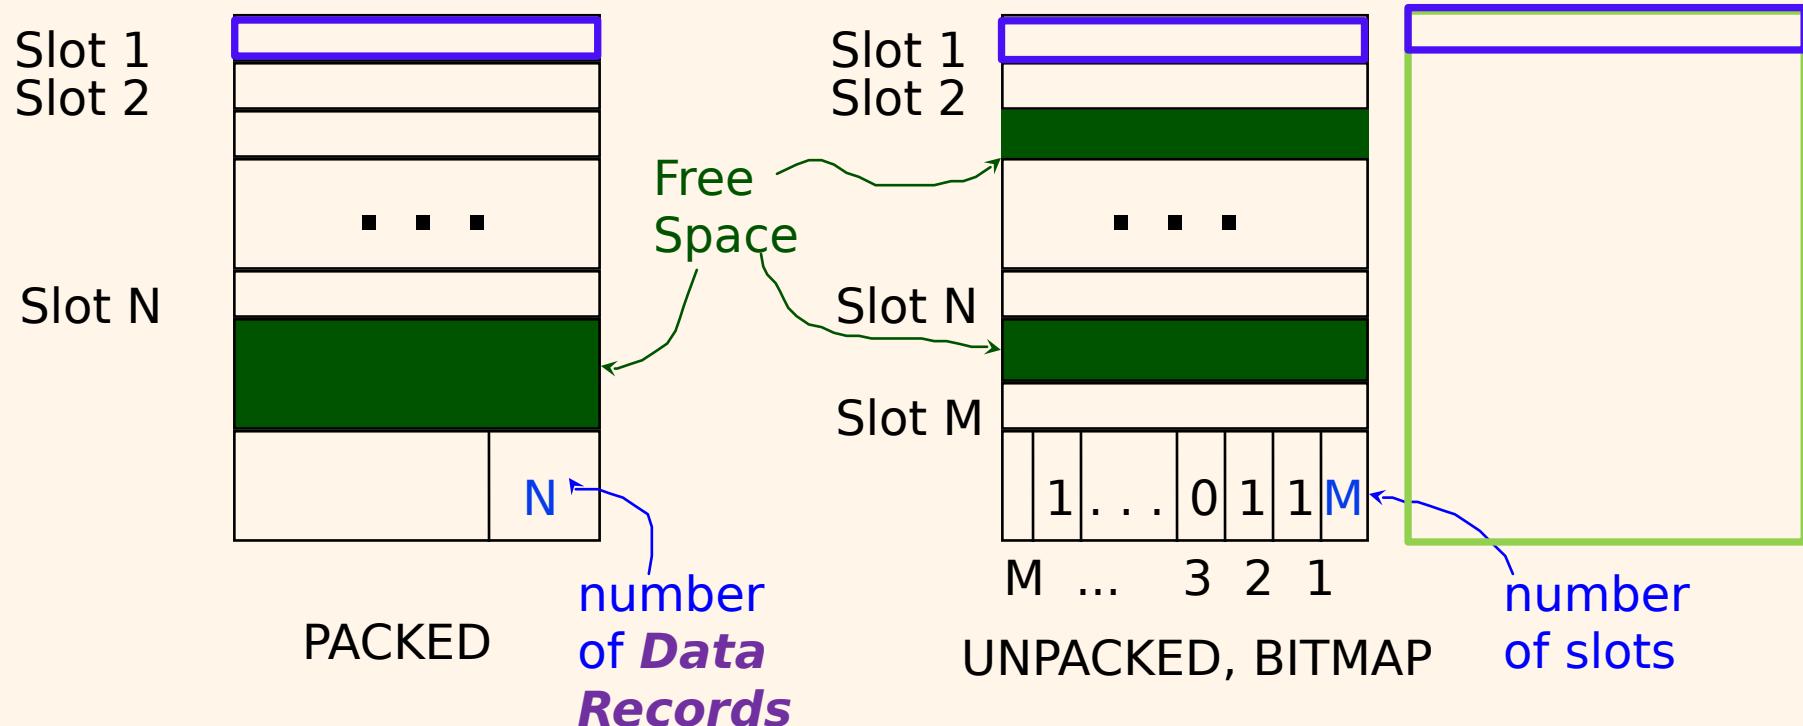
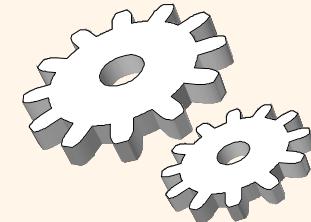
Page 1

234567	CS305	S1999	4.0
878787	MGT123	S1996	3.0

Page 2

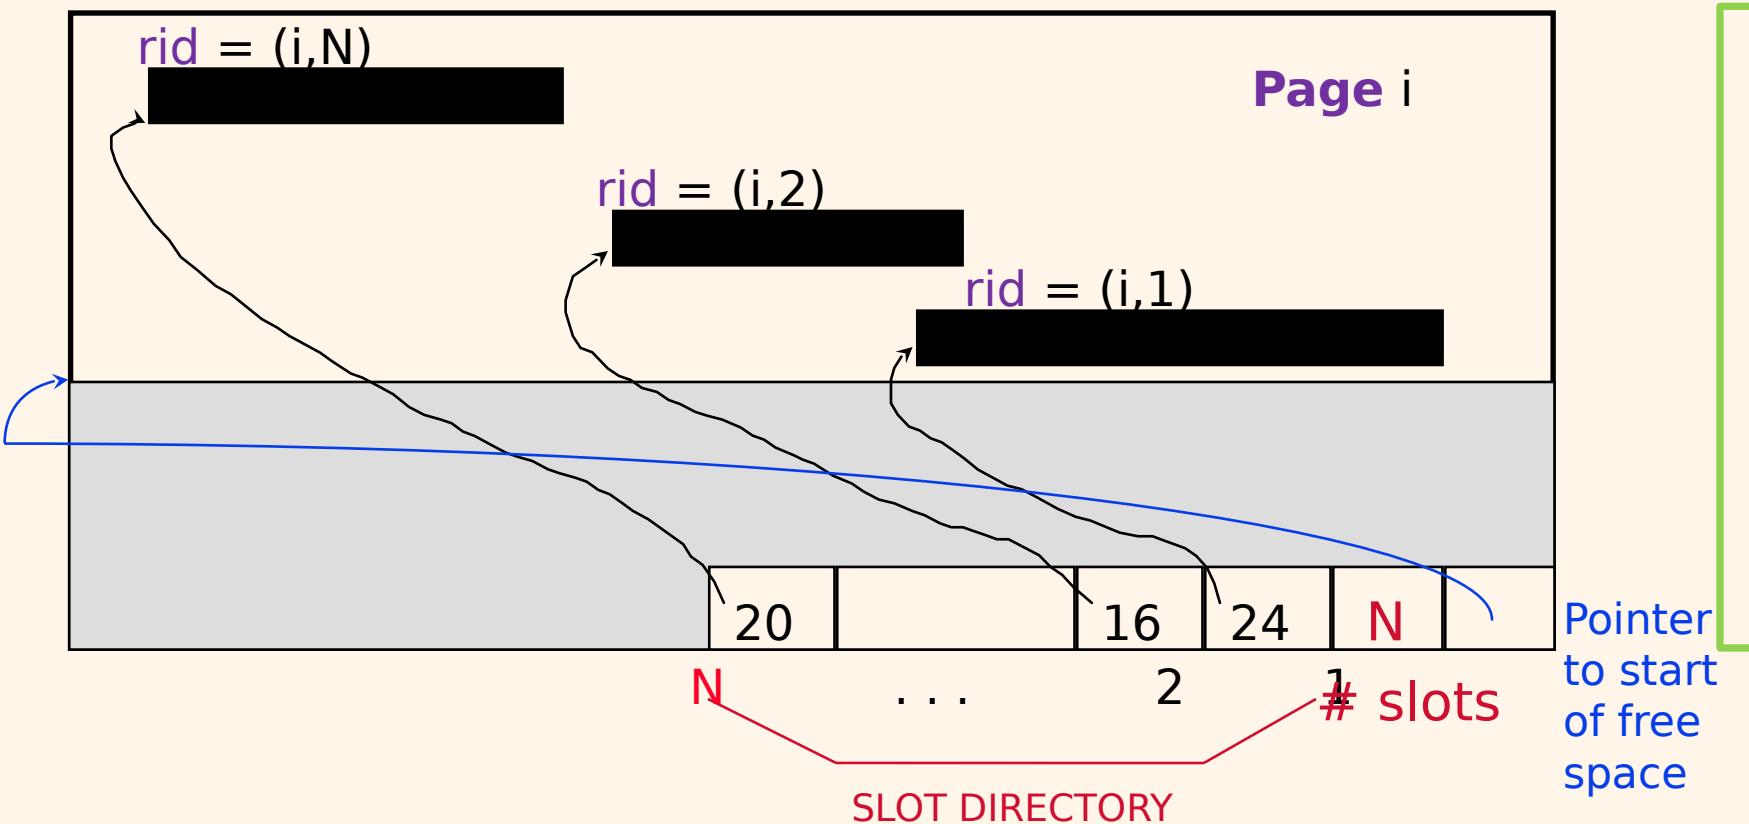
Page Formats?

# Page Formats: Fixed Length Records



- Record id (rid) = <Page id, slot #>. In first alternative, moving **Data Records** for free space management changes  $\text{rid}_{\text{ke}}$ ; may not be acceptable

# **Page Formats: Variable Length Records**



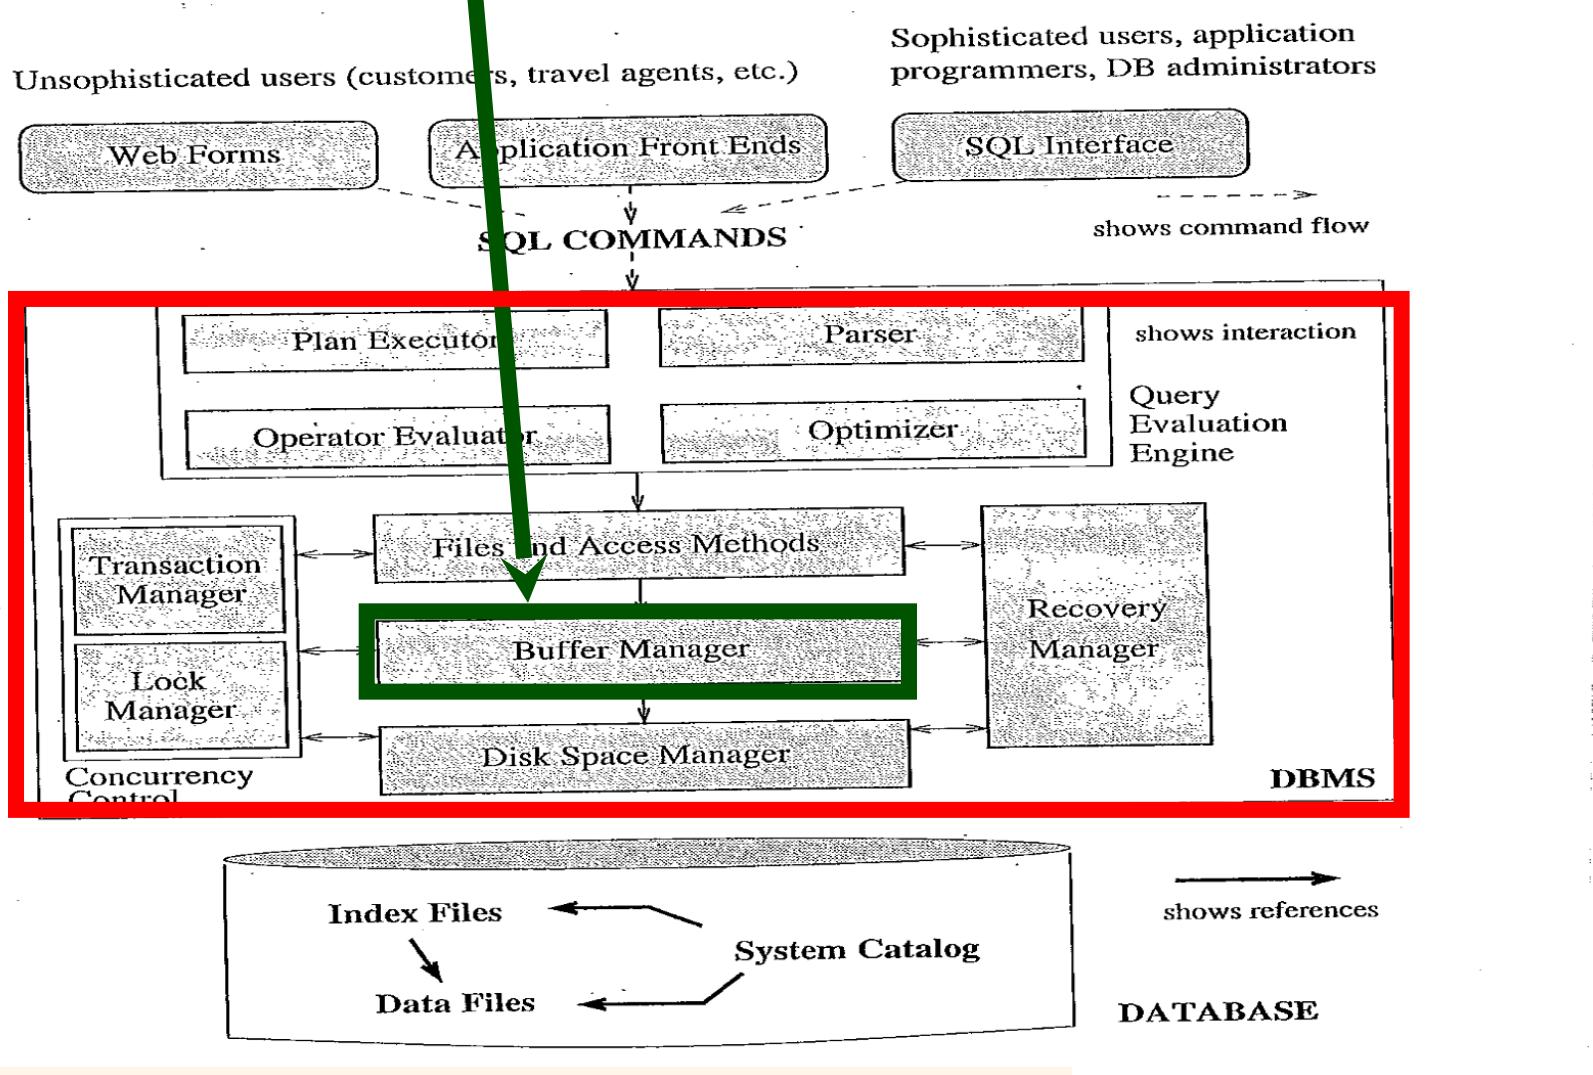
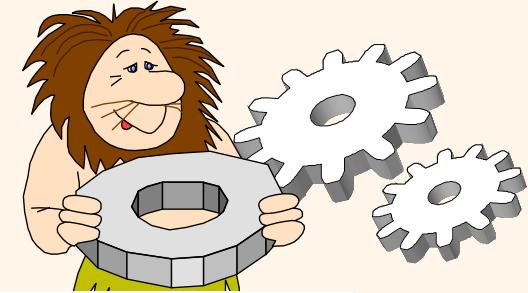
- Can move **Data Records** on **Page** without changing **rid**; so, attractive for **fixed-length**

**TA, Jordan (A – L).**

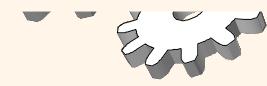
**TA, Fernando (M – Z).**

**Please compare CANVAS vs. TEAMS Attendance.  
Print screens of students in CANVAS but not in the TEAMS meeting.  
(3.20.2024 Attendance X missing LastName.docx)**

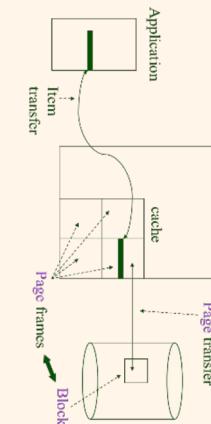
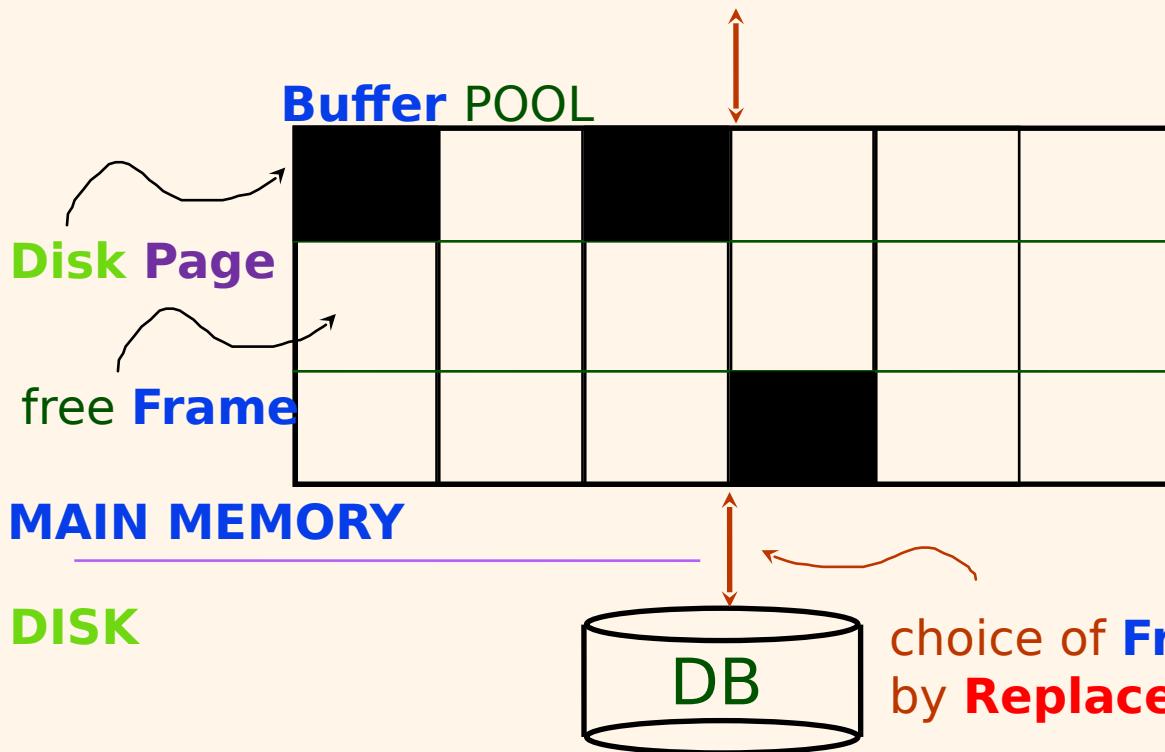
# A DBMS



# Buffer (Frame) Management In a DBMS



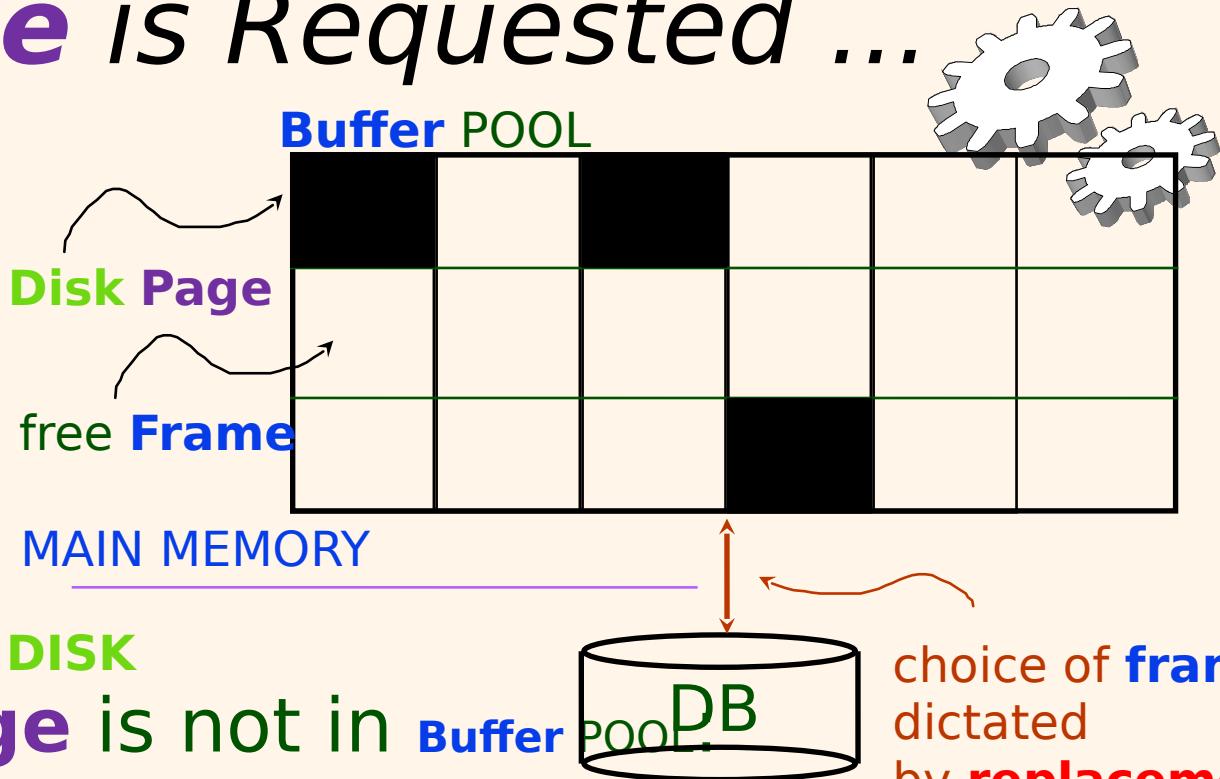
Page Requests from Higher Levels



- ❖ **Data must be in RAM for DBMS to operate on it!**
- ❖ **Table of <frame#, Pageid> pairs is**

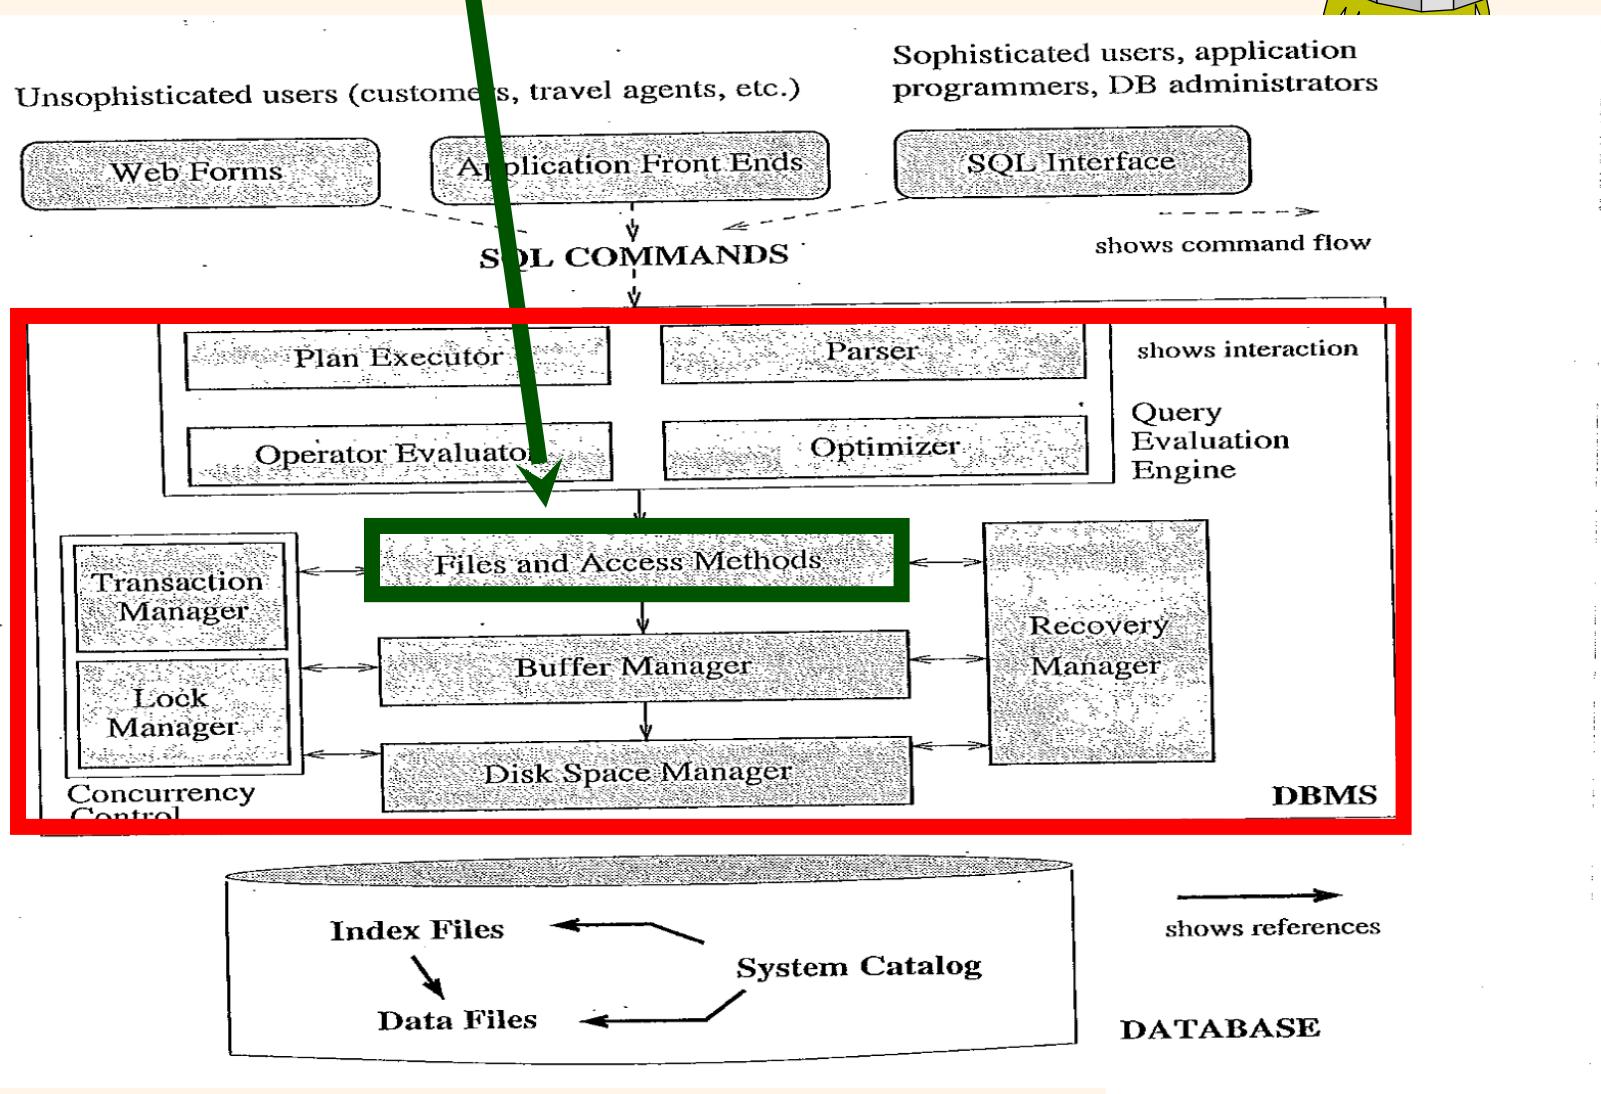
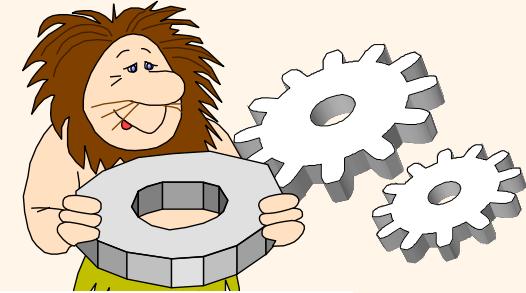


# When a **Page** is Requested ...

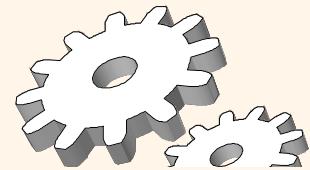


- ❖ If requested **Page** is not in **Buffer POOL**
  - Choose a **Frame** for *replacement*
  - If **Frame** is dirty, Write it to **Disk**
  - Read requested **Page** into chosen **Frame**
- ❖ *Pin the Page* and return its address.
  - If requests can be predicted (e.g., **sequential scans**) **Pages** can be pre-fetched several **Pages** at a time!

# A DBMS



# Files of Data Records (Tuples/Rows)



**Page** or Block is OK when doing I/O, but higher levels of **DBMS** operate on **Data Records** (Tuples), and **Files of Data Records** (Tuples).

**FILE**: A collection of **Pages** (B), each containing a **collection of Data Records (Tuples) (R)**. Must support:

- insert/delete/modify **Data Record**
- read a particular **Data Record** (specified using *record id rid*)
- scan all **Data Records** (possibly with some records to be retrieved)

sid	name	login	age	gpa
53831	Madayan	madayan@music	11	1.8
53832	Guldu	guldu@music	12	2.0
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	19	3.2
53650	Smith	smith@math	19	3.8

<1,1>  
<1,2>  
<1,3>  
<2,1>  
<2,2>



# *Unordered (**Heap**) Files of Data Records*

Simplest **File** structure contains **Data Records** in no particular order.

As **File** grows and shrinks, Disk **Pages** are allocated and deallocated.

To support **Data Record** level operations, we must:

- keep track of **the Pages** in a **File (B)**
- keep track of *free space* on **Pages**
- keep track of the **Data Records** on a **Page (R)**

sid	name	login	age	gpa
53831	Madayan	madayan@music	11	1.8
53832	Guldu	guldu@music	12	2.0
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	19	3.2
53650	Smith	smith@math	19	3.8

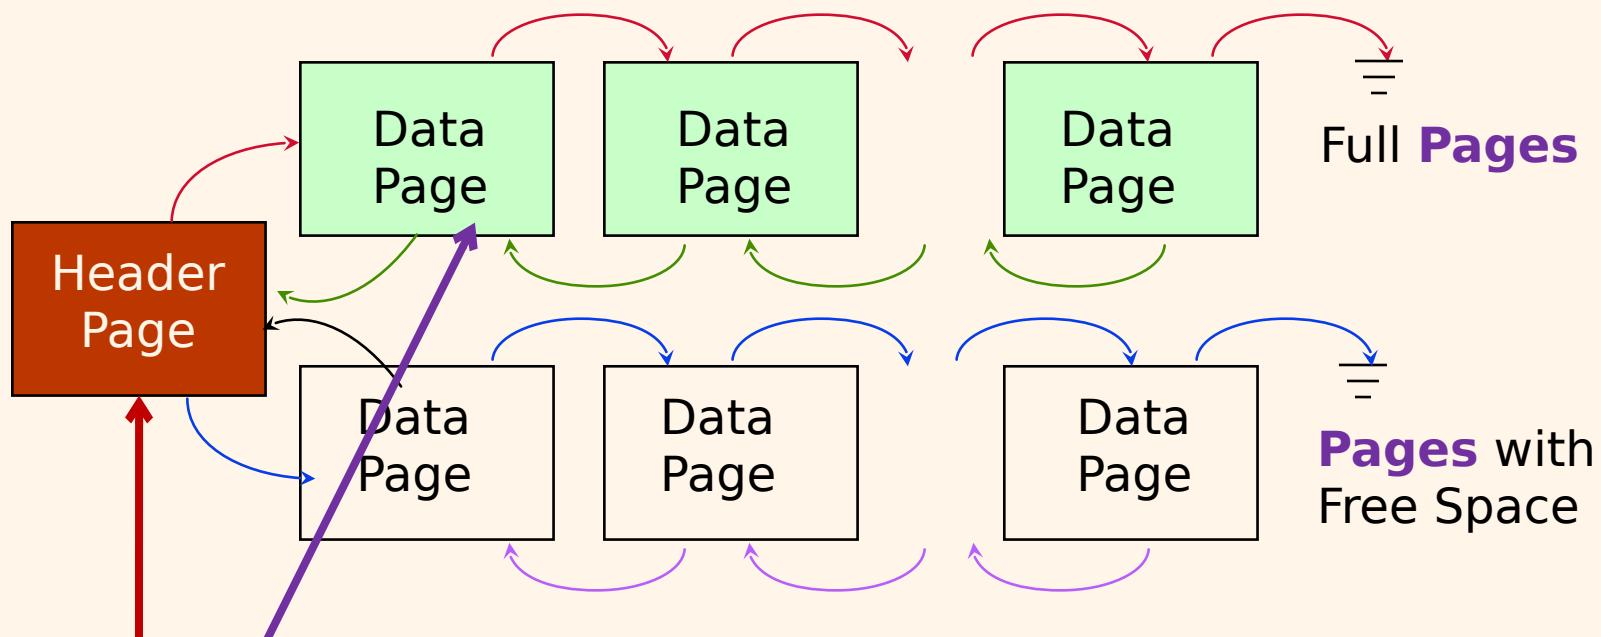
<1,1>  
<1,2>  
<1,3>  
<2,1>  
<2,2>

or keeping track of **Pages (B)**.

# **Heap File implemented as a Linked based LIST**

**(I loved my Data Structures**

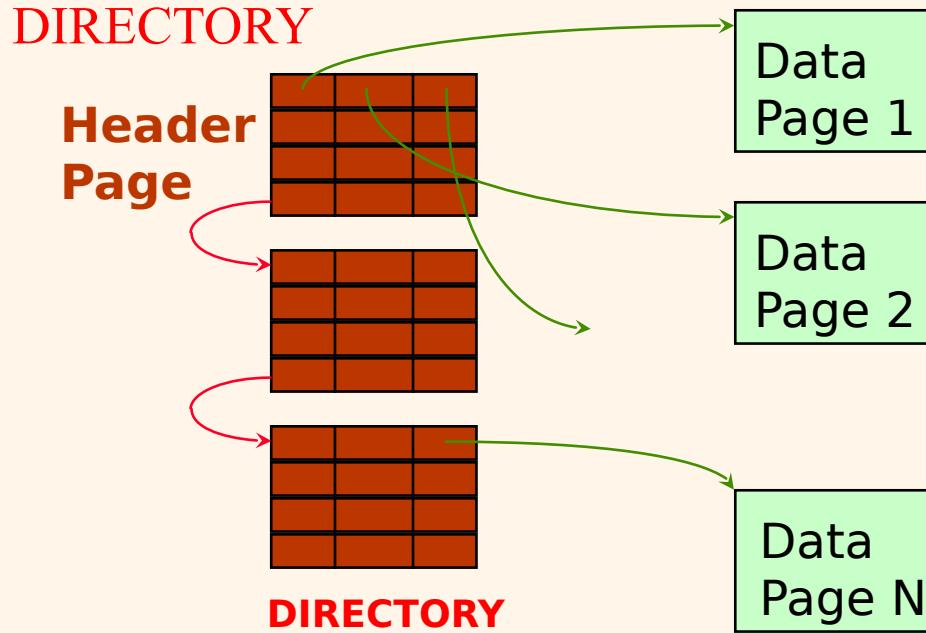
**Class!**)



The **Header Page id** and **Heap File name** must be stored someplace.

Each **Data Page** contains 2 `pointers' plus **Data Records**

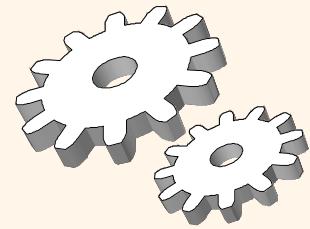
# Heap File Using a Page DIRECTORY



The **Header** entry for a **Data Page** can include the **number of free bytes** on the **Page**.

DIRECTORY is a collection of Pages; Linked LIST implementation is just one alternative.

- Much smaller than **Linked based List** of all **Heap File Pages**



# System **Catalogs**

- ❖ For each **Index**:
  - structure (e.g., **B+ tree**) and **search key** fields
- ❖ For each **Relation/Table**:
  - **name**, **File name**, **File** structure (e.g., **Heap File**)
  - attribute name and type, for each attribute
  - index name, for each **Index**
  - integrity constraints
- ❖ For each **View**:
  - view name and definition
- ❖ Plus Statistics, Authorization, **Buffer Pool Size**, etc.
- **Catalogs are themselves stored as Relations!**



## □ 14. SET 3: 2:STORAGE II



### □ 14.1 Physical design Hidden



### □ 14.2 Single-level indexes Hidden



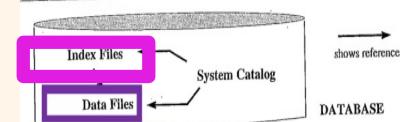
### □ 14.3 LAB - Create index and explain (Sakila) Hidden



### □ 14.4 LAB - Query execution plans (Sakila) Hidden



# Indexes - Data Structures



key,

$r_{id}$

- **Ordered** set of values that contains Index search key and pointers
- **More efficient** to use Index to access **Table** than to scan all Rows in **Table** sequentially

key,  $r_{id}$

STATE INDEX

Key	Row
AZ	1
....	....
....	....
FL	1
FL	7
FL	8
FL	13245
FL	14786
....	....
....	....

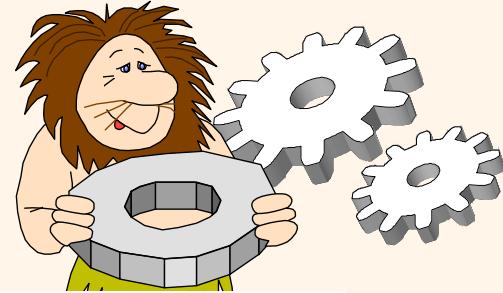
CUSTOMER TABLE  
(14,786 rows)

Row #	CUS_CODE	CUS_LNAME	CUS_FNAME	CUS_INITIAL	CUS_AREACODE	CUS_PHONE	CUS_STATE	CUS_BALANCE
1	10010	Ramas	Alfred	A	615	844-2573	FL	\$0.00
2	10011	Dunne	Leona	K	713	894-1238	AZ	\$0.00
3	10012	Smith	Kathy	W	615	894-2285	TX	\$345.86
4	10013	Olowski	Paul	P				
5	10014	Orlando	Myron	M				
6	10015	O'Brian	Amy	E				
7	10016	Brown	James	J				
8	10017	Williams	George	G				
9	10018	Farriss	Anne	A				
10	10019	Smith	Olette	O				
....	.....	.....	.....	.....	.....	.....	.....	.....
....	.....	.....	.....	.....	.....	.....	.....	.....
13245	23120	Veron	George	G				
....	.....	.....	.....	.....	.....	.....	.....	.....
....	.....	.....	.....	.....	.....	.....	.....	.....
14786	24560	Suarez	Victor	V				

▪ Indexes: Data structures to organize records via trees or hashing.

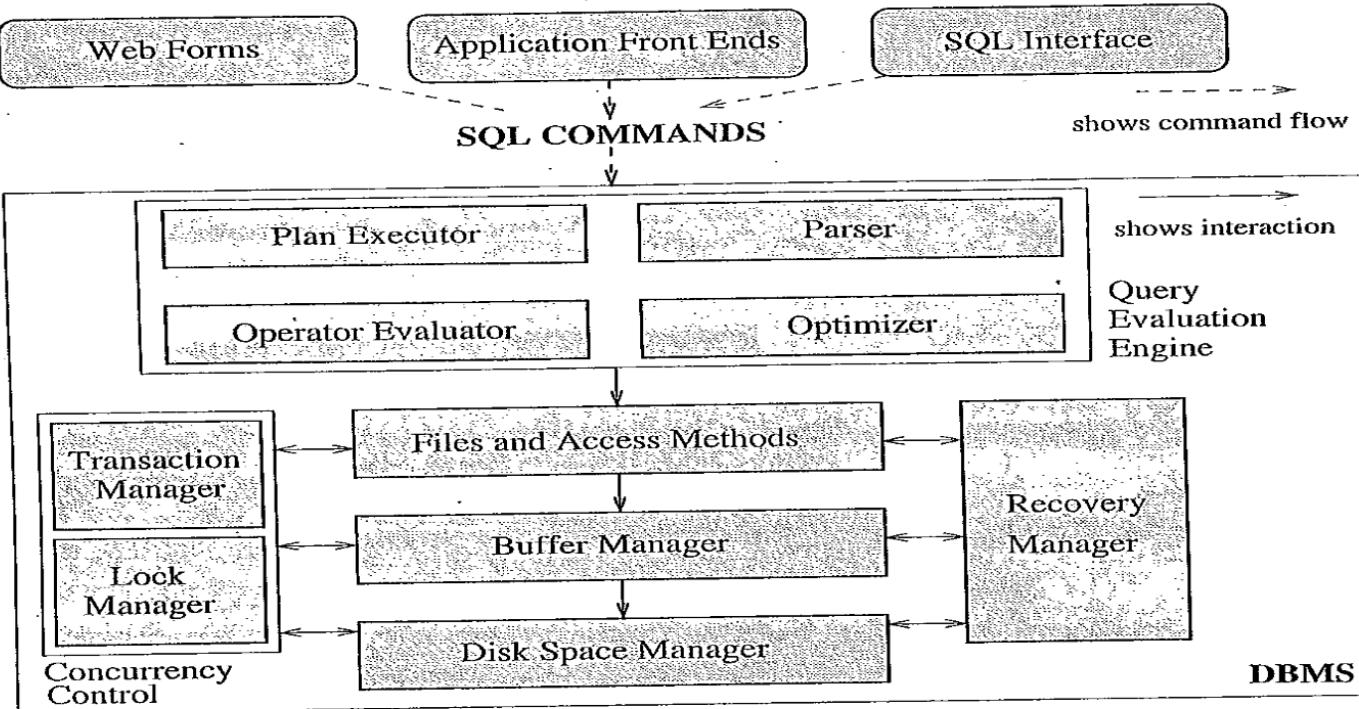
• Like Sorted Files, they speed up searches for a subset of records, based on values in certain ("search key") fields

# A DBMS



Unsophisticated users (customers, travel agents, etc.)

Sophisticated users, application  
programmers, DB administrators



## Dense Index

# index entries = # of data records

dense sorted  
index

table

B

N	block 1	ATL	40	Aer Lingus	10:30	DUB	Boeing 737
		DME	552	Lufthansa	18:00	FRA	MD 111
		DUB	666	United Airlines	6:00	ATL	Airbus 321
		DXB	698	Air India	3:15	DXB	Boeing 747
		FRA	702	American Airlines	5:05	ORD	Airbus 323
		HKG	739	Air China	3:20	DME	Boeing 777
	block 0		803	British Airways	13:50	PHL	Airbus 323
			971	Air India	19:35	LAX	Boeing 747
			1001	Southwest Airlines	5:05	SJC	Airbus 323
			1107	Air China	3:20	HKG	Boeing 767
			1154	American Airlines	13:50	ORD	MD 111
			1222	Delta Airlines	8:25	SEA	Boeing 747
			1291	Lufthansa	4:10	PIT	Boeing 777

## □ 14. SET 3: 2:STORAGE II



### □ 14.1 Physical design Hidden



### □ 14.2 Single-level indexes Hidden



### □ 14.3 LAB - Create index and explain (Sakila) Hidden



### □ 14.4 LAB - Query execution plans (Sakila) Hidden



# 14.1 Physical design

## MySQL storage engines

**Logical design** specifies tables, columns, and keys. The logical design process is described elsewhere in this material. **Physical design** specifies indexes, table structures, and partitions. Physical design affects query performance but never affects query results.

A **storage engine** or **storage manager** translates instructions generated by a query processor into low-level commands that access data on storage media. Storage engines support different index and table structures, so physical design is dependent on a specific storage engine.

MySQL can be configured with several different storage engines, including:

- InnoDB is the default storage engine installed with the MySQL download. InnoDB has full support for transaction management, foreign keys, referential integrity, and locking.
- MyISAM has limited transaction management and locking capabilities. MyISAM is commonly used for analytic applications with limited data updates.
- MEMORY stores all data in main memory. MEMORY is used for fast access with databases small enough to fit in main memory.

Different databases and storage engines support different table structures and index types. Ex:

- Table structure. Oracle Database supports heap, sorted, hash, and cluster tables. MySQL with InnoDB supports only heap and sorted tables.
- Index type. MySQL with InnoDB or MyISAM supports only B+tree indexes. MySQL with MEMORY supports both B+tree and hash indexes.

This section describes the physical design process and statements for MySQL with InnoDB. The process and statements can be adapted to other databases and storage engines, but details depend on supported index and table structures.

PARTICIPATION  
ACTIVITY

14.1.1: Logical and physical design.

B

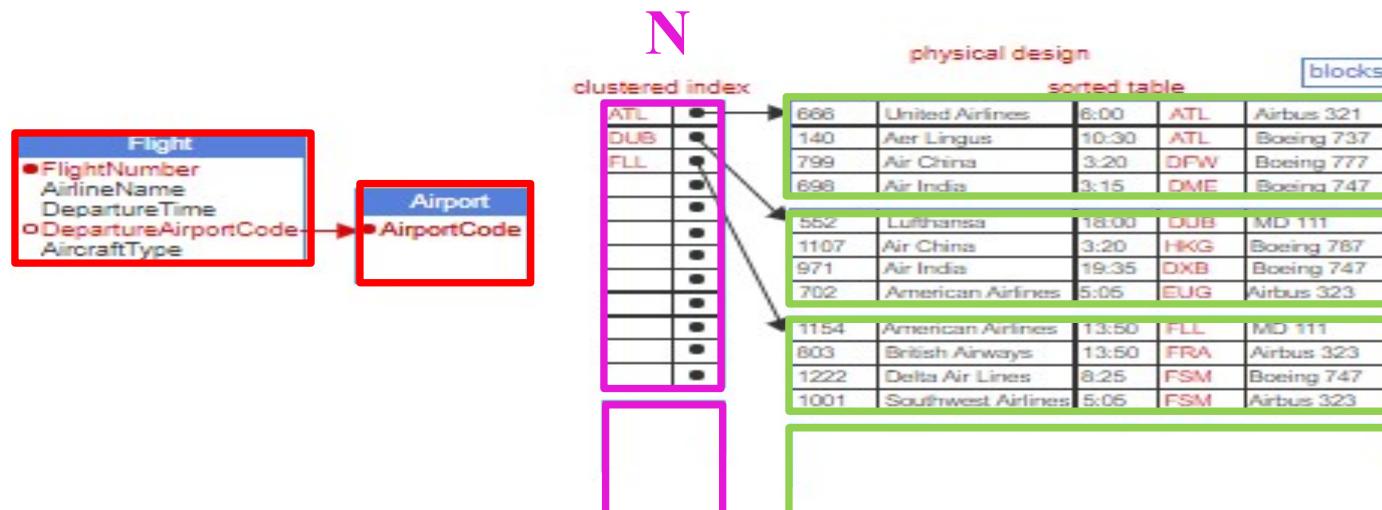
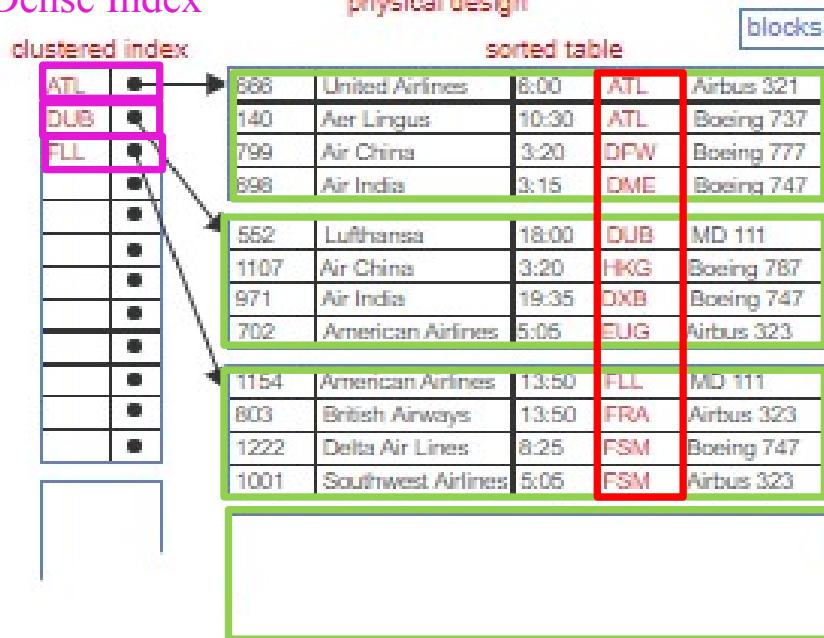


Table Sorted on search key

# index entries = # of blocks

### Dense Index



B

N

### Next two classes

The **CREATE INDEX** statement creates an index by specifying the index name and table columns that compose the index. Most indexes specify just one column, but a composite index specifies multiple columns.

The **DROP INDEX** statement deletes a table's index.

The **SHOW INDEX** statement displays a table's index. SHOW INDEX generates a result table with one row for each column of each index. A multi-column index has multiple rows in the result table.

The SQL standard includes logical design statements such as CREATE TABLE but not physical design statements such as CREATE INDEX. Nevertheless, CREATE INDEX and many other physical design statements are similar in most relational databases.

Table 14.1.1: INDEX statements.

Statement	Description	Syntax
CREATE INDEX	Create an index	<code>CREATE INDEX IndexName ON TableName (Column1, Column2, ..., ColumnN);</code>
DROP INDEX	Delete an index	<code>DROP INDEX IndexName ON TableName;</code>
SHOW INDEX	Show an index	<code>SHOW INDEX FROM TableName;</code>

## 14. SET 3: 2:STORAGE II



### 14.1 Physical design Hidden



### 14.2 Single-level indexes Hidden



#### ■ Participation activities

- |   |  |    |
|---|--|----|
| 14.2.1: Single-level index.                         |  | 0% |
| 14.2.2: Single-level indexes.                       |  | 0% |
| 14.2.3: Query processing with single-level index.   |  | 0% |
| 14.2.4: Query processing with single-level index... |  | 0% |
| 14.2.5: Binary search on a sorted index.            |  | 0% |
| 14.2.6: Binary search.                              |  | 0% |
| 14.2.7: Dense and sparse indexes.                   |  | 0% |
| 14.2.8: Primary, clustering, and secondary index... |  | 0% |
| 14.2.9: Insert with dense sorted index.             |  | 0% |
| 14.2.10: Inserts, updates, and deletes.             |  | 0% |

#### ■ Challenge activities

- |                               |  |    |
|-------------------------------|--|----|
| 14.2.1: Single-level indexes. |  | 0% |
|-------------------------------|--|----|

### 14.3 LAB - Create index and explain (Sakila) Hidden

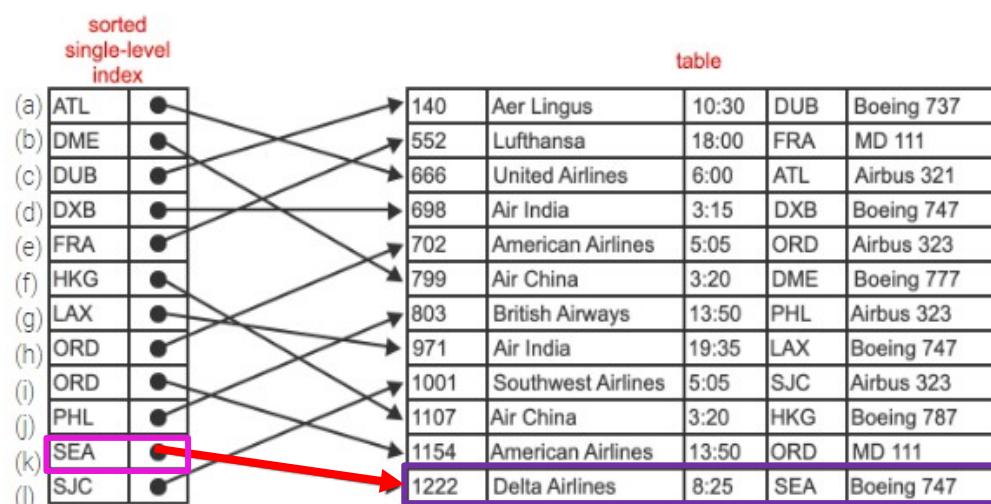


### 14.4 LAB - Query execution plans (Sakila) Hidden



CHALLENGE  
ACTIVITY

## 14.2.1: Single-level indexes.



Select the index entry that refers to the row for Delta Airlines flight 1222.

Pick ▾

k

1

?????

CHALLENGE  
ACTIVITY

## 14.2.1: Single-level indexes.

## Index

sorted single-level index

	1622 Alaska Airlines PEK	table
ATL	140	Aer Lingus
DME	552	Lufthansa
DUB	666	United Airlines
DXB	698	Air India
FRA	702	American Airlines
HKG	799	Air China
LAX	803	British Airways
ORD	971	Air India
ORD	1001	Southwest Airlines
PHL	1107	Air China
SEA	1154	American Airlines
SJC	1222	Delta Airlines

PEK

Alaska Airlines flight 1622 departing from PEK is inserted into the table. Where does the new index entry go?

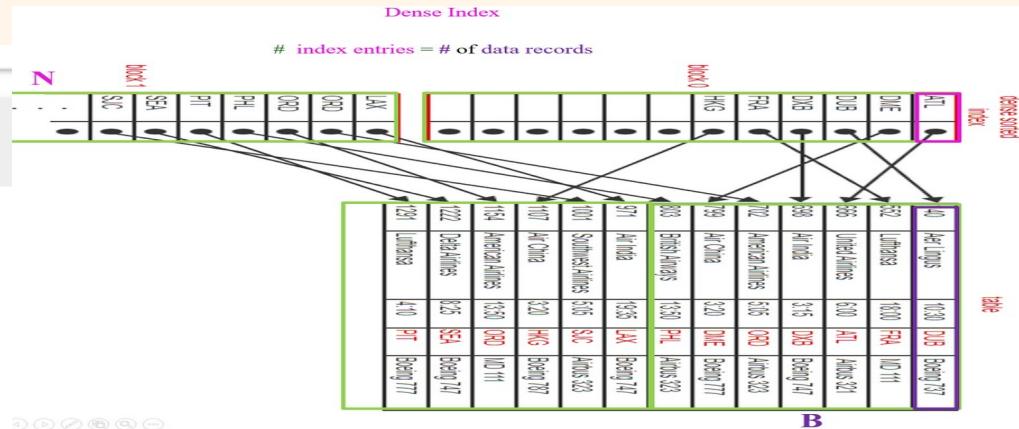
Pick Between ORD and PHL

2

?????

## CHALLENGE ACTIVITY

## 14.2.1: Single-level indexes.



$$B = 5,000 \text{ blocks}$$

A table occupies 5000 blocks. FlightNumber is the primary key. A single-level index on FlightNumber occupies 500 blocks. The WHERE clause of a SELECT specifies "FlightNumber = 2780".

$$N = 500 \text{ blocks}$$

$$500 + 1$$

What is the maximum number of blocks necessary to process the SELECT?

Ex: 5 **501** blocks

- 3
- 4
- 5
- 6

Assume the table has no index. What is the maximum number of blocks necessary to process the SELECT?

Ex: **5,000** blocks

B

CHALLENGE  
ACTIVITY

## 14.2.1: Single-level indexes.

Consider the following scenario:

- A table has 400,000,000 rows      400,000,000 rows
- Each row is 100 bytes      100 bytes / row
- Magnetic disk transfer rate is 0.1 gigabytes per second      0.1 gigabytes / sec
- Assume 1 gigabyte is approximately 1,000,000,000 bytes      1 gigabyte = 1,000,000,000 bytes

Assuming no free space, a table scan requires approximately how many seconds?

Table scan

Ex: 1234      seconds

$$400,000,000 \text{ rows} * 100 \text{ bytes / row} = 40000000000 \text{ bytes} = 40 \text{ gigabytes}$$

$$40 \text{ gigabytes} * 0.1 \text{ gigabytes / sec} = 400 \text{ sec}$$

4

750,000



B-blocks of data records

Consider the following scenario:

- A table is sorted on indexed column and has 10,000,000 rows 10,000,000 rows
- Each row is 300 bytes 300 bytes / row
- Table and index blocks are 4 kilobytes 1 block = 4,000 bytes / block
- Assume 1 kilobyte is approximately 1,000 bytes

Assuming the index is sparse how many entries are in the index?

Ex: 1234567 entries

# index entries = B blocks of data records/ rows for the Table

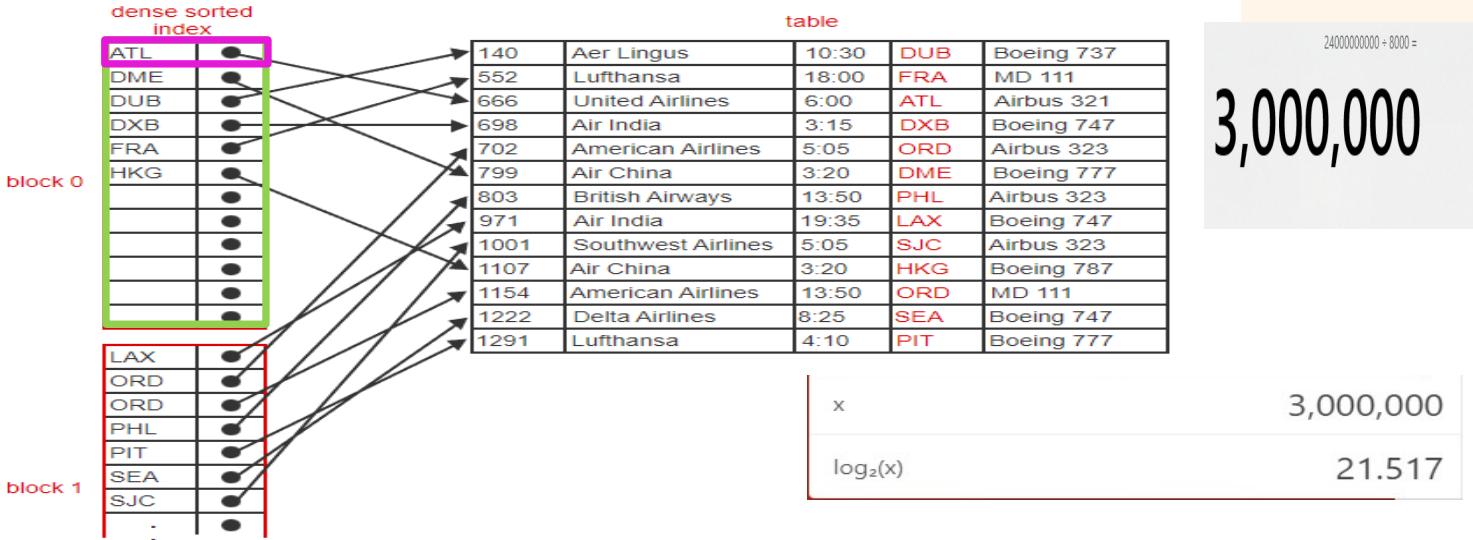
$$\text{Total number of bytes} = 10,000,000 \text{ rows} * 300 \text{ bytes / row} = 3,000,000,000 \text{ bytes}$$

$$\begin{aligned} \text{B-blocks of data records} &= 3,000,000,000 \text{ bytes} / 4,000 \text{ bytes / block} \\ &= 750,000 \text{ blocks} \end{aligned}$$

5

# index entries = 750,000

?????



N –blocks of data entries k\*

$$\log_2 3,000,000 = 22 \text{ blocks}$$

Consider the following scenario:

- A table has 600,000,000 rows. **600,000,000 rows**
- Each index block is 8 kilobytes. **1 block = 8,000 bytes / block**
- Each index entry is 40 bytes. **1 index entry = 40 bytes**
- Assume 1 kilobyte is approximately 1,000 bytes.

Assuming the index is dense and sorted, a binary search for one row reads approximately how many blocks?

$$\log_2 \text{Ex: 123456} \text{ blocks}$$

$$\# \text{ index entries} = \# \text{ of data records} = 600,000,000 \text{ index entries}$$

?????

$$\text{Total number of bytes} = 600,000,000 \text{ index entries} * 1 \text{ index entry} = 40 \text{ bytes} = 24,000,000,000 \text{ index bytes}$$

6

$$\text{ta N –blocks of data entries k*} = 24,000,000,000 \text{ index bytes} / 1 \text{ block} = 8,000 \text{ bytes / block} \\ = 3,000,000 \text{ blocks}$$

# TA time (Fernando)

## (CA 14.2.1 Step 1 – Single level index)

COSC 3380 23578 M & W 4 to 5:30 PM LECTURE (in TEAMS from 3:50 to 5:15 PM)

58:27

Take control Pop out Chat People Raise View Rooms Apps More Camera Mic Share Leave

Participants

Invite someone or dial a number Share invite

In this meeting (108) Mute all

CHALLENGE ACTIVITY 14.2.1: Single-level indexes.

544874.3144414.qx3zqy? Jump to level 1

sorted single-level index table

sorted single-level index	table
(a) ATL	140 Aer Lingus 10:30 DUB Boeing 737
(b) DME	552 Lufthansa 18:00 FRA MD 111
(c) DUB	666 United Airlines 6:00 ATL Airbus 321
(d) DXB	698 Air India 3:15 DXB Boeing 747
(e) FRA	702 American Airlines 5:05 ORD Airbus 323
(f) HKG	799 Air China 3:20 DME Boeing 777
(g) LAX	803 British Airways 13:50 PHL Airbus 323
(h) ORD	971 Air India 19:35 LAX Boeing 747
(i) ORD	1001 Southwest Airlines 5:05 SJC Airbus 323
(j) PHL	1107 Air China 3:20 HKG Boeing 787
(k) SEA	1154 American Airlines 13:50 ORD MD 111
(l) SJC	1222 Delta Airlines 8:25 SEA Boeing 747

Select the index entry that refers to the row for Lufthansa flight 552.

((e) ▾)

1 2 3 4 5 6

Check Next

✓ Expected: (e)

Index entry (e) contains the value FRA with a pointer to the row that refers to Lufthansa flight 552.

View solution (Instructors only)

Ramirez, Fernando

Type here to search

71°F Partly sunny 4:49 PM 3/20/2024

# TA time (Fernando)

## (CA 14.2.1 Step 2 – Single level index)

COSC 3380 23578 M & W 4 to 5:30 PM LECTURE (in TEAMS from 3:50 to 5:15 PM)

59:39

Take control Pop out Chat People 107 Raise View Rooms Apps More Camera Mic Share Leave

Participants

Invite someone or dial a number Share invite

In this meeting (107) Mute all

Hilford, Victoria Organizer

Adhikari, Rohit

Ahmed, Mohamed A

Akram, Ali

Akukwe, Benetta O

Alsayed, Sami H

Alvarez, Stephanie

Avcı, Hatice Kubra

Aysola, Riya

Banza, Sean Paolo B

Bui, Hieu

Burger, Jake

Carrillo-Zepeda, Victor E

CHALLENGE ACTIVITY 14.2.1: Single-level indexes.

544874\_3144414.qz0zy7

Jump to level 1

sorted single-level index

		table			
ATL	•	140	Aer Lingus	10:30	DUB
DME	•	552	Lufthansa	18:00	FRA
DUB	•	666	United Airlines	6:00	ATL
DXB	•	698	Air India	3:15	DXB
FRA	•	702	American Airlines	5:05	ORD
HKG	•	799	Air China	3:20	DME
LAX	•	803	British Airways	13:50	PHL
ORD	•	971	Air India	19:35	LAX
ORD	•	1001	Southwest Airlines	5:05	SJC
PHL	•	1107	Air China	3:20	HKG
SEA	•	1154	American Airlines	13:50	ORD
SJC	•	1222	Delta Airlines	8:25	MD 111

Ethiopian Airlines flight 1920 departing from PEK is inserted into the table. Where does the new index entry go?

between the ORD and PHL entries

1 2 3 4 5 6

Check Next

✓ Expected: between the ORD and PHL entries

Since the single-level index is sorted, an entry for PEK goes between the ORD and PHL entries. The database must create space for the new entry.

Ramirez, Fernando

Type here to search

View solution (Instructors only)

71°F Partly sunny 4:50 PM 3/20/2024

# TA time (Fernando)

## (CA 14.2.1 Step 3 – Single level index)

COSC 3380 23578 M & W 4 to 5:30 PM LECTURE (in TEAMS from 3:50 to 5:15 PM)

01:02:09

Take control Pop out Chat People Raise View Rooms Apps More Camera Mic Share Leave

Participants

Invite someone or dial a number Share invite

In this meeting (108) Mute all

CHALLENGE ACTIVITY 14.2.1: Single-level indexes.

544674\_3144414.qz3zqj7

Jump to level 1

A table occupies 1000 blocks. FlightNumber is the primary key. A single-level index on FlightNumber occupies 250 blocks. The WHERE clause of a SELECT specifies "FlightNumber = 2719".

What is the maximum number of blocks necessary to process the SELECT with an index scan?

251 blocks

Assume the table has no index. What is the maximum number of blocks necessary to process the SELECT with a table scan?

1000 blocks

1 2 3 4 5 6

Check Next

✓ Expected: 251, 1000

The database reads at most 250 index blocks, plus one table block containing the row for flight 2719. With no index, the database performs a table scan and reads at most all 1000 table blocks.

View solution (Instructors only)

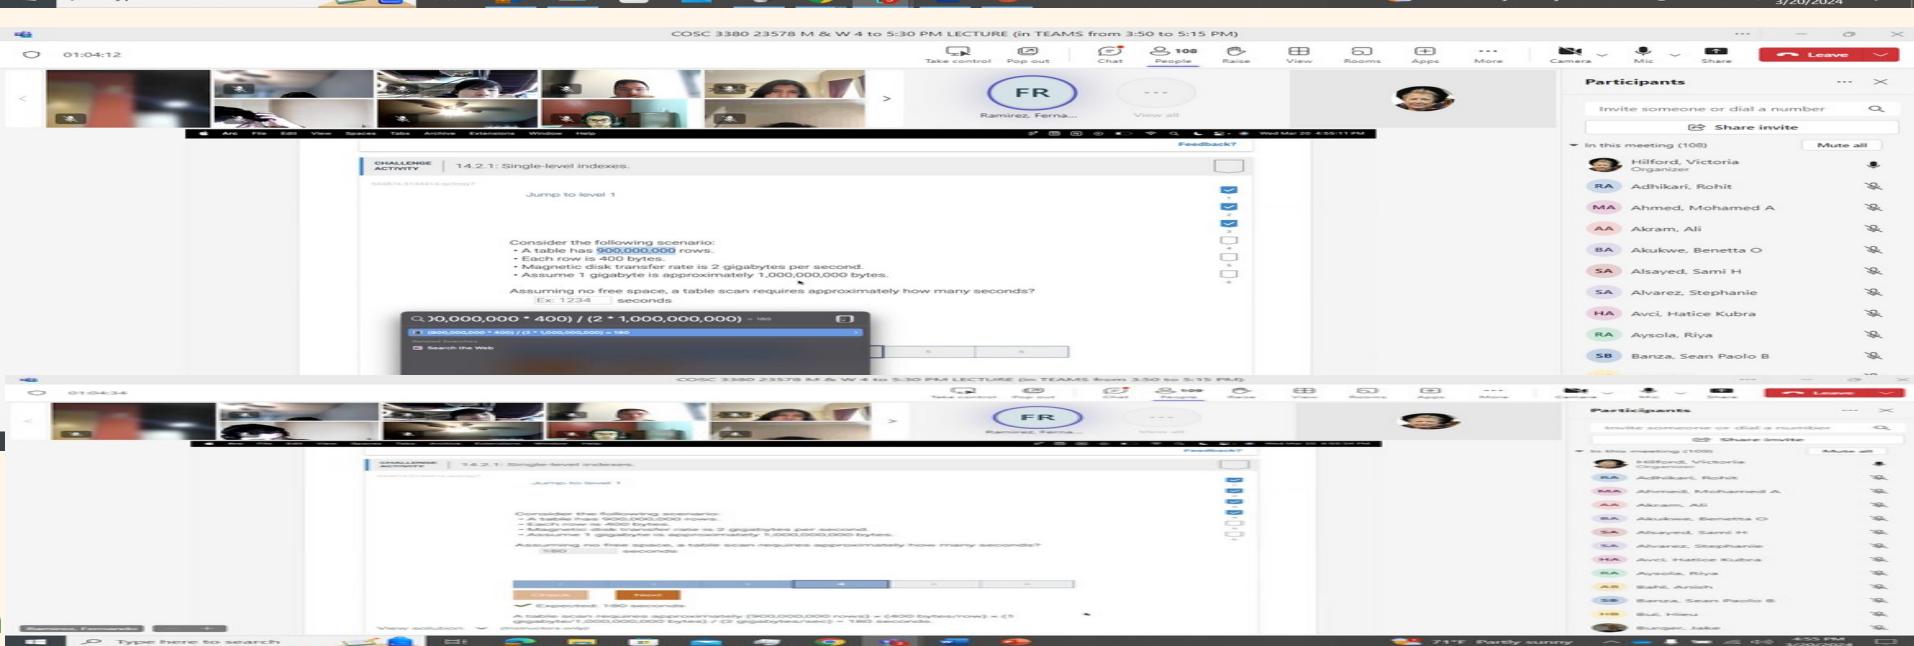
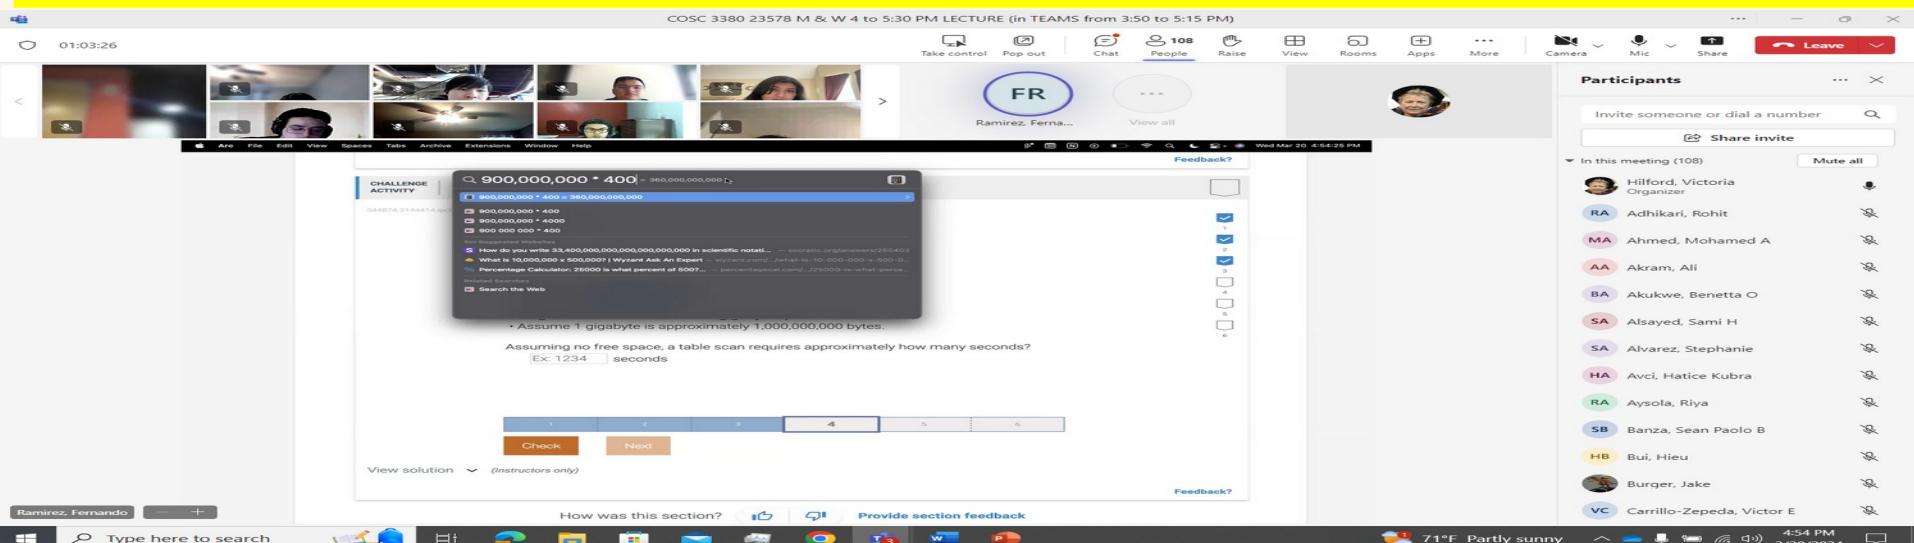
Ramirez, Fernando

Type here to search

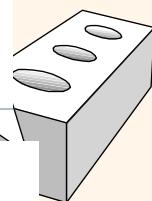
71°F Partly sunny 4:53 PM 3/20/2024

# TA time (Fernando)

## (CA 14.2.1 Step 4 – Single level index)



# ***Storage II Practice Questions***



# Which storage engine provides full support for transaction management and locking?

a. MyISAM

b. InnoDB

c. MEMORY

d. MySQL

## MySQL storage engines

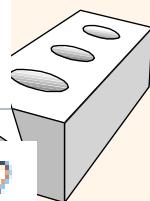
**Logical design** specifies tables, columns, and keys. The logical design process is described elsewhere in this material. **Physical design** specifies indexes, table structures, and partitions. Physical design affects query performance but never affects query results.

A **storage engine** or **storage manager** translates instructions generated by a query processor into low-level commands that access data on storage media. Storage engines support different index and table structures, so physical design is dependent on a specific storage engine.

MySQL can be configured with several different storage engines, including:

- InnoDB is the default storage engine installed with the MySQL download. InnoDB has full support for transaction management, foreign keys, referential integrity, and locking.
- MyISAM has limited transaction management and locking capabilities. MyISAM is commonly used for analytic applications with limited data updates.
- MEMORY stores all data in main memory. MEMORY is used for fast access with databases small enough to fit in main memory.

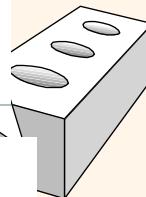
?????



In MySQL with InnoDB, how are primary and secondary indexes created?

- a. Primary indexes are automatically created for each primary key.  
Secondary indexes are created manually for all foreign keys.
- b. Primary indexes are manually created for each primary key. All secondary indexes are created automatically.
- c. Primary indexes are automatically created for each primary key.  
Secondary indexes are created automatically for all foreign keys.
- d. Primary indexes are manually created for each primary key. Secondary indexes are created manually for all foreign keys.

?????



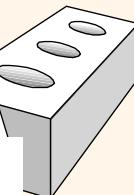
In MySQL with InnoDB, tables with a primary key have a \_\_\_ structure, while those without a primary key have a \_\_\_ structure.

a. heap, sorted

b. sorted, heap

c. sorted, hash

d. cluster, heap



What approach can a database administrator use to assess the effectiveness of indexes when investigating slow queries?

a. Run EXPLAIN on queries

b. Change the primary key

c. Inspect query logs

d. Partition large tables

?????

**At 5:00 PM .**

03.20.2024

zyBook SET 3 - 2

(17 - We)

Set 3

LECTURE 13 STORAGE II

- 
12. SET 3 Empty ▾
- 
14. SET 3: 2:STORAGE II Hidden 0% ▾
- 

VH work on  
SET 3 – 2: STORAGE II

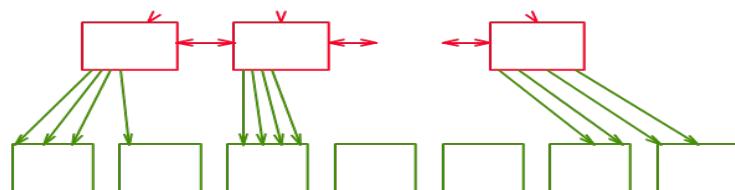
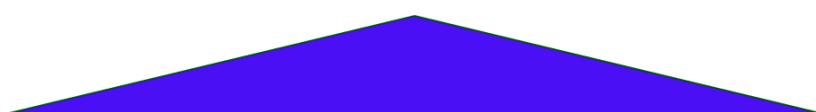
03.25.2024

ZyBook SET 3-3

Set 3

(18 - Mo)

LECTURE 14 STORAGE III B TREE INDEX



N blocks of Data Entries k\*

B blocks of Data Records

$$\text{Cost} = \lceil \log_F N \rceil + 1 = \text{Tree Level}$$

F = # pointers/ Index Page, N = # Leaf (Data Entry k\*) blocks

12. SET 3

Empty

15. SET 3 - 3:STORAGE III B TREE INDEX

Hidden 0% 0%

VH, unHIDE

# From 5:05 to 5:15 PM – 5 minutes.

## STORAGE II

This is a synchronous online class.

Attendance is required.

Recording or distribution of class materials is prohibited.

1. At the beginning of selected classes there is an assessment in the first 10 minutes. (beige BOX in the Detailed Syllabus)

2. At the end of selected classes there is an assessment in the last 10 minutes. (blue BOX in the Detailed Syllabus)

3. ZyBook sections will be downloaded and used for 30% of Total Score on the dates specified in the Detailed Syllabus.

4. EXAMS are in CANVAS. No late EXAMS.

5. I have to be present in TEAMS in order to take any graded assignment assigned during that class.

At 5:15 PM.

End Class 17

VH, unhide ZyBook Section 15.



15. SET 3 - 3:STORAGE III B TREE INDEX



0%



0%



VH, Download Attendance Report  
Rename it:  
**3.20.2024 Attendance Report FINAL**

VH, upload **Class 17** to CANVAS.