



UNIVERSITYof **HOUSTON**

DEPARTMENT OF COMPUTER SCIENCE

COSC 3380 Spring 2024

Database Systems

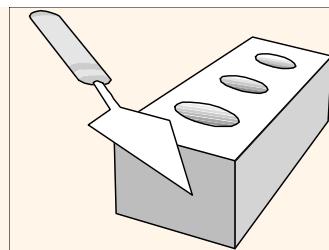
M & W 4:00 to 5:30 PM

Prof. Victoria Hilford

PLEASE TURN your webcam ON (must have)

NO CHATTING during LECTURE

VH, unhide Section 15: SET 3 STORAGE III – B PLUS TREE INDEX



COSC 3380

4 to 5:30

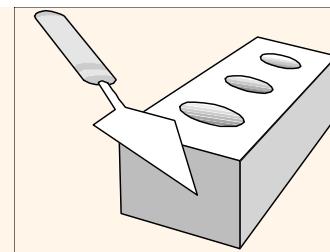
PLEASE

LOG IN

CANVAS

Please close all other windows.

03.25.2024 (18 – Mo)	ZyBook SET 3- 3	Set 3 LECTURE 14 STORAGE III B TREE INDEX
03.27.2024 (19 - We)	ZyBook SET 3- 4	Set 3 LECTURE 15 STORAGE IV HASH INDEX
04.01.2024 (20 - Mo)		EXAM 3 Practice (PART of 20 points)
04.03.2024 (21 – We)	IA Download ZyBook SET 3 Sections (4 PM) (PART of 30 points)	EXAM 3 Review (PART of 20 points)
04.08.2024 (22 - Mo)		EXAM 3 (PART of 50 points)



Class 18

03.25.2024

ZyBook SET 3-3

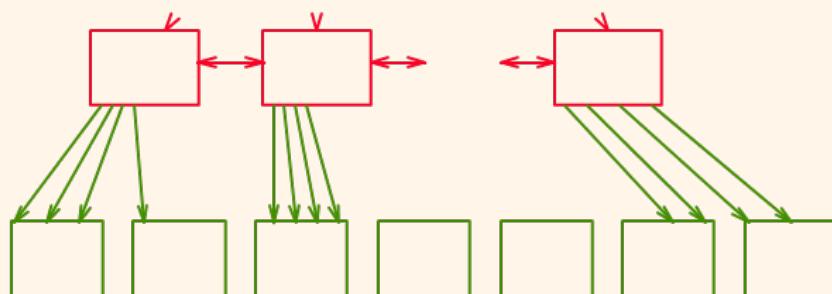
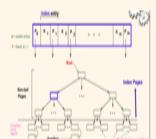
Set 3

(18 - Mo)

LECTURE 14 STORAGE III B TREE INDEX

$$\text{Cost} = \lceil \log_F N \rceil + 1 = \text{Tree Level}$$

$F = \# \text{ pointers}/\text{Index Page}$, $N = \# \text{ Leaf (Data Entry } k^*) \text{ blocks}$



N blocks of Data Entries k^*

B blocks of Data Records

From 4:00 to 4:07 PM – 5 minutes.

03.25.2024

ZyBook SET 3-3

(18 - Mo)

Set 3

LECTURE 14 STORAGE III B TREE INDEX

CLASS PARTICIPATION 20 points

20% of Total + :

B PLUS TREE

Class 18 BEGIN PARTICIPATION

Not available until Mar 25 at 4:00pm | Due Mar 25 at 4:07pm | 100 pts

VH, publish

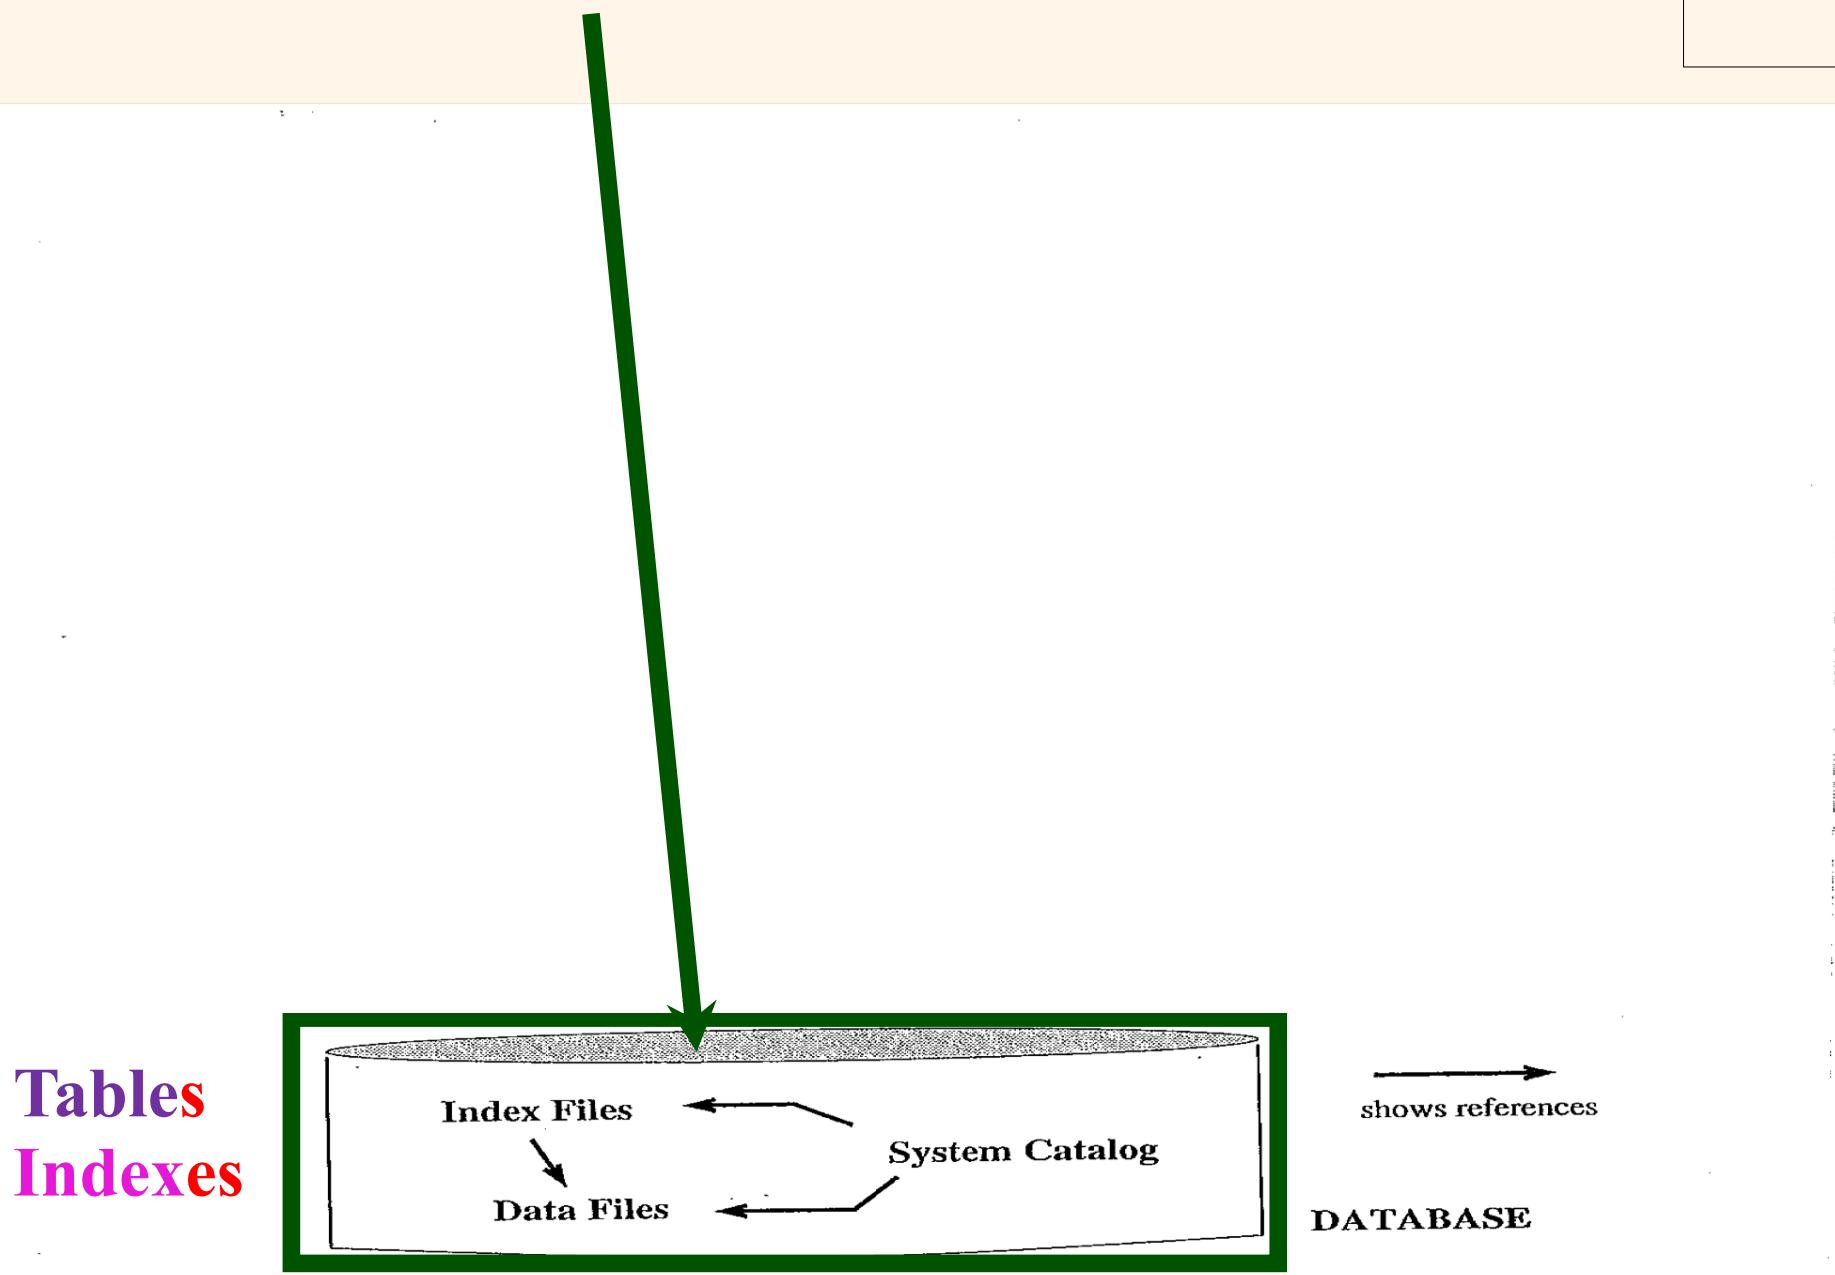
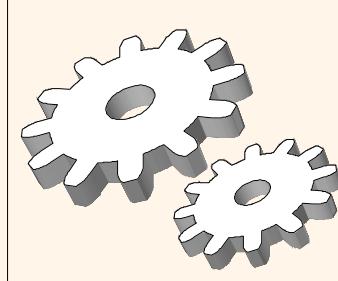
This is a synchronous online class.

Attendance is required.

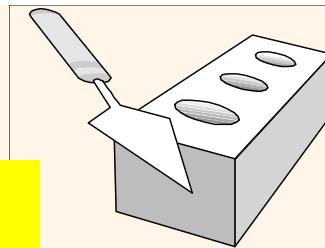
Recording or distribution of class materials is prohibited.

1. At the beginning of selected classes there is an assessment in the first 10 minutes. (beige BOX in the Detailed Syllabus)
2. At the end of selected classes there is an assessment in the last 10 minutes. (blue BOX in the Detailed Syllabus)
3. ZyBook sections will be downloaded and used for 30% of Total Score on the dates specified in the Detailed Syllabus.
4. EXAMS are in CANVAS. No late EXAMS.
5. I have to be present in TEAMS in order to take any graded assignment assigned during that class.

A DBMS



COSC 3380



ZyBook SET 3

10.24.2022 (18 – Mo)	ZyBook SET 3- 3	Set 3 LECTURE 14 STORAGE III B TREE INDEX
10.26.2022 (19 - We)	ZyBook SET 3- 4	Set 3 LECTURE 15 STORAGE IV HASH INDEX
10.31.2022 (20 - Mo)		EXAM 3 Practice (PART of 10 points)
11.02.2022 (21 – We)	Download ZyBook SET 3 Sections (4 PM) (PART of 30 points)	EXAM 3 Review (PART of 10 points)
11.07.2022 (22 - Mo)		EXAM 3 – 15 points



Lecture 14

Tree Structured Indexes

R - number of data records per page/block

B - number of data records blocks

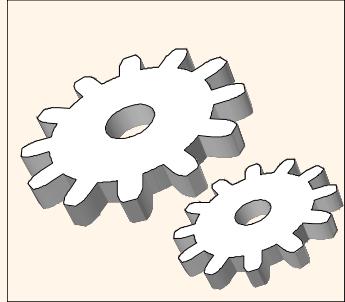
N - number of data entries blocks

F - fan out - number of data entries per page/block

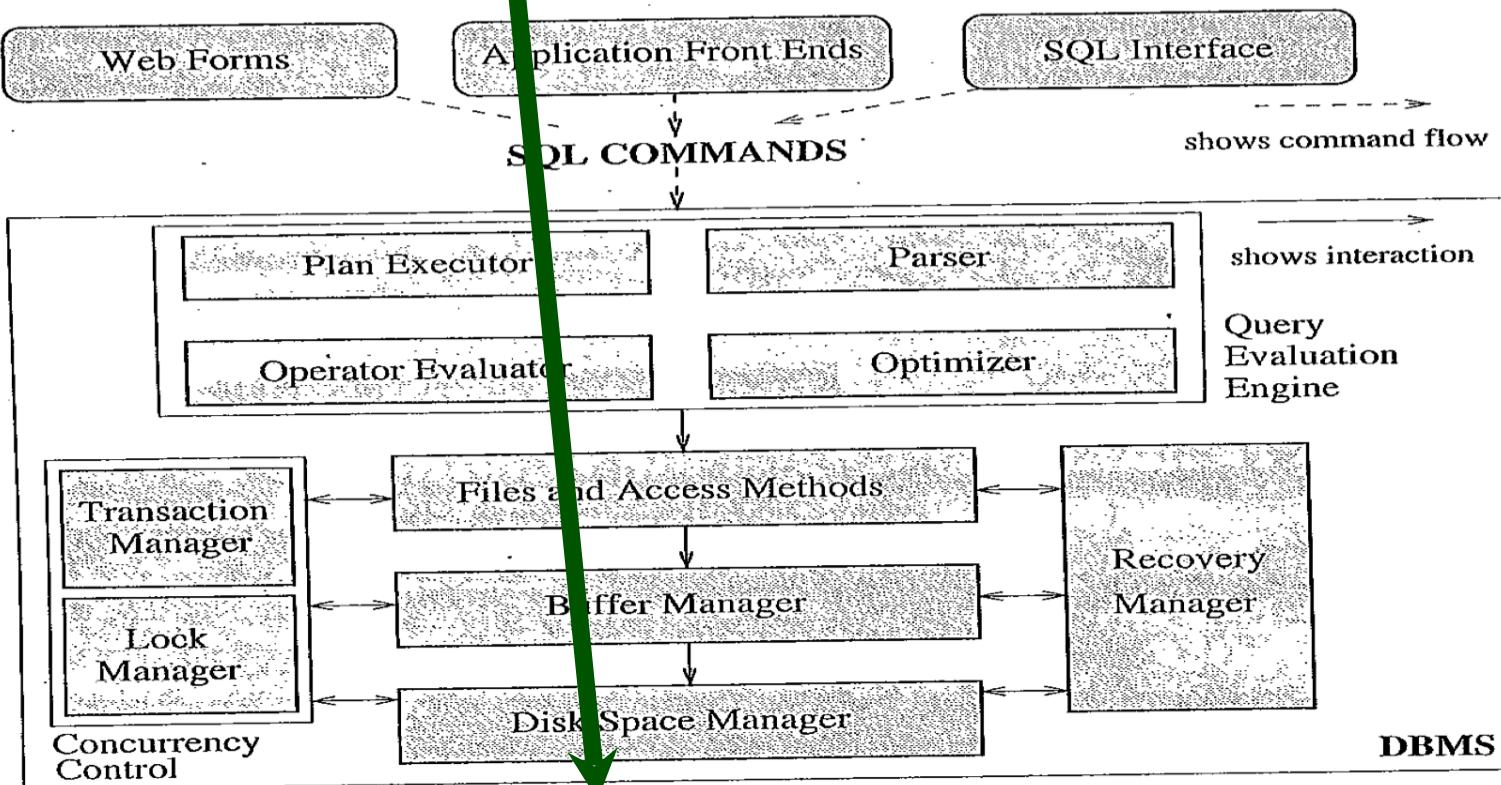
$$\text{Cost} = \lceil \log_F N \rceil + 1 = \text{Tree Level}$$

$F = \# \text{ pointers/ Index Page}$, $N = \# \text{ Leaf (Data Entry k*) blocks}$

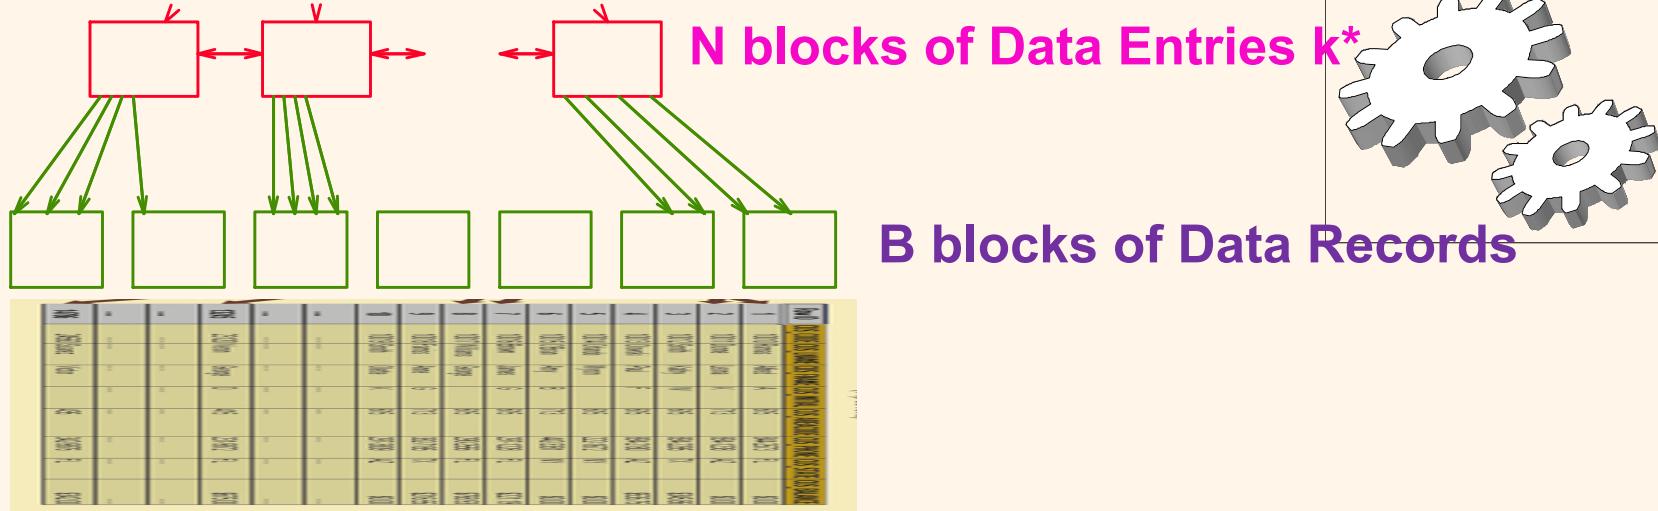
A DBMS



Unsophisticated users (customers, travel agents, etc.) Sophisticated users, application programmers, DB administrators

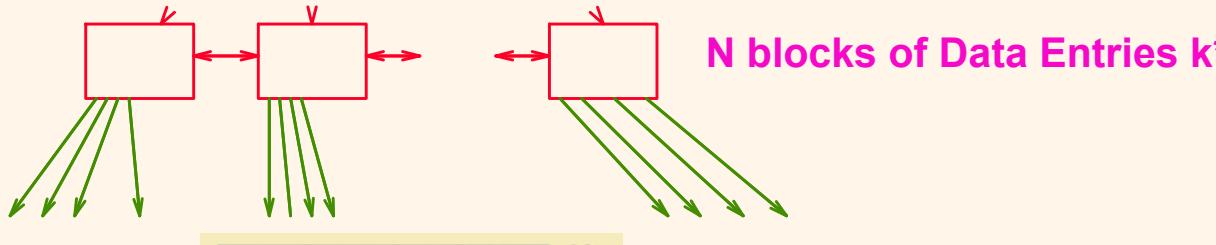


Index



Data Entries k*

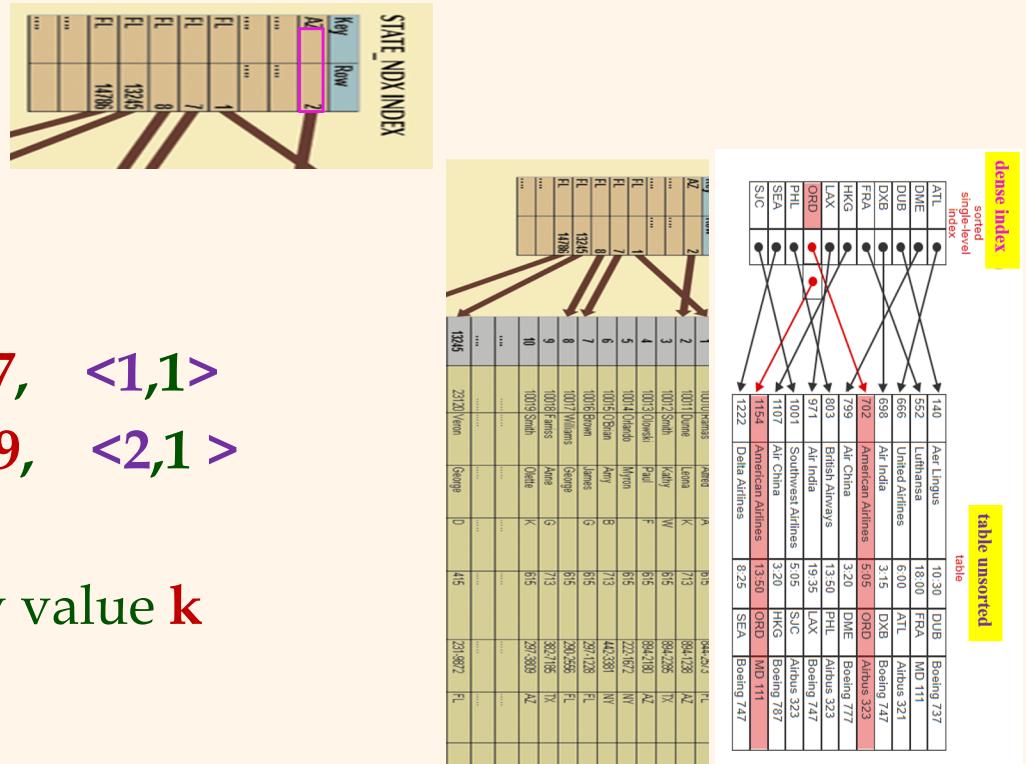
- smaller than **Data Records**
 - reside in the leaf node/**Page**

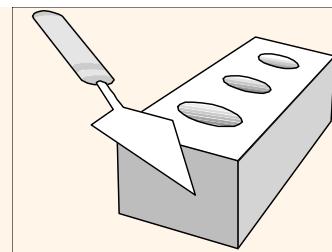


- Index rows contains
 - $k^* = \langle k, \text{rid} \rangle$ of the
 - Data Record k
 - $\text{rid} = \langle \text{pid}, \text{slot \#} \rangle$

Data Records k

- Data Record (**Record**, **Tuple**) with key value **k**





15. SET 3 - 3: STORAGE III B TREE INDEX

█ 0% █ 0% ^

15.1 Single-level indexes Hidden

█ 0% █ 0% ▼

Single-level indexes

A **single-level index** is a file containing column values, along with pointers to rows containing the column value. The pointer identifies the block containing the row. In some indexes, the pointer also identifies the exact location of the row within the block. Indexes are created by database designers with the CREATE INDEX command, described elsewhere in this material.

Single-level indexes are normally sorted on the column value. A sorted index is not the same as an index on a sorted table. Ex: An index on a heap table is a sorted index on an unsorted table.

If an indexed column is unique, the index has one entry for each column value. If an indexed column is not unique, the index may have multiple entries for some column values, or one entry for each column value, followed by multiple pointers.

An index is usually defined on a single column, but an index can be defined on multiple columns. In a **multi-column index**, each index entry is a composite of values from all indexed columns. In all other respects, multi-column indexes behave exactly like indexes on a single column.

PARTICIPATION
ACTIVITY

15.1.1: Single-level index.



Indexes – Data Structures

key,

r_id

- **Ordered set of values that contains Index search key and pointers**
- **More efficient to use Index to access Table than to scan all Rows in Table sequentially**

Index representation for the CUSTOMER table

key, r_id

STATE-INDEX INDEX

Key	Row
AZ	2
FL	1
FL	7
FL	8
FL	13245
FL	14786
...	...
...	...

CUSTOMER TABLE
(14,786 rows)

Row #	CUS_CODE	CUS_LNAME	CUS_FNAME	CUS_INITIAL	CUS_AREACODE	CUS_PHONE	CUS_STATE	CUS_BALANCE
1	10010	Ramas	Alfred	A	615	844-2573	FL	\$0.00
2	10011	Dunne	Leona	K	713	894-1238	AZ	\$0.00
3	10012	Smith	Kathy	W	615	894-2285	TX	\$345.86
4	10013	Olowski	Paul	F				
5	10014	Orlando	Myron					
6	10015	O'Brian	Amy	E				
7	10016	Brown	James	G				
8	10017	Williams	George					
9	10018	Farris	Anne	C				
10	10019	Smith	Olette	H				
...
13245	23120	Veron	George	D				
...
14786	24560	Suarez	Victor	I				

▪ **Indexes:** Data structures to organize records via trees or hashing.

▪ Like Sorted Files, they speed up searches for a subset of records, based on values in certain ("search key") fields

dense indexsorted
single-level
index

ATL	●
DME	●
DUB	●
DXB	●
FRA	●
HKG	●
LAX	●
ORD	●
PHL	●
SEA	●
SJC	●

table unsorted

table

140	Aer Lingus	10:30	DUB	Boeing 737
552	Lufthansa	18:00	FRA	MD 111
666	United Airlines	6:00	ATL	Airbus 321
698	Air India	3:15	DXB	Boeing 747
702	American Airlines	5:05	ORD	Airbus 323
799	Air China	3:20	DME	Boeing 777
803	British Airways	13:50	PHL	Airbus 323
971	Air India	19:35	LAX	Boeing 747
1001	Southwest Airlines	5:05	SJC	Airbus 323
1107	Air China	3:20	HKG	Boeing 787
1154	American Airlines	13:50	ORD	MD 111
1222	Delta Airlines	8:25	SEA	Boeing 747

Refer to the index and table in the animation above.

1) Which index entry refers to Air China flight 1107?

- The first entry
- The sixth entry
- The eleventh entry

Correct

The sixth entry contains the value HKG, with a pointer to the row for Air China flight 1107.

?????

sparse clustered index

physical design

table sorted

blocks

clustered index

sorted table

666	United Airlines	6:00	ATL	Airbus 321
140	Aer Lingus	10:30	ATL	Boeing 737
799	Air China	3:20	DFW	Boeing 777
698	Air India	3:15	DME	Boeing 747

552	Lufthansa	18:00	DUB	MD 111
1107	Air China	3:20	HKG	Boeing 787
971	Air India	19:35	DXB	Boeing 747
702	American Airlines	5:05	EUG	Airbus 323

1154	American Airlines	13:50	FLL	MD 111
803	British Airways	13:50	FRA	Airbus 323
1222	Delta Air Lines	8:25	FSM	Boeing 747
1001	Southwest Airlines	5:05	FSM	Airbus 323

N blocks of data entries

B blocks of data records

Query processing

To execute a SELECT query, the database can perform a table scan or an index scan:

- A **table scan** is a database operation that reads table blocks directly, without accessing an index.
- An **index scan** is a database operation that reads index blocks sequentially, in order to locate the needed table blocks.

Hit ratio, also called **filter factor** or **selectivity**, is the percentage of table rows selected by a query. When a SELECT query is executed, the database examines the WHERE clause and estimates hit ratio. If hit ratio is high, the database performs a table scan. If hit ratio is low, the query needs only a few table blocks, so a table scan would be inefficient. Instead, the database:

1. Looks for an indexed column in the WHERE clause.
2. Scans the index.
3. Finds values that match the WHERE clause.
4. Reads the corresponding table blocks.

If the WHERE clause does not contain an indexed column, the database must perform a table scan.

Since a column value and pointer occupy less space than an entire row, an index requires fewer blocks than a table. Consequently, index scans are much faster than table scans. In some cases, indexes are small enough to reside in main memory, and index scan time is insignificant. When hit ratio is low, additional time to read the table blocks containing selected rows is insignificant.

PARTICIPATION
ACTIVITY

15.1.3: Query processing with single-level index.

B blocks of data records

table				
140	Aer Lingus	10:30	DUB	Boeing 737

B blocks

971	Air India	19:35	LAX	Boeing 747
1001	Southwest Airlines	5:05	SJC	Airbus 323
1107	Air China	3:20	HKG	Boeing 787
1154	American Airlines	13:50	ORD	MD 111
1222	Delta Airlines	8:25	SEA	Boeing 747

Query processing

To execute a SELECT query, the database can perform a table scan or an index scan:

- A **table scan** is a database operation that reads table blocks directly, without accessing an index.
- An **index scan** is a database operation that reads index blocks sequentially, in order to locate the needed table blocks.

Hit ratio, also called **filter factor** or **selectivity**, is the percentage of table rows selected by a query. When a SELECT query is executed, the database examines the WHERE clause and estimates hit ratio. If hit ratio is high, the database performs a table scan. If hit ratio is low, the query needs only a few table blocks, so a table scan would be inefficient. Instead, the database:

1. Looks for an indexed column in the WHERE clause.
2. Scans the index.
3. Finds values that match the WHERE clause.
4. Reads the corresponding table blocks.

If the WHERE clause does not contain an indexed column, the database must perform a table scan.

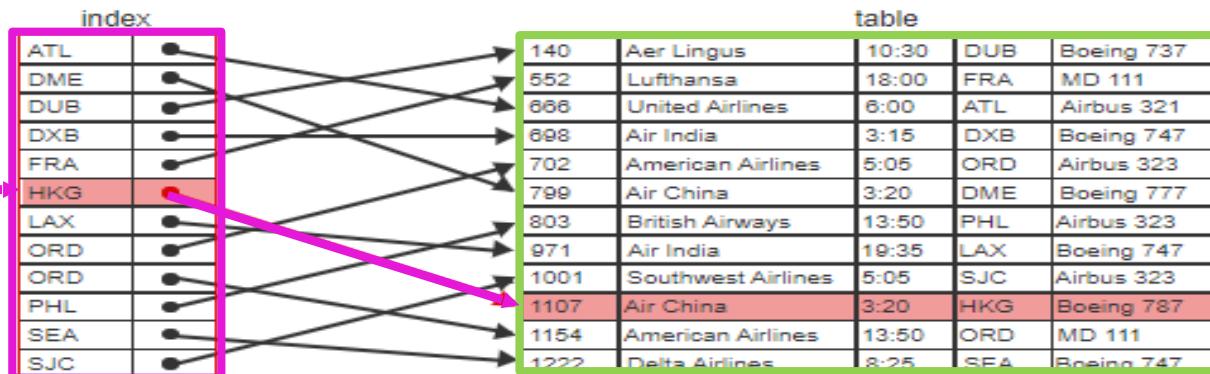
Since a column value and pointer occupy less space than an entire row, an index requires fewer blocks than a table. Consequently, index scans are much faster than table scans. In some cases, indexes are small enough to reside in main memory, and index scan time is insignificant. When hit ratio is low, additional time to read the table blocks containing selected rows is insignificant.

PARTICIPATION
ACTIVITY

15.1.3: Query processing with single-level index.



INDEX
~~Binary Search~~



N blocks of data entries

1

```
SELECT FlightNumber, AirlineName
FROM Flight
WHERE DepartureAirport = 'HKG';
```

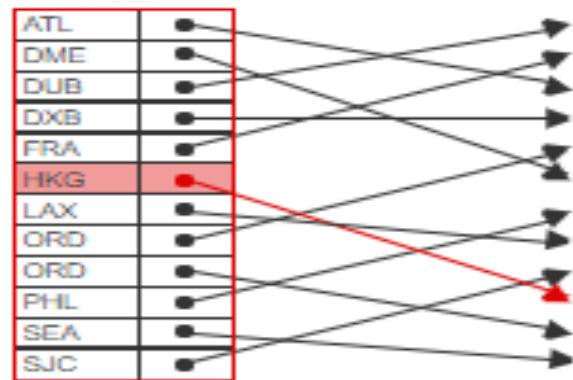
N = 200**INDEX****Binary Search**

table				
140	Aer Lingus	10:30	DUB	Boeing 737
552	Lufthansa	18:00	FRA	MD 111
666	United Airlines	6:00	ATL	Airbus 321
698	Air India	3:15	DXB	Boeing 747
702	American Airlines	5:05	ORD	Airbus 323
799	Air China	3:20	DME	Boeing 777
803	British Airways	13:50	PHL	Airbus 323
971	Air India	19:35	LAX	Boeing 747
1001	Southwest Airlines	5:05	SJC	Airbus 323
1107	Air China	3:20	HKG	Boeing 787
1154	American Airlines	13:50	ORD	MD 111
1222	Delta Airlines	8:25	SEA	Boeing 747

B = 2,000

SELECT FlightNumber, AirlineName
FROM Flight;

1

FlightNumber	AirlineName
140	Aer Lingus
552	Lufthansa
666	United Airlines
698	Air India
702	American Airlines
799	Air China
803	British Airways
971	Air India
1001	Southwest Airlines
1107	Air China
1154	American Airlines
1222	Delta Airlines

SELECT FlightNumber, AirlineName
FROM Flight
WHERE FlightNumber = 803;

FlightNumber	AirlineName
803	British Airways

SELECT FlightNumber, AirlineName
FROM Flight
WHERE DepartureAirport = "HKG";

FlightNumber	AirlineName
1107	Air China

Refer to the following scenario:

- A table occupies 2,000 blocks. **B = 2,000**
- FlightNumber is the primary key.
- An index on FlightNumber occupies 200 blocks. **N = 200**

1) What is the maximum number of blocks necessary to process the
SELECT?

- 2
 201
 2,000
—

200 + 1

Correct

The database reads at most 200 index blocks, plus one table block containing the row for flight 3988.

?????

Primary, clustering, and secondary indexes

Indexes on a sorted table may be primary, clustering, or secondary:

- A **primary index** is an index on a unique sort column.
- A **clustering index** is an index on a non-unique sort column.
- A **secondary index** is an index that is not on the sort column.

A sorted table can have only one sort column, and therefore only one primary or clustering index. Tables can have many secondary indexes. All indexes of a heap or hash table are secondary, since heap and hash tables have no sort column.

Indexes may also be dense or sparse:

- A **dense index** contains an entry for every table row.

When a table is sorted on an index column, the index may be sparse, as illustrated in the animation below. Primary and clustering indexes are on sort columns and usually sparse. Secondary indexes are on non-sort columns and therefore are always dense.

Sparse indexes are much faster than dense indexes since sparse indexes have fewer entries and occupy fewer blocks. Consider the following scenario:

- A table has 10 million rows.
- Each row is 100 bytes.
- Table and index blocks are 4 kilobytes.
- Each index entry is 10 bytes.

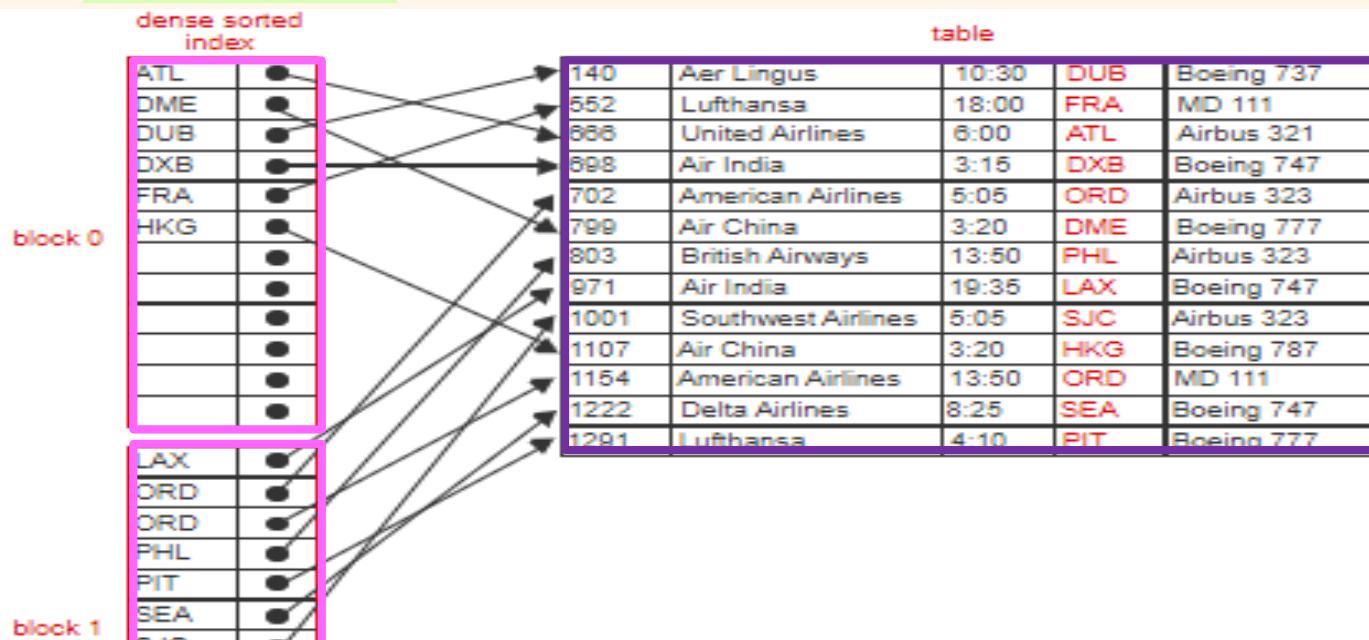
B = 250,000 blocks

N = 625 blocks

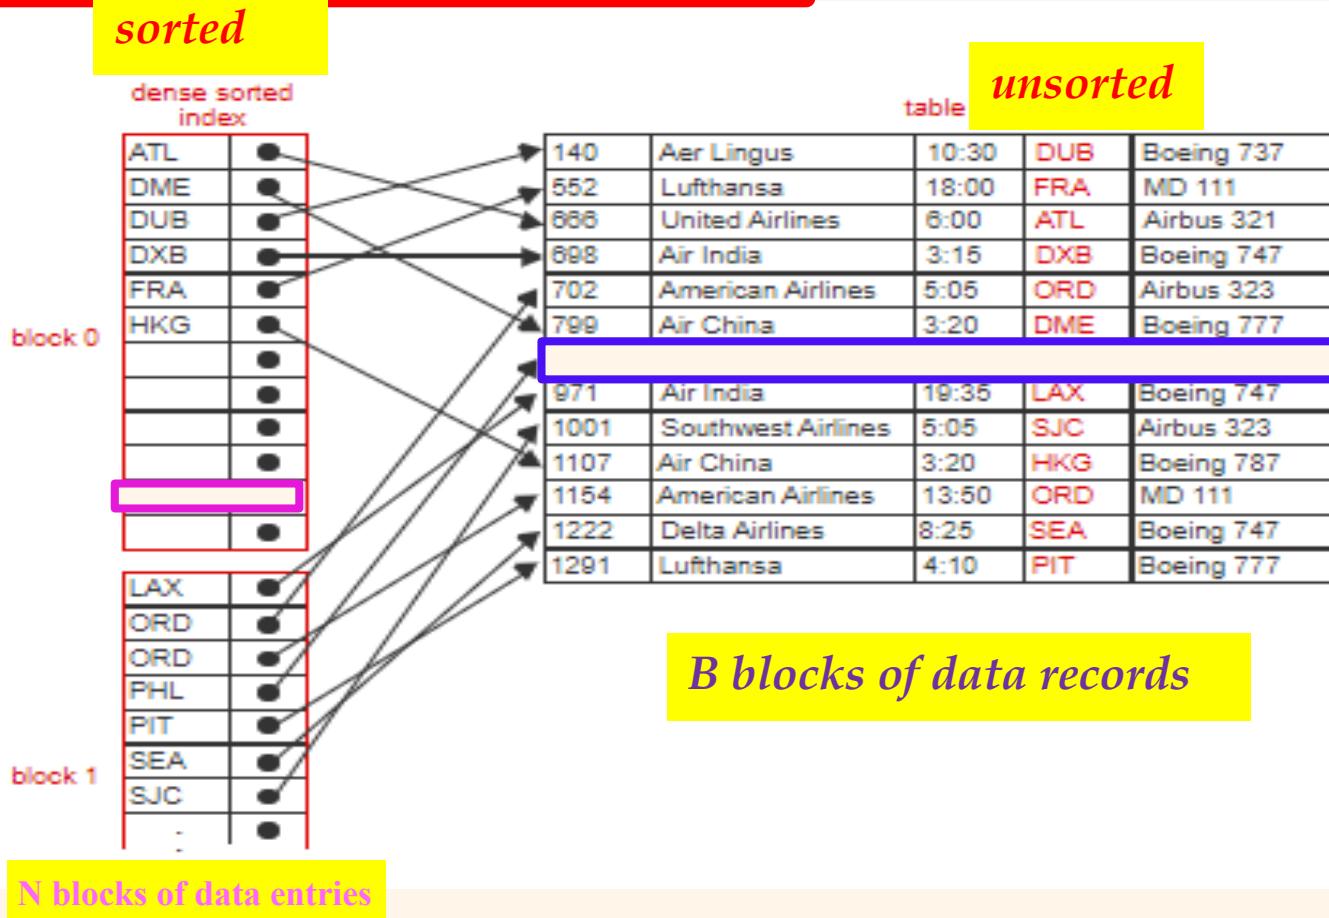
The table occupies 250,000 blocks ($10 \text{ million rows} \times 100 \text{ bytes/row} / 4 \text{ kilobytes/block}$). A sparse index requires 250,000 entries (one entry per table block) and occupies 625 blocks ($250,000 \text{ entries} \times 10 \text{ bytes/entry} / 4 \text{ kilobytes/block}$). The sparse index can easily be retained in main memory.

Primary and clustering indexes are always dense. Secondary and sparse indexes are fast. As a result, database designers usually create a primary or clustering index on large tables.

Dense index



index entries = # data records



- 1) Inserts to a sorted table are always faster when the table has no indexes.

- True
 False

Correct

Inserting a row may generate a new index entry, which slows down the insert. However, if the index is on the sort column, the index helps locate the table block for the insert, which speeds up the query.

?????

Binary search

When hit ratio is low, index scans are always faster than table scans. Consider the following scenario:

- A table has 10 million rows.
 - Each row is 100 bytes.
 - Each block is 4 kilobytes.
 - Each index entry is 10 bytes, including a 6-byte value and a 4-byte pointer.
 - Magnetic disk transfer rate is 0.1 gigabytes per second.

If hit ratio is low, an index scan is roughly 10 times faster than a table scan.

- **Table scan:** The table contains 1 billion bytes ($10 \text{ million rows} \times 100 \text{ bytes/row}$). The table scan takes around 10 seconds (1 gigabyte / 0.1 gigabytes/sec).
 - **Index scan:** The index contains 100 million bytes ($10 \text{ million rows} \times 1 \text{ entry/row} \times 10 \text{ bytes/entry}$). The index scan takes around 1 second (0.1 gigabyte / 0.1 gigabytes/sec). The index scan returns pointers to blocks containing rows selected by the query. When hit ratio is low, additional time to read table blocks is insignificant.

Although index scans are faster than table scans, index scans are too slow in many cases. If a single-level index is sorted, each value can be located with a **binary search**. In a **binary search**, the database repeatedly splits the index in two until it finds the entry containing the search value:

1. The database first compares the search value to an entry in the middle of the index.
 2. If the search value is less than the entry value, the search value is in the first half of the index. If not, the search value is in the second half.
 3. The database now compares the search value to the entry in the middle of the selected half, to narrow the search to one quarter of the index.
 4. The database continues in this manner until it finds the index block containing the search value.

For an index with N entries, a binary search examines $\log_2 N$ values, rounded up to the nearest integer. In the example above, the index has 10 million entries. The binary search reads at most $\log_2 10,000,000 = 24$ blocks and requires about 0.001 seconds (24 blocks \times 4 kilobytes/block / 0.1 gigabytes/sec).

PARTICIPATION ACTIVITY

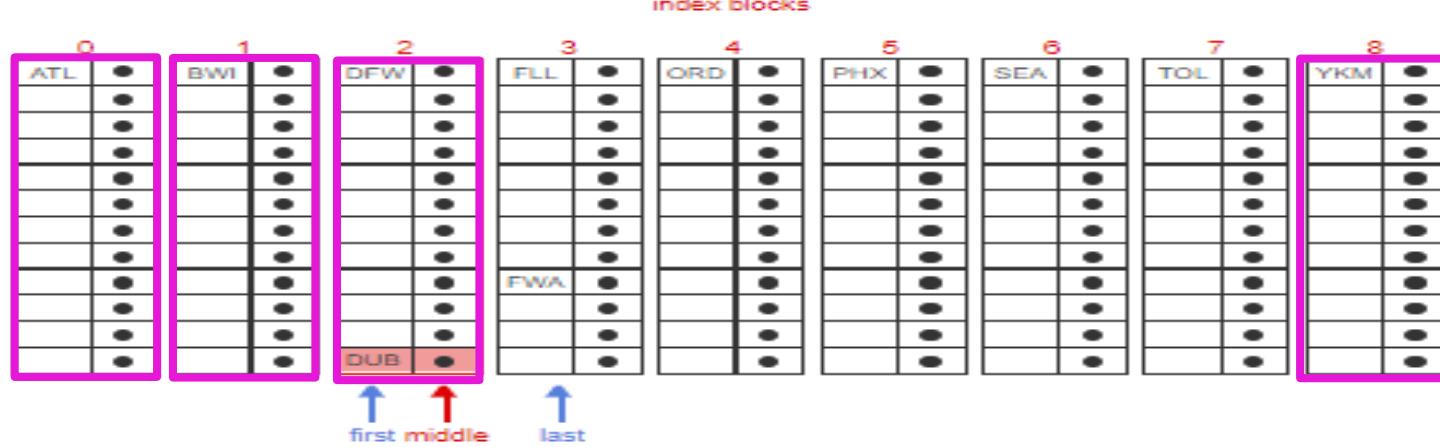
15.1.5: Binary search on a sorted index.

INDEX

Binary Search

$$\log_2 9$$

N = 9



```
SELECT FlightNumber, AirlineName  
FROM Flight  
WHERE DepartureAirport = 'DUB';
```



Refer to the following scenario:

- A table has 100 million rows.
- Each row is 400 bytes.
- Each block is 8 kilobytes.
- Each index entry is 20 bytes
- Magnetic disk transfer rate is 0.1 gigabytes per second.

$$R = 8,000 / 400 = \\ 20 \text{ data records / page}$$

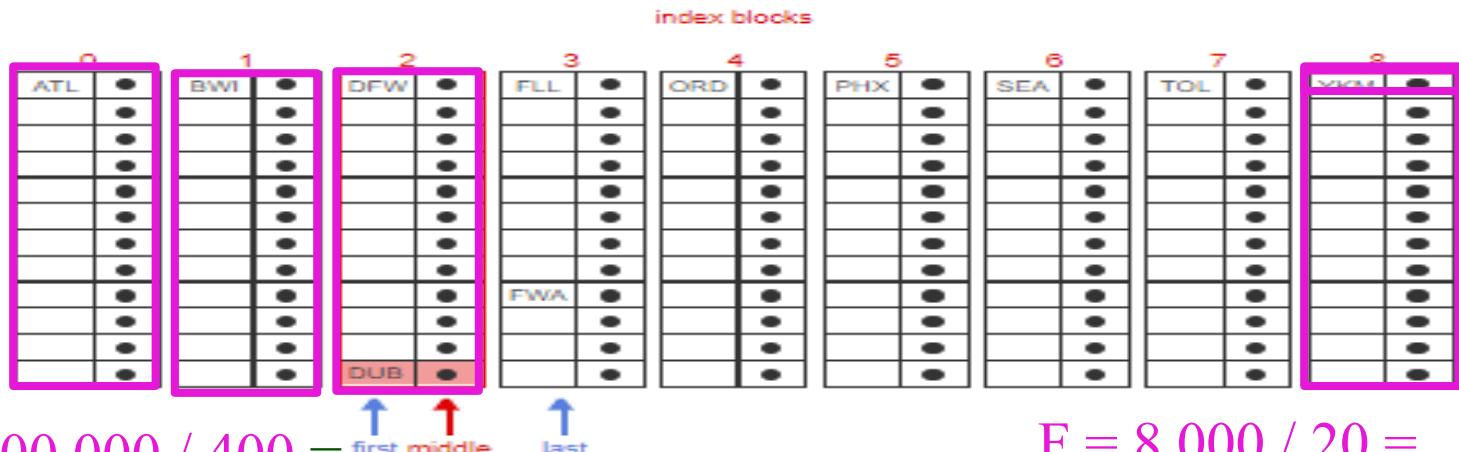
$$8\text{KB} = \\ 8,000 \text{ B}$$

- 1) Assuming no free space, a table scan requires approximately how many seconds?

- 0.4
- 40
- 400

Correct

$100,000,000 \text{ rows} \times 400 \text{ bytes/row} / 0.1 \text{ gigabytes/sec} = \\ 400 \text{ seconds.}$

INDEX**Binary Search** $\log_2 N$ 

```
SELECT FlightNumber, AirlineName
FROM Flight
WHERE DepartureAirport = "DUB";
```

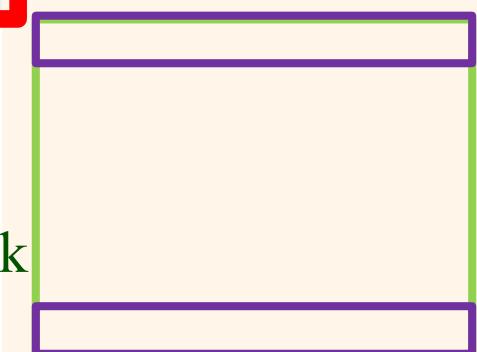
?????

Refer to the following scenario:

- A table has 100 million rows.
- Each row is 400 bytes.
- Each block is 8 kilobytes.
- Each index entry is 20 bytes
- Magnetic disk transfer rate is 0.1 gigabytes per second.

$$R = 8,000 / 400 = 20 \text{ data records / block}$$

8KB =
8,000 B



2) Assuming the index is sorted, a binary search for one row reads approximately how many blocks?

- $\log_2 250,000$
- $\log_2 5,000,000$
- $\log_{10} 250,000$

Correct

Each index block contains 8,000 bytes / 20 bytes/entry = 400 entries. The index has 100,000,000 entries / 400 entries/block = 250,000 blocks. A binary search requires $\log 250,000$ (base 2).

$$B = 100,000,000 / 20 = 50,000,000 \text{ blocks}$$

Primary, clustering, and secondary indexes

Indexes on a sorted table may be primary, clustering, or secondary:

- A **primary index** is an index on a unique sort column.
- A **clustering index** is an index on a non-unique sort column.
- A **secondary index** is an index that is not on the sort column.

A sorted table can have only one sort column, and therefore only one primary or clustering index. Tables can have many secondary indexes. All indexes of a heap or hash table are secondary, since heap and hash tables have no sort column.

Indexes may also be dense or sparse:

- A **dense index** contains an entry for every table row.
- A **sparse index** contains an entry for every table block.

When a table is sorted on an index column, the index may be sparse, as illustrated in the animation below. Primary and clustering indexes are on sort columns and usually sparse. Secondary indexes are on non-sort columns and therefore are always dense.

Sparse indexes are much faster than dense indexes since sparse indexes have fewer entries and occupy fewer blocks. Consider the following scenario:

- A table has 10 million rows.
- Each row is 100 bytes.
- Table and index blocks are 4 kilobytes.
- Each index entry is 10 bytes.

$$B = 250,000 \text{ blocks}$$

The table occupies 250,000 blocks ($10 \text{ million rows} \times 100 \text{ bytes/row} / 4 \text{ kilobytes/block}$). A sparse index requires 250,000 entries (one entry per table block) and occupies 625 blocks ($250,000 \text{ entries} \times 10 \text{ bytes/entry} / 4 \text{ kilobytes/block}$). The sparse index can easily be retained in main memory.

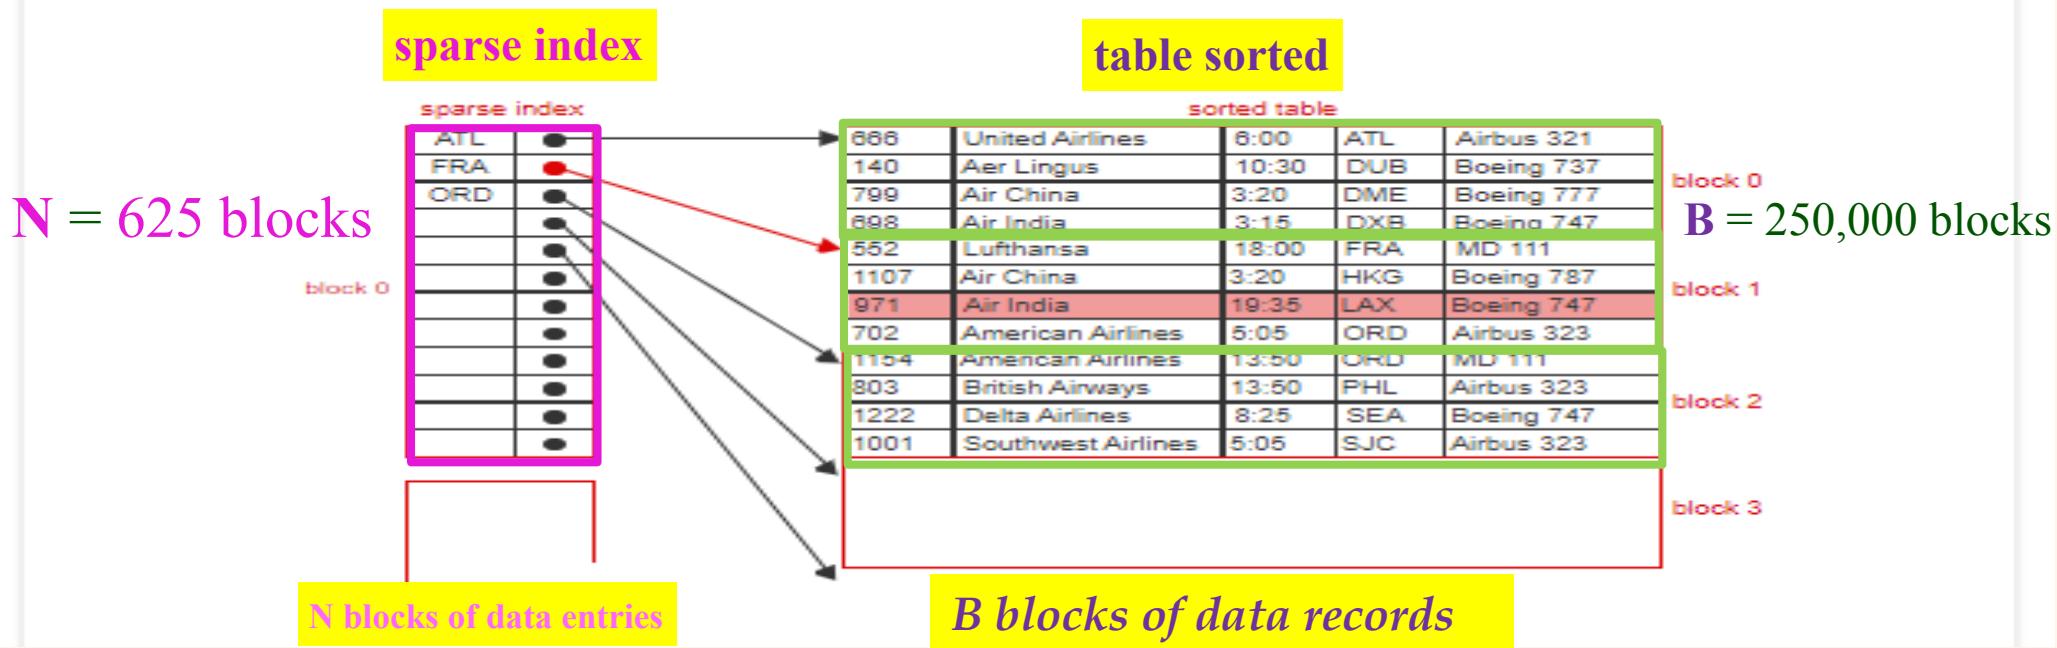
$$N = 625 \text{ blocks}$$

Primary and clustering indexes are usually sparse and sparse indexes are fast. As a result, database designers usually create a primary or clustering index on large tables.

PARTICIPATION
ACTIVITY

15.1.7: Dense and sparse indexes.

$N = 625 \text{ blocks}$



Inserts, updates, and deletes

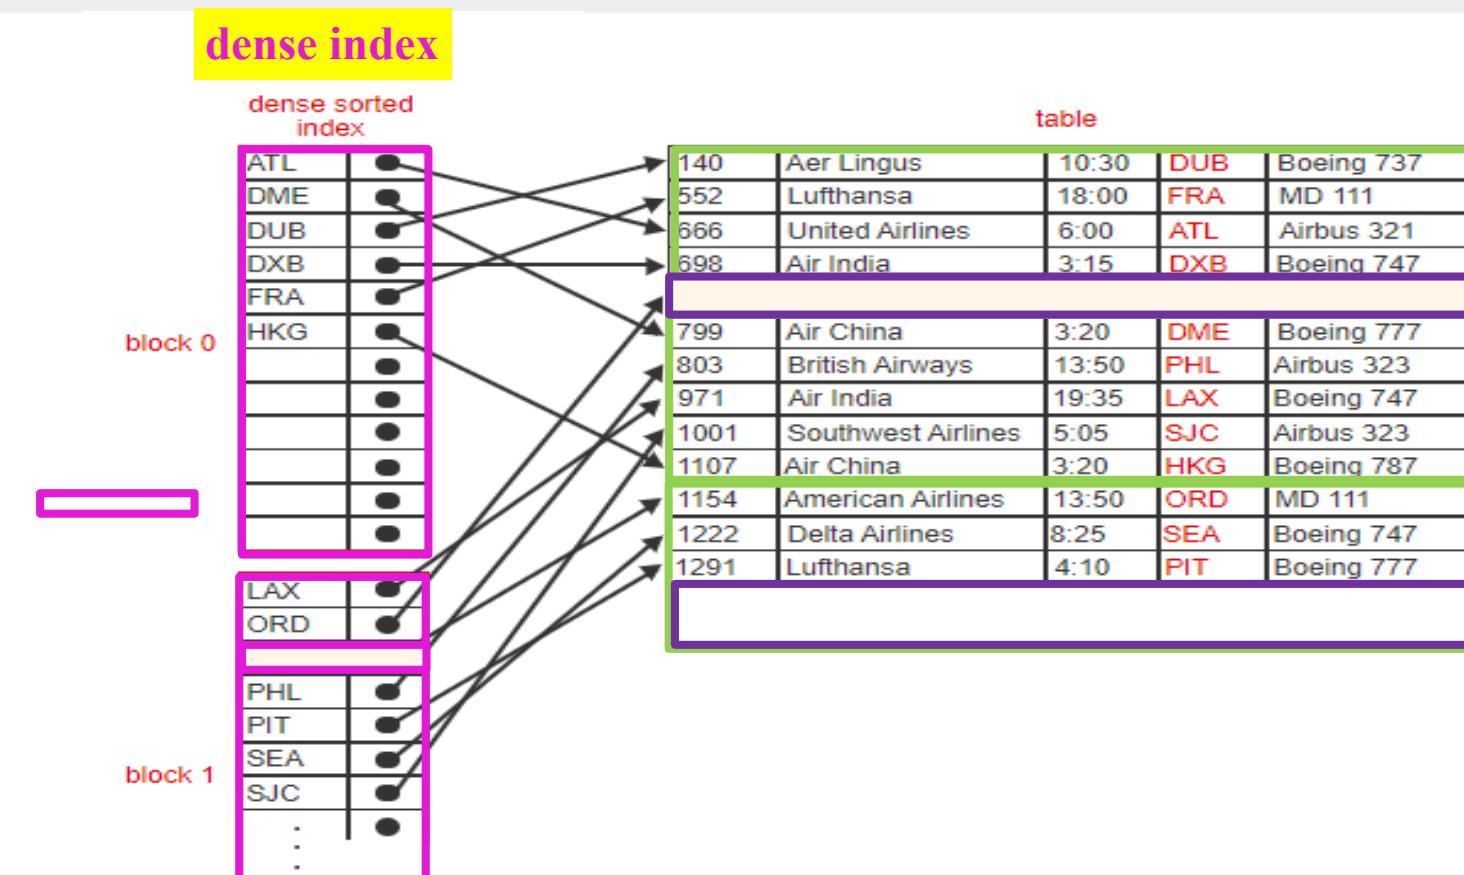
Inserts, updates, and deletes to tables have an impact on single-level indexes. Consider the behavior of dense indexes:

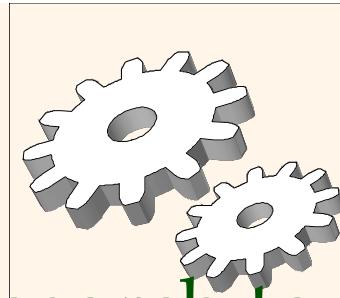
- **Insert.** When a row is inserted into a table, a new index entry is created. Since single-level indexes are sorted, the new entry must be placed in the correct location. To make space for the new entry, subsequent entries must be moved, which is too slow for large tables. Instead, the database splits an index block and reallocates entries to the new block, creating space for the new entry.
- **Delete.** When a row is deleted, the row's index entry must be deleted. The deleted entry can be either physically removed or marked as 'deleted'. Since single-level indexes are sorted, physically removing an entry requires moving all subsequent entries, which is slow. For this reason, index entries are marked as 'deleted'. Periodically, the database may reorganize the index to remove deleted entries and compress the index.
- **Update.** An update to a column that is not indexed does not affect the index. An update to an indexed column is like a delete followed by an insert. The index entry for the initial value is deleted and an index entry for the updated value is inserted.

With a sparse index, each entry corresponds to a table block rather than a table row. Index entries are inserted or deleted when blocks split or merge. Since blocks contain many rows, block splits and mergers occur less often than row inserts and deletes. Aside from frequency, however, the behavior of sparse and dense indexes is similar.

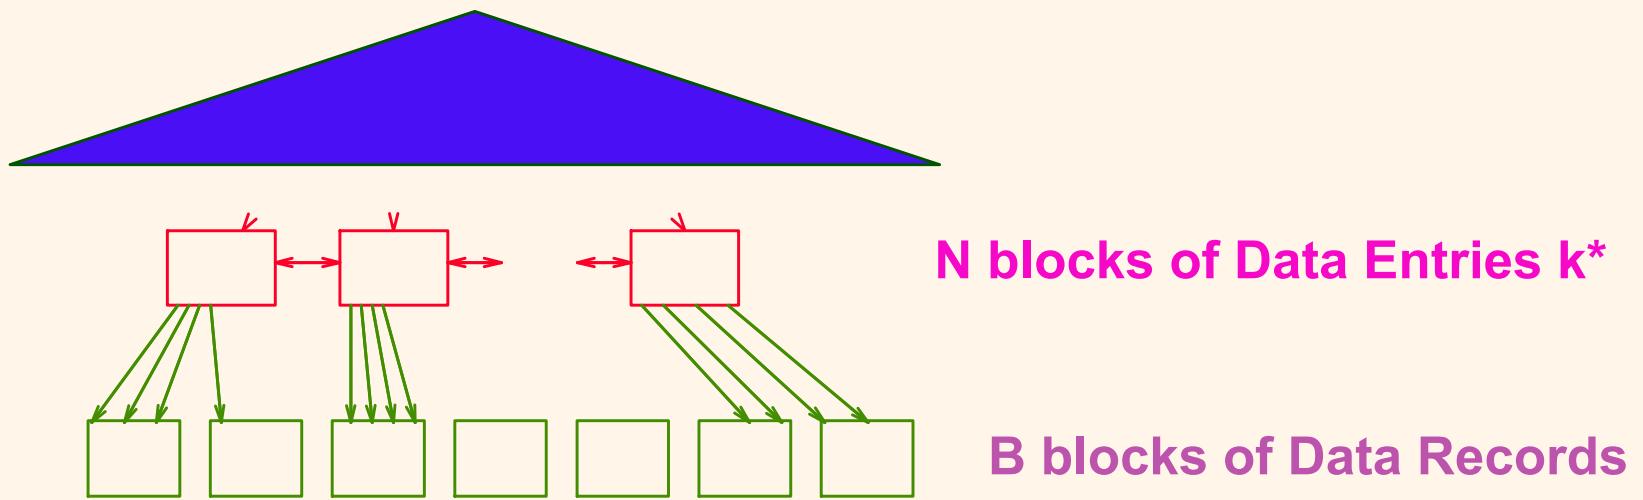
PARTICIPATION
ACTIVITY

15.1.9: Insert with dense sorted index.



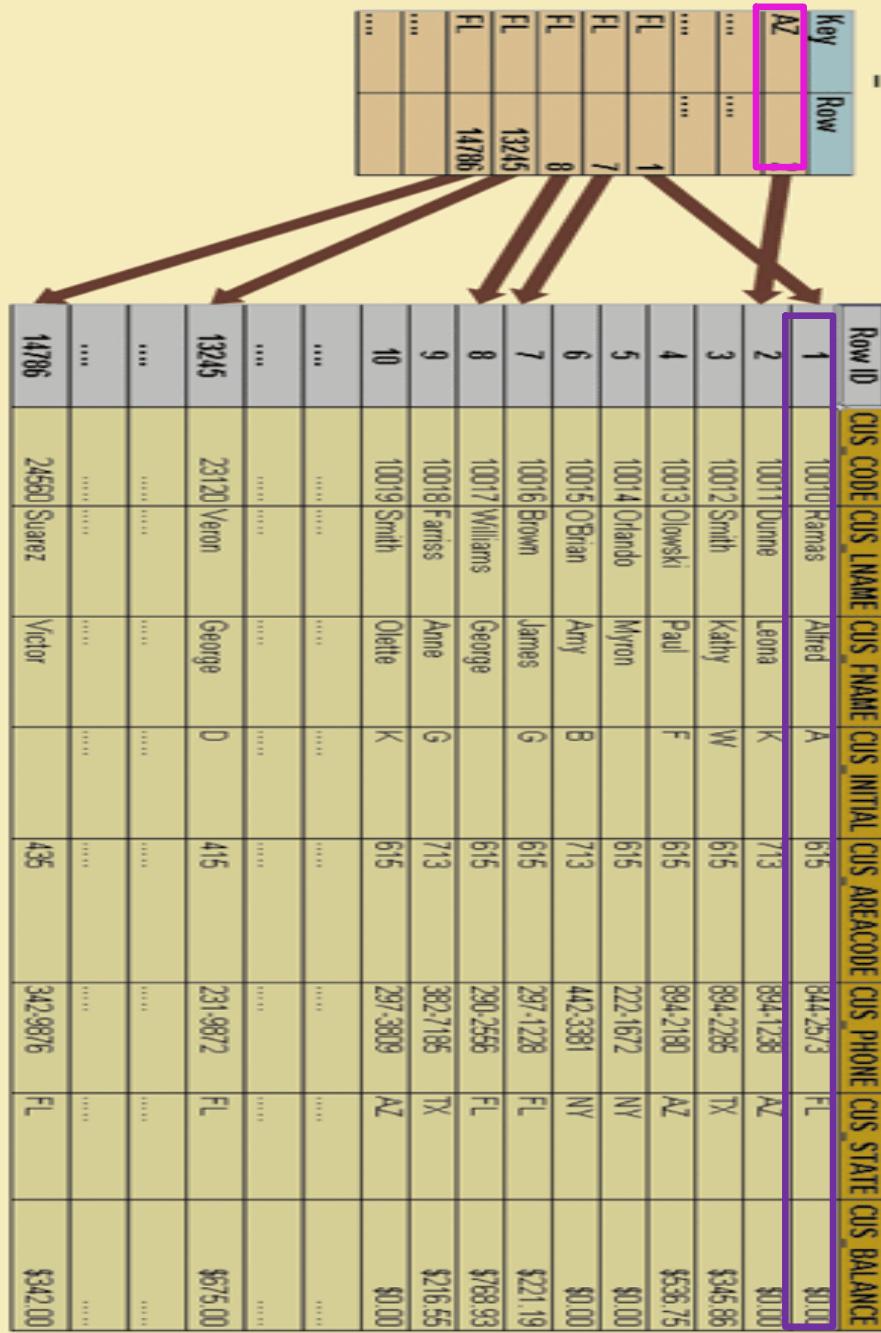
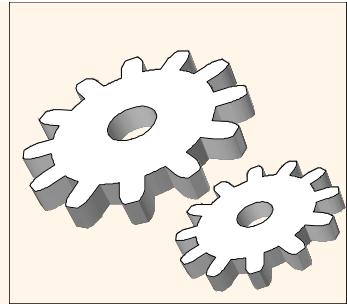


- ❖ Auxiliary information used to quickly direct the search to Data Entries k^* : **Tree Structured Indexing** techniques support both *range searches* and *equality searches*.
- ❖ ISAM: static structure;
- ❖ B⁺ Tree: dynamic, adjusts gracefully under **inserts** and **deletes**.



Indexes

blocks of Search Tree



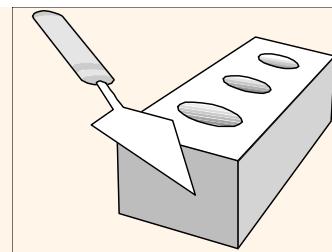
B blocks of Data Records

CUSTC

11

N blocks of Data Entries k*

Index rep



15. SET 3 - 3: STORAGE III B TREE INDEX

█ 0% █ 0% ^

15.2 Multi-level indexes Hidden

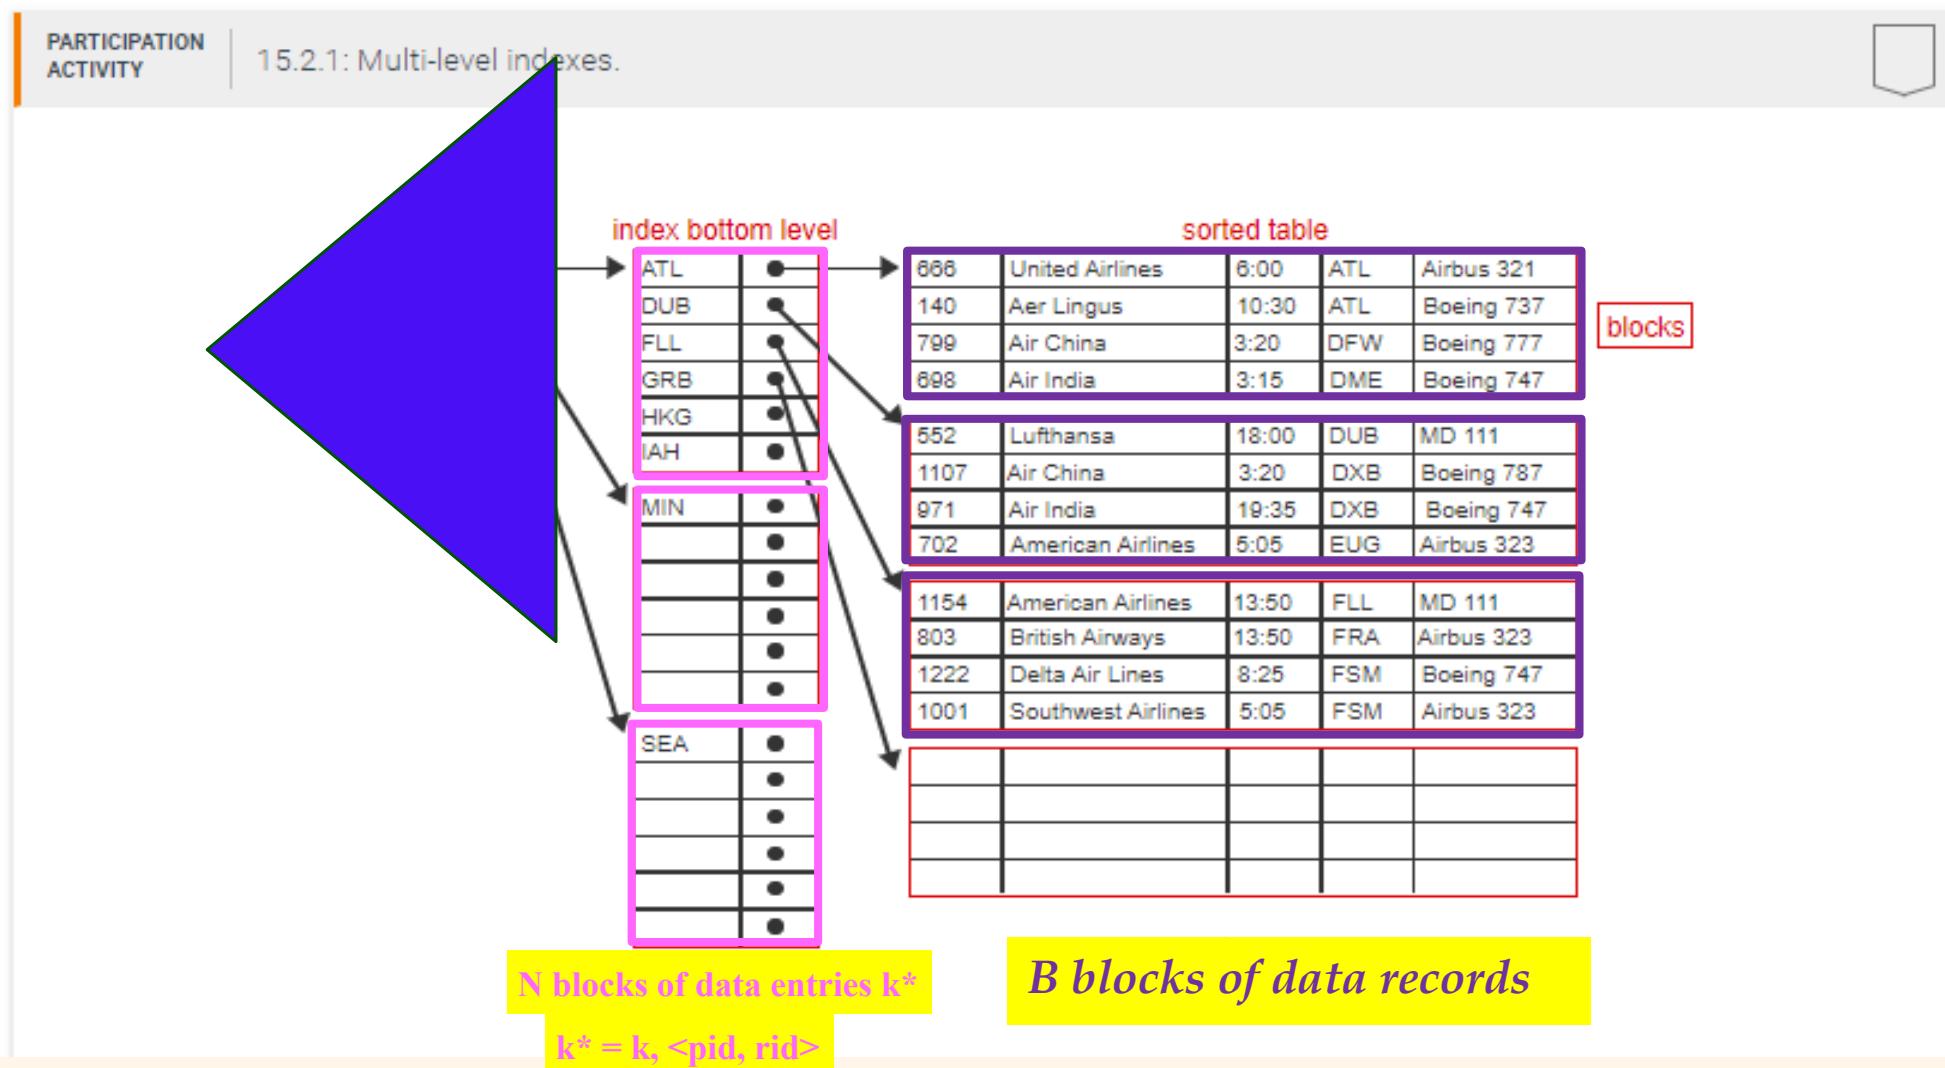
█ 0% █ 0% ▼

Multi-level indexes

A **multi-level index** stores column values and row pointers in a hierarchy. The bottom level of the hierarchy is a sorted single-level index. The bottom level is sparse for primary and clustering indexes, or dense for secondary indexes.

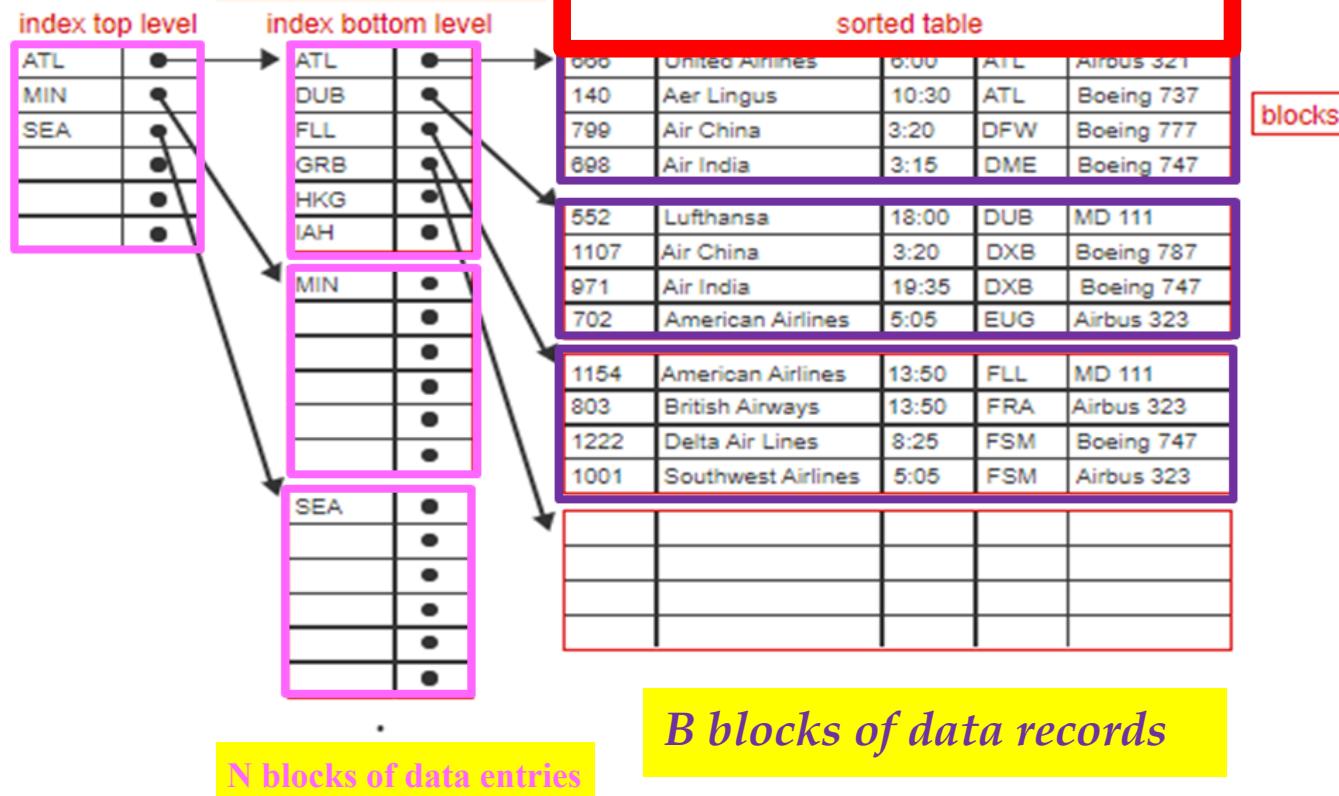
Each level above the bottom is a sparse sorted index to the level below. Since all levels above the bottom are sparse, levels rapidly become smaller. The top level always fits in one block.

To locate a row containing an indexed value, the database first reads the top-level block. The database compares the indexed value to entries in the block and locates the next level block containing the value. Continuing in this manner, the database eventually locates the bottom-level block containing the value. The bottom-level block contains a pointer to the correct table block.





Sparse index



- 1) The bottom level of a multi-level index is always sparse.

- True
 False

Correct

The bottom level of a multi-level index is like a single-level index. The bottom level is usually sparse when the table is sorted on the indexed column. The bottom level is dense when the table is not sorted on the indexed column.

?????

Query processing

- A table has 10 million rows.
 - Each row is 100 bytes.
 - Each index entry is 10 bytes.
 - Table and index blocks are 4 kilobytes.

The table contains 250,000 blocks = 10 million rows \times 100 bytes per row / 4,000 bytes per block.

A dense, single-level index contains 25,000 blocks = 10 million entries \times 10 bytes per entry / 44,000 bytes per block.

A dense, multi-level index contains 3 levels = log₂ 400 entries per index block 10,000,000 rows, rounded up.

A query searches for rows containing a specific value of an indexed column. Assuming query hit ratio is low:

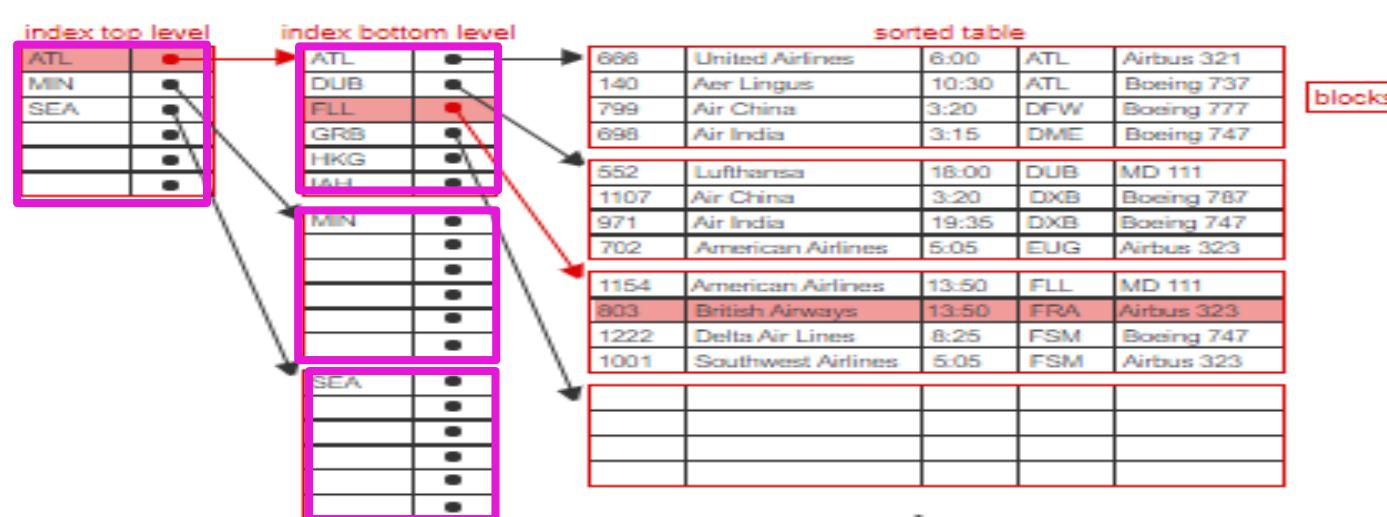
- A table scan reads at most 250,000 table blocks.
 - A single-level index scan reads at most 25,000 index blocks plus a few table blocks.
 - A binary search of a sorted, single-level index reads at most 24 index blocks plus a few table blocks.
 - A multi-level index search reads 3 index blocks plus a few table blocks.

The multi-level index search reads one index block per level. Usually the top two levels are small and retained in memory. Since the index has three levels, the query reads just one index block from storage media.

Because multi-level indexes are faster than single-level indexes on most queries, databases commonly use multi-level rather than single-level indexes.

PARTICIPATION ACTIVITY

15.2.4: Query processing with a multi-level index.



```
SELECT AirlineName  
FROM Flight  
WHERE DepartureAirport = 'FRA'
```

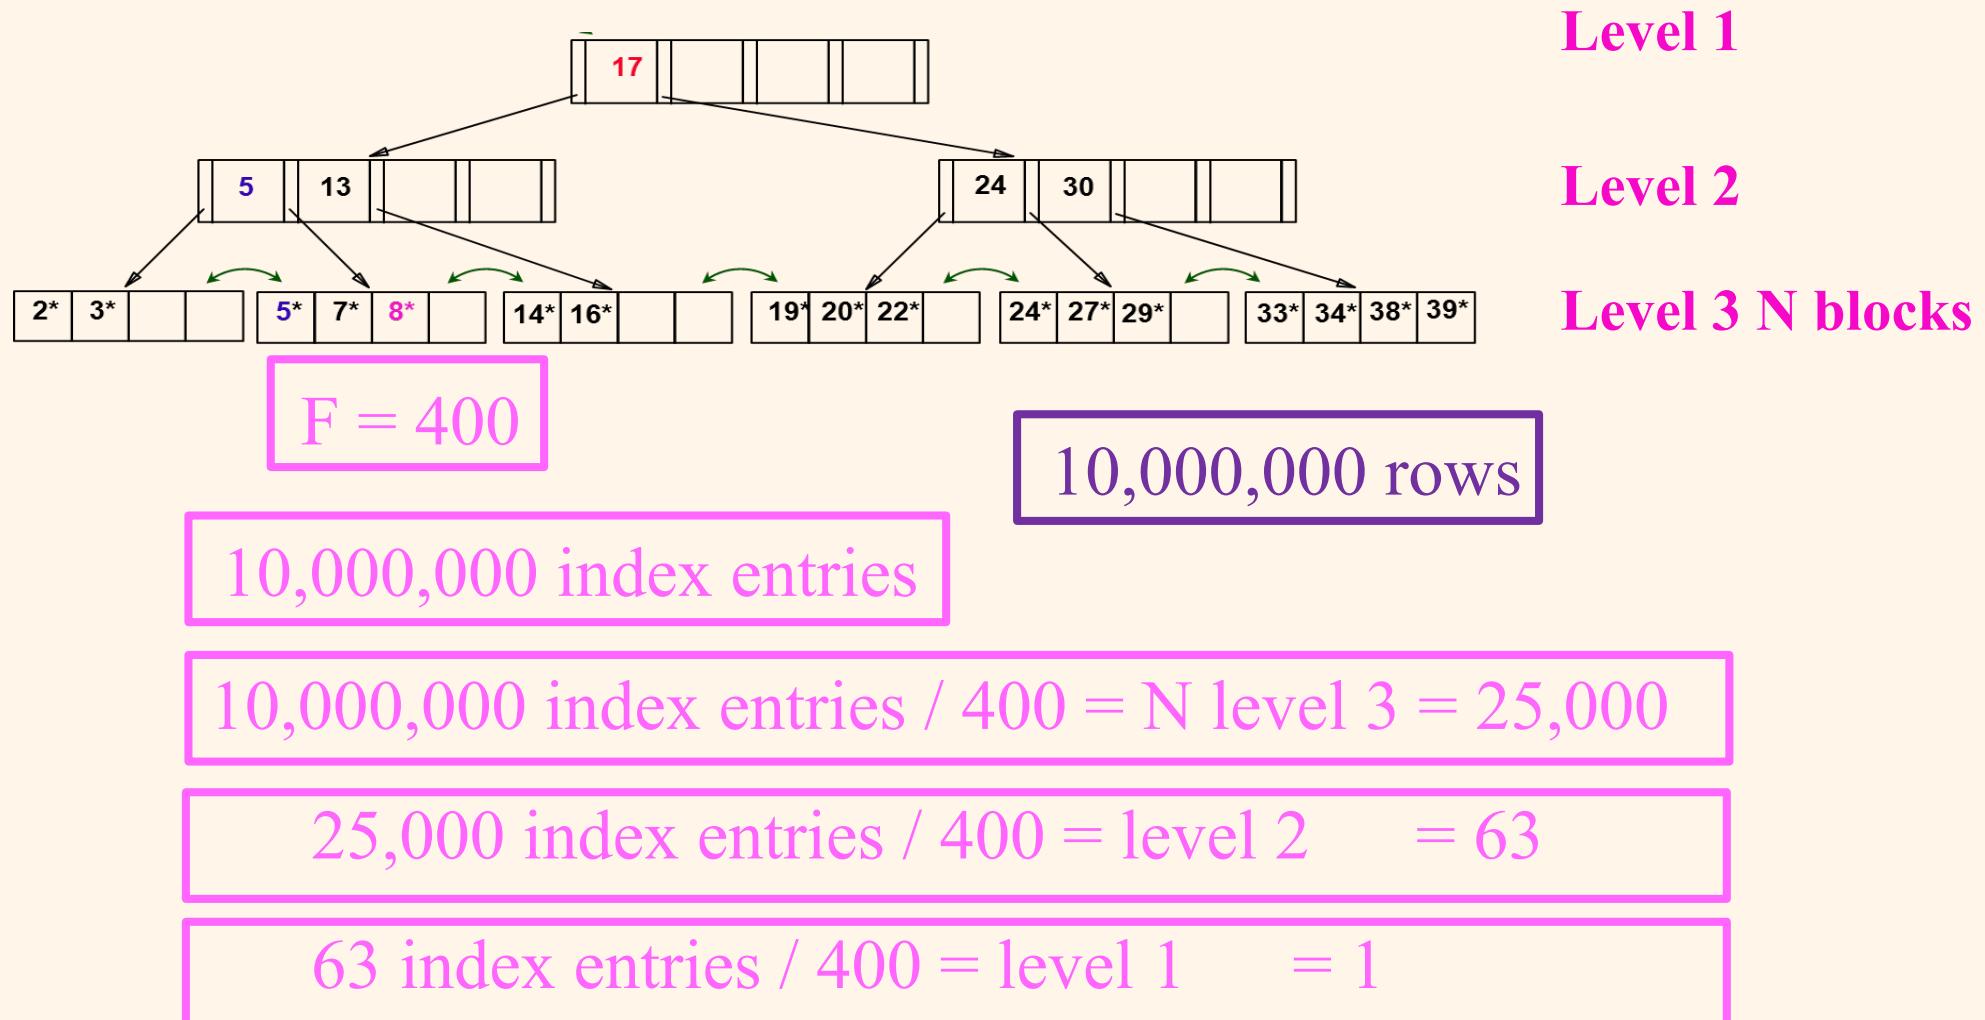
Multilevel Index



Number of levels

A dense index has more bottom-level entries than a sparse index, and may have more levels. Assuming a table with 10 million rows and 400 index entries per block, a dense index has three levels:

- Level 3 is dense and has 25,000 blocks = $10 \text{ million rows} / 400 \text{ index entries per block}$
- Level 2 is sparse and has 63 blocks = $25,000 \text{ level 3 blocks} / 400 \text{ index entries per block}$
- Level 1 is sparse and has one block containing 63 index entries.



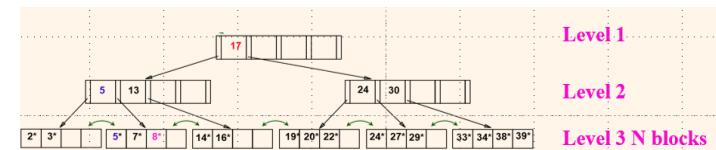
Number of levels

A dense index has more bottom-level entries than a sparse index, and may have more levels. Assuming a table with 10 million rows and 400 index entries per block, a dense index has three levels:

- Level 3 is dense and has 25,000 blocks = 10 million rows / 400 index entries per block
- Level 2 is sparse and has 63 blocks = 25,000 level 3 blocks / 400 index entries per block
- Level 1 is sparse and has one block containing 63 index entries.

The number of index entries per block is called the **fan-out** of a multi-level index. The number of levels in a multi-level index can be computed from fan-out, number of rows, and rows per block

- For a dense index, number of levels = $\log_{\text{fan-out}}(\text{number of rows})$
- For a sparse index, number of levels = $\log_{\text{fan-out}}(\text{number of rows} / \text{rows per block})$



In both cases, log is a fractional number and must be rounded up to the nearest integer. Both formulas assume minimal free space in the index.

Dense indexes usually have four levels or less. Sparse indexes usually have three levels or less.

Table 15.2.1: Number of levels.

dense index		rows	
		1 million	1 billion
index entries per block	100	3	5
	400	3	4

sparse index		rows	
		1 million	1 billion
index entries per block	100	3	4
	400	2	3

2*	3*			5*	7*	8*		14*	16*			19*	20*	22*		24*	27*	29*		33*	34*	38*	3
----	----	--	--	----	----	----	--	-----	-----	--	--	-----	-----	-----	--	-----	-----	-----	--	-----	-----	-----	---

$$F = 8,000 / 26 \\ = 308 \sim 300$$

5*	7*	8*	
----	----	----	--

8,000 bytes

5*	7*	8*	
----	----	----	--

26 bytes

PARTICIPATION ACTIVITY

15.2.3: Number of levels.

Assume a table has 1,000,000 rows. Index entries are 26 bytes, and index blocks are 8 kilobytes.

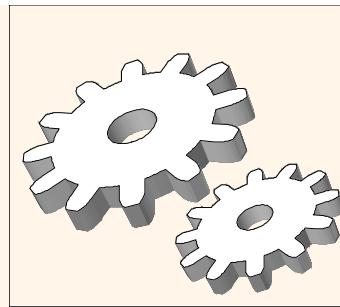
1) What is the fan-out for a multi-level index?

- Approximately 200
- Approximately 300
- Approximately 400

Correct

Fan-out is the number of entries per block = 8,000 bytes per index block / 26 bytes per index entry, or roughly 300.

?????



ISAM (*Indexed Sequential Access Method*)

14.1 Physical design

Present

Note

MySQL storage engines

Logical design specifies tables, columns, and keys. The logical design process is described elsewhere in this material. **Physical design** specifies indexes, table structures, and partitions. Physical design affects query performance but never affects query results.

A **storage engine** or **storage manager** translates instructions generated by a query processor into low-level commands that access data on storage media. Storage engines support different index and table structures, so physical design is dependent on a specific storage engine.

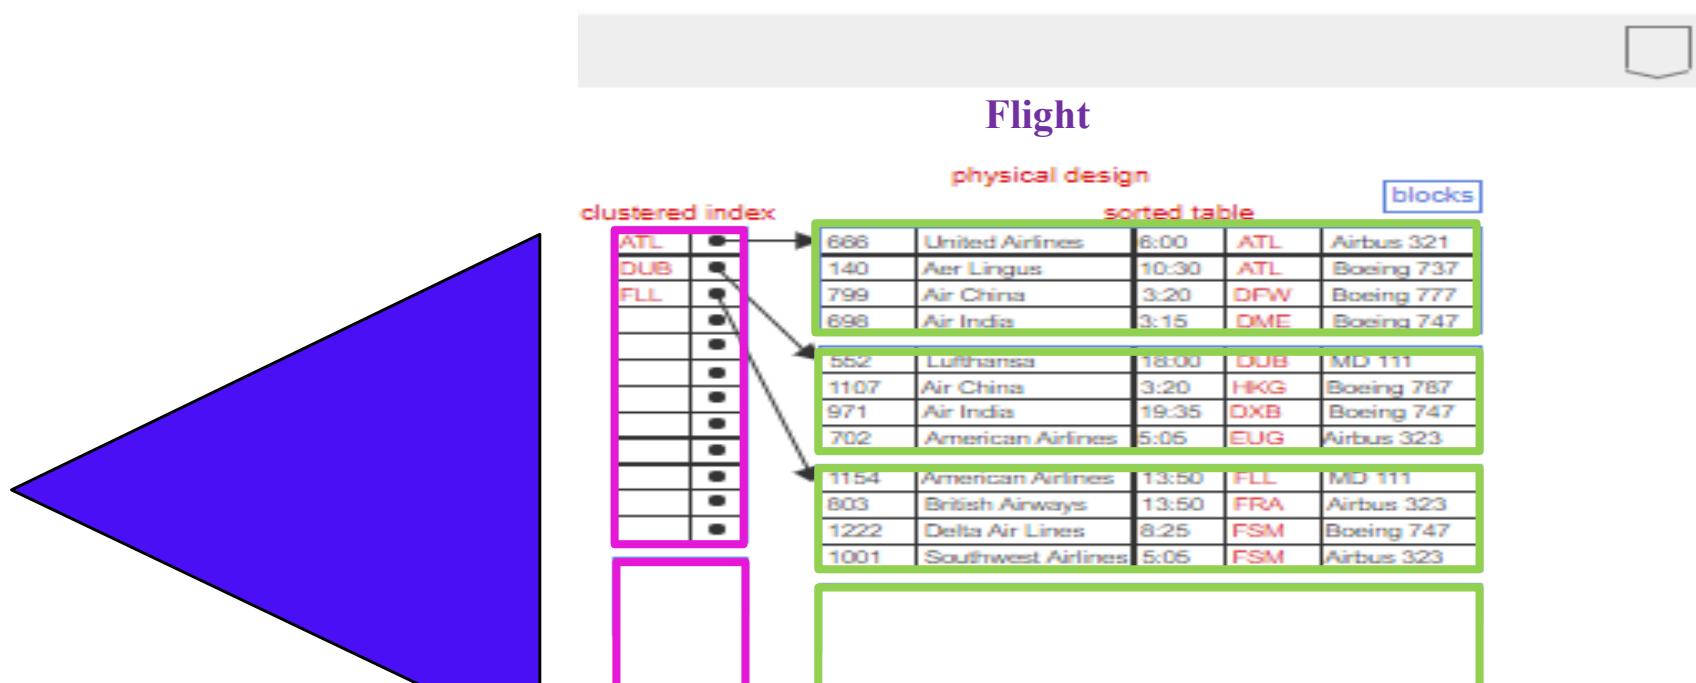
MySQL can be configured with several different storage engines, including:

- InnoDB is the default storage engine installed with the MySQL download. InnoDB has full support for transaction management, foreign keys, referential integrity, and locking.
- MyISAM has limited transaction management and locking capabilities. MyISAM is commonly used for analytic applications with limited data updates.
- MEMORY stores all data in main memory. MEMORY is used for fast access with databases small enough to fit in main memory.

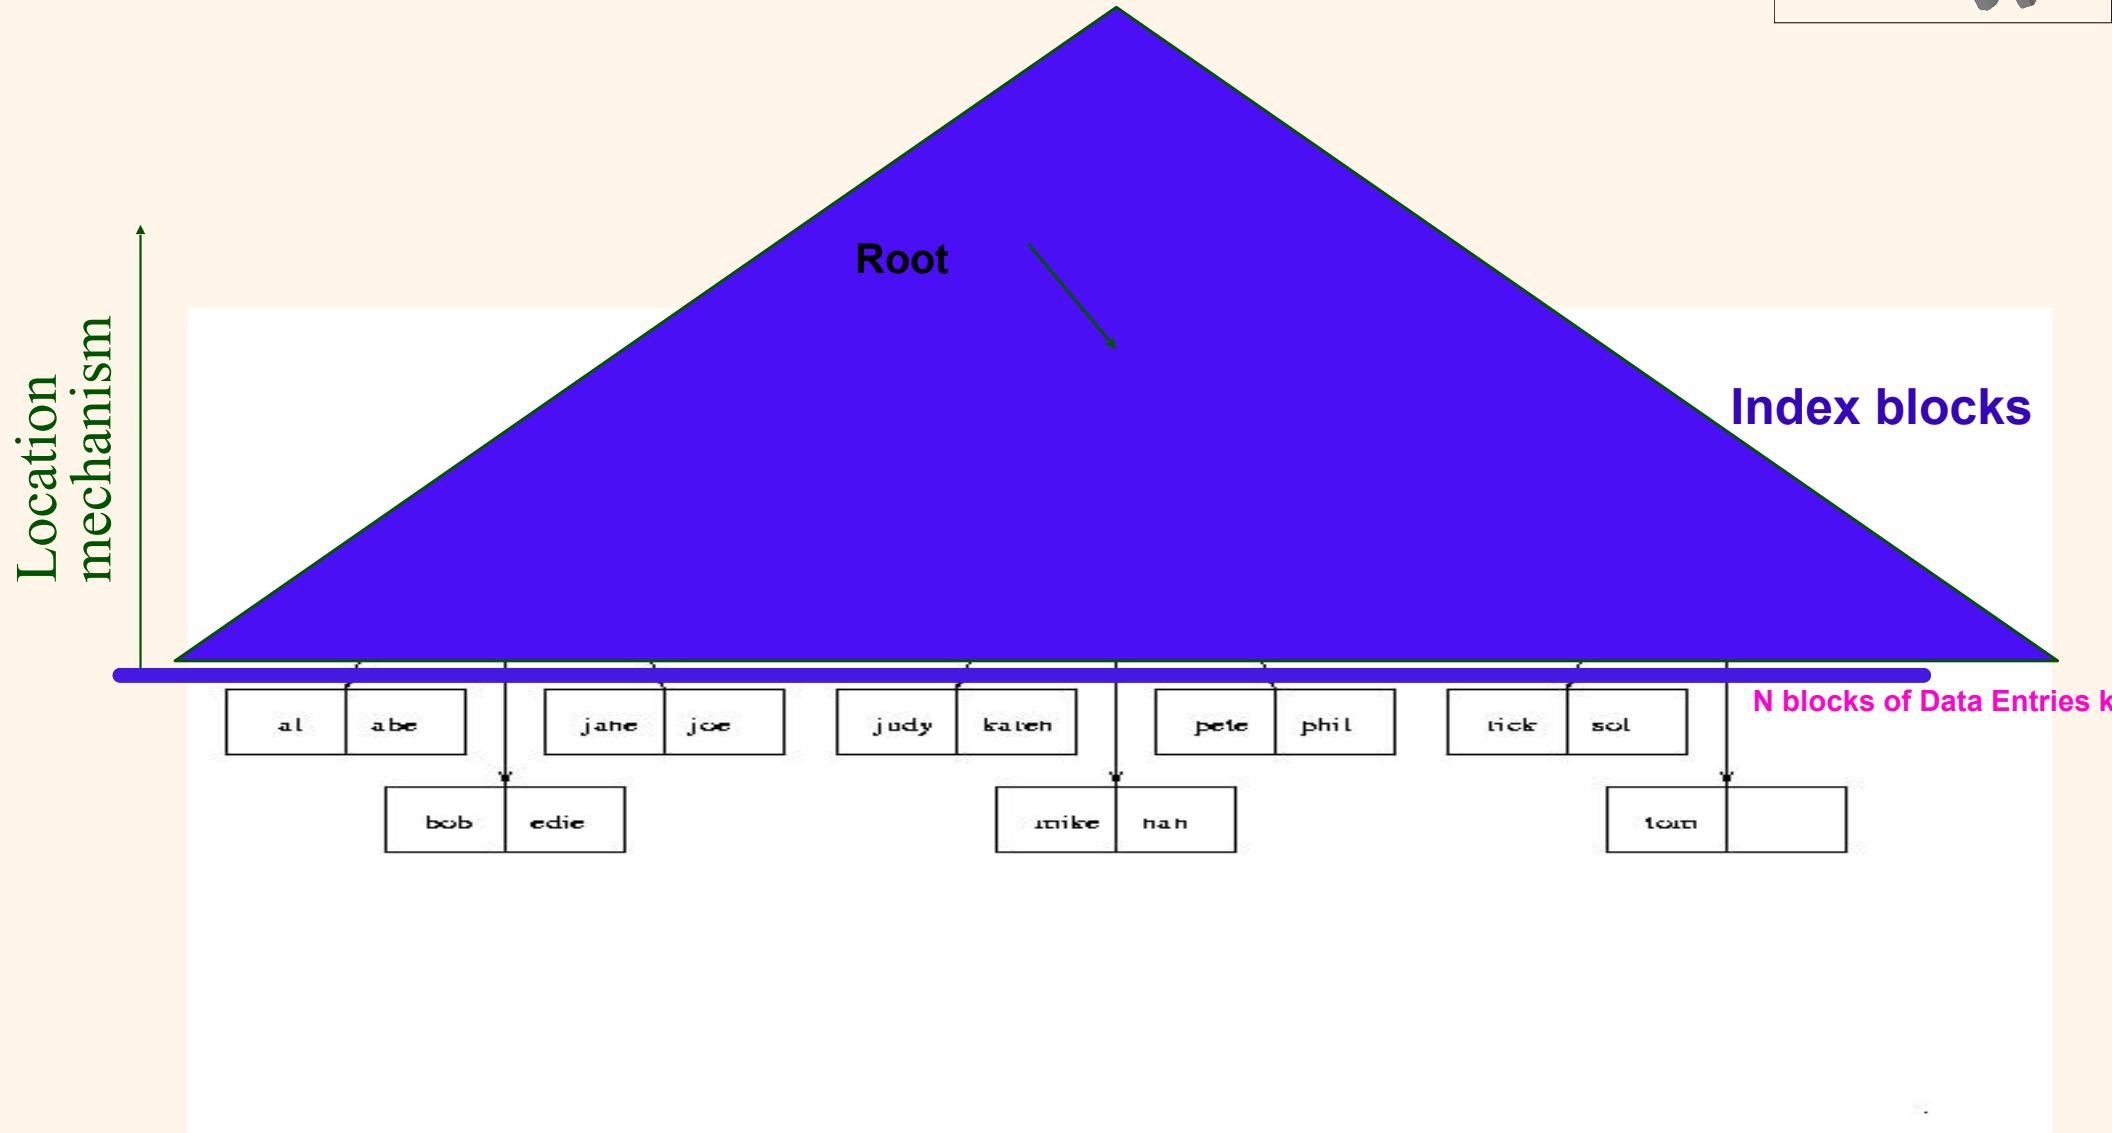
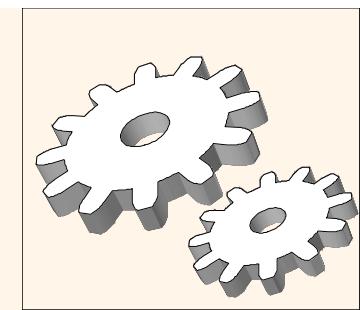
Different databases and storage engines support different table structures and index types. Ex:

- Table structure. Oracle Database supports heap, sorted, hash, and cluster tables. MySQL with InnoDB supports only heap and sorted tables.
- Index type. MySQL with InnoDB or MyISAM supports only B-tree indexes. MySQL with MEMORY supports both B-tree and hash indexes.

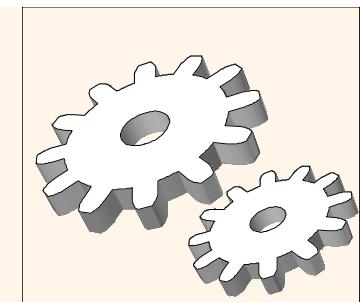
This section describes the physical design process and statements for MySQL with InnoDB. The process and statements can be adapted to other databases and storage engines, but details depend on supported index and table structures.



Index Sequential Access Method - ISAM



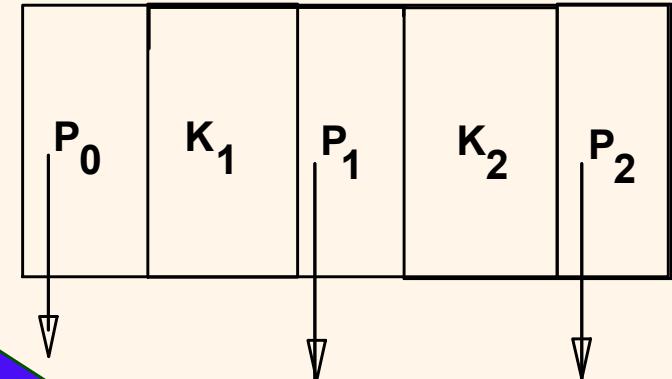
Example ISAM Tree



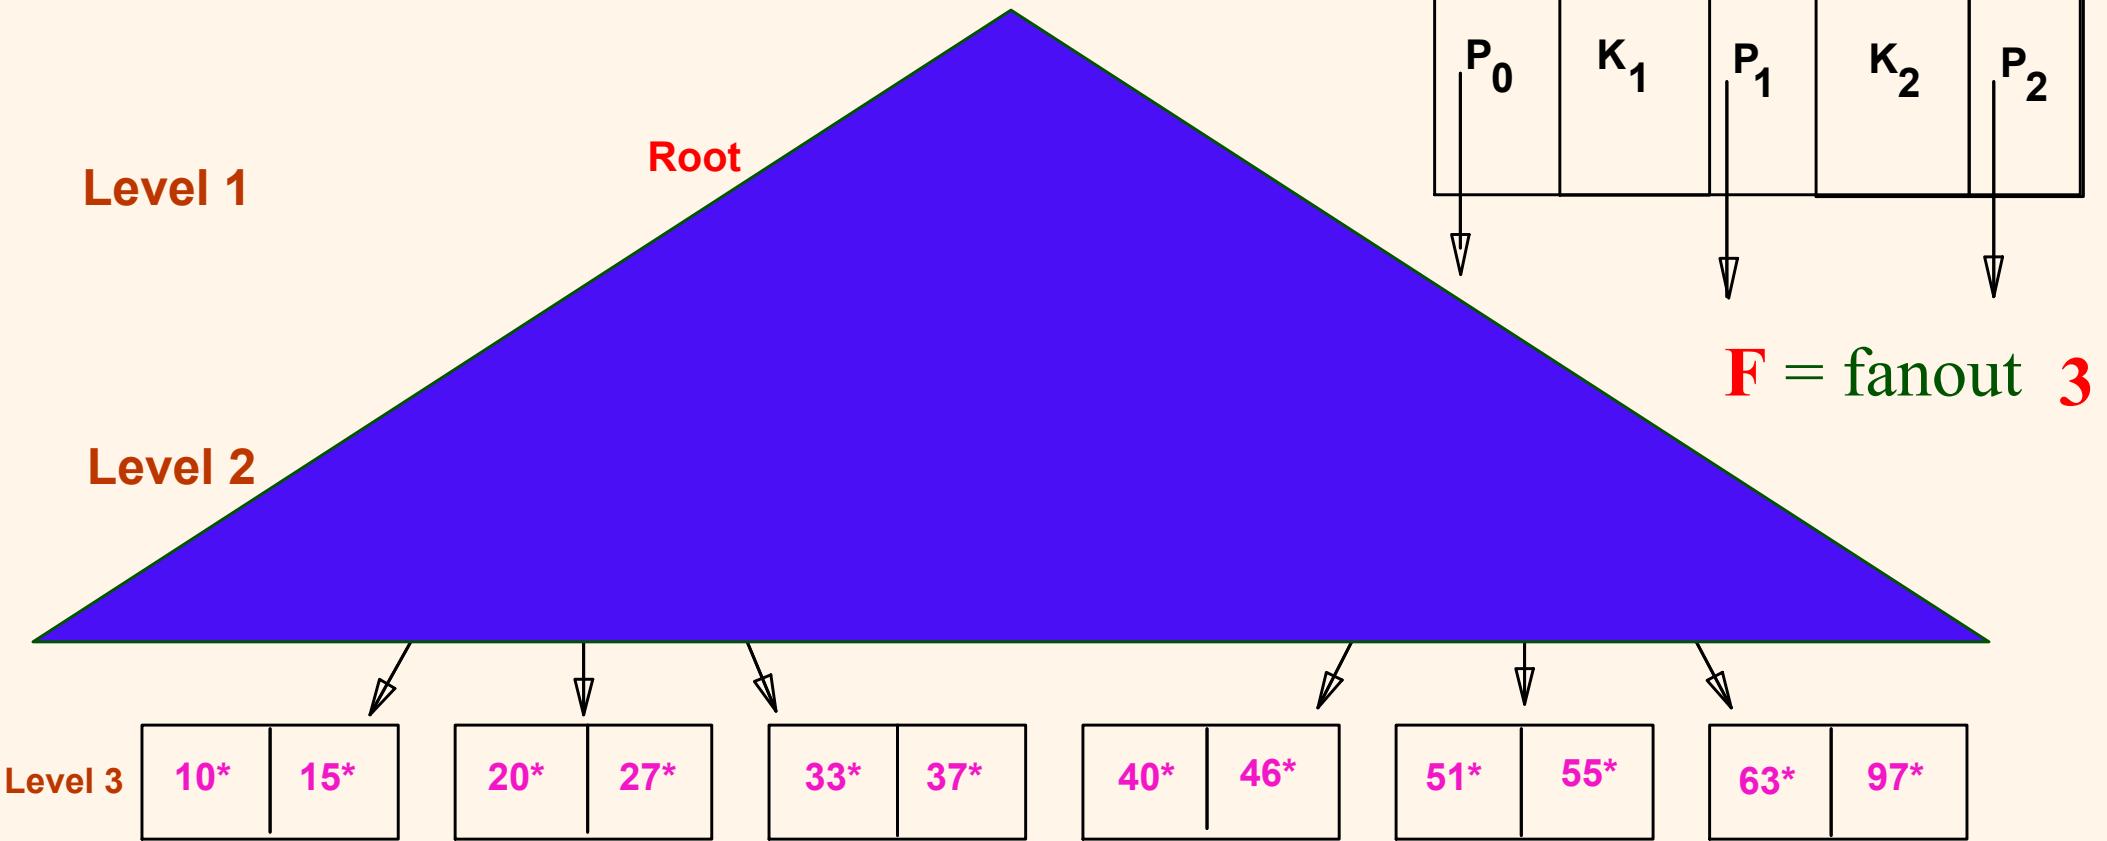
- Each node/page(Index k & Data Entry k^*) can hold 2 entries.

$m = \text{number of keys } 2$

Index entry



$F = \text{fanout } 3$



$N = 6 \text{ blocks of Data Entries } k^*$

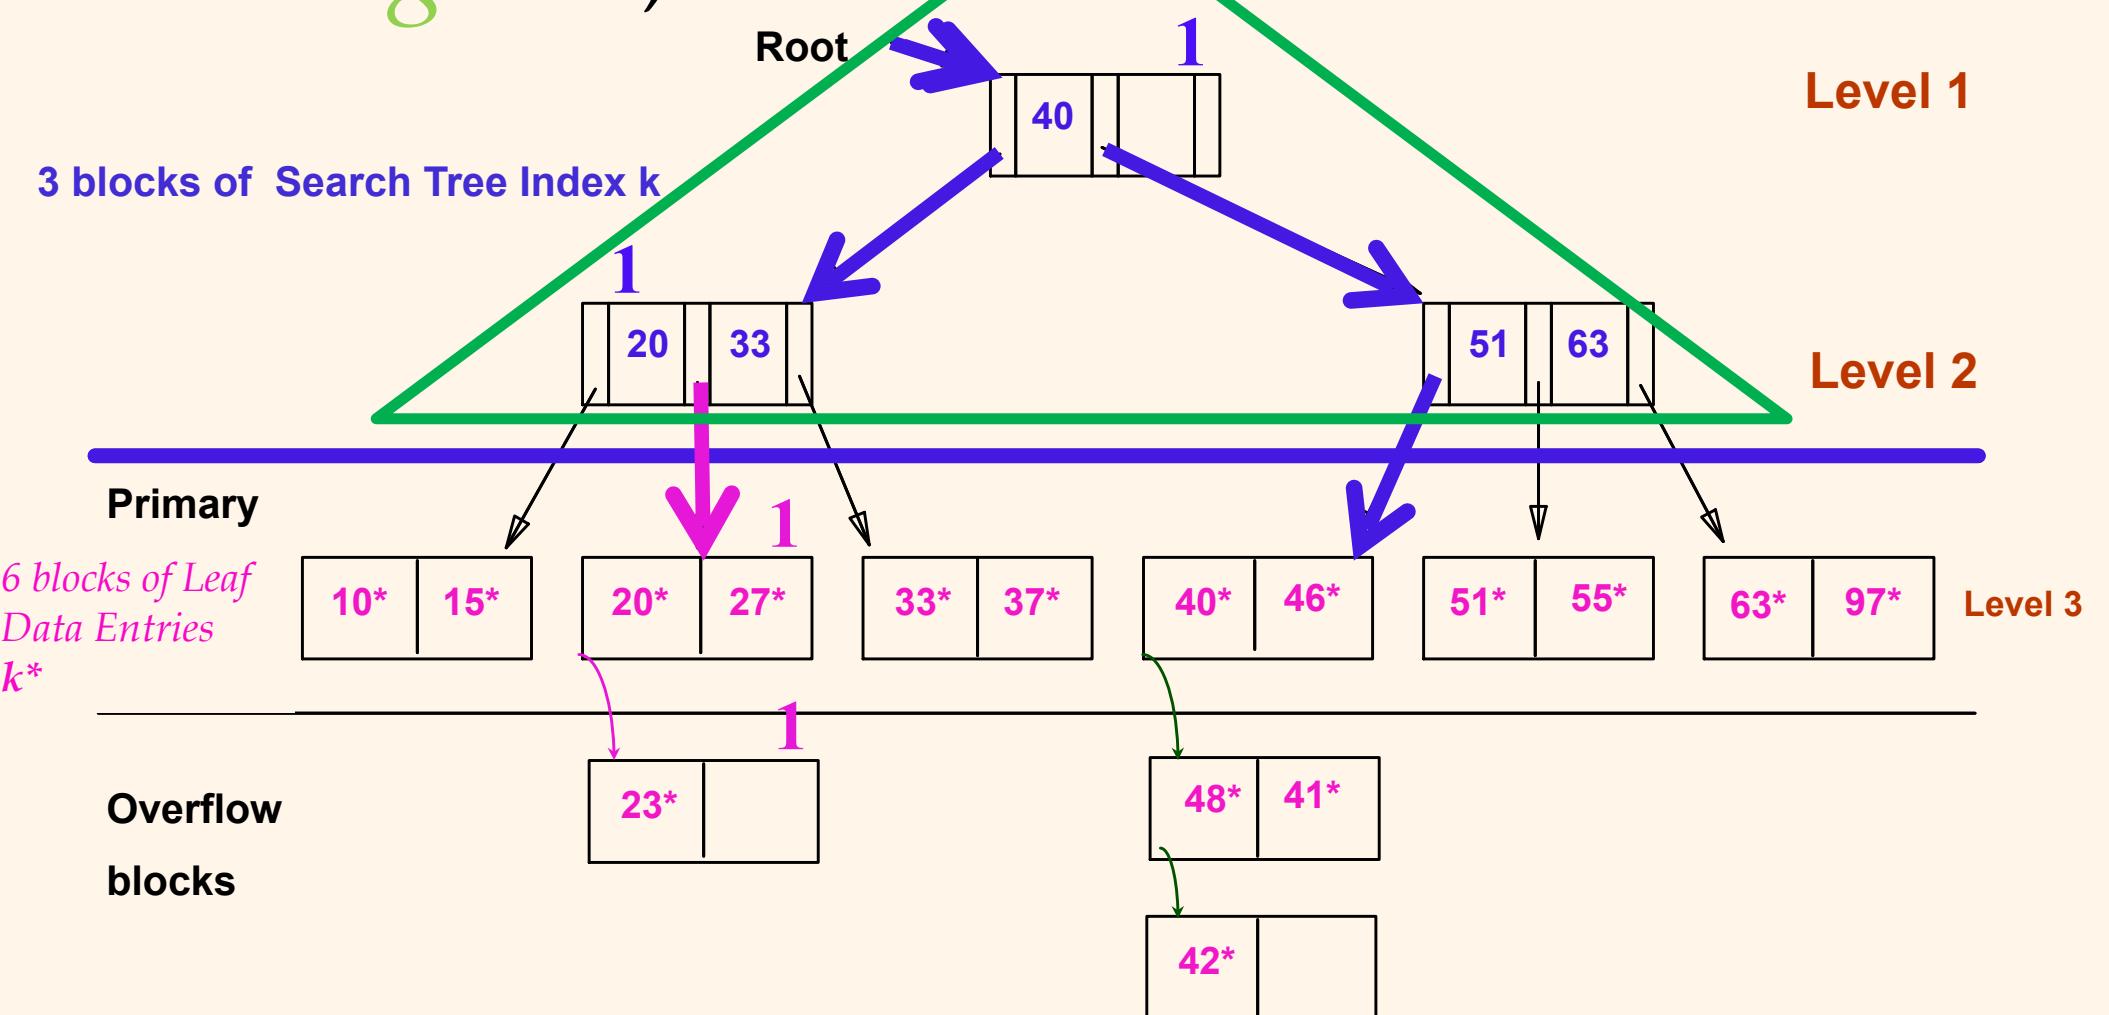
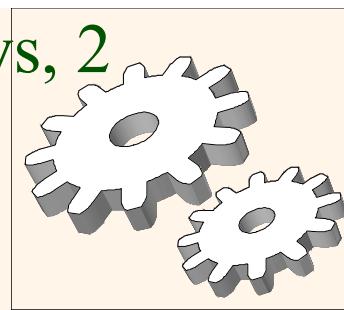
Inserting 23*

Inserting 48*

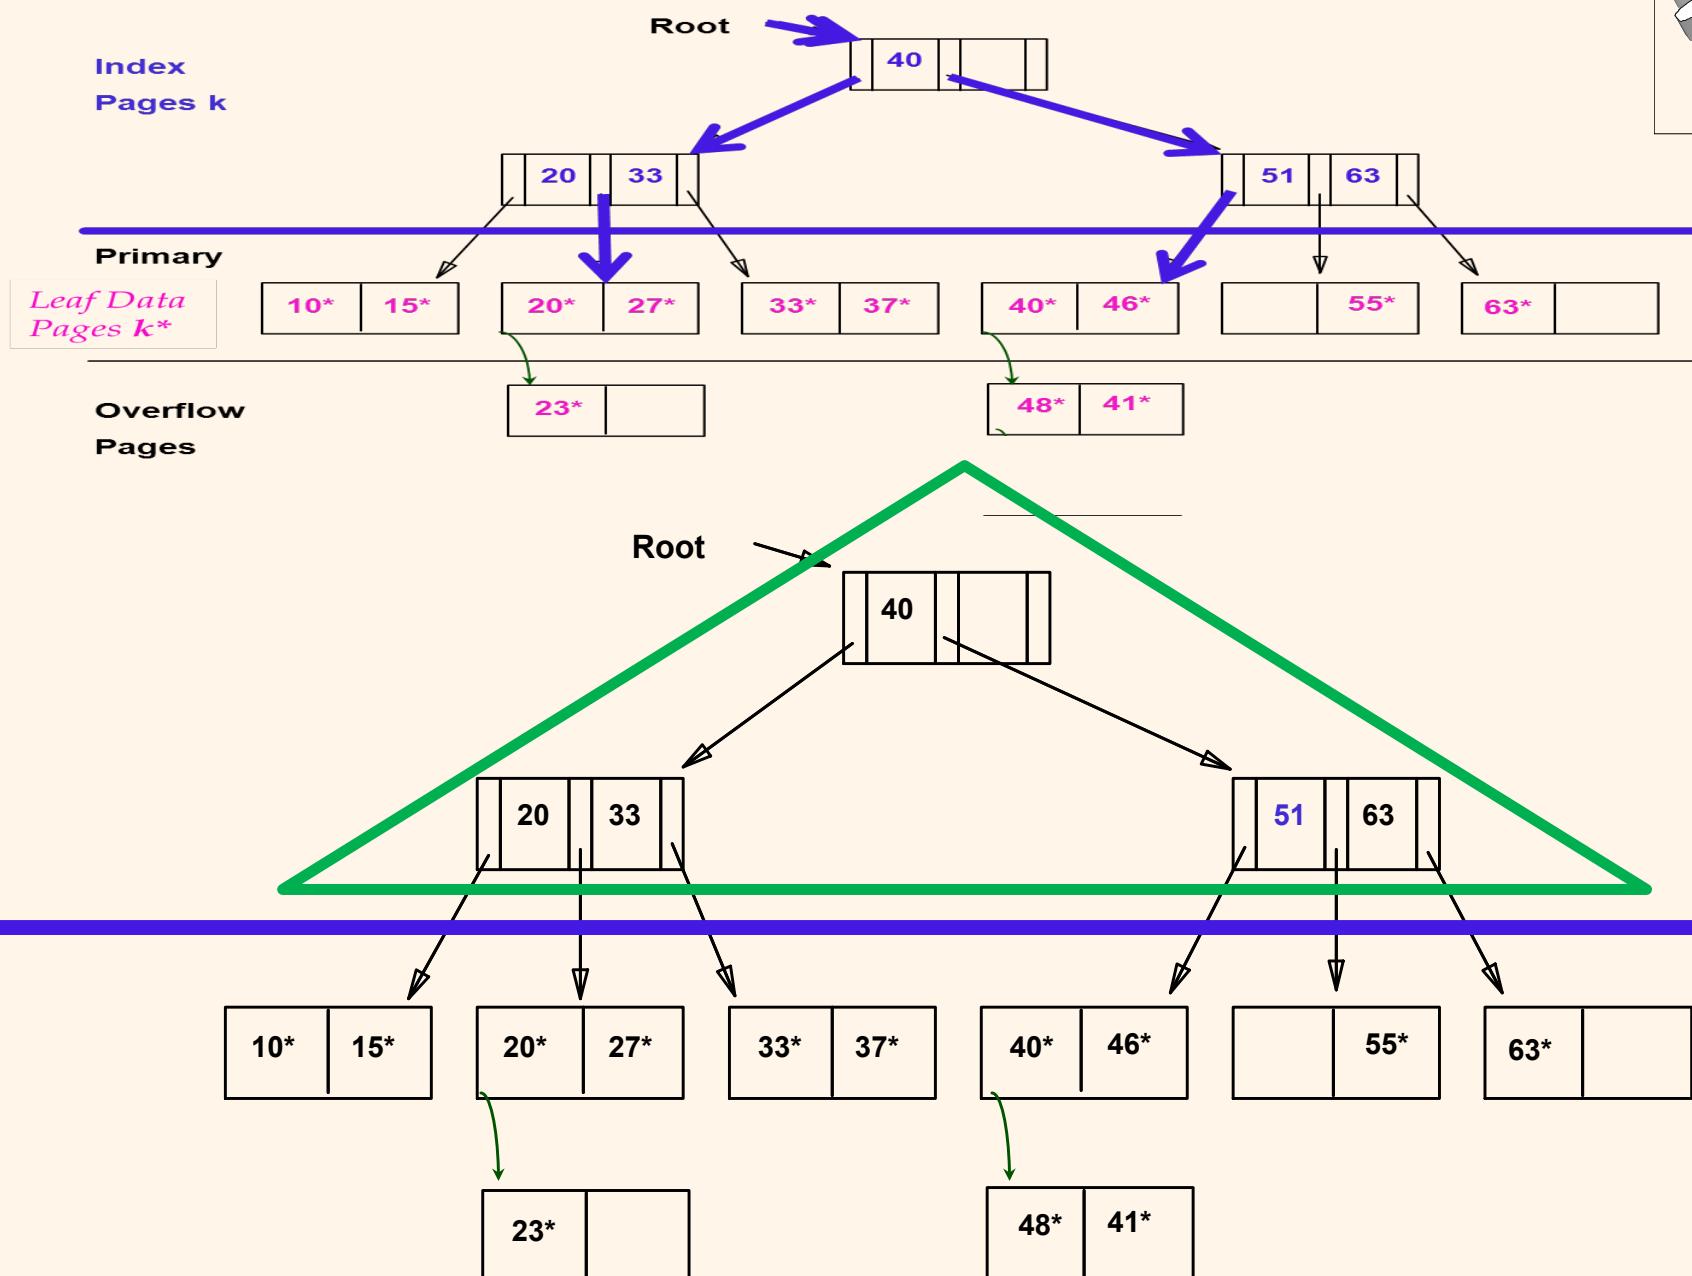
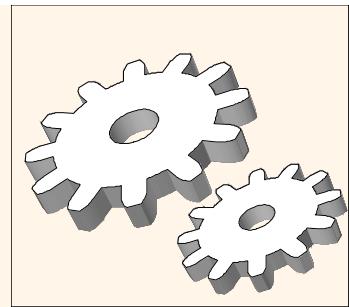
Inserting 41*, 42*

m = number of keys, 2

F = fanout, 3



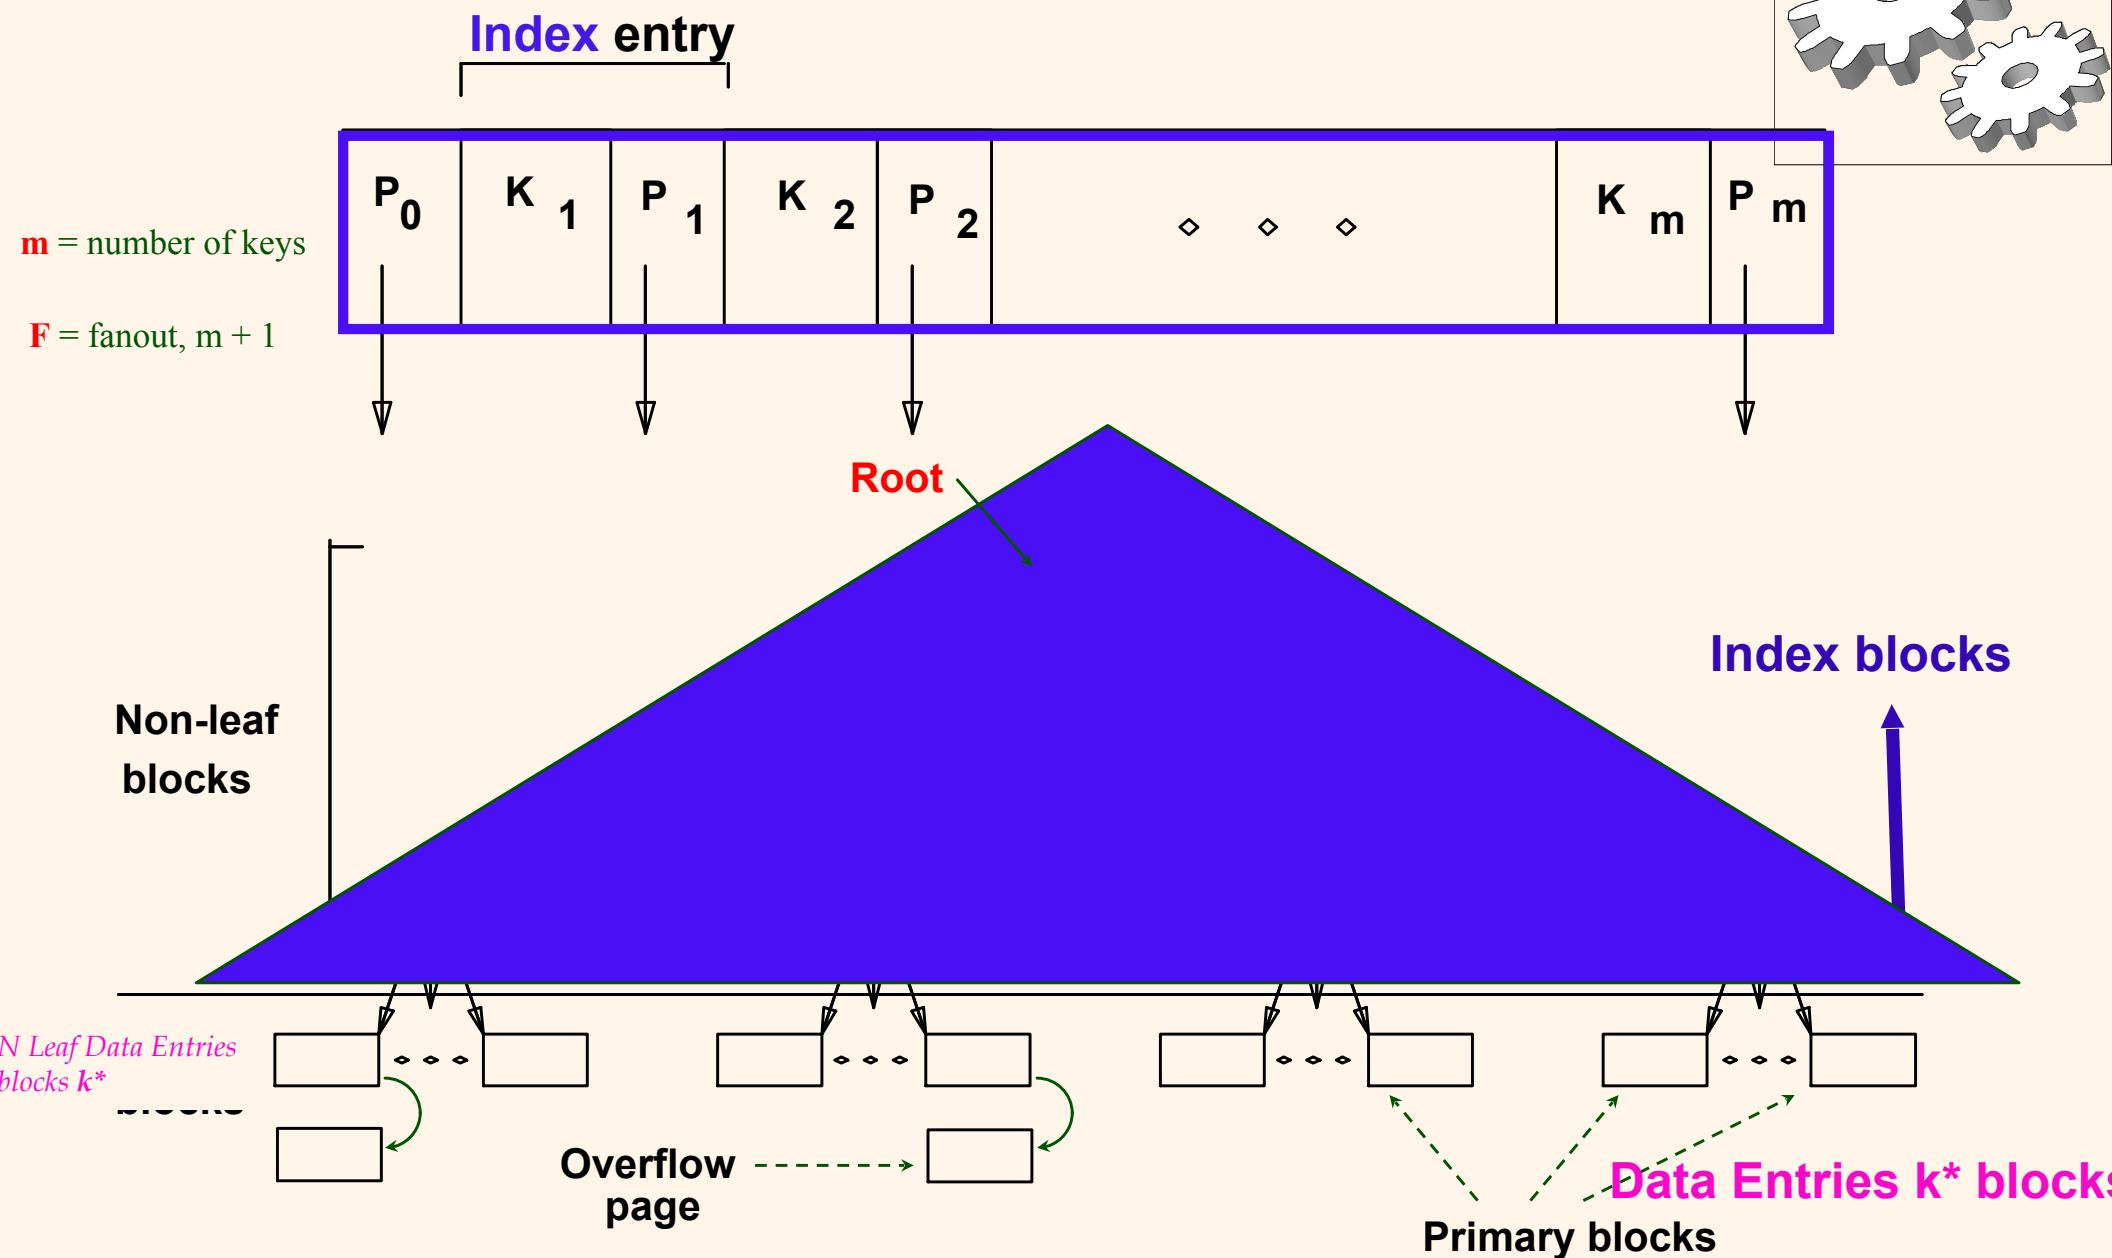
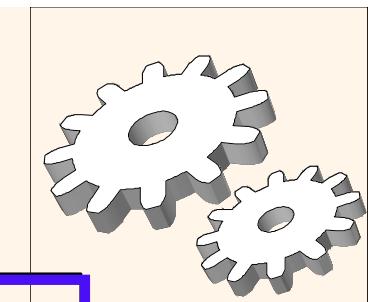
... Then Deleting 42*, 51*, 97*



?

Note that 51 appears in Index levels, but not in leaf!

ISAM (Indexed Sequential Access Method)



?

N Leaf Data Entries blocks contain Data Entries k^ .*

Comments on ISAM

File creation: Leaf ($\text{Data Entries } k^*$) blocks allocated sequentially,
Sorted by search key; then **Index blocks** allocated,
then space for **Overflow blocks**.

Index blocks k: <search key value, pid>;
`direct' search for *Leaf blocks*.

Leaf blocks k*: <search key value, rid>;
`points' to *Data Record Page*.

Search k^* : Start at **Root**; use key comparisons to go to **Leaf Page**.

$$\text{Cost} = \lceil \log_F N \rceil + 1 = \text{Tree Level}$$

$F = \# \text{ pointers/ Index Page}$, $N = \# \text{ Leaf (Data Entry } k^*$

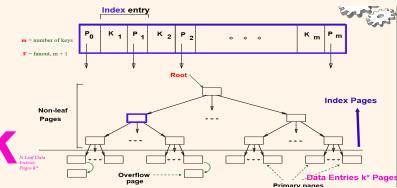
Insert k^* : Find **Leaf (Data) Page** belongs to, and put it there.

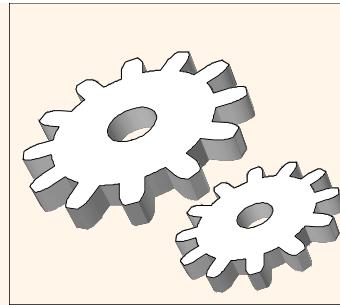
Delete k^* : Find and remove from **Leaf Page**; if empty **Overflow Page**, de-allocates

Leaf (Data) blocks

Index blocks

Overflow blocks





B⁺ Tree: Most Widely Used Index

14.1 Physical design

Present

Note

MySQL storage engines

Logical design specifies tables, columns, and keys. The logical design process is described elsewhere in this material. **Physical design** specifies indexes, table structures, and partitions. Physical design affects query performance but never affects query results.

A **storage engine** or **storage manager** translates instructions generated by a query processor into low-level commands that access data on storage media. Storage engines support different index and table structures, so physical design is dependent on a specific storage engine.

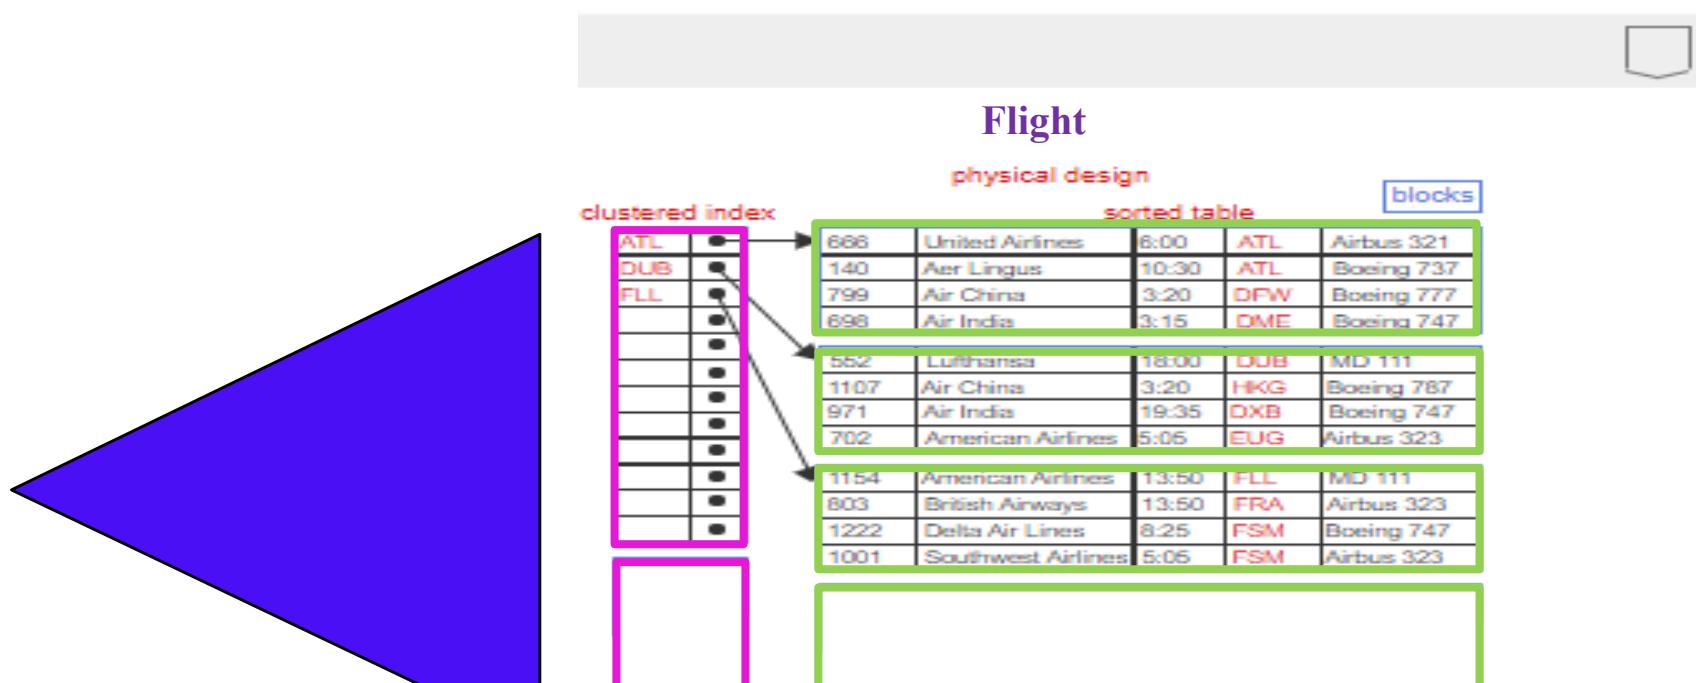
MySQL can be configured with several different storage engines, including:

- InnoDB is the default storage engine installed with the MySQL download. InnoDB has full support for transaction management, foreign keys, referential integrity, and locking.
- MyISAM has limited transaction management and locking capabilities. MyISAM is commonly used for analytic applications with limited data updates.
- MEMORY stores all data in main memory. MEMORY is used for fast access with databases small enough to fit in main memory.

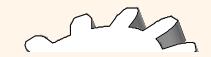
Different databases and storage engines support different table structures and index types. Ex:

- Table structure. Oracle Database supports heap, sorted, hash, and cluster tables. MySQL with InnoDB supports only heap and sorted tables.
- Index type. MySQL with InnoDB or MyISAM supports only B-tree indexes. MySQL with MEMORY supports both B-tree and hash indexes.

This section describes the physical design process and statements for MySQL with InnoDB. The process and statements can be adapted to other databases and storage engines, but details depend on supported index and table structures.



B⁺ Tree: Most Widely Used Index



Cost to get to the Leaf Page is the Level of the B⁺ Tree

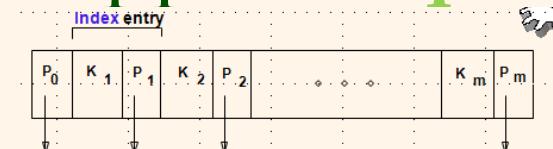
$$\lceil \log_F N \rceil + 1$$

Keeps Tree *height-balanced*. (F = fanout, N = # leaf blocks)

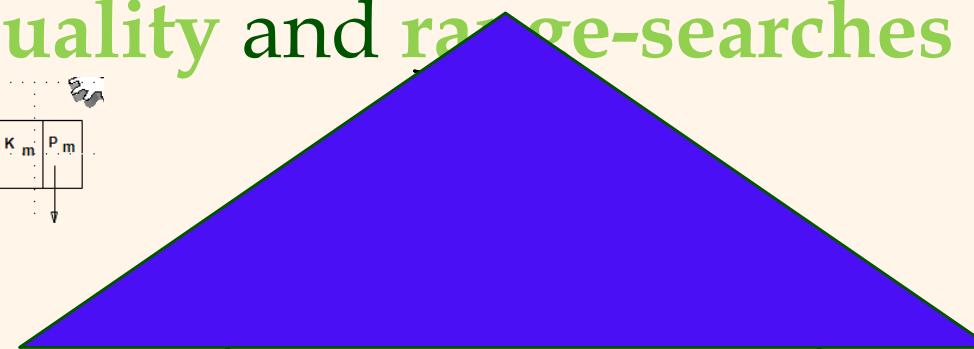
Minimum 50% occupancy (except for Root Node/Page). Each Node contains $d \leq m \leq 2d$ entries. The parameter d is called the *order* of the Tree.

$F = m+1$ - number of Pointers P_0 to P_m to other Index blocks

Supports equality and range-searches efficiently.

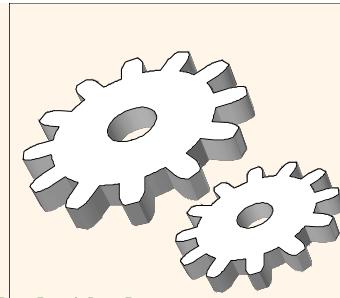


Index Page
(Direct search)



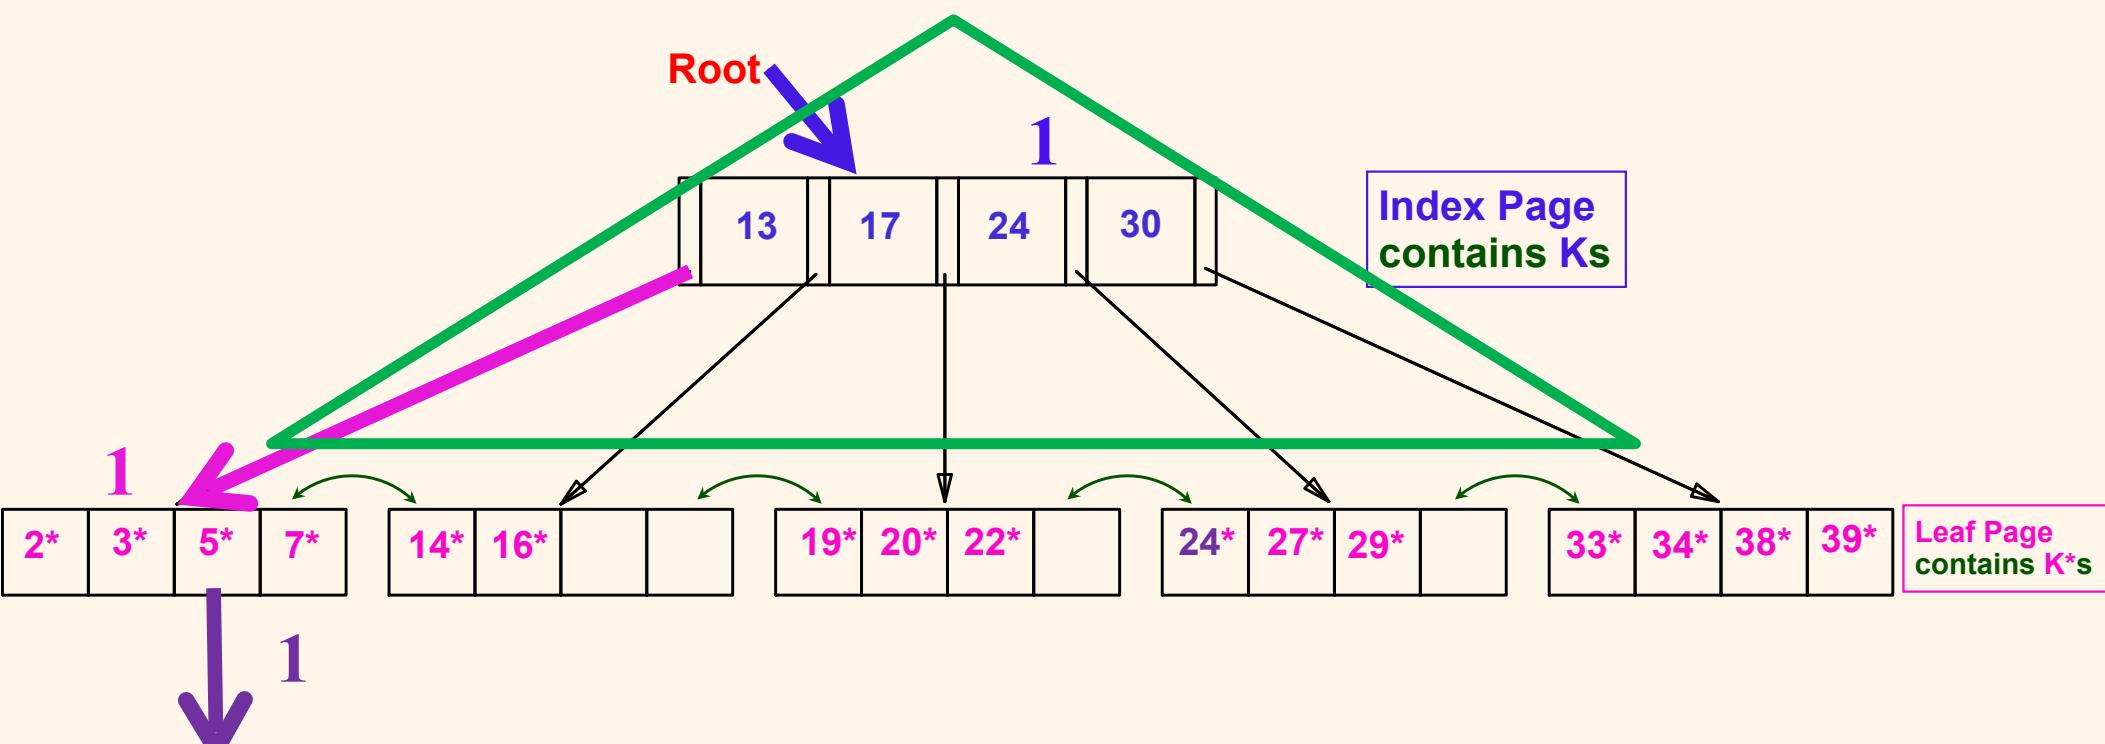
N blocks of Data Entries k^*
("Sequence set")

Example B^+ Tree Index



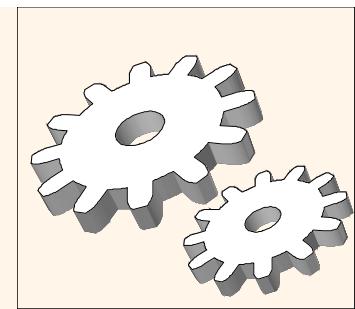
- ❖ Search for k^* begins at **Root**, and **key** comparisons direct it to a **leaf Page of Data Entries k^*** (as in **ISAM**).
- ❖ Search for 5^*

Search for 15^* , all data entries $\geq 24^*$...



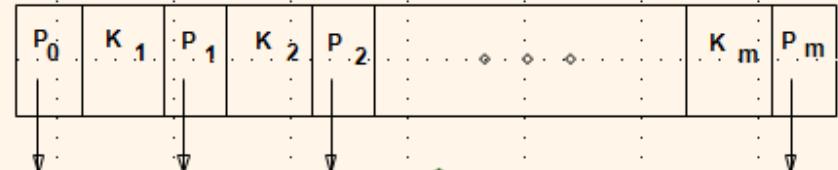
Based on the search for 15^* , we know it is not in the Tree!

B⁺ Tree Indexes in Practice



Typical order: $d = 100$. Typical fill-factor: 67%.

- average fanout $F = 133 = m + 1$

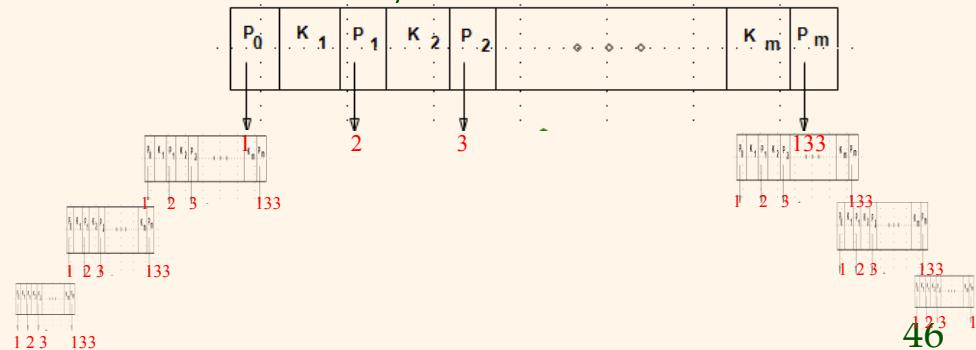


Typical capacities:

- Height 4: $133^4 = 312,900,700$ k*
- Height 3: $133^3 = 2,352,637$ k*

Can often hold top levels in **Buffer Pool**:

- Level 1 = 1 block = 8 KBytes
- Level 2 = 133 blocks = 133×8 KByte = 1064 KB ~ 1 MByte
- Level 3 = $133 \times 133 = 17,689$ blocks ~ 133 MBytes

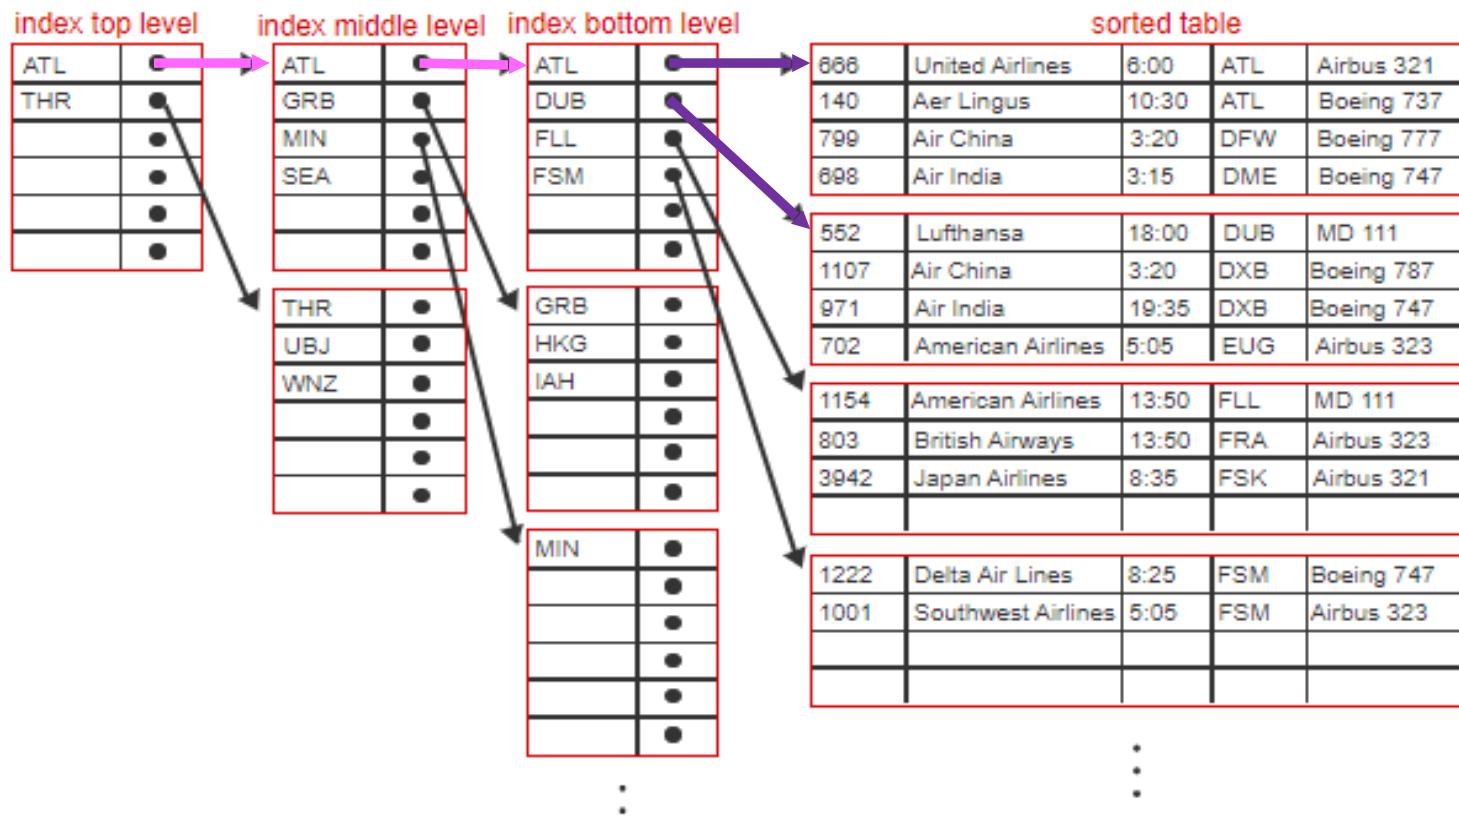


* Root is at Level 1!

Balanced indexes



Flight



```
INSERT INTO Flight (FlightNumber, AirlineName, DepartureTime, DepartureAirport, AircraftType)
VALUES (3942, 'Japan Airlines', '8:35', 'FSK', 'Airbus 321')
```

PARTICIPATION ACTIVITY

15.2.7: Balanced indexes.

?????

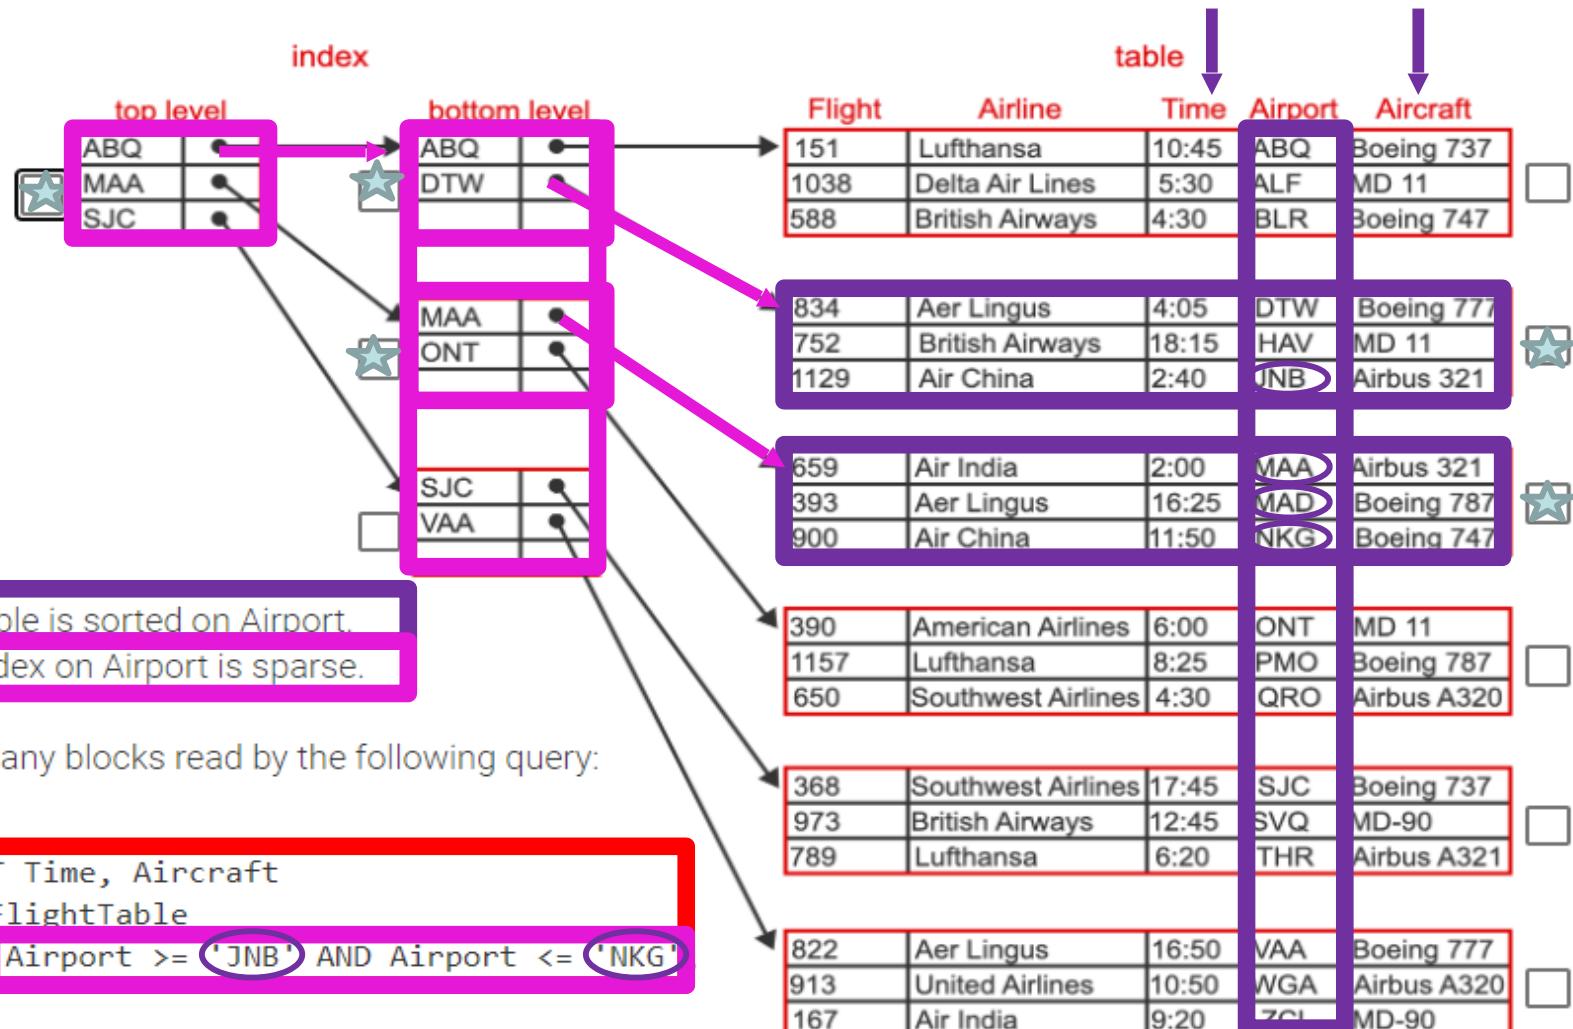
A table has a multi-level index with no free space in any index blocks.

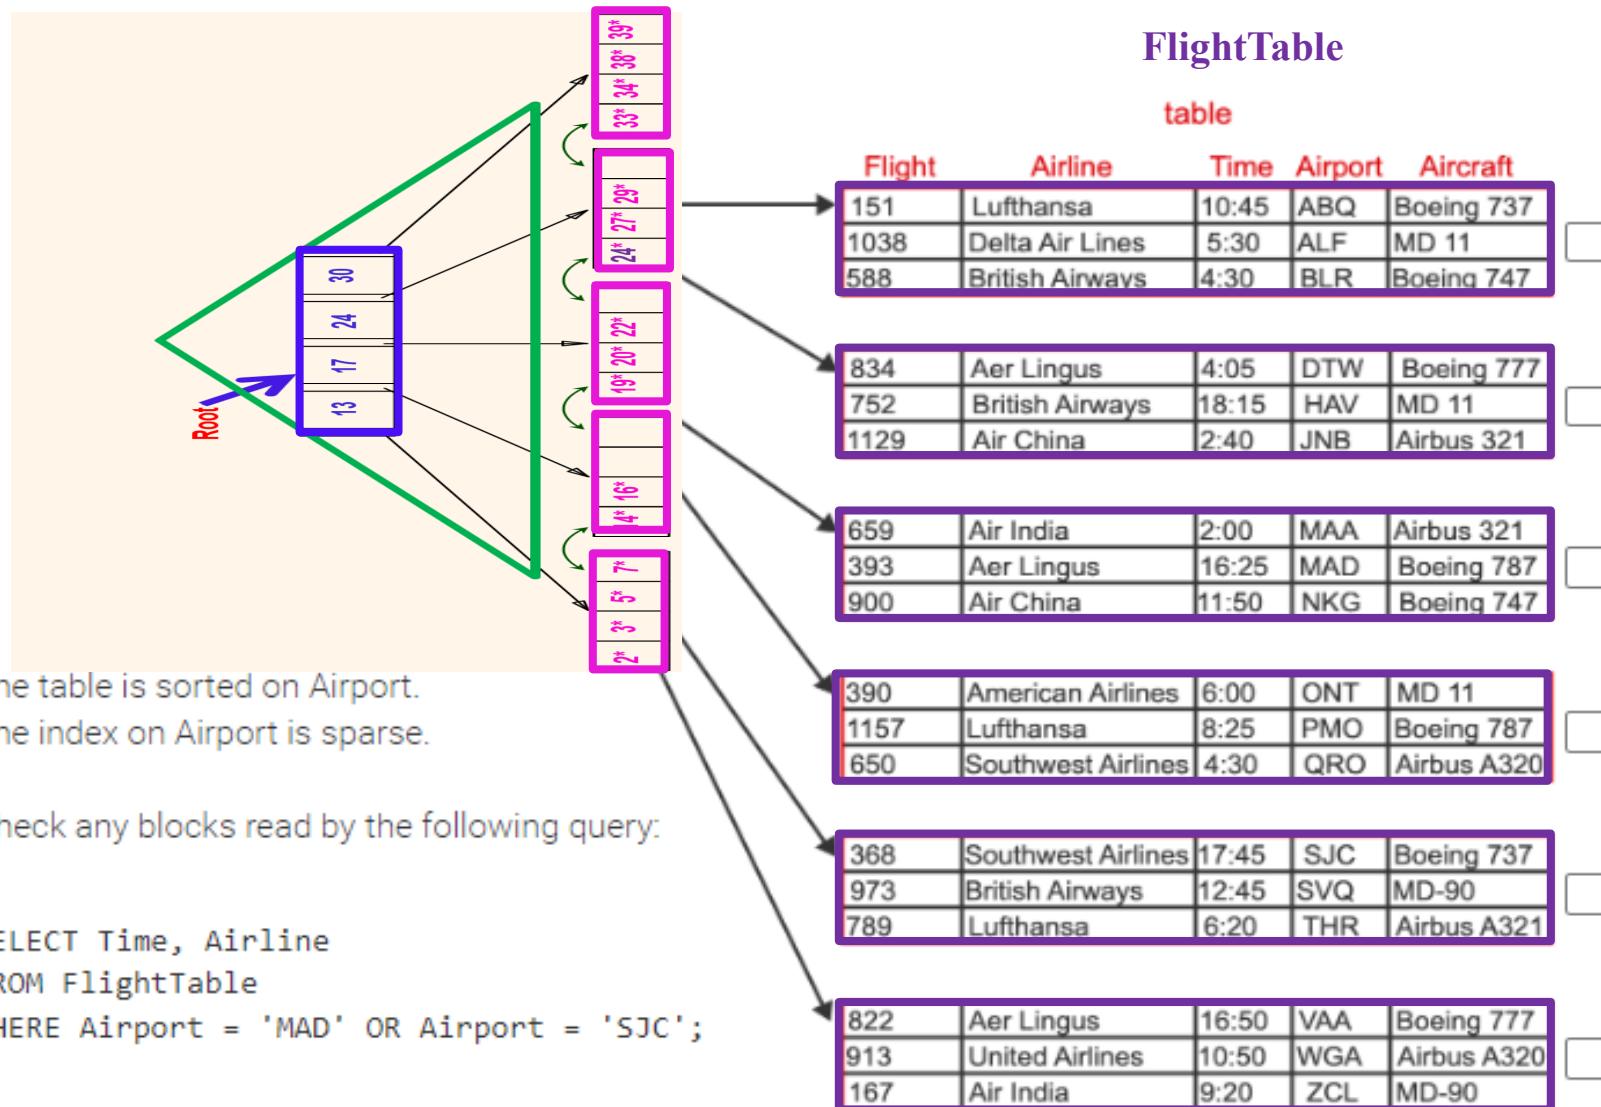
- If the index is sparse, an insert to the table always generates a new index level.

- True
 False

Correct

Since the index is sparse, a new index entry is generated only when the insert causes a table block split.

FlightTable



CHALLENGE
ACTIVITY

15.1.1: Multi-level indexes.

D

index

top level	
ABQ	●
MAA	●
SJC	●

bottom level	
ABQ	●
DTW	●

MAA	●
ONT	●

SJC	●
VAA	●

The table is sorted on Airport.
The index on Airport is sparse

Check any blocks read by the following query:

```
SELECT Aircraft, Time
FROM FlightTable
WHERE Airport = 'WGA' OR Airport = 'ZCL';
```

1

FlightTable
table

Flight	Airline	Time	Airport	Aircraft
151	Lufthansa	10:45	ABQ	Boeing 737
1038	Delta Air Lines	5:30	ALF	MD 11
588	British Airways	4:30	BLR	Boeing 747

834	Aer Lingus	4:05	DTW	Boeing 777
752	British Airways	18:15	HAV	MD 11
1129	Air China	2:40	JNB	Airbus 321

659	Air India	2:00	MAA	Airbus 321
393	Aer Lingus	16:25	MAD	Boeing 787
900	Air China	11:50	NKG	Boeing 747

390	American Airlines	6:00	ONT	MD 11
1157	Lufthansa	8:25	PMO	Boeing 787
650	Southwest Airlines	4:30	QRO	Airbus A320

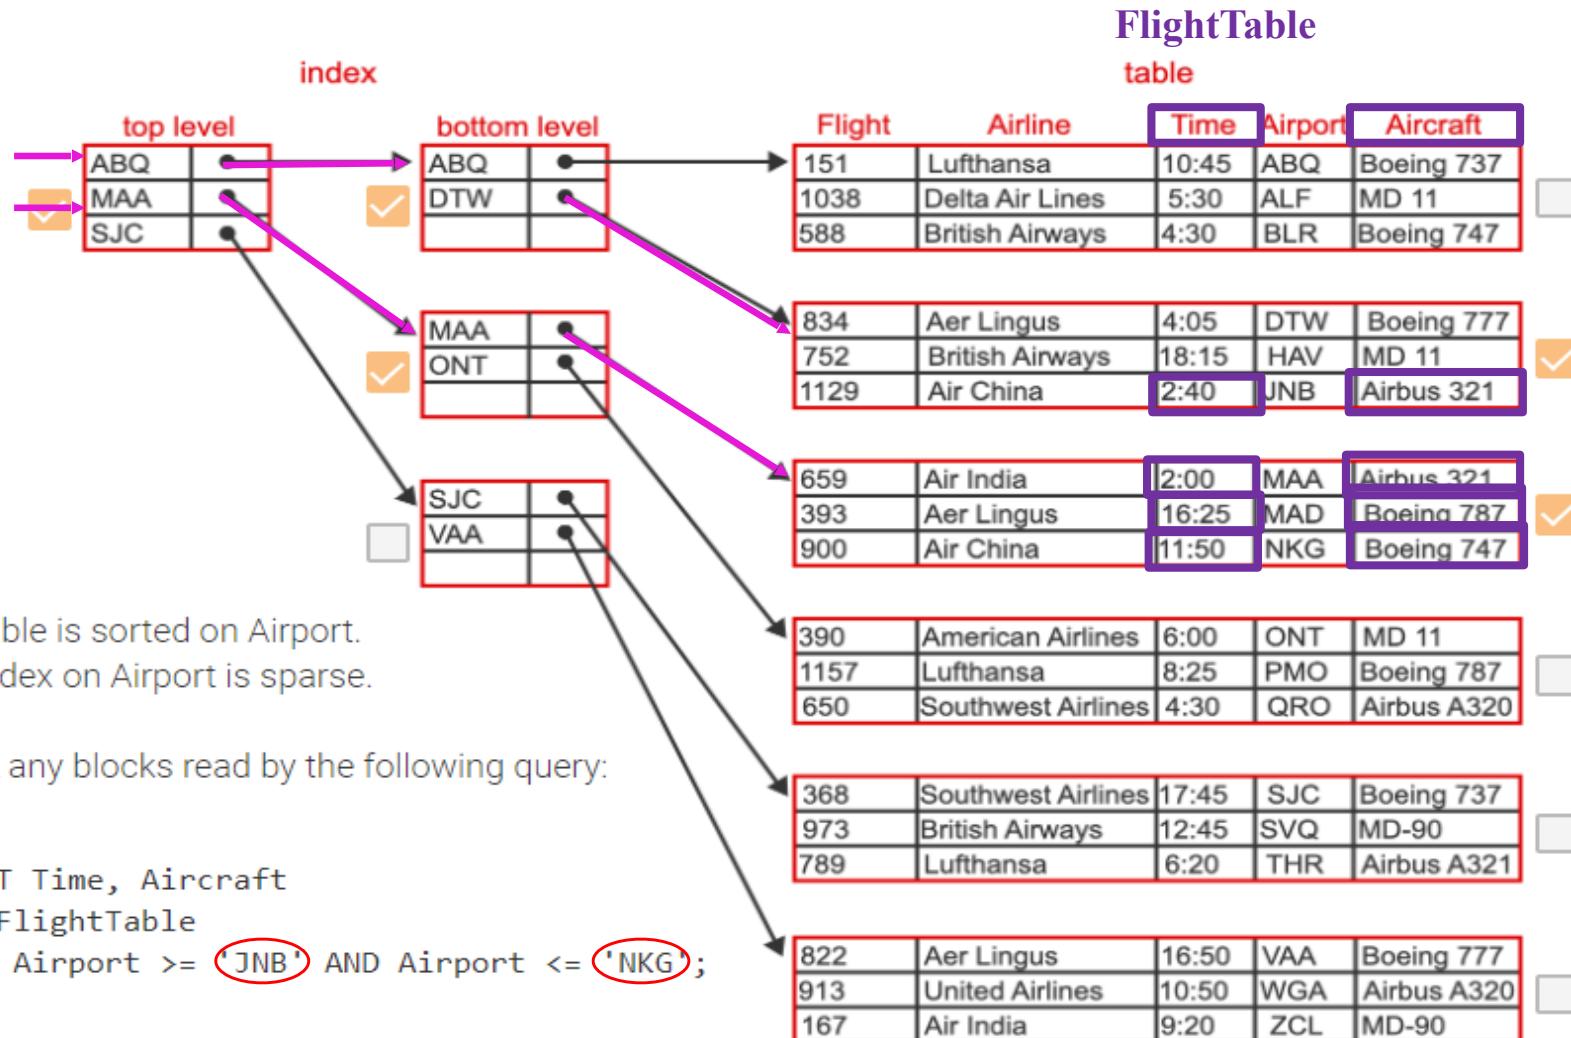
368	Southwest Airlines	17:45	SJC	Boeing 737
973	British Airways	12:45	SVQ	MD-90
789	Lufthansa	6:20	THR	Airbus A321

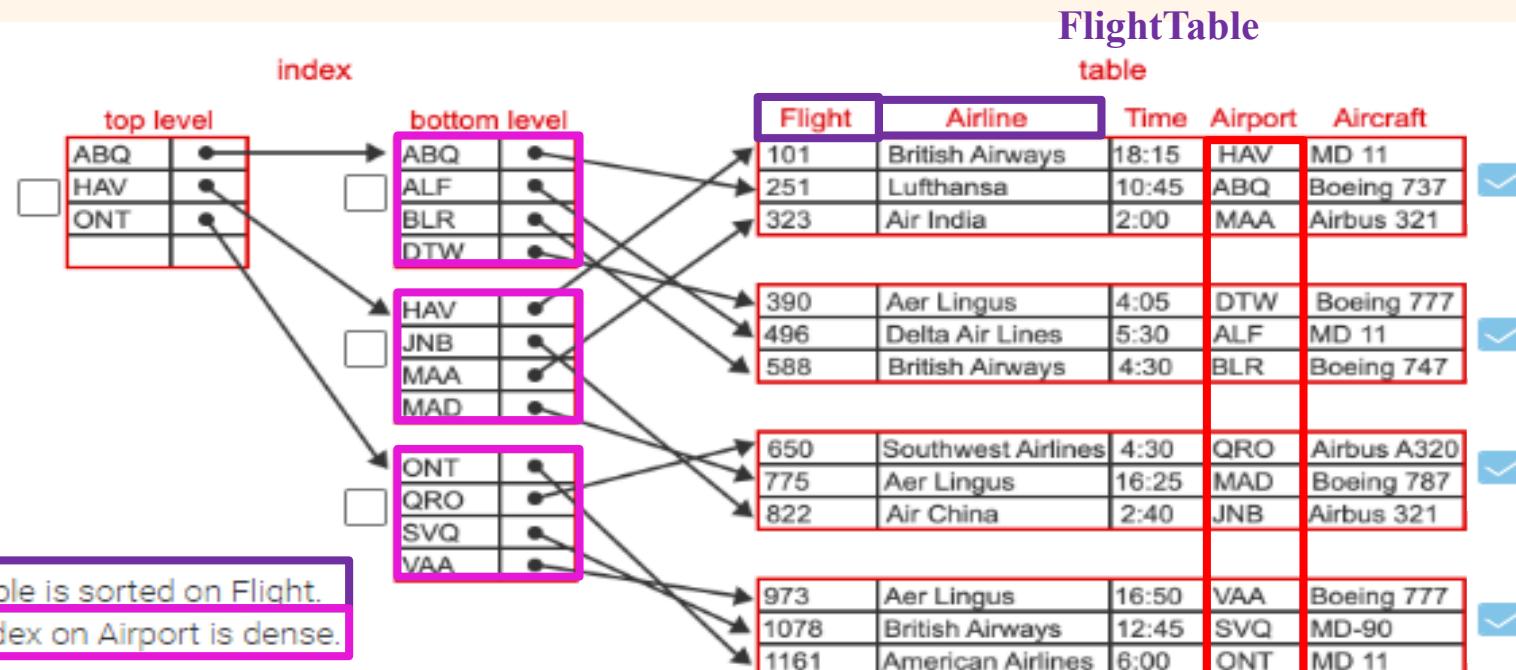
822	Aer Lingus	16:50	VAA	Boeing 777
913	United Airlines	10:50	WGA	Airbus A320
167	Air India	9:20	ZCL	MD-90

Aircraft	Time
Airbus A320	10:50
MD-90	9:20

D

?????



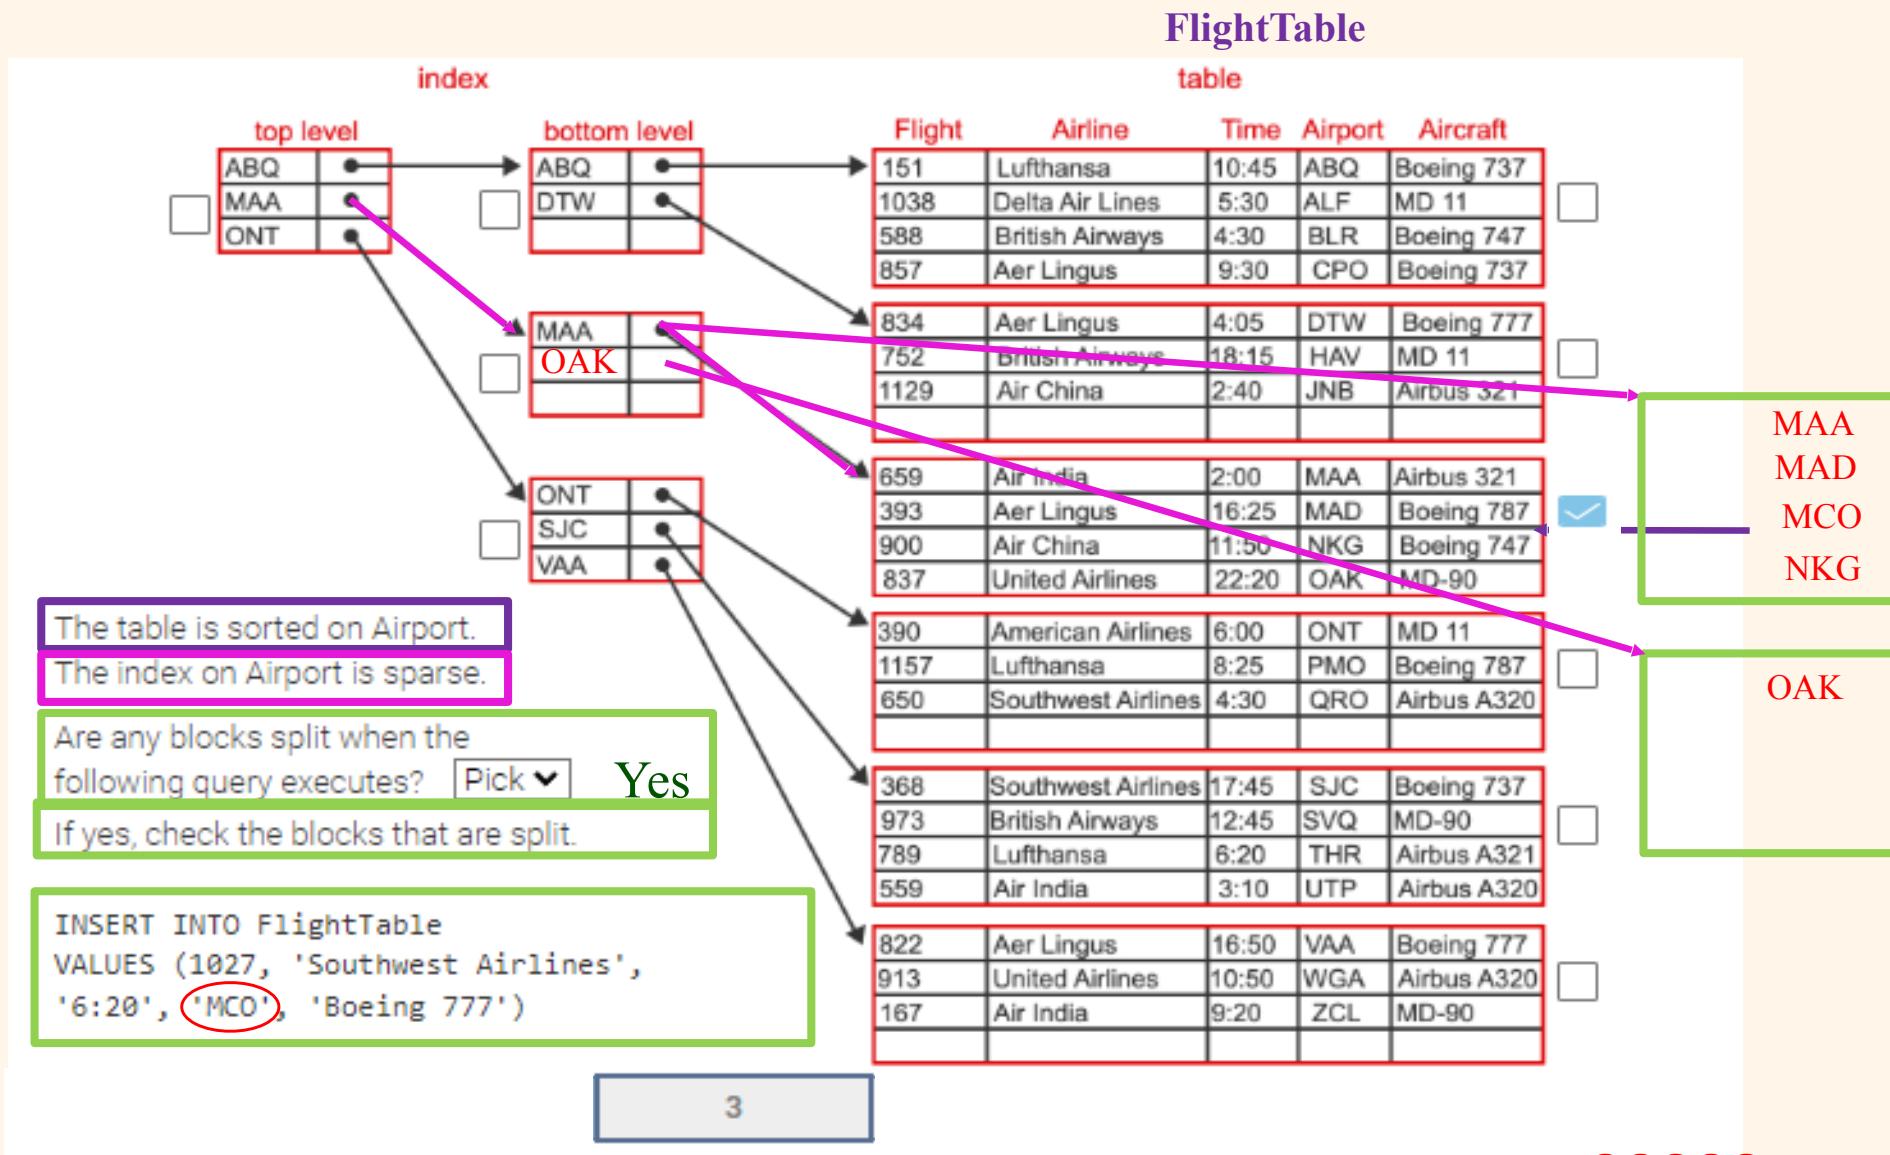


Check any blocks read by the following query:

```
SELECT Airline, Flight
FROM FlightTable
WHERE Airline > 'Lufthansa';
```

Index on Airport

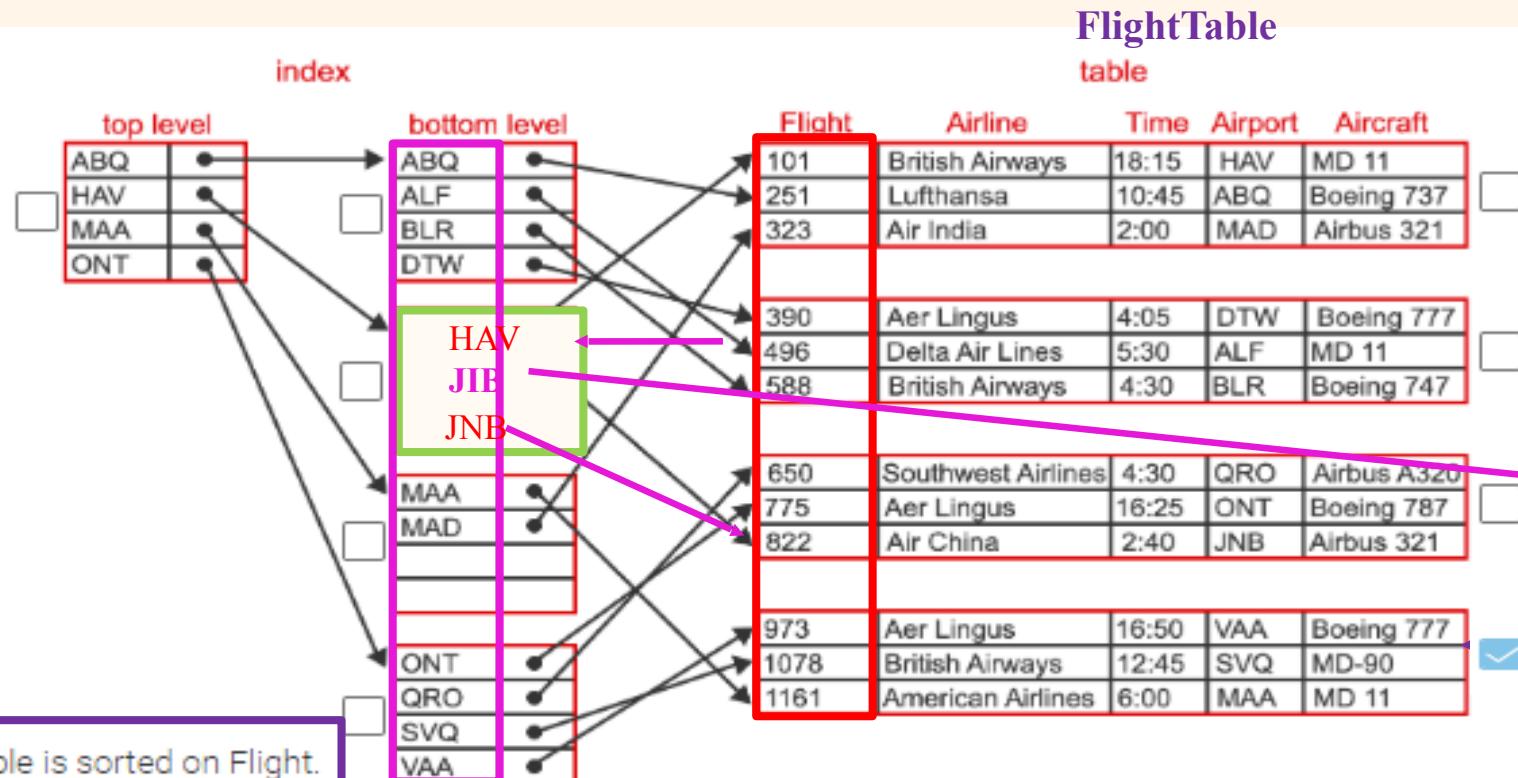
All table blocks B



CHALLENGE
ACTIVITY

15.1.1: Multi-level indexes.

995, 'American Airlines', '2:15', 'JIB', 'Boeing 747'



Check any blocks that split when the following query executes:

```
INSERT INTO FlightTable
VALUES (995, 'American Airlines', '2:15', 'JIB', 'Boeing 747')
```

4

VH

?????

55

TA time (Alvaro)
(CA 15.1.1 Step 1– Multi-level indexes)

TA time (Alvaro)
(CA 15.1.1 Step 2 – Multi-level indexes)

TA time (Alvaro)
(CA 15.1.1 Step 3– Multi-level indexes)

TA time (Alvaro)
(CA 15.1.1 Step 4— Multi-level indexes)

B+Tree Index Practice Questions



Branches that are similar in length in an index hierarchy are ____.

a. dense

b. balanced

c. sparse

d. primary

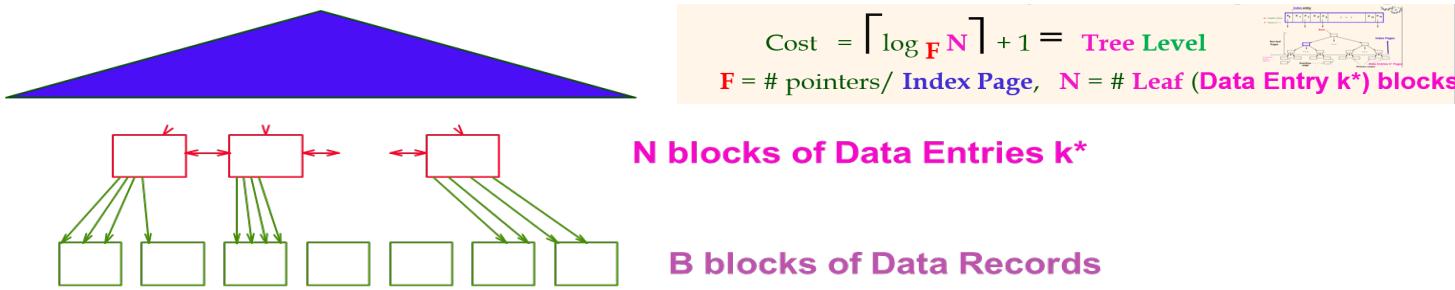
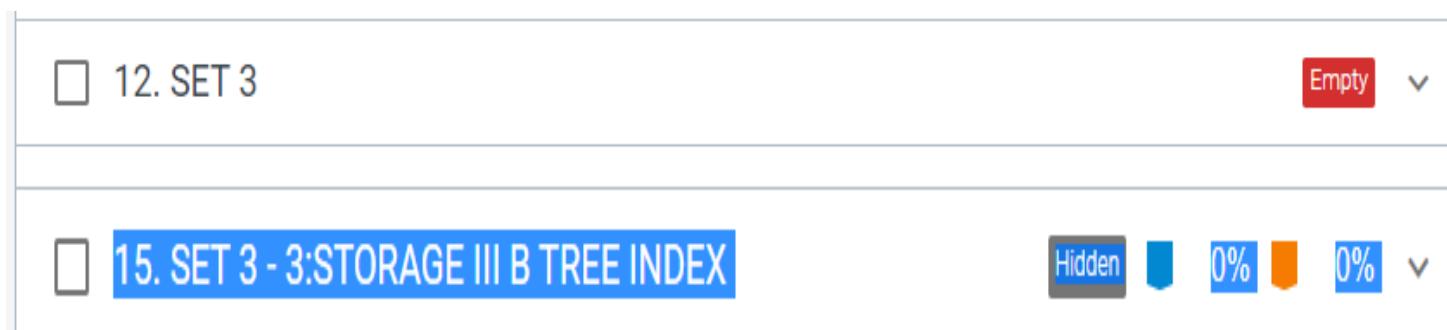
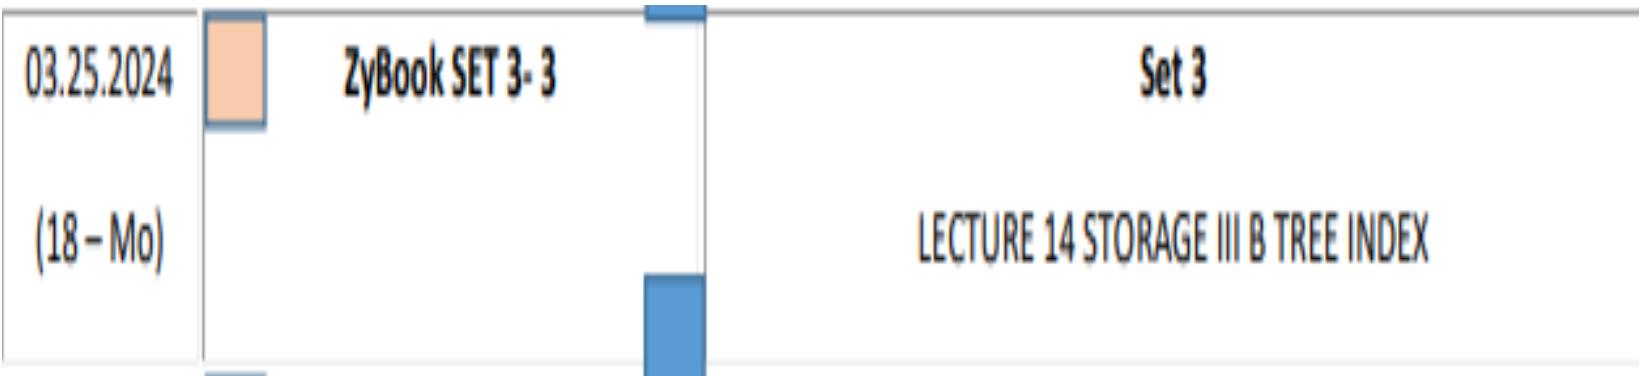
????

How are blocks read with a single-level index scan?

- a. A scan reads all index blocks to find table blocks that contain selected rows. The table blocks are then read.
- b. A scan is initiated to read all table blocks. Index blocks are then read.
- c. The search reads one index block plus selected table blocks.
- d. A scan is initiated to read all table blocks that contain index blocks with selected data.

????

At 5:00 PM .



VH work on
SET 3 – 3: Storage III B Tree Index

03.27.2024

ZyBook SET 3-4

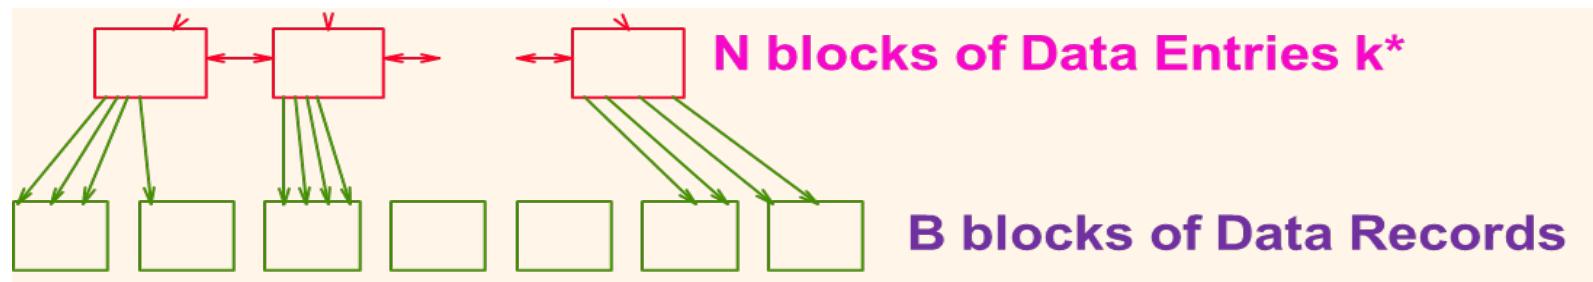
(19 - We)

Set 3

LECTURE 15 STORAGE IV HASH INDEX

HASH

$O(1)$



12. SET 3

Empty ▼

16. SET 3 - 4:STORAGE IV HASH INDEX

Hidden 0% 0% ▼

VH, unHIDE

From 5:05 to 5:15 PM – 5 minutes.

The screenshot shows a ZyBook interface. On the left, there is a date and time indicator: "03.25.2024 (18 - Mo)". Next to it is a box labeled "ZyBook SET 3-3". To the right of the box is the text "Set 3". Below these elements is the title "LECTURE 14 STORAGE III B TREE INDEX". A yellow square icon is positioned near the title.

This screenshot shows a Canvas assignment titled "CLASS PARTICIPATION" worth 20 points. It is set to be 20% of the total grade. There are three small icons at the top left of the assignment box.

B PLUS TREE

This screenshot shows a Canvas assignment titled "Class 18 END PARTICIPATION" with a due date of Mar 25 at 5:15pm and 100 points available. The assignment is marked as "VH, publish". There are three small icons at the top right of the assignment box.

This is a synchronous online class.

Attendance is required.

Recording or distribution of class materials is prohibited.

1. At the beginning of selected classes there is an assessment in the first 10 minutes. (beige BOX in the Detailed Syllabus)
2. At the end of selected classes there is an assessment in the last 10 minutes. (blue BOX in the Detailed Syllabus)
3. ZyBook sections will be downloaded and used for 30% of Total Score on the dates specified in the Detailed Syllabus.
4. EXAMS are in CANVAS. No late EXAMS.
5. I have to be present in TEAMS in order to take any graded assignment assigned during that class.

At 5:15 PM.

End Class 18

VH, unhide ZyBook Section 16.



**VH, Download Attendance Report
Rename it:
3.25.2024 Attendance Report FINAL**

VH, upload Class 18 to CANVAS.