

COSC 3380 Spring 2024

Database Systems

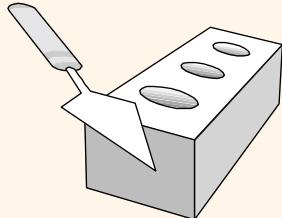
M & W 4:00 to 5:30 PM

Prof. **Victoria Hilford**

PLEASE TURN your webcam ON (must have)

NO CHATTING during LECTURE

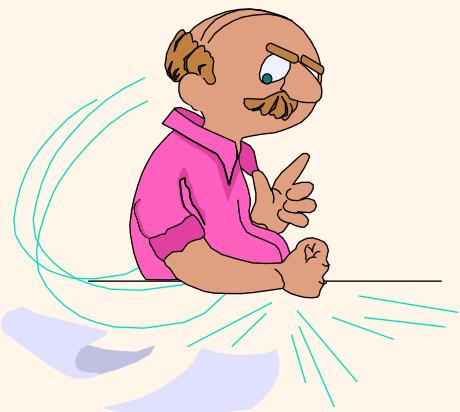
VH, unhide Section 16: SET 3 STORAGE IV – HASH INDEX



COSC 3380

4 to 5:30

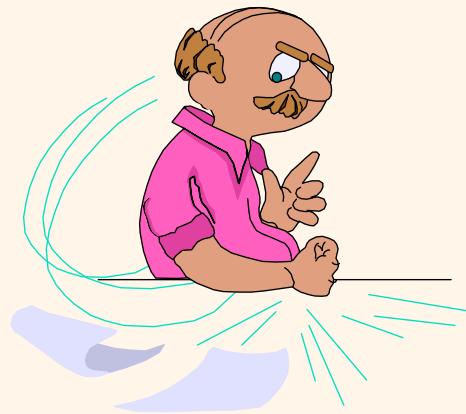
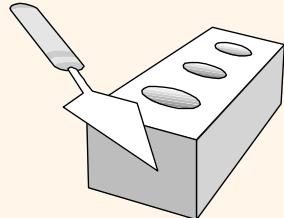
**PLEASE
LOG IN
CANVAS**



Please close all other windows.

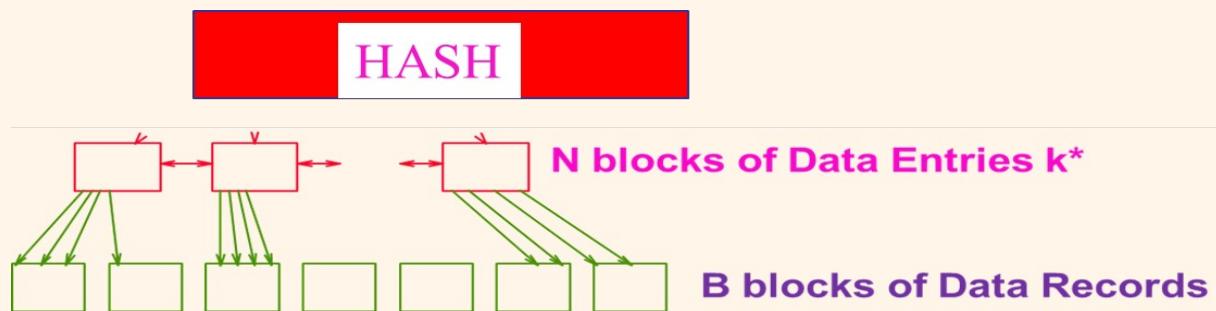
03.27.2024 (19 - We)	ZyBook SET 3- 4	Set 3 LECTURE 15 STORAGE IV HASH INDEX
04.01.2024 (20 - Mo)		EXAM 3 Practice (PART of 20 points)
04.03.2024 (21 – We)	IA Download ZyBook SET 3 Sections (4 PM) (PART of 30 points)	EXAM 3 Review (PART of 20 points)
04.08.2024 (22 - Mo)		EXAM 3 (PART of 50 points)

COSC 3380

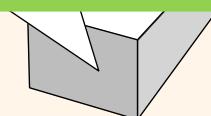


Class 19

09.27.2024 ZyBook SET 3-4 Set 3
(19 - We) LECTURE 15 STORAGE IV HASH INDEX



From 4:00 to 4:07 PM – 5 minutes.



03.27.2024

ZyBook SET 3-4

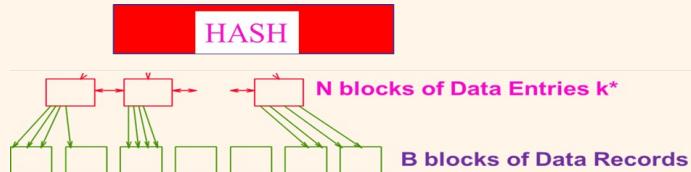
(19 - We)

Set 3

LECTURE 15 STORAGE IV HASH INDEX

CLASS PARTICIPATION 20 points

20% of Total + :



HASH INDEX



Class 19 BEGIN PARTICIPATION

Not available until Mar 27 at 4:00pm | Due Mar 27 at 4:07pm | 100 pts

VH, publish ⊖ :

This is a synchronous online class.

Attendance is required.

Recording or distribution of class materials is prohibited.

1. At the beginning of selected classes there is an assessment in the first 10 minutes. (beige BOX in the Detailed Syllabus)

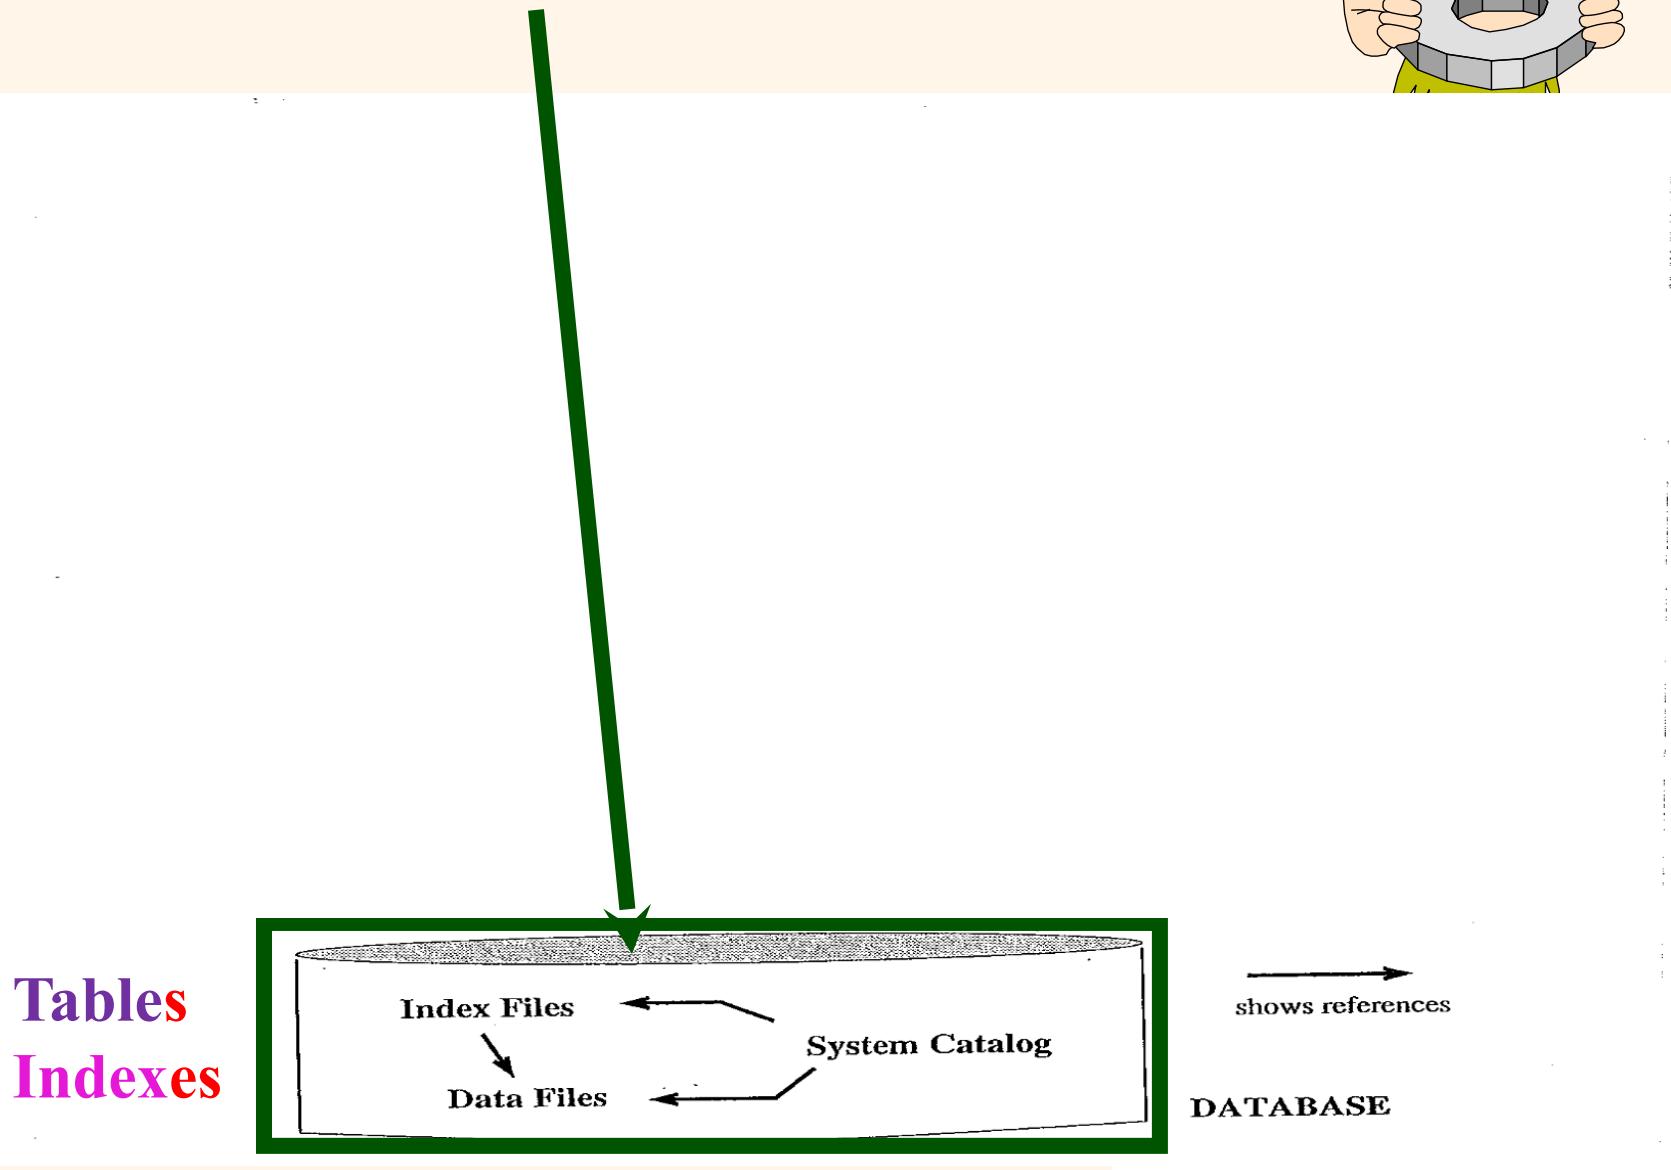
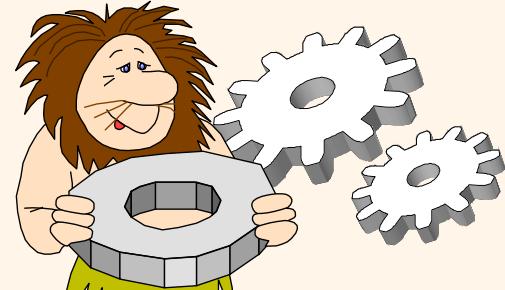
2. At the end of selected classes there is an assessment in the last 10 minutes. (blue BOX in the Detailed Syllabus)

3. ZyBook sections will be downloaded and used for 30% of Total Score on the dates specified in the Detailed Syllabus.

4. EXAMS are in CANVAS. No late EXAMS.

5. I have to be present in TEAMS in order to take any graded assignment assigned during that class.

A *DBMS*



Tables
Indexes

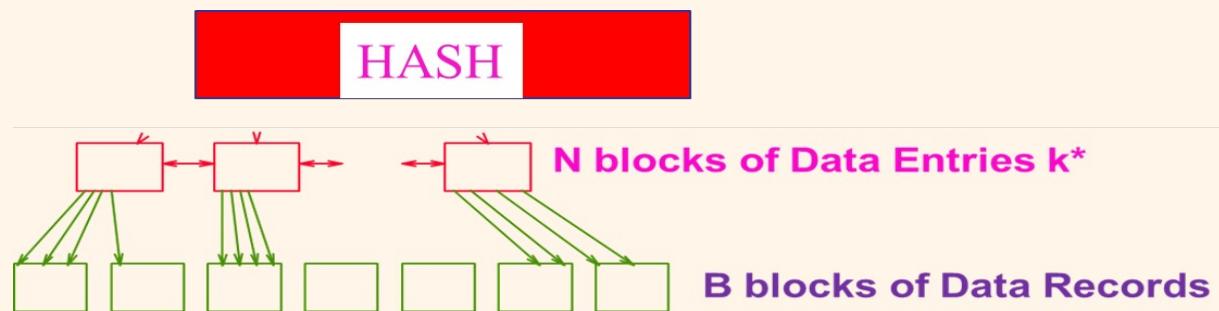
From 4:07 to 5:00 PM – 53 minutes.

COSC 3380



Lecture 15

Hash Based Indexes



$$h + 1 + 1$$

Leaf (Data Entry k^*) blocks; 1 Table Block

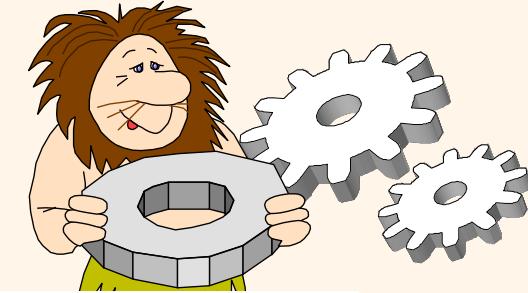
$$h + 1 + 1 + 1 + 1$$

2 (read and write back) Leaf (Data Entry k^*) blocks; 2 (read and write back) Table Blocks



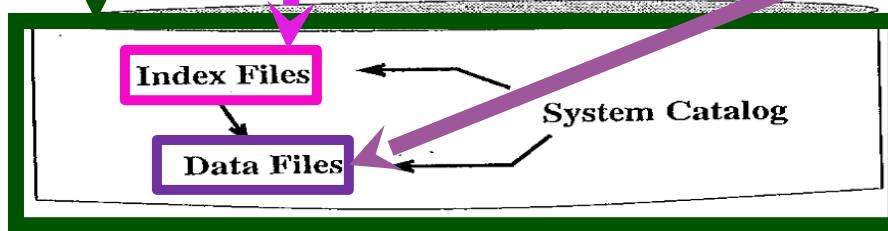
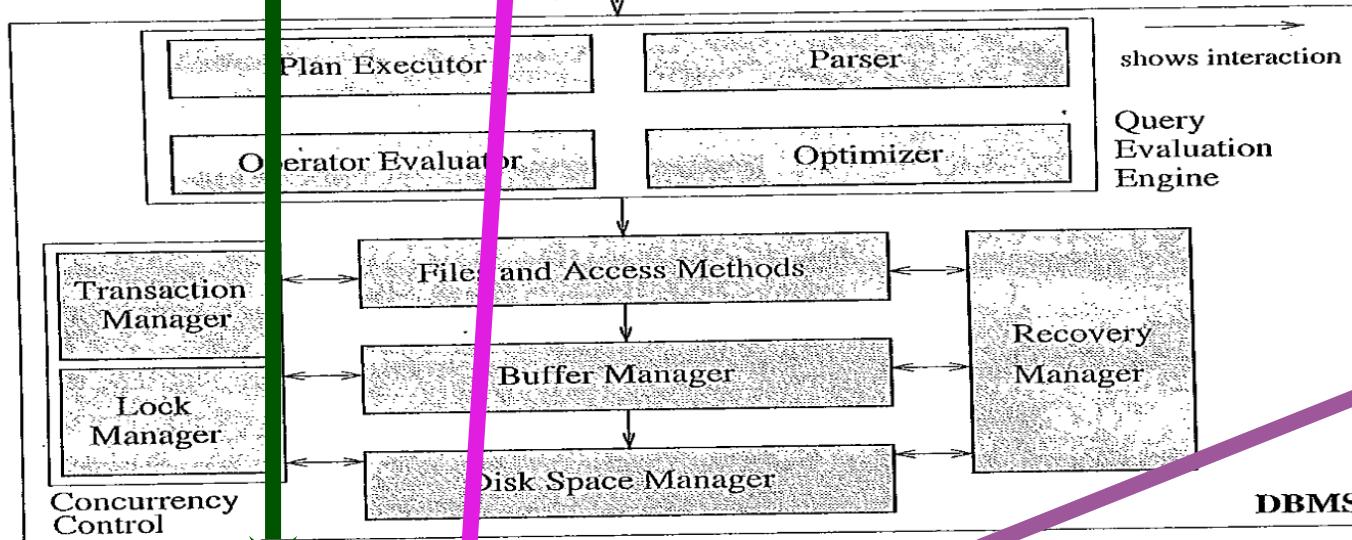
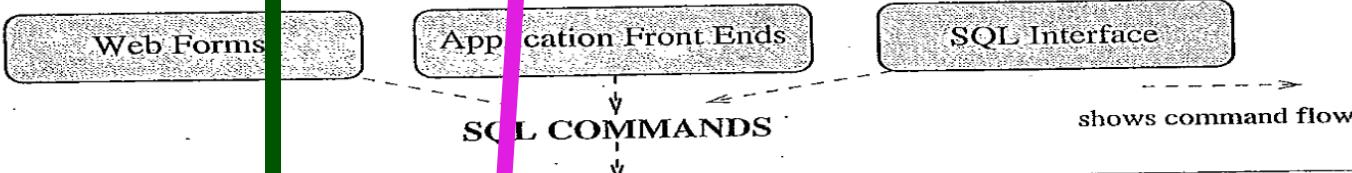
A **DBMS**

Index File Data Entries k*



Unsophisticated users (customers, travel agents, etc.)

Sophisticated users, application programmers, DB administrators

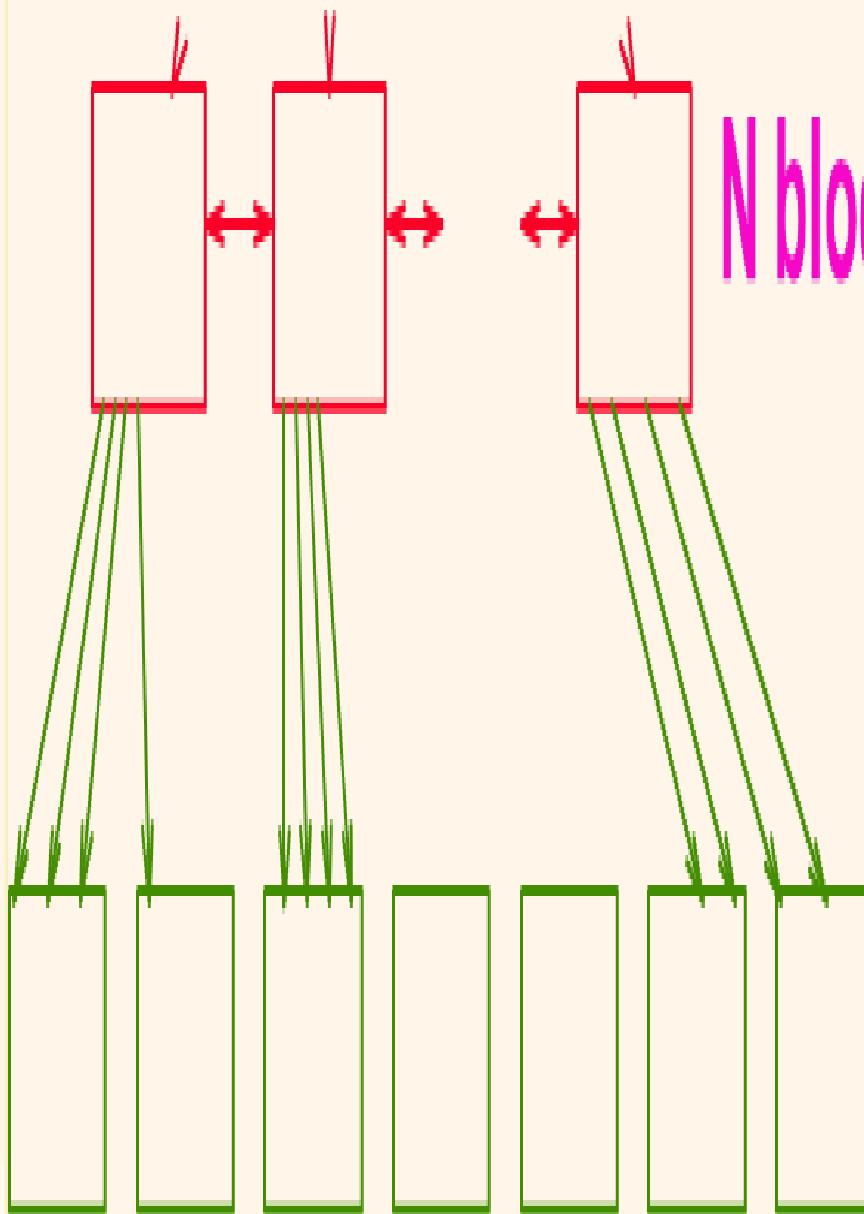
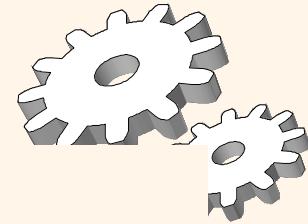


**FILE/TABLE
Data Records**

DATABASE

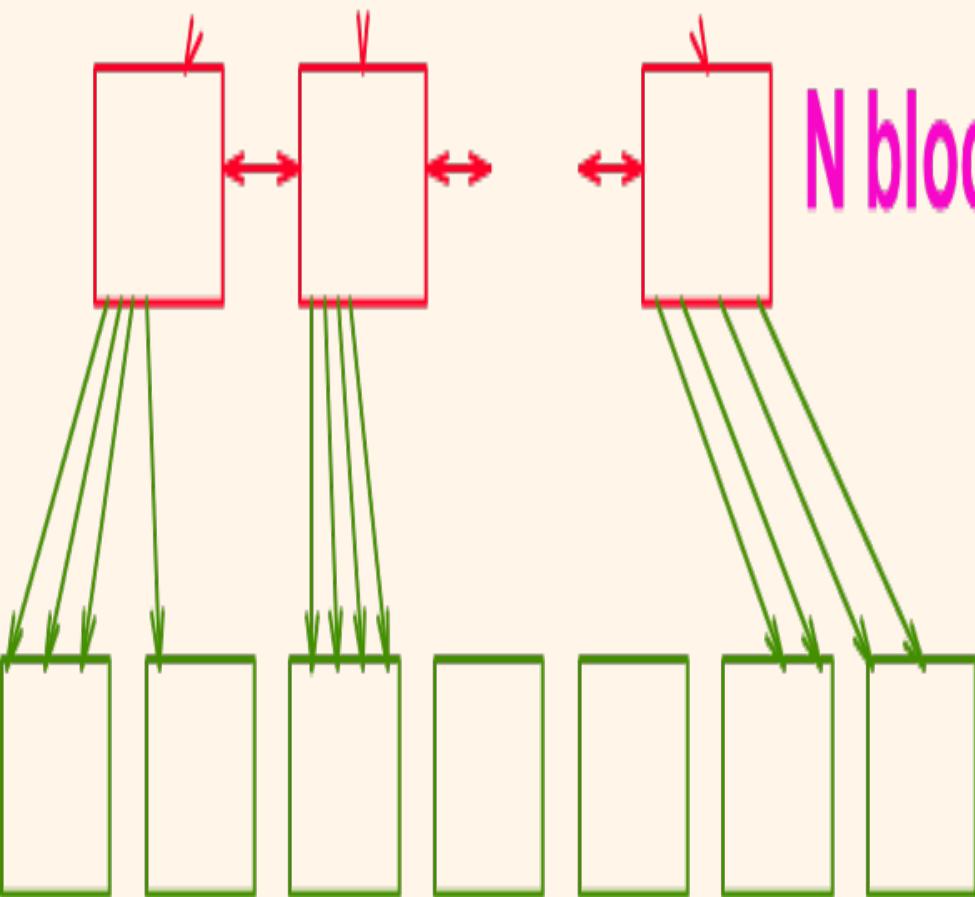
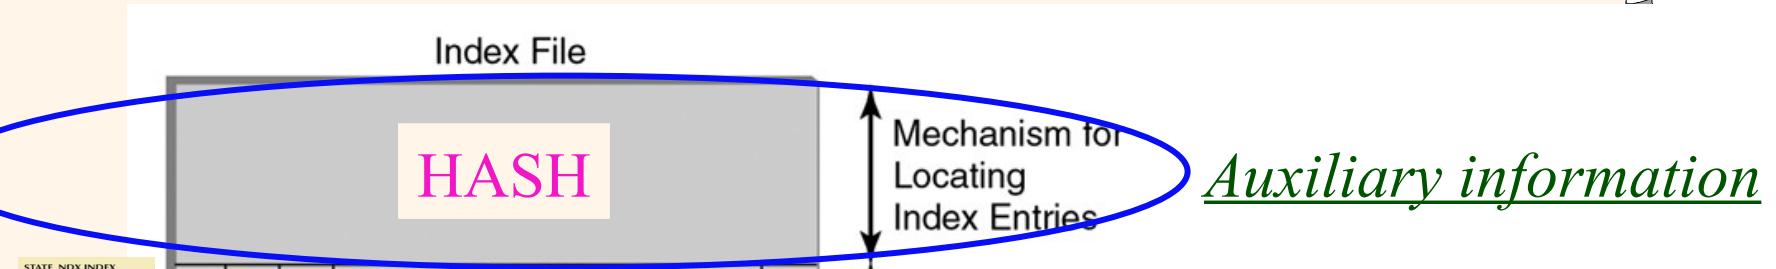
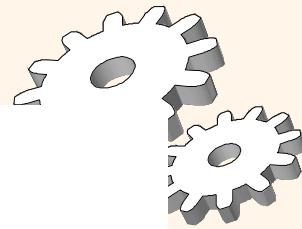
Indexes

HASH

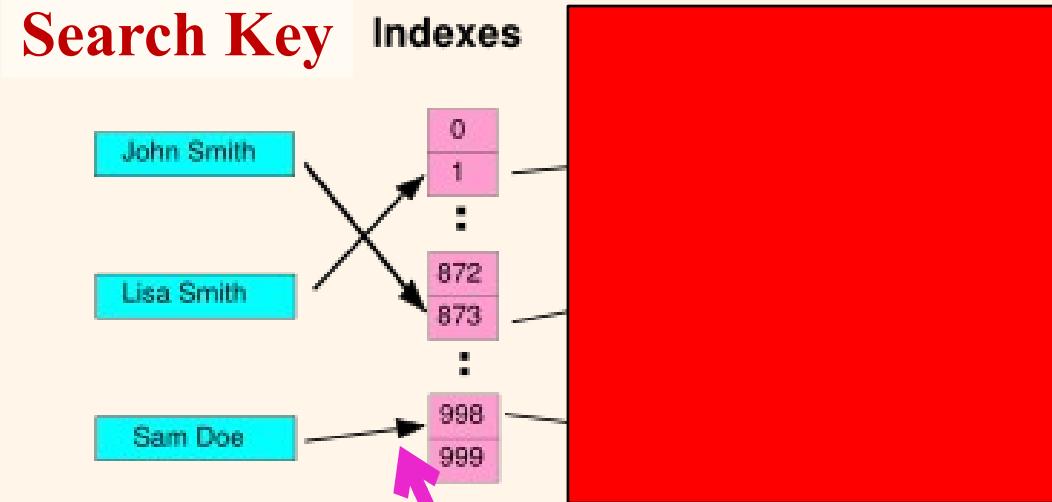


N blocks of Data Entries k^*

B blocks of Data Records



Hash Table (Hash Map)

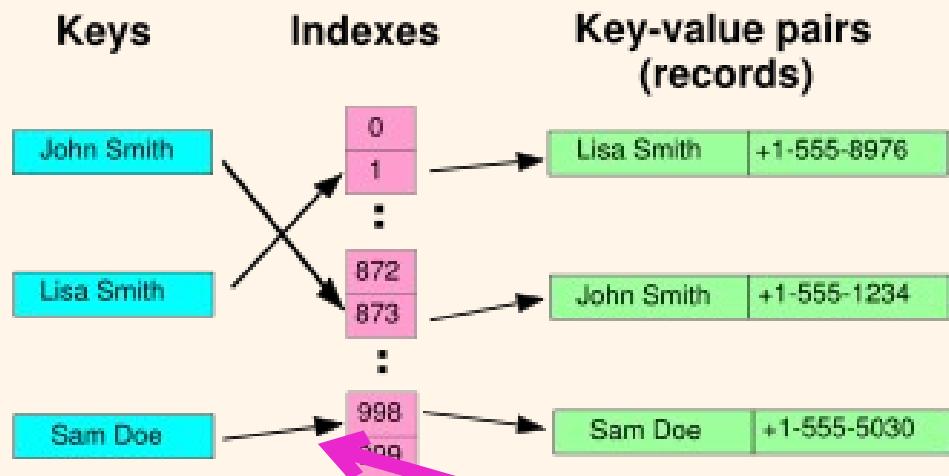
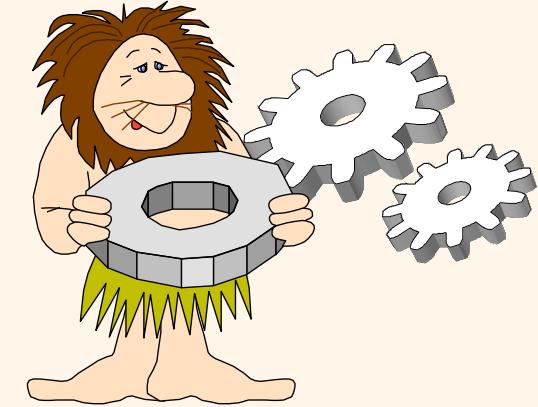


A small phone book as a **Hash Table** (Array).

A **Hash Table** associates **keys** with **values**. Given a **key** (e.g., a *person's name*), find the corresponding **value** (e.g. that *person's telephone number*).

Transform the **key** using a **hash function** into a **hash**, a number that is used as an **index** in an **Hash Table** (Array) to locate the desired location ("**bucket**") where the **value** should be.

Hash Table (Hash Map)

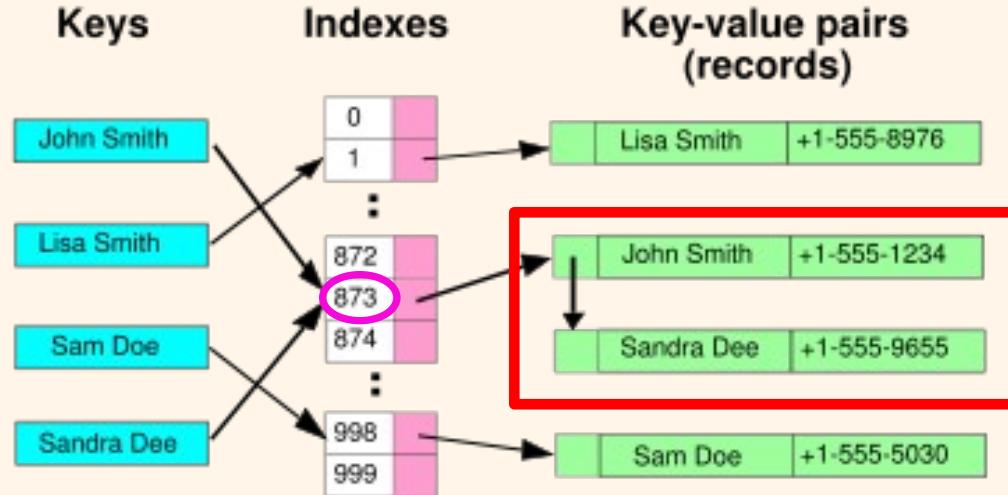


*A small phone book as a **Hash Table** (array).*

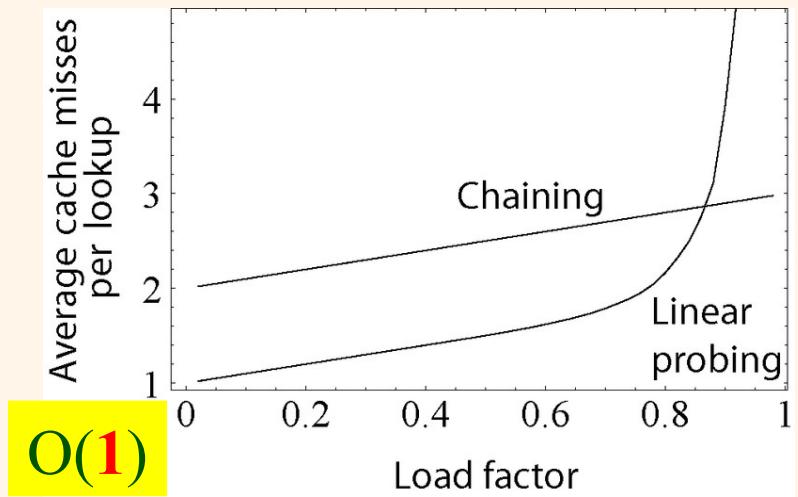
Collision: when two **keys** **hash** to the same **Hash Table** (Array) **index**, an alternate location must be determined because two **keys** cannot be stored in the same **Hash Table** (Array) location (same **index**).

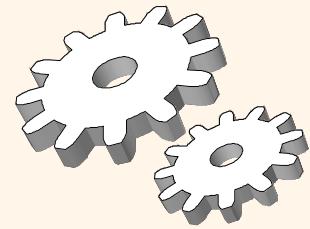
Collision resolution: **Separate Chaining & Open**

Separate Chaining



Bucket
(Linked)





Hash-Based Indexes are best for **equality** **O(1)** selections. **Cannot** support **range** searches.

Static and **Dynamic Hashing Index** techniques exist;

- trade-offs similar to **ISAM** vs. **B⁺ Trees Index**.

Hash table

In a **hash table**, rows are assigned to buckets. A **bucket** is a block or group of blocks containing rows. Initially, each bucket has one block. As a table grows, some buckets eventually fill up with rows, and the database allocates additional blocks. New blocks are linked to the initial block, and the bucket becomes a chain of linked blocks.

The bucket containing each row is determined by a hash function and a hash key. The **hash key** is a column or group of columns, usually the primary key. The **hash function** computes the bucket containing the row from the hash key.

Hash functions are designed to scramble row locations and evenly distribute rows across blocks. The **modulo function** is a simple hash function with four steps:

1. Convert the hash key by interpreting the key's bits as an integer value.
2. Divide the integer by the number of buckets.
3. Interpret the division remainder as the bucket number.
4. Convert the bucket number to the physical address of the block containing the row.

As tables grow, a fixed hash function allocates more rows to each bucket, creating deep buckets consisting of long chains of linked blocks. Deep buckets are inefficient since a query may read several blocks to access a single row. To avoid deep buckets, databases may use dynamic hash functions. A **dynamic hash function** automatically allocates more blocks to the table, creates additional buckets, and distributes rows across all buckets. With more buckets, fewer rows are assigned to each bucket and, on average, buckets contain fewer linked blocks.

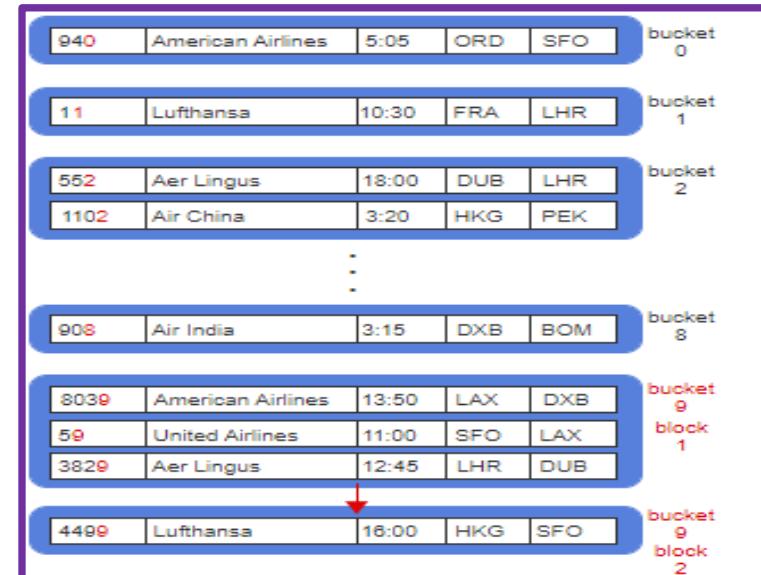
Hash tables are optimal for inserts and deletes of individual rows, since row location is quickly determined from the hash key. For the same reason, hash tables are optimal for selecting a single row when the hash key value is specified in the WHERE clause. Hash tables are slow on queries that select many rows with a range of values, since rows are randomly distributed across many blocks.

PARTICIPATION ACTIVITY

13.2.5: Assigning rows to buckets with a modulo function.

Flight				
Flight Number	AirlineName	Depart Time	Depart Airport	Arrive Airport
11	Lufthansa	10:30	FRA	LHR
552	Aer Lingus	18:00	DUB	LHR
908	Air India	3:15	DXB	BOM
940	American Airlines	5:05	ORD	SFO
1102	Air China	3:20	HKG	PEK
8039	American Airlines	13:50	LAX	DXB
59	United Airlines	11:00	SFO	LAX
3829	Aer Lingus	12:45	LHR	DUB
4499	Lufthansa	16:00	HKG	SFO

hash key



Hash indexes

The multi-level index is the most commonly used index type. Several additional index types are used less often but supported by many databases:

- Hash index
- Bitmap index
- Logical index
- Function index

Hash indexes

The multi-level index is the most commonly used index type. Several additional index types are used less often but supported by many databases:

- Hash index
- Bitmap index
- Logical index
- Function index

In a **hash index**, index entries are assigned to buckets. A **bucket** is a block or group of blocks containing index entries. Initially, each bucket has one block. As an index grows, some buckets eventually fill up, and additional blocks are allocated and linked to the initial block.

The bucket containing each index entry is determined by a **hash function**, which computes a bucket number from the value of the indexed column. To locate a row containing a column value, the database:

1. Applies the hash function to the column value to compute a bucket number.
2. Reads the index blocks for the bucket number.
3. Finds the index entry for the column value and reads the table block pointer.
4. Reads the table block containing the row.

A hash index is similar to a hash table, described in another section. However, a hash index stores **index entries** in each bucket, while a hash table stores **table rows** in each bucket.

PARTICIPA
ACTIVITY

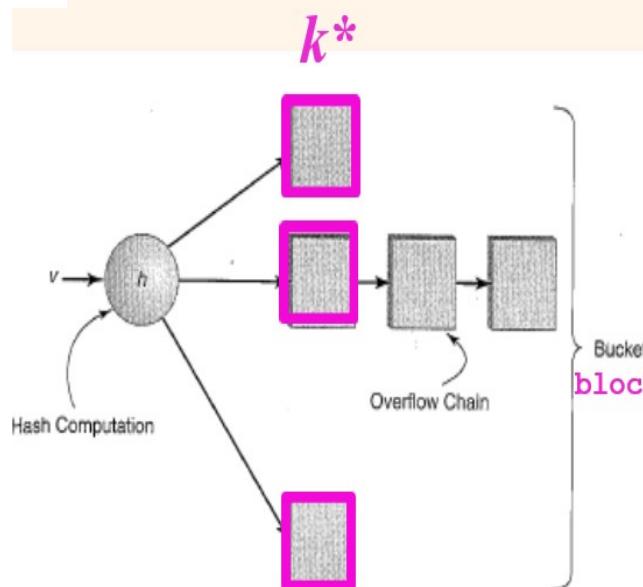
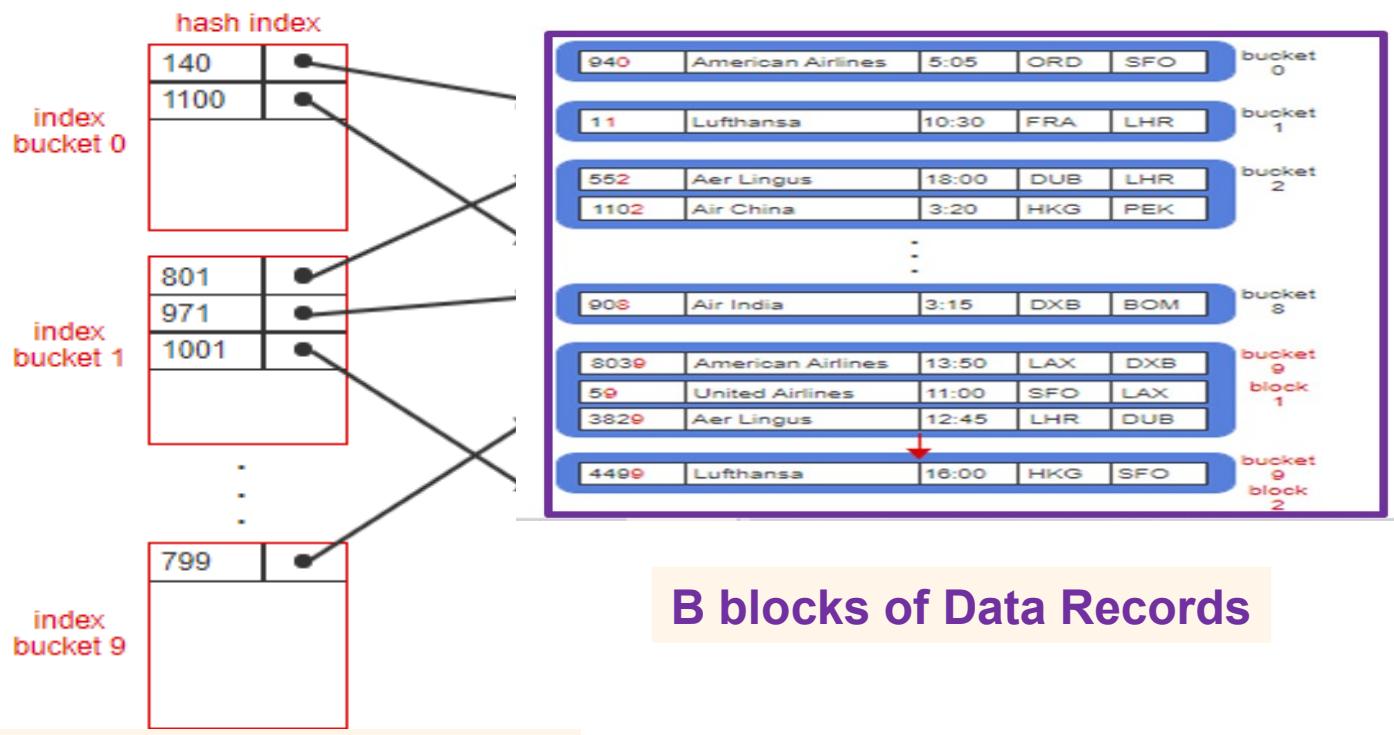


table				
666	United Airlines	6:00	ATL	Airbus 321
140	Aer Lingus	10:30	DUB	Boeing 737
803	Air China	3:20	DME	Boeing 777
801	Air India	3:15	DXB	Boeing 747
552	Lufthansa	18:00	FRA	MD 111
1100	Air China	3:20	HKG	Boeing 787
971	Air India	19:35	LAX	Boeing 747
702	American Airlines	5:05	ORD	Airbus 323
1154	American Airlines	13:50	ORD	MD 111
799	British Airways	13:50	PHL	Airbus 323
1222	Delta Airlines	8:25	SEA	Boeing 747
1001	Southwest Airlines	5:05	SJC	Airbus 323

blocks

N blocks of Data Entries k^*

B blocks of Data Records



N blocks of Data Entries k*

Terminology

Sometimes the term *hash index* is used to mean a *hash key*, but the two terms are different. A *hash index* is an index that is structured using a hash function. A *hash key* is a column that determines the physical location of rows in a hash table.

PARTICIPATION ACTIVITY

16.1.2: Hash indexes.

- 1) A hash table can have a hash index.

- True
 False

Correct

Hash tables cannot have a primary or clustering index because rows of a hash table are not stored in sort order. However, hash tables can have secondary indexes. The secondary indexes can be any structure, including hash.

?????

TA time (Jordan)

(CA 16.1.1 Step 1 – Hash Index)

CHALLENGE
ACTIVITY

16.1.1: Other indexes.

Hash Index

Given that the hash index is on FlightNumber, and the hash function is modulo 10, select the row(s) with a hash index in bucket 3.

Ending in 3

table

<input type="checkbox"/>	581	United Airlines	6:00	ATL	Airbus 321
<input type="checkbox"/>	696	Aer Lingus	10:30	DUB	Boeing 737
<input type="checkbox"/>	126	Air China	3:20	DME	Boeing 777
<input checked="" type="checkbox"/>	183	Air India	3:15	DXB	Boeing 747
<input type="checkbox"/>	165	Lufthansa	18:00	FRA	MD 111
<input type="checkbox"/>	783	Air China	3:20	HKG	Boeing 787
<input type="checkbox"/>	721	Air India	19:35	LAX	Boeing 747
<input type="checkbox"/>	924	American Airlines	5:05	ORD	Airbus 323
<input type="checkbox"/>	936	American Airlines	13:50	ORD	MD 111
<input type="checkbox"/>	686	British Airways	13:50	PHL	Airbus 323
<input type="checkbox"/>	714	Delta Airlines	8:25	SEA	Boeing 747
<input type="checkbox"/>	791	Southwest Airlines	5:05	SJC	Airbus 323

183

783



35:42

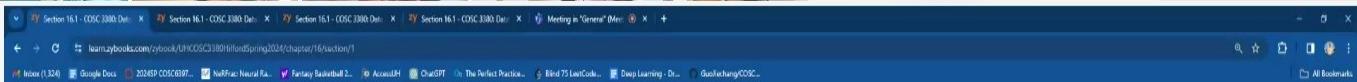
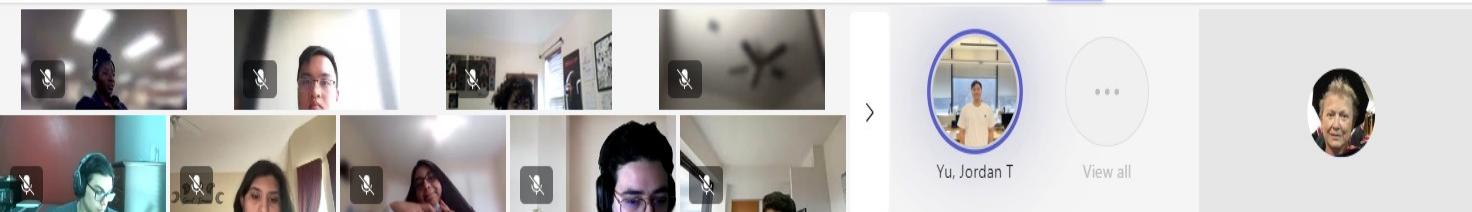


table				
<input checked="" type="checkbox"/>	733	United Airlines	6:00	ATL
<input type="checkbox"/>	707	Aer Lingus	10:30	DUB
<input checked="" type="checkbox"/>	863	Air China	3:20	DME
<input type="checkbox"/>	126	Air India	3:15	DXB
<input checked="" type="checkbox"/>	983	Lufthansa	18:00	FRA
<input checked="" type="checkbox"/>	753	Air China	3:20	HKG
<input checked="" type="checkbox"/>	313	Air India	19:35	LAX
<input type="checkbox"/>	226	American Airlines	5:05	ORD
<input type="checkbox"/>	828	American Airlines	13:50	MD 111
<input type="checkbox"/>	888	British Airways	13:50	PHL
<input type="checkbox"/>	419	Delta Airlines	8:25	SEA
<input type="checkbox"/>	478	Southwest Airlines	5:05	SJC
				Airbus 323

1 2 3 4

✓ Expected: Flight numbers 733, 863, 983, 753, and 313

The hash function is modulo 10, so index entries for flight numbers ending in 3 go in bucket 3.

733, 863, 983, 753, and 313 are the flight numbers that end in 3.

[View solution](#) • [\(instructors only\)](#)

teams.microsoft.com is sharing your screen

64°F Cloudy 4:22 PM 3/27/2024

Participants

Invite someone or dial a number

Share invite

Waiting in lobby (1)

Martinez, Valerio



In this meeting (111)

Hilford, Victoria
Organizer



Adhikari, Rohit



Ahmed, Mohamed A



Akram, Ali



Akukwe, Benetta O



Alsayed, Sami H



Altaf, Sameer



Alvarez, Stephanie



Anayor-Achu, Ogochukwu E



Avci, Hatice Kubra



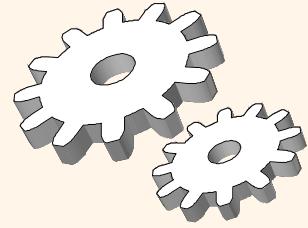
Yu, Jordan T



TA, Jordan (A – L).

TA, Alvaro (M – Z).

**Please compare CANVAS vs. TEAMS Attendance.
Print screens of students in CANVAS but not in the TEAMS meeting.
(3.27.2024 Attendance X missing LastName.docx)**



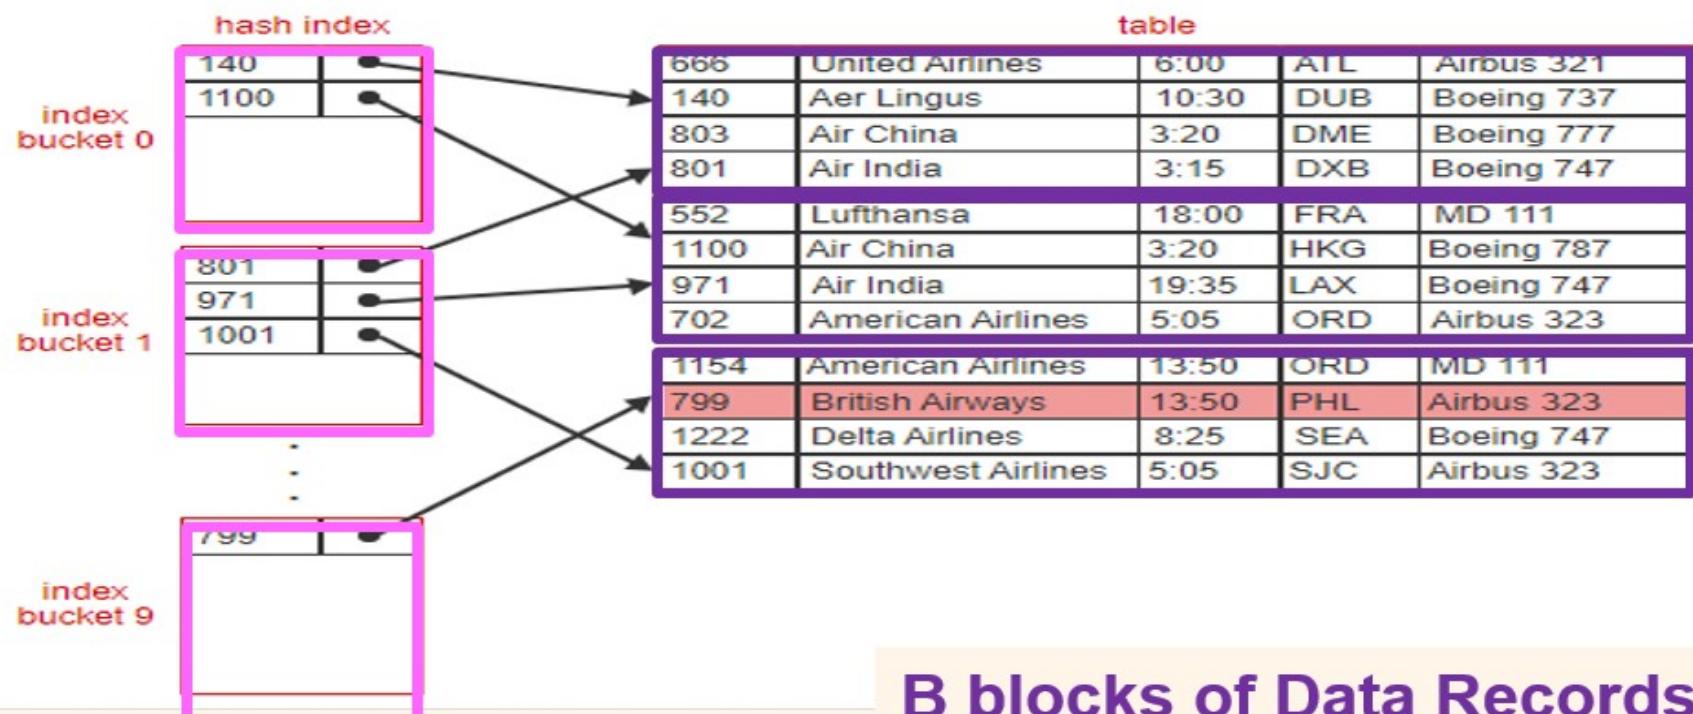
Static Hashing

Indexes

Static Hashing Indexes



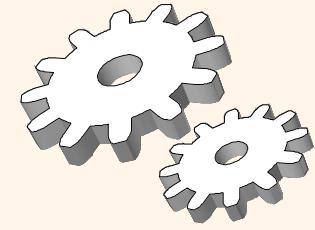
- # primary **blocks fixed**, allocated sequentially, never de-allocated; **overflow blocks if needed**.
- $h(\text{key}) \bmod N = \text{Bucket/Page}$ to which data entry k^* with key key belongs. ($N = \# \text{ of blocks}$)



B blocks of Data Records

N blocks of Data Entries k*

Static Hashing Indexes



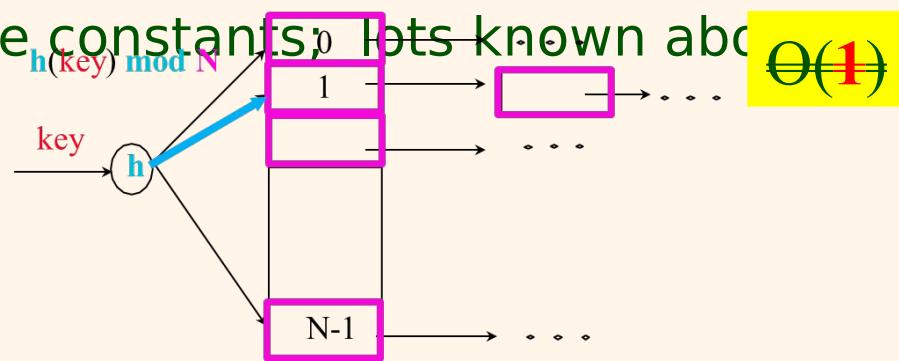
Buckets/blocks contain *Data Entries*

Data
Entry k*
k*blocks

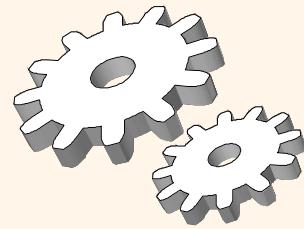
Hash function h works on search **key** field of **Data Record r**.

Must distribute values over range 0 ... N-1.

- $h(\text{key}) = (a * \text{key} + b)$ usually works well.
- a and b are constants; lots known about how to tune h .



Long overflow **Chains** can develop and degrade performance



Dynamic Extendable Hashing

Indexes

Dynamic Extendable Hashing Indexes

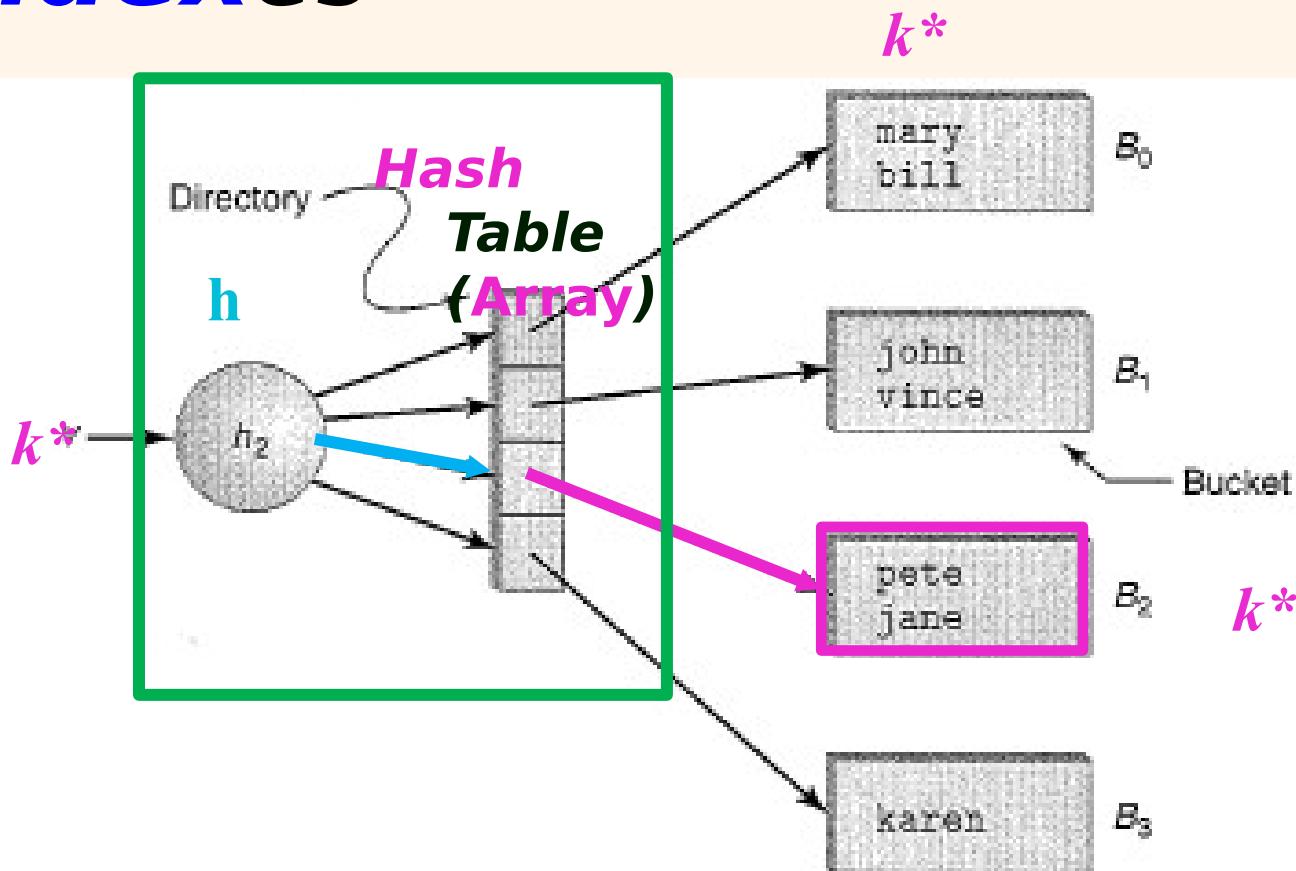
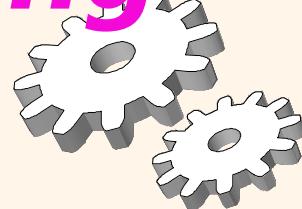


FIGURE 9.25 With extendable hashing, the data entry k^* is mapped to a bucket through a directory.

Data
Entry k^*
blocks

Example

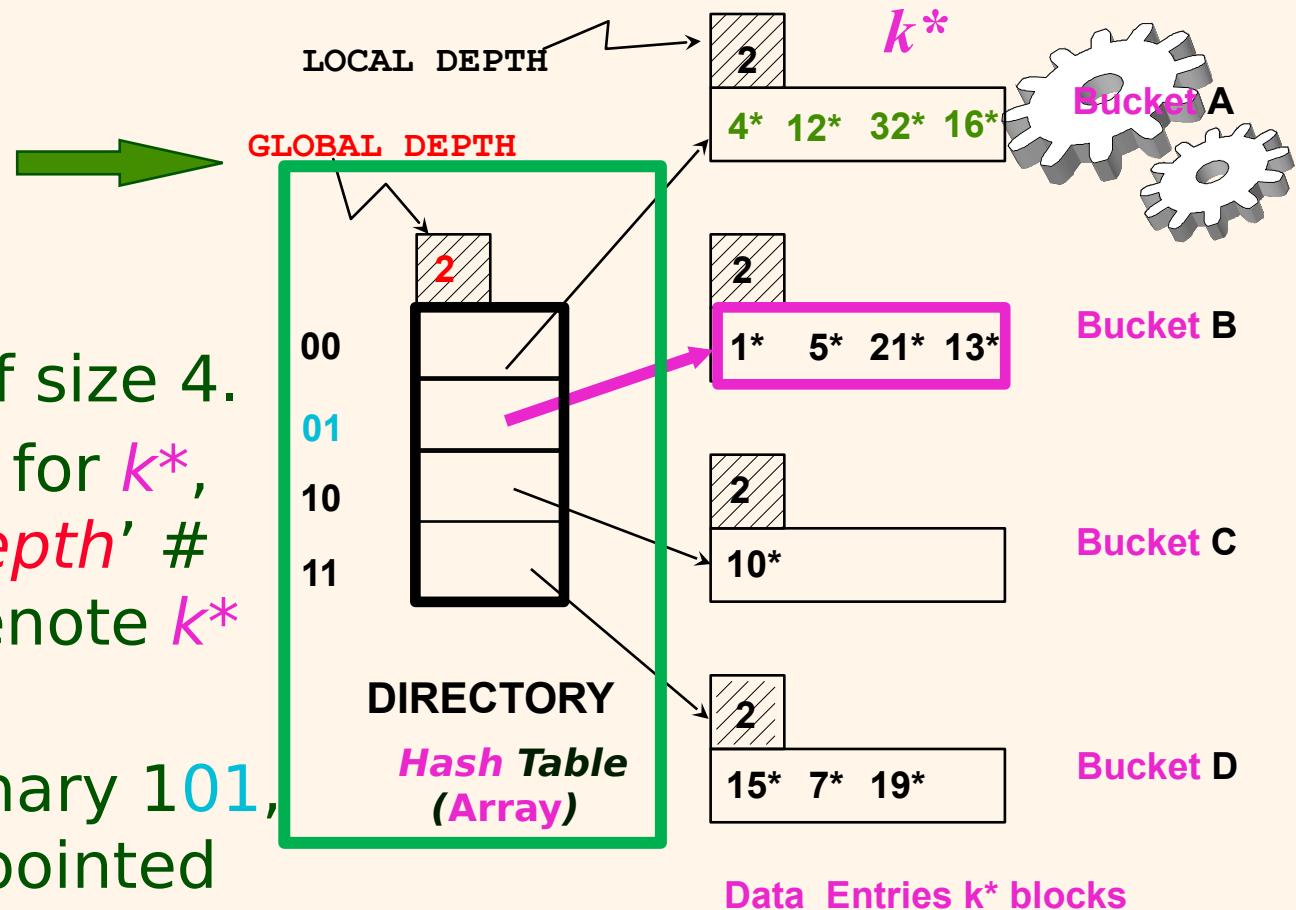
DIRECTORY is Array of size 4.

To Search Bucket for k^* , take last '*global depth*' # bits of $h(k^*)$; we denote k^* by $h(k^*)$.

- If $h(k^*) = 5 = \text{binary } 101$, it is in Bucket pointed to by 01.

insert k^* : If Bucket is full, split it (allocate **new Page**, re-distr)

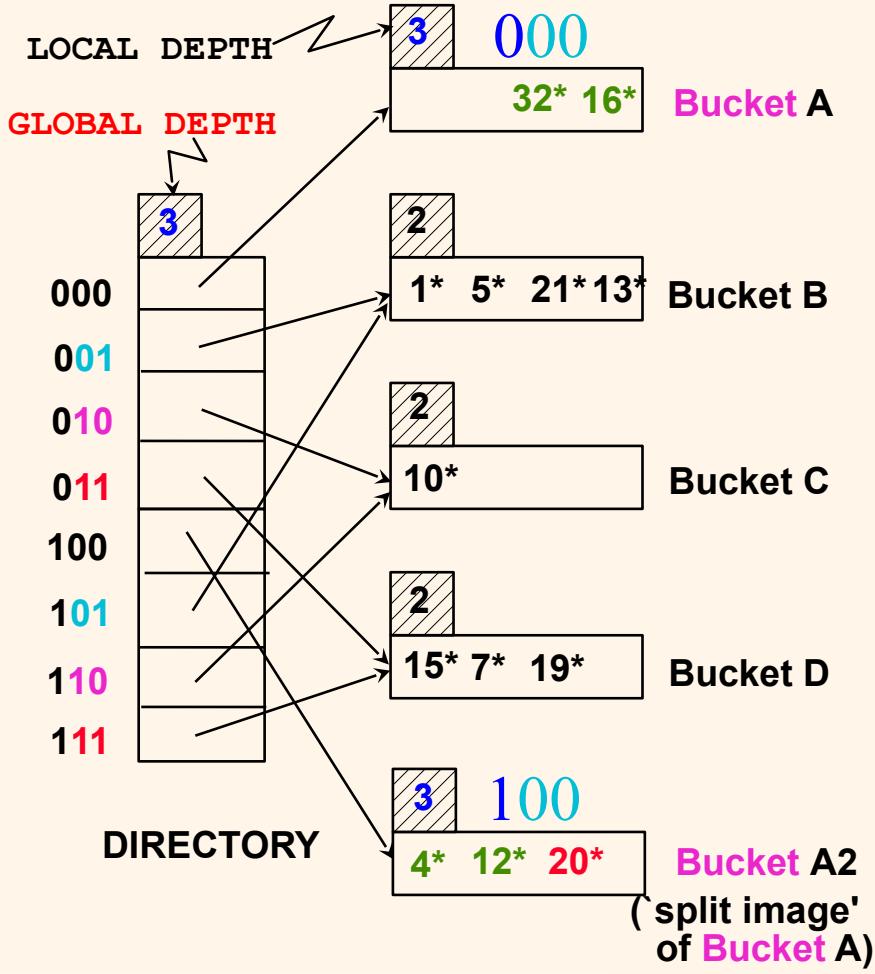
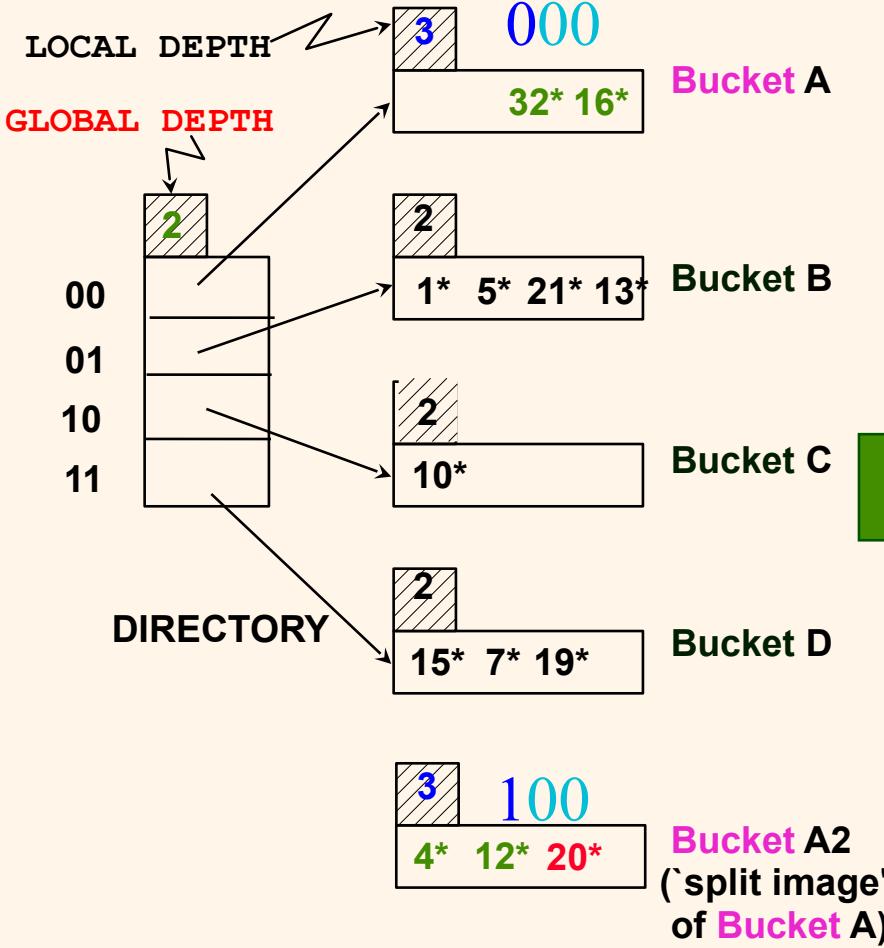
- ❖ If necessary, double the DIRECTORY. (As we will see, splitting a Bucket does not always require doubling; we can tell by comparing GLOBAL DEPTH with LOCAL DEPTH for the



Insert $n(K^*) = 20$ (Causes Doubling)

(10100)

Insert: If Bucket is full, split it (allocate new Page, re-distribute).



Dynamic Extendable Hashing Indexes



Situation: **Bucket** (primary **Page**) becomes full

Why not re-organize **Index** by *doubling* # of **Buckets**?

D per Page

- Reading and writing all **blocks** is expensive!
- Idea: Use **Directory Hash Table** of pointers (pid) to **Buckets/blocks**, *doubling* the **Directory**, **splitting just the Bucket that overflowed!**
- **Directory Hash Table** much smaller than **File**. so doubling it is muc Entries k* is split.

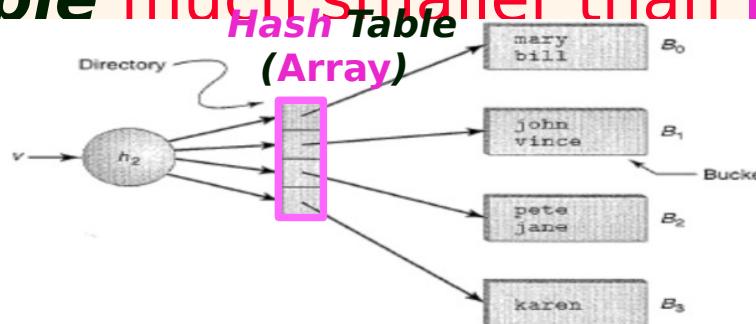
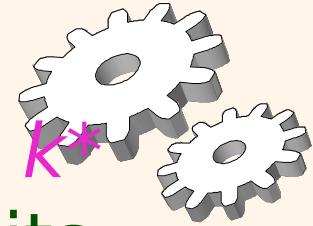


FIGURE 9.25 With extendable hashing, the hash result is mapped to a bucket through a directory.

Points to Note

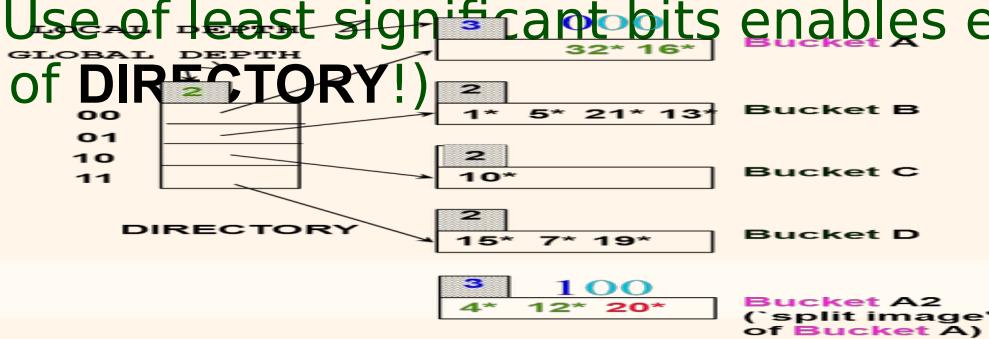


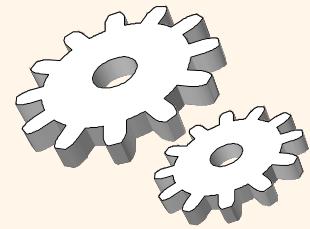
20 = binary 10100. Last **2** bits (00) tell us k^* belongs in **Bucket A** or **Bucket A2**. Last **3** bits needed to tell which.

- **GLOBAL DEPTH of DIRECTORY:** Max # of bits needed to tell which **Bucket** an Data Entry k^* belongs to.
- **LOCAL DEPTH of a Bucket:** # of bits used to determine if an Data Entry k^* belongs to this **Bucket**.

When does **Bucket** split cause **DIRECTORY doubling**

- Before **Insert**, **LOCAL DEPTH** of **Bucket** = **GLOBAL DEPTH**. **Insert** causes **LOCAL DEPTH** to become $>$ **GLOBAL DEPTH**; **DIRECTORY** is doubled by *copying it over* and `fixing' pointer to split image **Page**. (Use of least significant bits enables efficient doubling via copying of **DIRECTORY**!)





Dynamic Linear Hashing

Indexes

Dynamic Linear Hashing Indexes

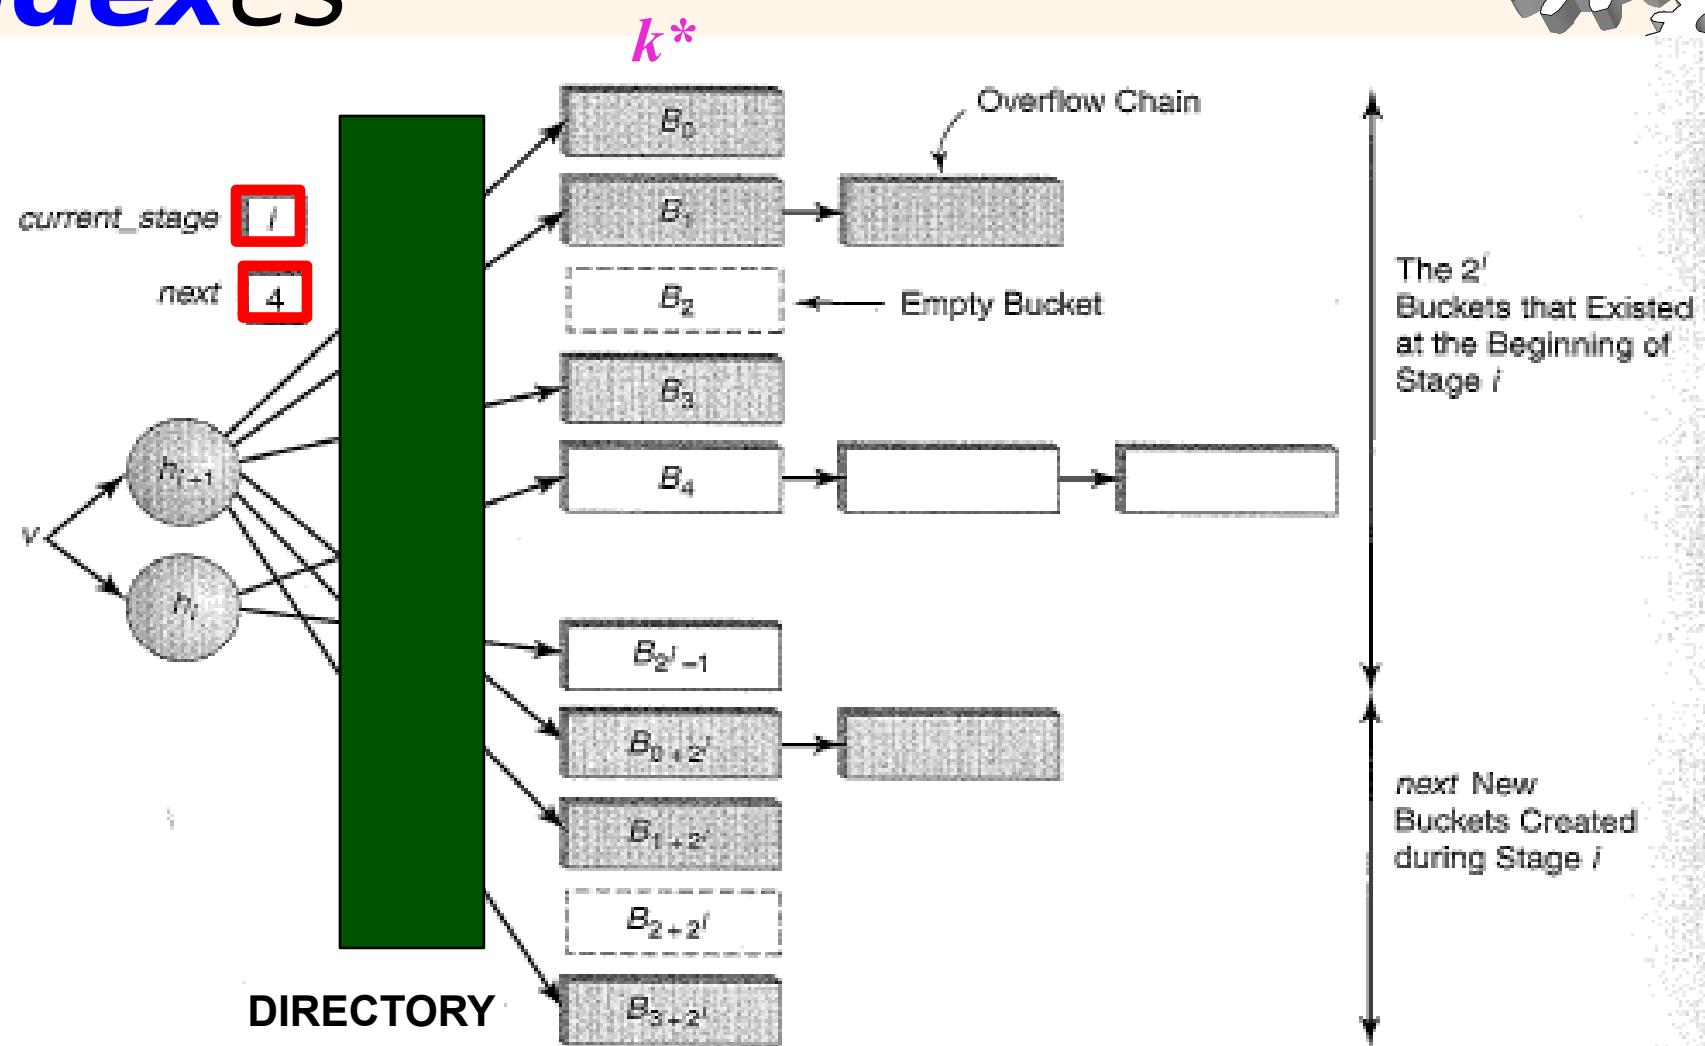
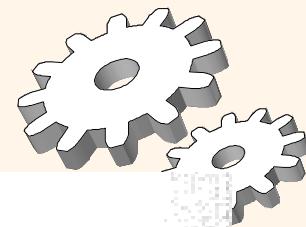


FIGURE 9.28: Linear hashing spans k^* blocks sequentially. The shaded buckets are accessed through h_{i+1} .

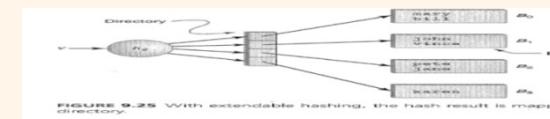
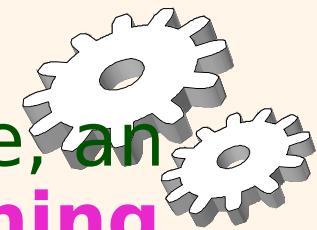
What is missing???

What is place of DIRECTORY???

Dynamic Linear Hashing

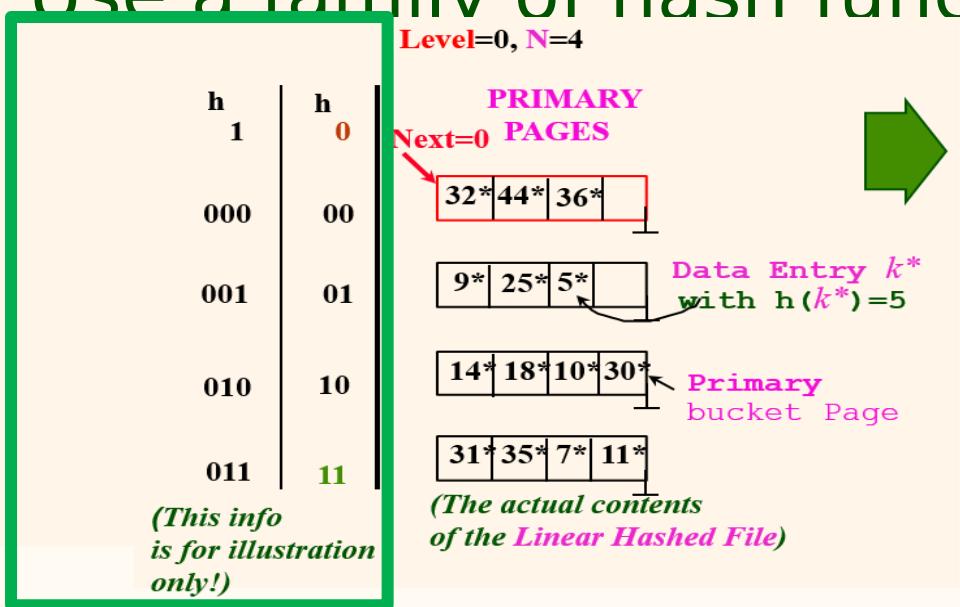
Indexes

This is another Dynamic Hashing scheme, an alternative to Dynamic Extendable Hashing Indexes.



LH handles the problem of long overflow chains without using a DIRECTORY, and handles duplicates.

Idea: Use a family of hash functions h_0, h_1, h_2, \dots



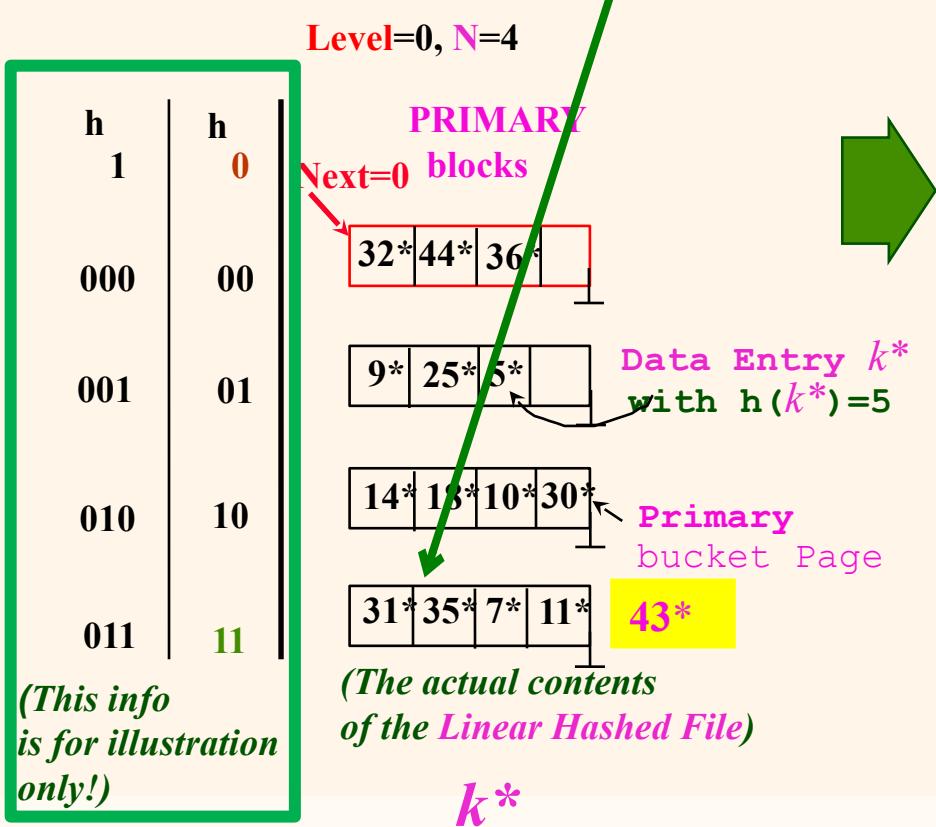
Example of Dynamic Linear Hashing Indexes



(Insert 43*) (101011)

- ❖ Insert: Find Bucket by applying $h_{Level} / h_{Level+1}$:
 - If Bucket to insert into is full:
 - Add overflow Page and insert data entry k^* .
 - Split Next Bucket and increment Next.

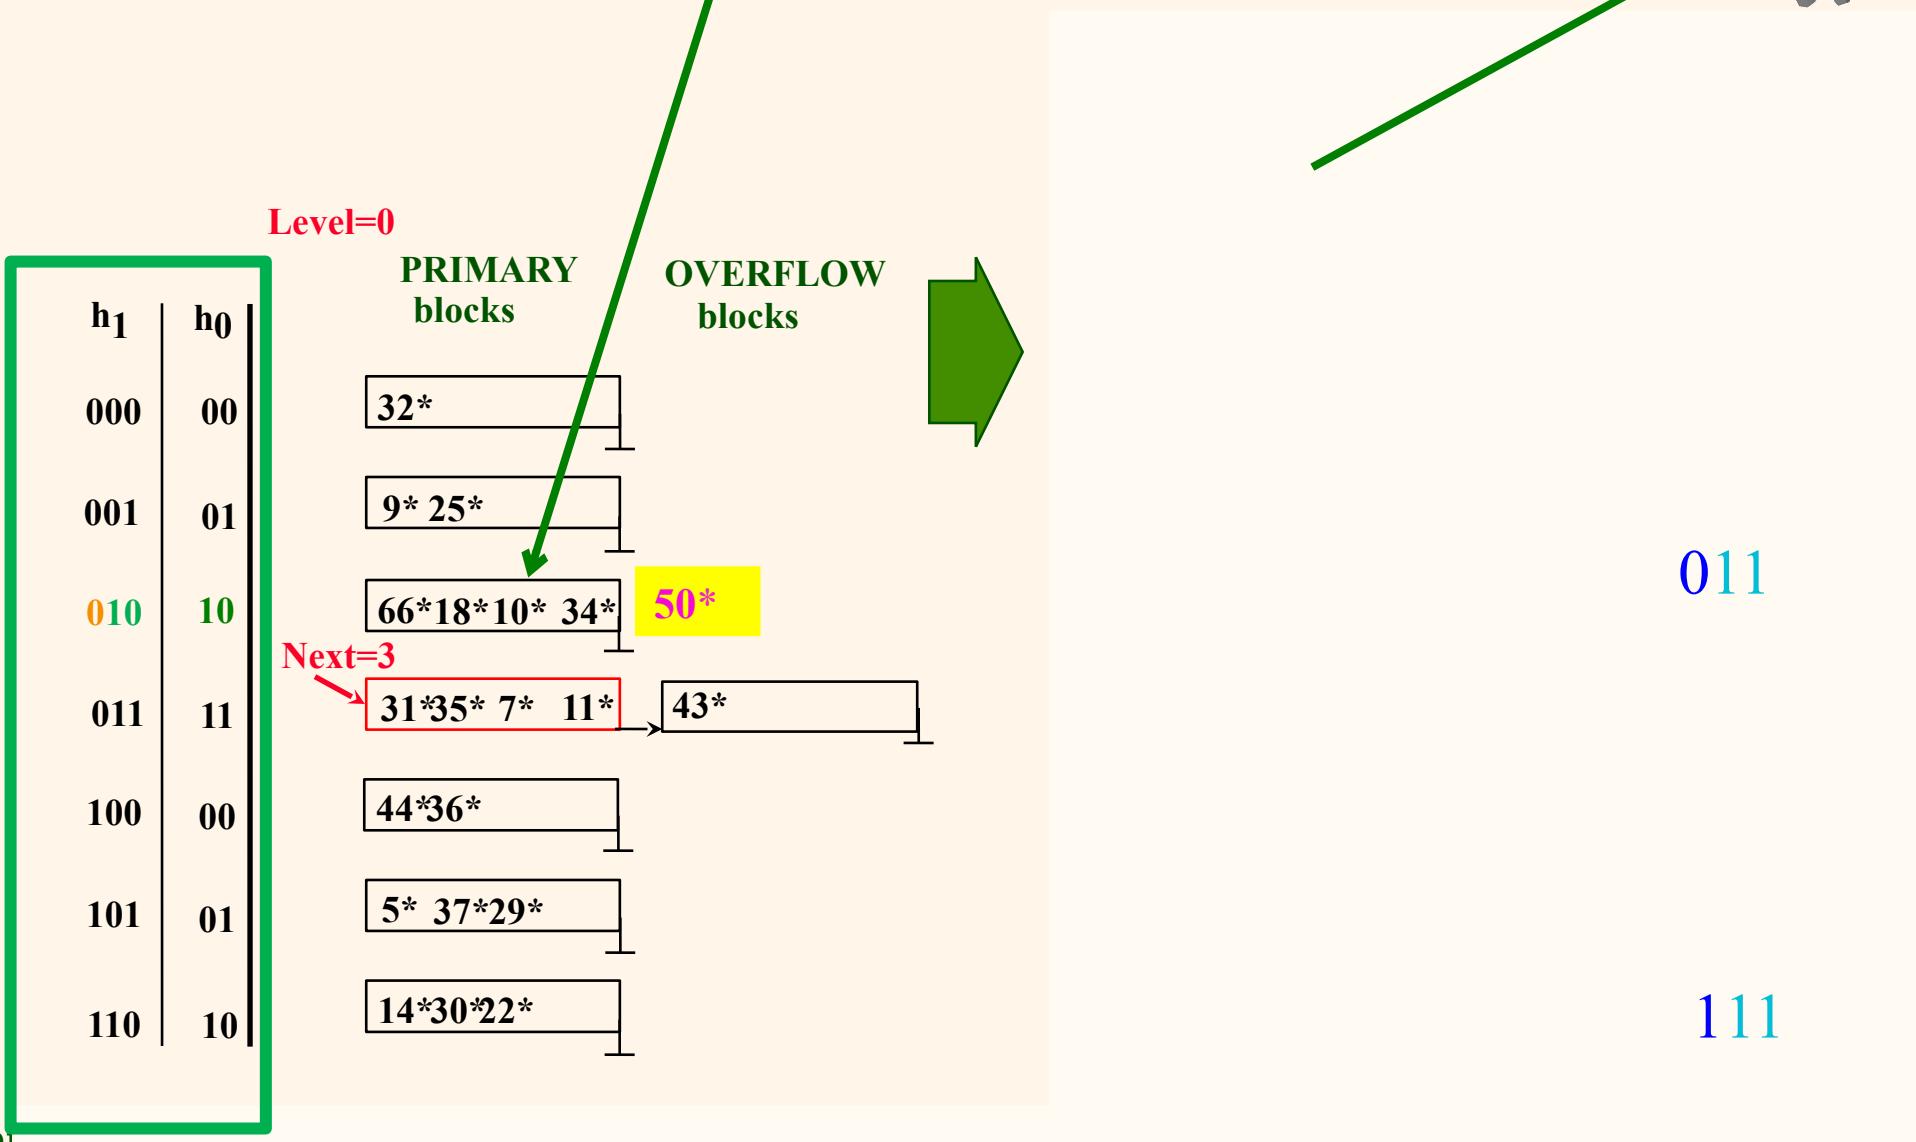
- ❖ On split, $h_{Level+1}$ is used to **re-distribute Data Entries k^*** .



Example: End of a Rour

Insert 50* (110010)

- ❖ **Insert:** Find Bucket by applying $h_{Level} / h_{Level+1}$:
- If Bucket to insert into is full:
 - Add **overflow Page** and insert data entry k*.
 - Split **Next Bucket** and increment **Next**.

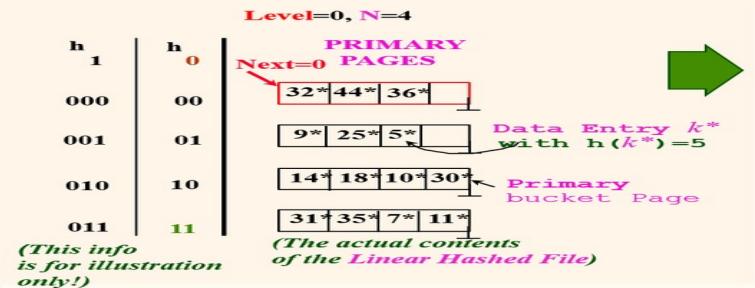


Dynamic Linear Hashing



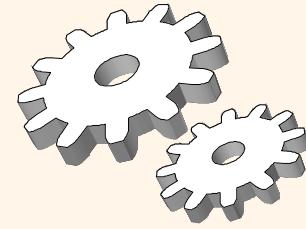
Indexes

DIRECTORY avoided in LH by using overflow **blocks**, and choosing **Bucket/Page** to split round-robin.

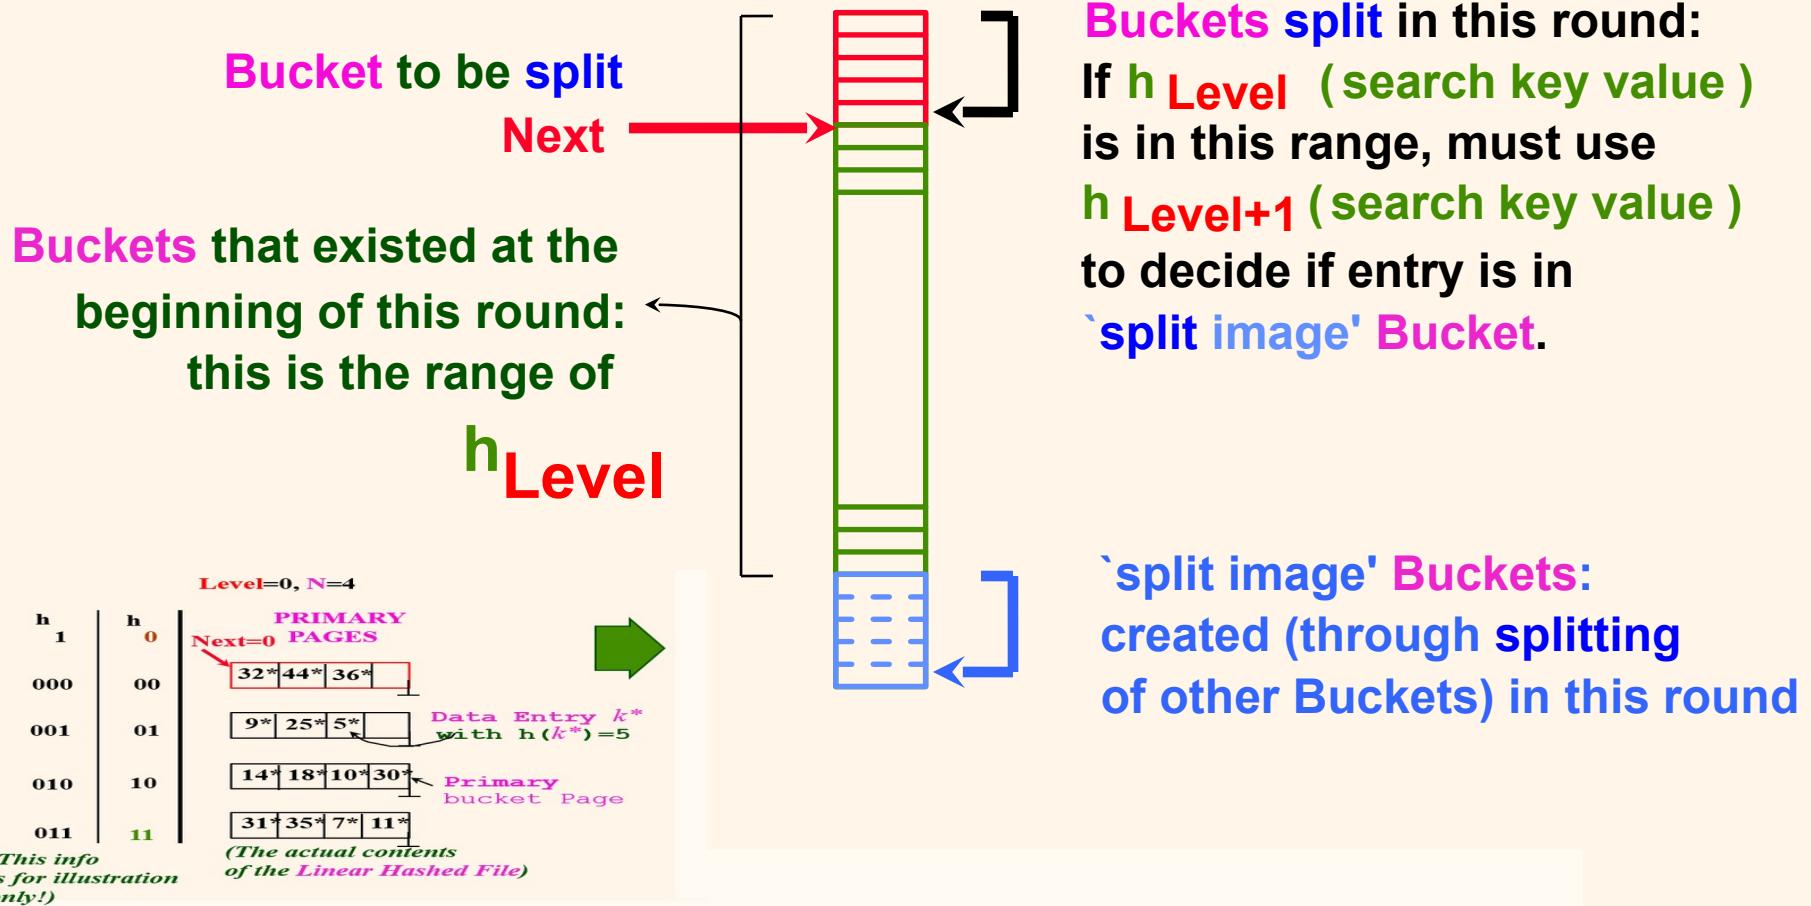


- Splitting proceeds in 'Rounds'. 'Round' ends when all N_R initial (for round R) **Buckets** are split. **Buckets** 0 to $Next-1$ (**next**) have been split; $Next$ to N_R yet to be split.
- Current 'Round' number is *Level*. (**current_stage**)
- Search k^* : To find **Bucket** for Data Entry k^* , find $h_{Level}(k^*)$:
 - If $h_{Level}(k^*)$ in range ' $Next$ to N_R ', k^* belongs here.
 - Else (i.e., if $h_{Level}(k^*)$ in range ' 0 to $Next-1$ ') **Bucket** $h_{Level}(k^*)$ is split.

Overview of *LH* File Indexes



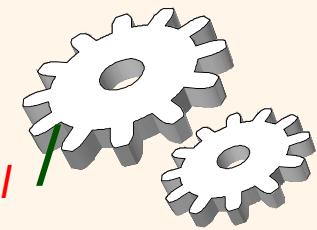
- ❖ In the middle of a 'Round'.



Dynamic Linear Hashing

Indexes

❖ Insert k^* : Find **Bucket** by applying h_{Level}



$h_{Level+1}$:

- If **Bucket** to insert into is full:
 - Add **overflow Page** and insert Data Entry k^* .
 - Split **Next Bucket** and increment **Next**.

❖ Since **Buckets** are split **round-robin**, long overflow chains don't develop!

❖ Doubling of **DIRECTORY** in **Extendable Hashing** is similar; switching of hash function shows how the # of bits examined is

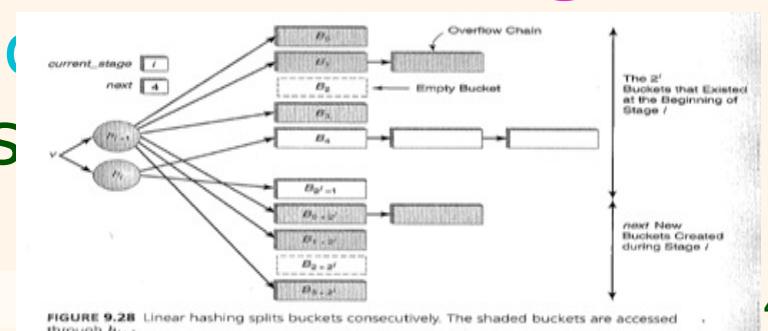


FIGURE 9.28 Linear hashing splits buckets consecutively. The shaded buckets are accessed through h_{i+1} .

LH Described as a Variant of EH

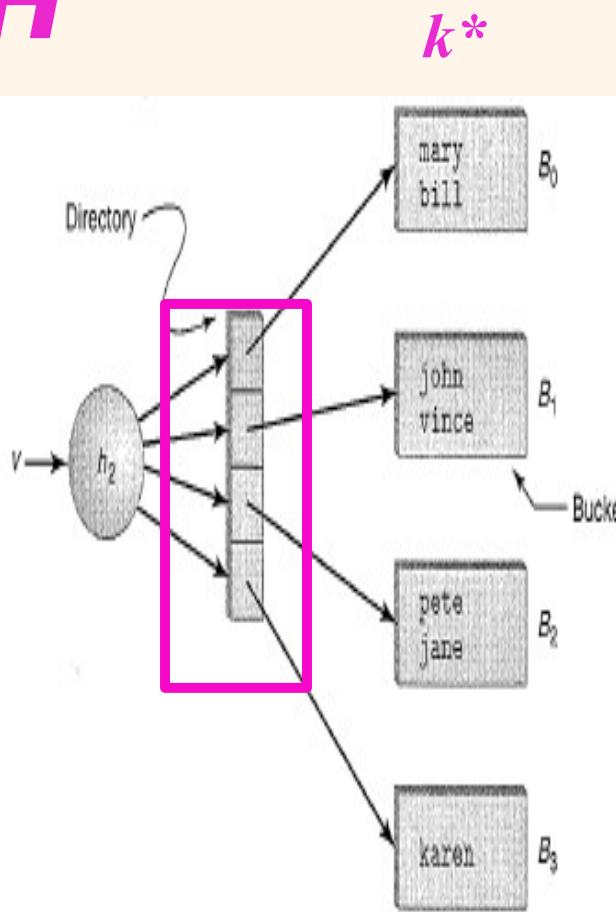


FIGURE 9.25 With extendable hashing, the hash result is mapped to a bucket through a directory.

EH

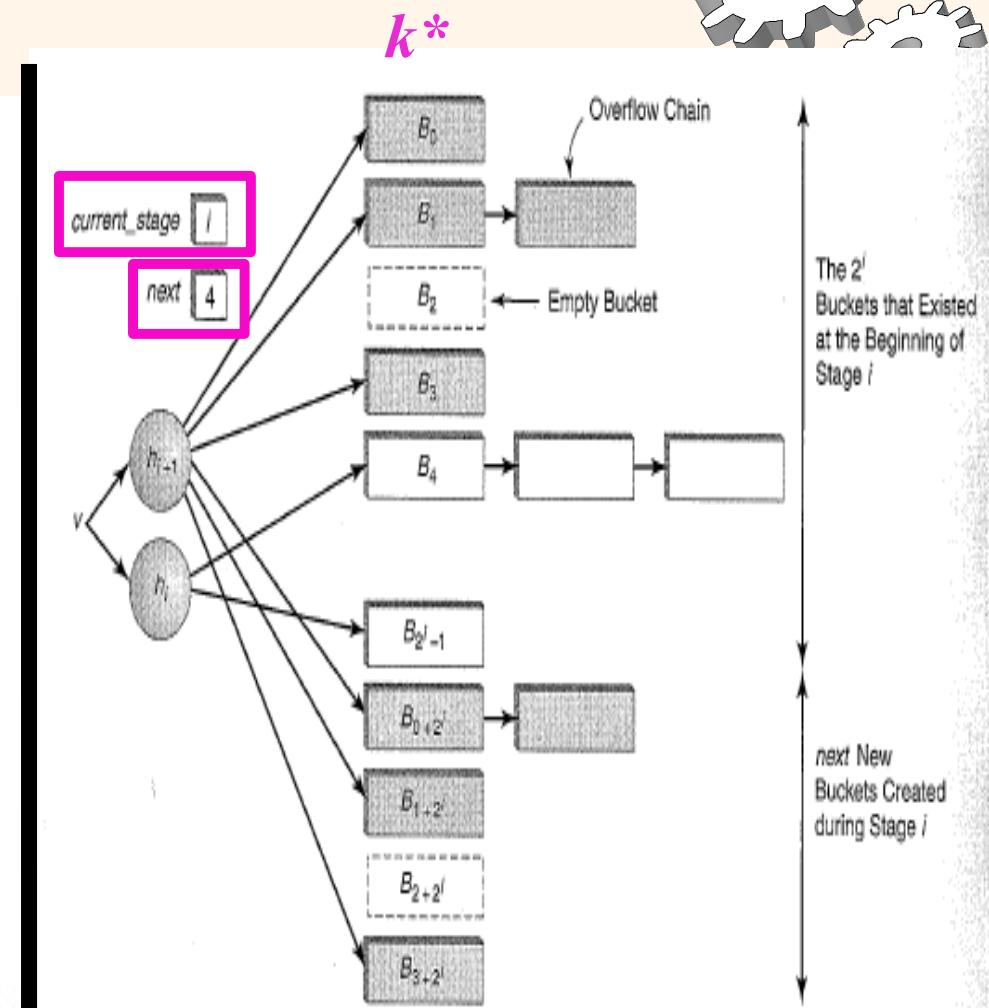
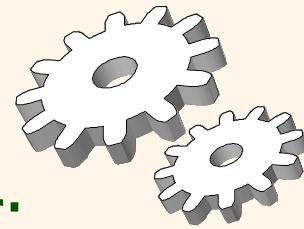


FIGURE 9.28 Linear hashing splits buckets consecutively. The shaded buckets are accessed through h_{i+1} .

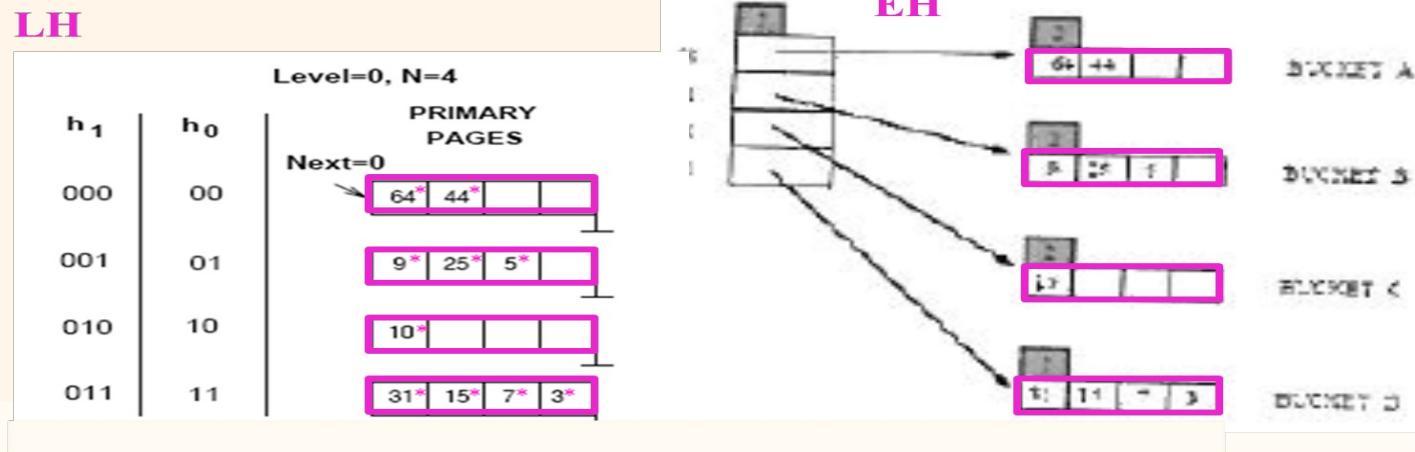
LH

LH Described as a Variant of EH

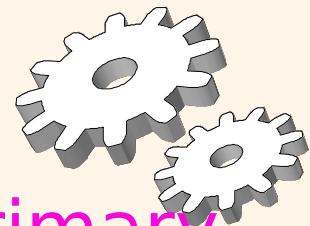


The two schemes are actually quite similar:

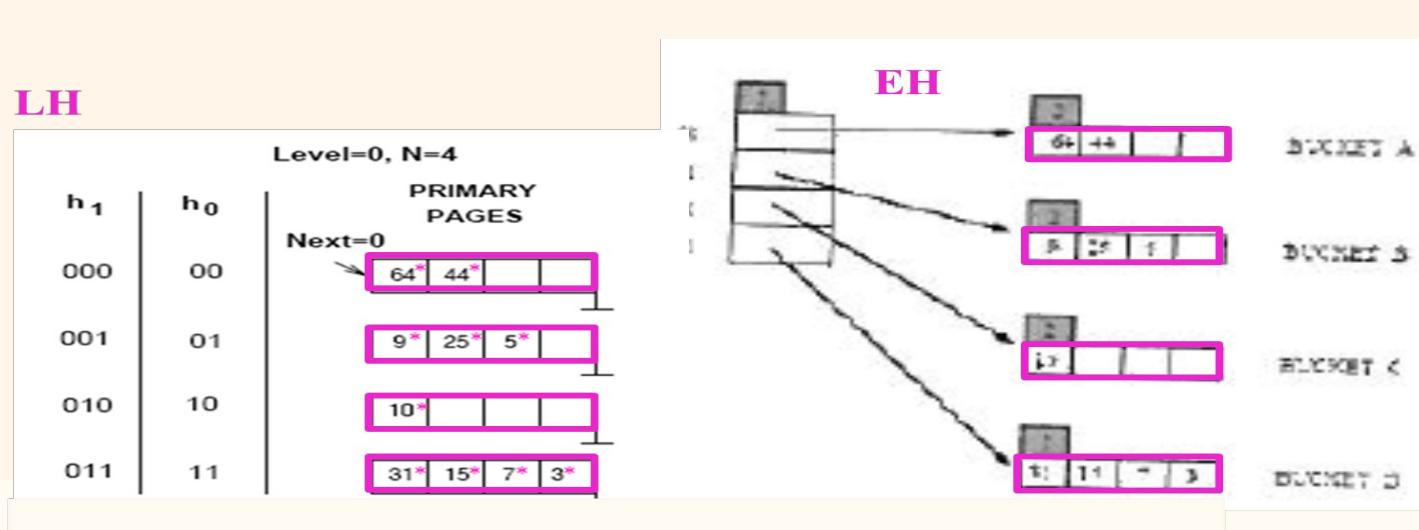
- Begin with an **EH Index** where **DIRECTORY** has **N** elements.
- Use overflow **blocks**, split **Buckets** round-robin.
- First split is at **Bucket 0**. (Imagine **DIRECTORY** being doubled at this point.) But elements $\langle 1, N+1 \rangle$, $\langle 2, N+2 \rangle$, ... are the same. So, need only create **DIRECTORY** element **N**, which differs from 0, now.
 - When **Bucket 1 splits**, create **DIRECTORY** element **N+1**, etc.

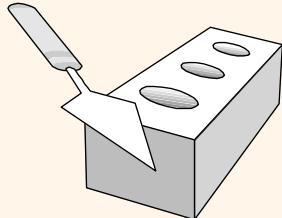


LH Described as a Variant of EH



So, DIRECTORY can double gradually. Also, primary Bucket blocks are created in order. If they are *allocated* in sequence too (so that finding i'th is easy), **we actually don't need a DIRECTORY!** Voila LH.





16. SET 3 - 4: STORAGE IV HASH INDEX



0%



0%



16.1 Other indexes Hidden



0%



0%



Hash indexes

The multi-level index is the most commonly used index type. Several additional index types are used less often but supported by many databases:

- Hash index
- Bitmap index
- Logical index
- Function index

Bitmap indexes

A **bitmap index** is a grid of bits:

- Each index row corresponds to a unique table row. If the table's primary key is an integer, the index row number might be the primary key value. Alternatively, the index row number might be an internal table row number, maintained by the database.
- Each index column corresponds to a distinct table value. Ex: If the index is on AirportCode, each index column corresponds to a different three-letter airport code. The mapping of index column numbers to table values is computed with a function or stored in an internal 'lookup' table.

Bitmap indexes contain ones and zeros. 'One' indicates that the table row corresponding to the index row number contains the table value corresponding to the index column number. Zero' indicates the row does not contain the value.

To locate rows containing a table value, the database:

1. Determines the index column corresponding to the table value.
2. Reads the index column and finds index rows that are set to 'one'.
3. Determines table rows corresponding to the index rows.
4. Determines pointers to blocks containing the table rows.
5. Reads the blocks containing the table rows.

An efficient bitmap index has two characteristics:

- The database can quickly determine the block containing a table row from the index row number (steps 3 and 4). Ex: The index row number is the hash key for a hash table. The block is determined by applying the hash function to the row number. Ex: The index row number is the table primary key. The block is determined with a primary index.
- The indexed column contains relatively few distinct values, typically tens or hundreds. If the indexed column contains thousands of distinct values, the bitmap index is large and inefficient.

Bitmap indexes with the above characteristics enable fast reads. Ex: A table with 10 million rows has a bitmap index on a column with 100 distinct values. The index contains 125 million bytes (10 million rows × 100 column values × 1 bit per index entry / 8 bits per byte) and can easily be retained in memory.

PARTICIPATION ACTIVITY

16.1.3: Bitmap index.



	ATL	DUB	LAX	MSN	ORD	SEA	SJC
0	1	0	0	0	0	0	0
1	0	1	0	0	0	0	0
2	0	0	1	0	0	0	0
3	0	0	0	1	0	0	0
4	0	0	0	0	1	0	0
5	1	0	0	0	0	0	0
6	0	0	0	0	0	1	0
7	0	0	0	1	0	0	0
8	0	0	0	0	1	0	0
9	1	0	0	0	0	0	0
10	0	0	1	0	0	0	0
11	0	0	0	0	0	0	1

	table				blocks
0	688	United Airlines	8:00	ATL	Airbus 321
1	140	Aer Lingus	10:30	DUB	Boeing 737
2	803	Air China	3:20	LAX	Boeing 777
3	801	Air India	3:15	MSN	Boeing 747
4	552	Lufthansa	18:00	MSN	MD 111
5	1100	Air China	3:20	ATL	Boeing 787
6	971	Air India	19:35	SEA	Boeing 747
7	702	American Airlines	5:05	MSN	Airbus 323
8	1154	American Airlines	13:50	ATL	MD 111
9	799	British Airways	13:50	ATL	Airbus 323
10	1222	Delta Airlines	8:25	LAX	Boeing 747
11	1001	Southwest Airlines	5:05	SJC	Airbus 323



bitmap index

	ATL	DUB	LAX	MSN	ORD	SEA	SJC
0	1	0	0	0	0	0	0
1	0	1	0	0	0	0	0
2	0	0	1	0	0	0	0
3	0	0	0	1	0	0	0
4	0	0	0	1	0	0	0
5	1	0	0	0	0	0	0
6	0	0	0	0	0	1	0
7	0	0	0	1	0	0	0
8	0	0	0	0	1	0	0
9	1	0	0	0	0	0	0
10	0	0	1	0	0	0	0
11	0	0	0	0	0	0	1

table

	Flight Number	Airlines	Time	Origin	Destination
0	666	United Airlines	6:00	ATL	Airbus 321
1	140	Aer Lingus	10:30	DUB	Boeing 737
2	803	Air China	3:20	LAX	Boeing 777
3	801	Air India	3:15	MSN	Boeing 747
4	552	Lufthansa	18:00	MSN	MD 111
5	1100	Air China	3:20	ATL	Boeing 787
6	971	Air India	19:35	SEA	Boeing 747
7	702	American Airlines	5:05	MSN	Airbus 323
8	1154	American Airlines	13:50	ORD	MD 111
9	799	British Airways	13:50	ATL	Airbus 323
10	1222	Delta Airlines	8:25	LAX	Boeing 747
11	1001	Southwest Airlines	5:05	SJC	Airbus 323

blocks

1) Refer to the bitmap index in the above animation. The three 1's in the MSN column means:

- The value MSN appears in three table blocks.
- The value MSN appears in three table rows.
- The value MSN appears in three buckets.

Correct

Each 1 in the MSN column indicates MSN appears in a distinct row. Since the MSN column contains three 1's, the value MSN appears in three rows.

?????

TA time (Jordan)

(CA 16.1.1 Step 2 – Bitmap Index)

CHALLENGE
ACTIVITY

16.1.1: Other indexes.

Bitmap Index

	LAX	DUB	MSN	SEA	ORD	ATL	SJC
0	0	0	0	0	1	1	0
1	0	0	0	0	0	0	1
2	0	0	0	0	0	1	0
3	0	1	0	0	0	0	0
4	0	0	0	1	0	0	0
5	0	0	0	0	0	0	0
6	0	0	1	0	0	0	0
7	0	0	0	0	0	0	0
8	0	0	1	0	0	0	0
9	0	1	1	0	0	0	0
10	1	0	0	0	0	0	1
11	0	0	0	0	0	1	0

row 3
row 9

Table row number(s) containing DUB: Ex: **3,9** Enter as comma-separated list.

2

VH

?????



COSC 3380 23578 M & W 4 to 5:30 PM LECTURE (in TEAMS from 3:50 to 5:15 PM)



01:00:26

Pop out

Chat

114
People

Raise

View

Rooms

Apps

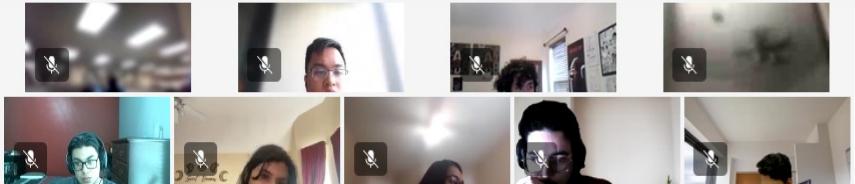
More

Camera

Mic

Share

Leave



Yu, Jordan T



View all



Participants

Invite someone or dial a number

Share invite

In this meeting (114)

Mute all

Hilford, Victoria
Organizer

RA Adhikari, Rohit

MA Ahmed, Mohamed A

AA Akram, Ali

BA Akukwe, Benetta O

SA Alsayed, Sami H

SA Altaf, Sameer

SA Alvarez, Stephanie

OA Anayor-Achu, Ogochukwu E

HA Avci, Hatice Kubra

RA Aysola, Riya

AB Bahl, Anish

zyBooks My library > COSC 3380: Database Systems home > 16.1: Other indexes

Feedback?

CHALLENGE ACTIVITY 16.1.1: Other indexes.

Jump to level 1

bitmap index

ORD	MSN	SEA	ATL	DUB	SJC	LAX
0	0	0	0	0	1	0
1	0	0	1	1	0	0
2	0	0	0	0	0	1
3	0	0	0	0	0	0
4	0	0	0	0	0	0
5	0	1	0	1	1	0
6	0	0	0	0	0	0
7	0	0	0	0	1	0
8	1	0	1	1	0	0
9	1	0	0	0	0	0
10	0	0	0	0	1	0
11	0	0	0	0	0	0

Table row number(s) containing LAX: Enter as comma-separated list.

1 2 3 4

✓ Expected: 2, 5

In the LAX column, each 1 indicates the table row number containing LAX.

LAX has a 1 in row 2, so 2 is a table row number containing LAX. Same for 5.

[View solution](#) (Instructors only) learns.oreilly.com is sharing your screen 441 PM 3/27/2024

Hash indexes

The multi-level index is the most commonly used index type. Several additional index types are used less often but supported by many databases:

- Hash index
- Bitmap index
- Logical index
- Function index

Logical indexes

A single- or multi-level index normally contains pointers to table blocks and is called a **physical index**.

A **logical index** is a single- or multi-level index in which pointers to table blocks are replaced with primary key values. Logical indexes are always secondary indexes and require a separate primary index on the same table. To locate a row containing a column value, the database:

1. Looks up the column value in the logical index to find the primary key value.
2. Looks up the primary key value in the primary index to find the table block pointer.
3. Reads the table block containing the row.

Logical indexes change only when primary key values are updated, which occurs infrequently. Physical indexes change whenever a row moves to a new block, which occurs in several ways:

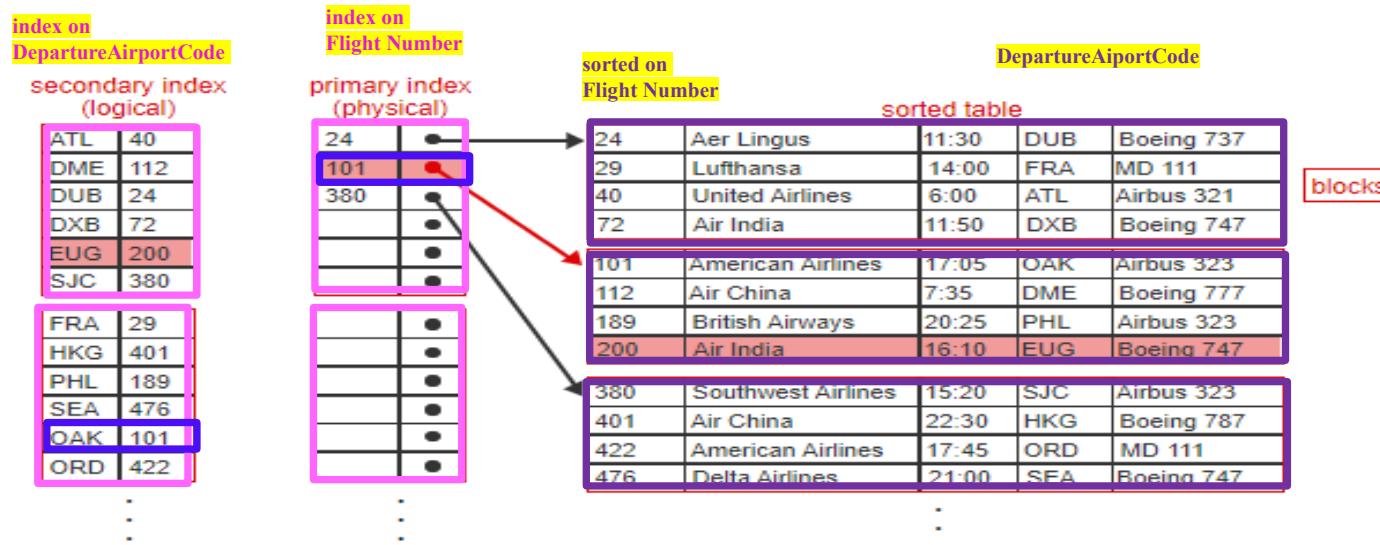
- A row is inserted into a full block. To create space for the new row, the block splits and some rows move to a new block.
- A clustering column is updated. When a clustering column is updated, the row may move to a new block to maintain sort order.
- The table is reorganized. Occasionally, a database administrator may physically reorganize a table to recover deleted space or order blocks contiguously on magnetic disk.

If a table has several indexes, the time required to update physical indexes is significant, and logical indexes are more efficient.

On read queries, a logical index requires an additional read of the primary index and is slower than a physical index. However, the primary index is often retained in memory, mitigating the cost of the additional read.

PARTICIPATION ACTIVITY

16.1.5: Logical index.



Logical indexes

A single- or multi-level index normally contains pointers to table blocks and is called a **physical index**.

A **logical index** is a single- or multi-level index in which pointers to table blocks are replaced with primary key values. Logical indexes are always secondary indexes and require a separate primary index on the same table. To locate a row containing a column value, the database:

1. Looks up the column value in the logical index to find the primary key value.
2. Looks up the primary key value in the primary index to find the table block pointer.
3. Reads the table block containing the row.

Logical indexes change only when primary key values are updated, which occurs infrequently. Physical indexes change whenever a row moves to a new block, which occurs in several ways:

- A row is inserted into a full block. To create space for the new row, the block splits and some rows move to a new block.
- A clustering column is updated. When a clustering column is updated, the row may move to a new block to maintain sort order.
- The table is reorganized. Occasionally, a database administrator may physically reorganize a table to recover deleted space or order blocks contiguously on magnetic disk.

If a table has several indexes, the time required to update physical indexes is significant, and logical indexes are more efficient.

On read queries, a logical index requires an additional read of the primary index and is slower than a physical index. However, the primary index is often retained in memory, mitigating the cost of the additional read.

PARTICIPATION ACTIVITY

16.1.6: Logical indexes.

An index on a unique sort column.

An index on a non-sort column.

An index with primary key values rather than block pointers.

Logical index

A logical index contains primary index values rather than table block pointers. Usually the primary index is on the primary key of the table, so logical indexes usually contain primary key values.

?????

An index with table block pointers.

TA time (Jordan)

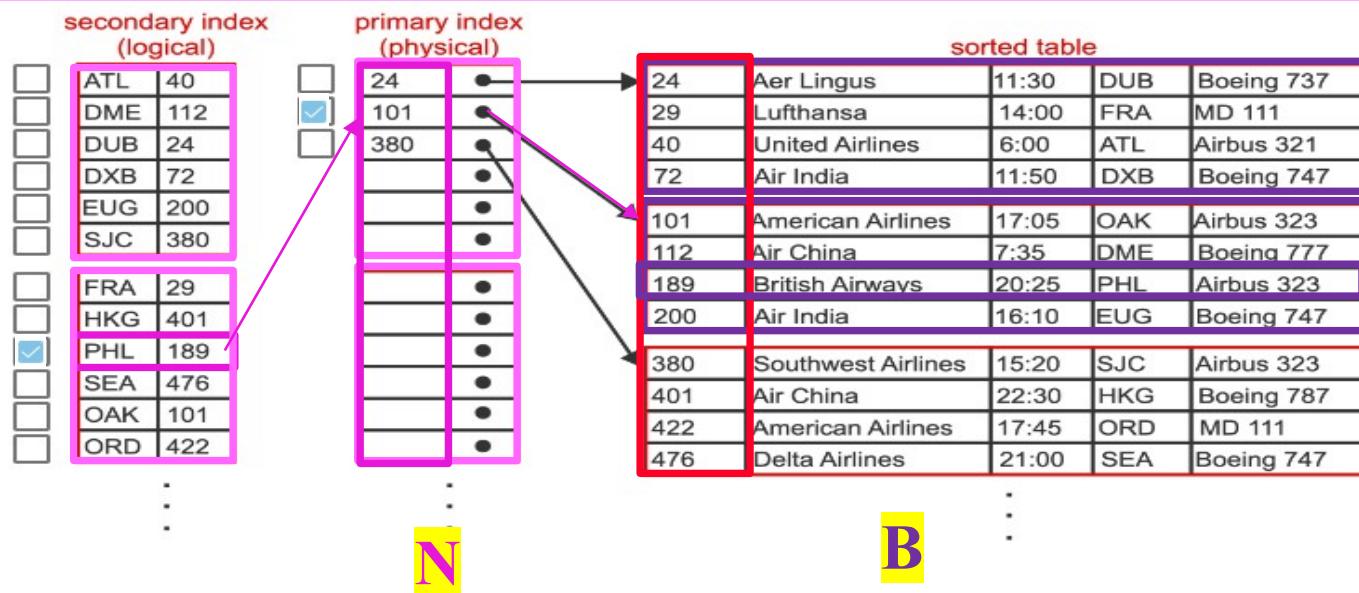
(CA 16.1.1 Step 3 – Logical Index)

CHALLENGE ACTIVITY

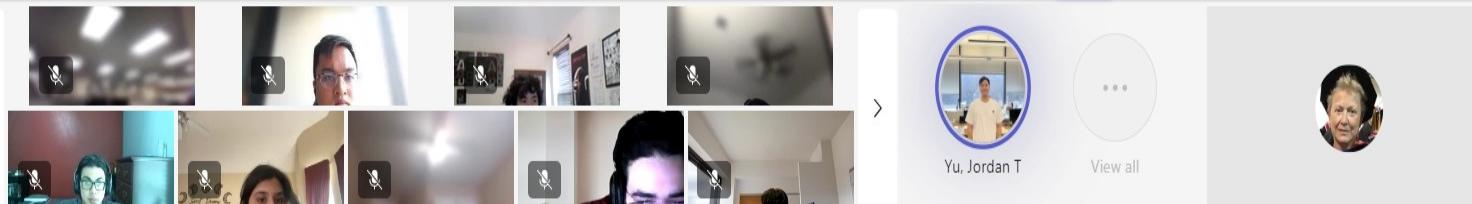
16.1.1: Other indexes.

Logical Index

The values in the primary index and the table are sorted on flight number. Select the indexes to find the row containing PHL.



01:07:22



Participants

Invite someone or dial a number

 Share invite

▼ In this meeting (113)

Mute all

Section 16.1 - COSC 3380: Databases

Meeting in "General" (1 person)

Index (1324) Google Docs 2024P-COSC3380... NumPy for Neural Ra... Fantasy Basketball 2... AccessDB ChatGPT The Perfect Practice... Blend 75 LearnCode... Deep Learning Du... GitHub for COSC... All Bookmarks

zyBooks My library > COSC 3380: Database Systems home > 16.1: Other indexes

zyBooks catalog Help/FAQ Jordan Yu

CHALLENGE ACTIVITY 16.1.1: Other indexes.

Jump to level 1

The values in the primary index and the table are sorted on flight number. Select the indexes to find the row containing DME.

secondary index (logical)

<input type="checkbox"/> ATL 40
<input checked="" type="checkbox"/> DME 112
<input type="checkbox"/> DUB 24
<input type="checkbox"/> DXB 72
<input type="checkbox"/> EUG 200
<input type="checkbox"/> SJC 380
<input type="checkbox"/> FRA 29
<input type="checkbox"/> HKG 401
<input type="checkbox"/> PHL 189
<input type="checkbox"/> SEA 476
<input type="checkbox"/> OAK 101
<input type="checkbox"/> ORD 422
...
...
...

primary index (physical)

<input checked="" type="checkbox"/> 24
<input type="checkbox"/> 101
<input type="checkbox"/> 380
...
...
...

sorted table

24	Aer Lingus	11:30	DUB	Boeing 737
29	Lufthansa	14:00	FRA	MD 111
40	United Airlines	6:00	ATL	Airbus 321
72	Air India	11:50	DXB	Boeing 747
101	American Airlines	17:05	OAK	Airbus 323
112	Air China	7:35	DME	Boeing 777
189	British Airways	20:25	PHL	Airbus 323
200	Air India	16:10	EUG	Boeing 747
380	Southwest Airlines	15:20	SJC	Airbus 323
401	Air China	22:30	HKG	Boeing 787
422	American Airlines	17:45	ORD	MD 111
476	Delta Airlines	21:00	SEA	Boeing 747

Check Next

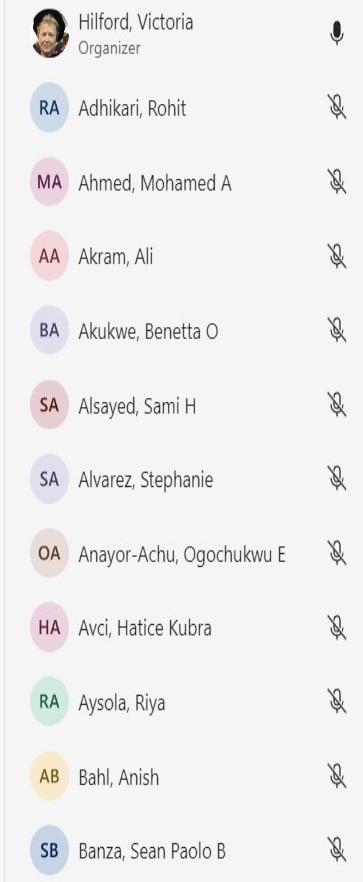
✓ Expected: Secondary index 112, primary index 101

The values in both the primary index and the table are sorted on flight number.

First, the database searches the secondary index to find the primary key for DME, which is 112. Next, the database searches the primary index for the largest key that is a primary key, which is 101. The table block pointer points to the table block that contains the row.

Type here to search View solution (Instructors only) teams.microsoft.com is sharing your screen Stop sharing Hide

65° Cloudy 3/27/2024



Hash indexes

The multi-level index is the most commonly used index type. Several additional index types are used less often but supported by many databases:

- Hash index
- Bitmap index
- Logical index
- Function index

Function indexes

In some cases, values specified in a WHERE clause may be in a different format or units than values stored in the column. Ex:

- The WHERE clause specifies values in upper case, but the column contains mixed upper and lower case characters.
- The WHERE clause specifies values as percentages, from 0 to 100, but the column contains values from 0 to 1.

In the above examples, index entries do not match values in the WHERE clause, so the database cannot use the index to execute the query.

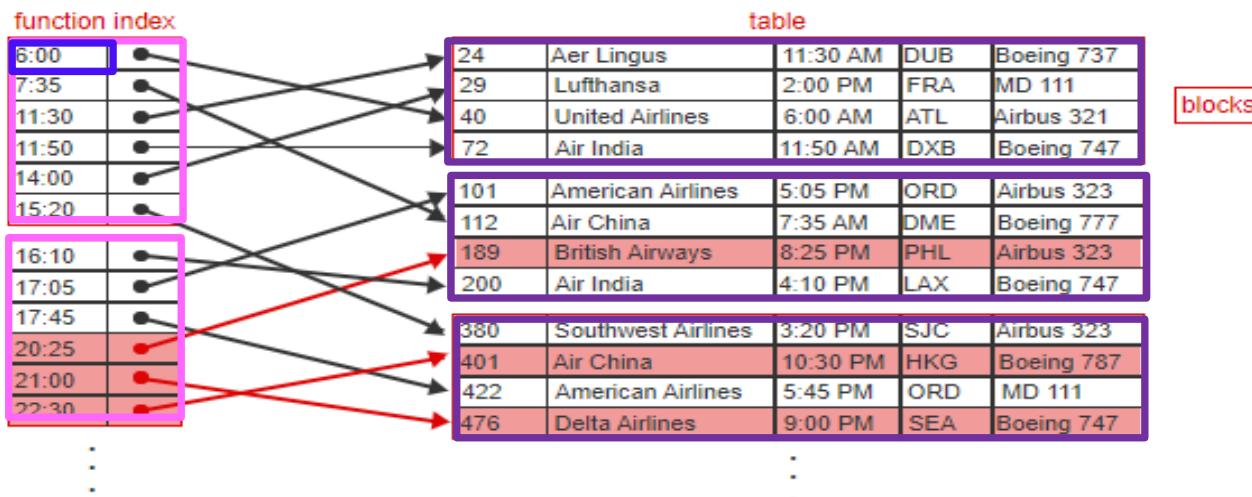
To address this problem, some databases support function indexes. In a **function index**, the database designer specifies a function on the column value. Index entries contain the result of the function applied to column values, rather than the column values.

Ex: Column values are stored as decimal numbers between 0 and 1, but users specify percentages as integers between 0 and 100 in queries. The database designer specifies a function index that multiplies column values by 100 and converts the result to an integer. The index contains integers between 0 and 100, so the database can use the function index to process queries.

In principle, functions can be used with any index type, including single-level, multi-level, hash, bitmap, and logical indexes. In practice, support varies by database.

PARTICIPATION ACTIVITY

16.1.7: Function index.



```
SELECT FlightNumber, AirlineName  
FROM Flight  
WHERE DepartureTime > "20:00"
```



function index

6:00
7:35
11:30
11:50
14:00
15:20
16:10
17:05
17:45
20:25
21:00
22:30

table

24	Aer Lingus	11:30 AM	DUB	Boeing 737
29	Lufthansa	2:00 PM	FRA	MD 111
40	United Airlines	6:00 AM	ATL	Airbus 321
72	Air India	11:50 AM	DXB	Boeing 747
101	American Airlines	5:05 PM	ORD	Airbus 323
112	Air China	7:35 AM	DME	Boeing 777
189	British Airways	8:25 PM	PHL	Airbus 323
200	Air India	4:10 PM	LAX	Boeing 747
380	Southwest Airlines	3:20 PM	SJC	Airbus 323
401	Air China	10:30 PM	HKG	Boeing 787
422	American Airlines	5:45 PM	ORD	MD 111
476	Delta Airlines	9:00 PM	SEA	Boeing 747

blocks

```
SELECT FlightNumber, AirlineName
FROM Flight
WHERE DepartureTime > "20:00"
```

- 1) Refer to the animation above. Which table blocks does the following query read?

```
SELECT FlightNumber, AirlineName
FROM Flight
WHERE DepartureTime > "12:00" AND DepartureTime < "16:00";
```

Correct

The query selects departure times of 14:00 and 15:20 from the index. 14:00 refers to flight 29 in block 0, and 15:20 refers to flight 380 in block 2.

?????

- Block 0
- Block 1
- Blocks 0 and 2

TA time (Jordan)

(CA 16.1.1 Step 4 – Function Index)

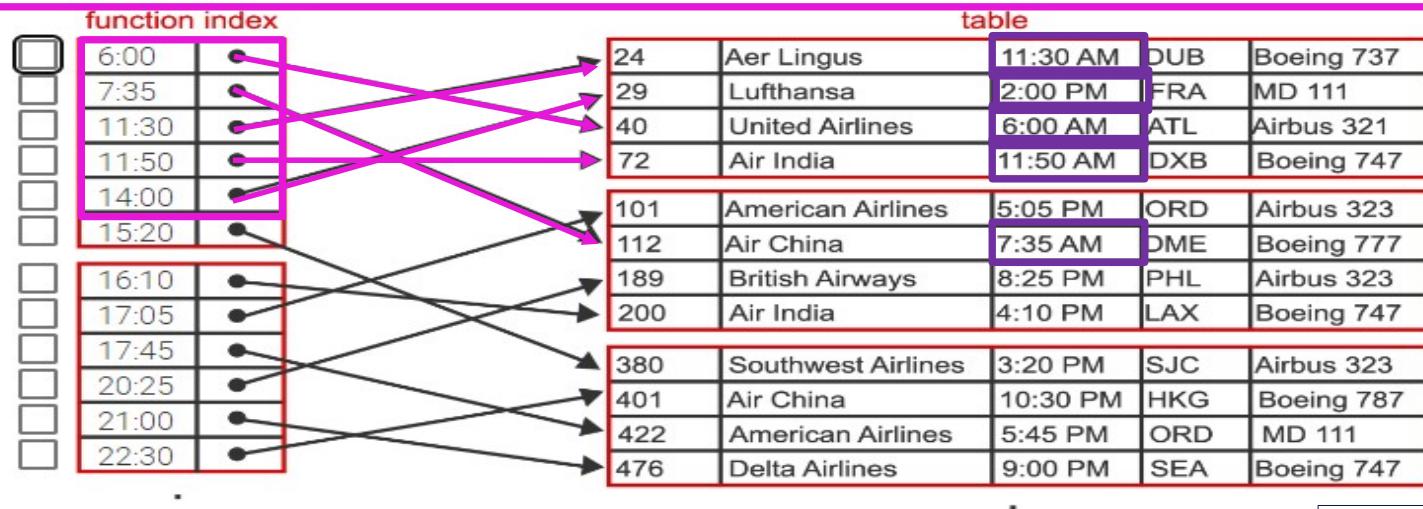
CHALLENGE
ACTIVITY

16.1.1: Other indexes.

Function Index

The table contains times in AM/PM format. The function index converts AM/PM format to 24-hour format.

Select the indexes that match the expression: DepartureTime < "15:00", given that the expression specifies time in 24-hour format.



VH

?????



01:14:30

Pop out Chat 113 People Raise View Rooms Apps More Camera Mic Share Leave

Section 16.1 - COSC 3380: Database Systems home > 16.1: Other indexes

invite (132) Google Docs 2024SP COSC6987... YouTube Neural Ra... Fantasy Basketball 2... AcousUR ChatGPT The Perfect Practice... Blend 75 Lat Codes... Deep Learning - Dr... GaoJicheng/COSC...

zyBooks catalog Help/FAQ Jordan Yu ▾

CHALLENGE ACTIVITY 16.1.1: Other indexes.

Jump to level 1

The table contains times in AM/PM format. The function index converts AM/PM format to 24-hour format.

Select the indexes that match the expression: DepartureTime > '17:30', given that the expression specifies time in 24-hour format:

function index	table
6:00	24 Aer Lingus 11:30 AM DUB Boeing 737
7:35	29 Lufthansa 2:00 PM FRA MD 111
11:30	40 United Airlines 6:00 AM ATL Airbus 321
11:50	72 Air India 11:50 AM DXB Boeing 747
14:00	
15:20	
16:10	101 American Airlines 5:05 PM ORD Airbus 323
17:05	112 Air China 7:35 AM DME Boeing 777
17:45	189 British Airways 8:25 PM PHL Airbus 323
20:25	200 Air India 4:10 PM LAX Boeing 747
21:00	
22:30	380 Southwest Airlines 3:20 PM SJC Airbus 323
	401 Air China 10:30 PM HKG Boeing 787
	422 American Airlines 5:45 PM ORD MD 111
	476 Delta Airlines 9:00 PM SEA Boeing 747

Feedback?

Check Next Done Click any level to practice more. Completion is preserved.

✓ Expected: 17:45, 20:25, 21:00, and 22:30

The departure times 17:45, 20:25, 21:00, and 22:30 match the expression DepartureTime > '17:30'.

[View solution](#) (instructors only)

learns.microsoft.com sharing your screen Stop sharing [Feedback](#)

Participants

Invite someone or dial a number

Share invite

In this meeting (113)

Mute all

Hilford, Victoria Organizer

RA Adhikari, Rohit

MA Ahmed, Mohamed A

AA Akram, Ali

BA Akukwe, Benetta O

SA Alsayed, Sami H

SA Alvarez, Stephanie

OA Anayor-Achu, Ogochukwu E

HA Avci, Hatice Kubra

RA Aysola, Riya

AB Bahl, Anish

SB Banza, Sean Paolo B

Hash Index Practice Questions

Hash index entries are assigned to _____.

- a. values
- b. clusters
- c. buckets
- d. blocks

????

What are two characteristics of an efficient bitmap index?

- a. Each bucket initially has one block. Any additional blocks are allocated and linked to the initial block.
- b. The WHERE clause may specify any values in upper case. The column used in a WHERE clause may contain mixed upper and lower case characters.
- c. The database can quickly determine the block containing a table row from the index row number. Any indexed column contains relatively few distinct values.
- d. Indexes change only when primary key values are updated. Physical indexes change whenever a row moves to a new block.

????

At 5:00 PM .

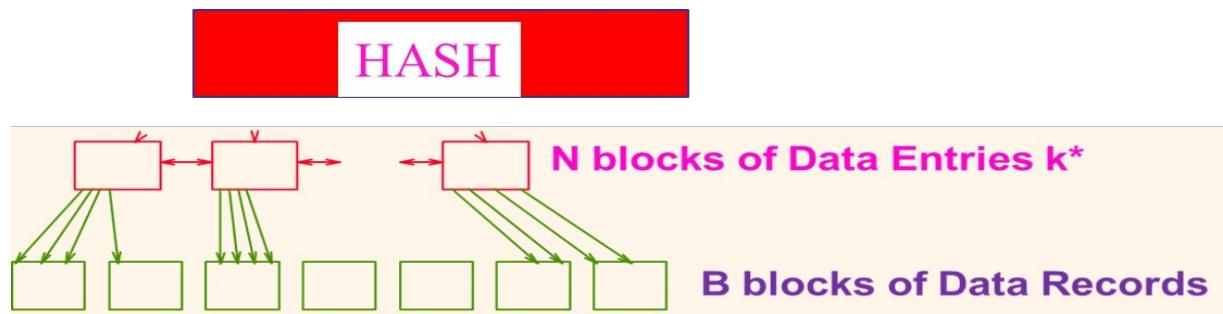
03.27.2024

ZyBook SET 3-4

Set 3

(19 - We)

LECTURE 15 STORAGE IV HASH INDEX



12. SET 3

Empty

16. SET 3 - 4:STORAGE IV HASH INDEX

0% 0% 0%

VH work on
SET 3 – 16

Next

04.01.2024

(20 - Mo)

EXAM 3 Practice

(PART of 20 points)

12. SET 3

Empty

13. SET 3 - 1:STORAGE I

Hidden

0%

0%

14. SET 3: 2:STORAGE II

Hidden

0%

15. SET 3 - 3:STORAGE III B TREE INDEX

Hidden

0%

0%

16. SET 3 - 4:STORAGE IV HASH INDEX

Hidden

0%

0%

From 5:05 to 5:15 PM – 5 minutes.



03.27.2024

ZyBook SET 3-4

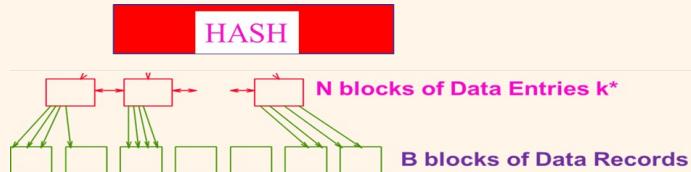
(19 - We)

Set 3

LECTURE 15 STORAGE IV HASH INDEX

CLASS PARTICIPATION 20 points

20% of Total + :



HASH INDEX

Class 19 END PARTICIPATION

Not available until Mar 27 at 5:05pm | Due Mar 27 at 5:15pm | 100 pts

VH, publish ⊖ :

This is a synchronous online class.

Attendance is required.

Recording or distribution of class materials is prohibited.

1. At the beginning of selected classes there is an assessment in the first 10 minutes. (beige BOX in the Detailed Syllabus)

2. At the end of selected classes there is an assessment in the last 10 minutes. (blue BOX in the Detailed Syllabus)

3. ZyBook sections will be downloaded and used for 30% of Total Score on the dates specified in the Detailed Syllabus.

4. EXAMS are in CANVAS. No late EXAMS.

5. I have to be present in TEAMS in order to take any graded assignment assigned during that class.

At 5:15 PM.

End Class 19

VH, Download Attendance Report

Rename it:

3.27.2024 Attendance Report FINAL TEAMS



EXAM 3 Practice



CLASS PARTICIPATION [20 points] Module | Not available until Apr 1 at 4:10pm | Due Apr 1 at 5:05pm | 80 pts

VH, upload Class 19 to CANVAS.