

COSC 3380 Spring 2024

Database Systems

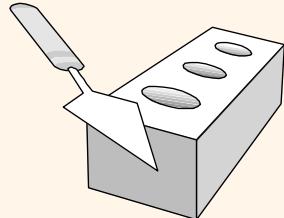
M & W 4:00 to 5:30 PM

Prof. **Victoria Hilford**

PLEASE TURN your webcam ON (must have)

NO CHATTING during LECTURE

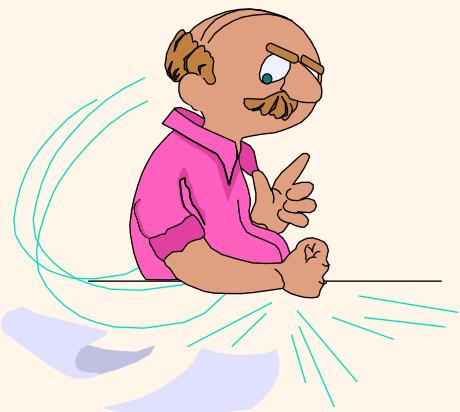
VH, UNhide Section 18



COSC 3380

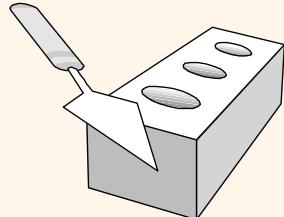
4 to 5:30

**PLEASE
LOG IN
CANVAS**

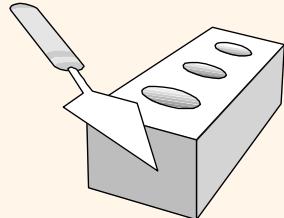


Please close all other windows.

COSC 3380

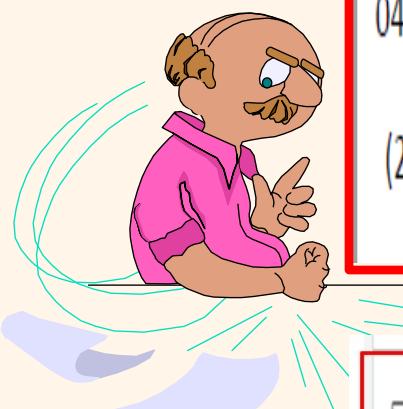


04.15.2024 (24 – Mo)	ZyBook SET 4 - 2	Set 4 LECTURE 17 CONCURRENCY CONTROL
04.17.2024 (25 – We)	ZyBook SET 4 - 3	Set 4 LECTURE 18 CRASH RECOVERY LECTURE 19 SECURITY and AUTHORIZATION
04.22.2024 (26 – Mo)		EXAM 4 Practice (PART of 20 points)
04.24.2024 (27 – We)	TA Download ZyBook SET 4 Sections (4 PM) (PART of 30 points)	EXAM 4 Review (PART of 20 points)
04.29.2024 (28 – Mo)		EXAM 4 (PART of 50 points)



COSC 3380

Class 24



04.15.2024

ZyBook SET 4 - 2

(24 - Mo)

Set 4

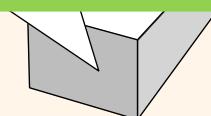
LECTURE 17 CONCURRENCY CONTROL

19. SET 4-2:CONCURRENCY CONTROL

Hidden | 0% v

LECTURE 17 CONCURRENCY CONTROL

From 4:00 to 4:07 PM – 5 minutes.



04.15.2024

ZyBook SET 4 · 2

(24-Mo)

Set 4

LECTURE 17 CONCURRENCY CONTROL

CLASS PARTICIPATION 20 points

20% of Total

+

⋮

CONCURRENCY



Class 24 BEGIN PARTICIPATION

Not available until Apr 15 at 4:00pm | Due Apr 15 at 4:07pm

VH, publish



This is a synchronous online class.

Attendance is required.

Recording or distribution of class materials is prohibited.

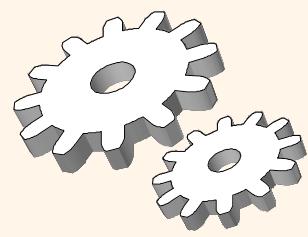
1. At the beginning of selected classes there is an assessment in the first 10 minutes. (beige BOX in the Detailed Syllabus)

2. At the end of selected classes there is an assessment in the last 10 minutes. (blue BOX in the Detailed Syllabus)

3. ZyBook sections will be downloaded and used for 30% of Total Score on the dates specified in the Detailed Syllabus.

4. EXAMS are in CANVAS. No late EXAMS.

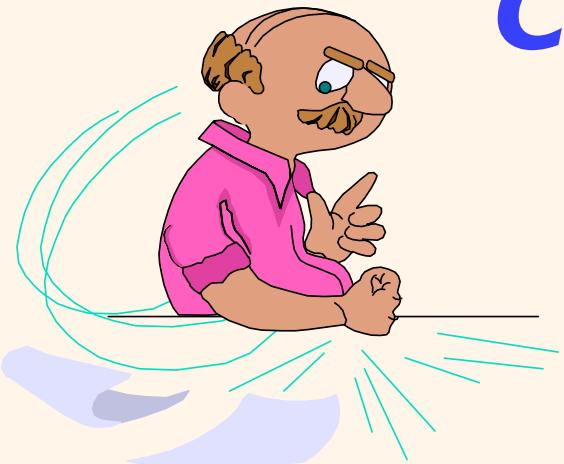
5. I have to be present in TEAMS in order to take any graded assignment assigned during that class.



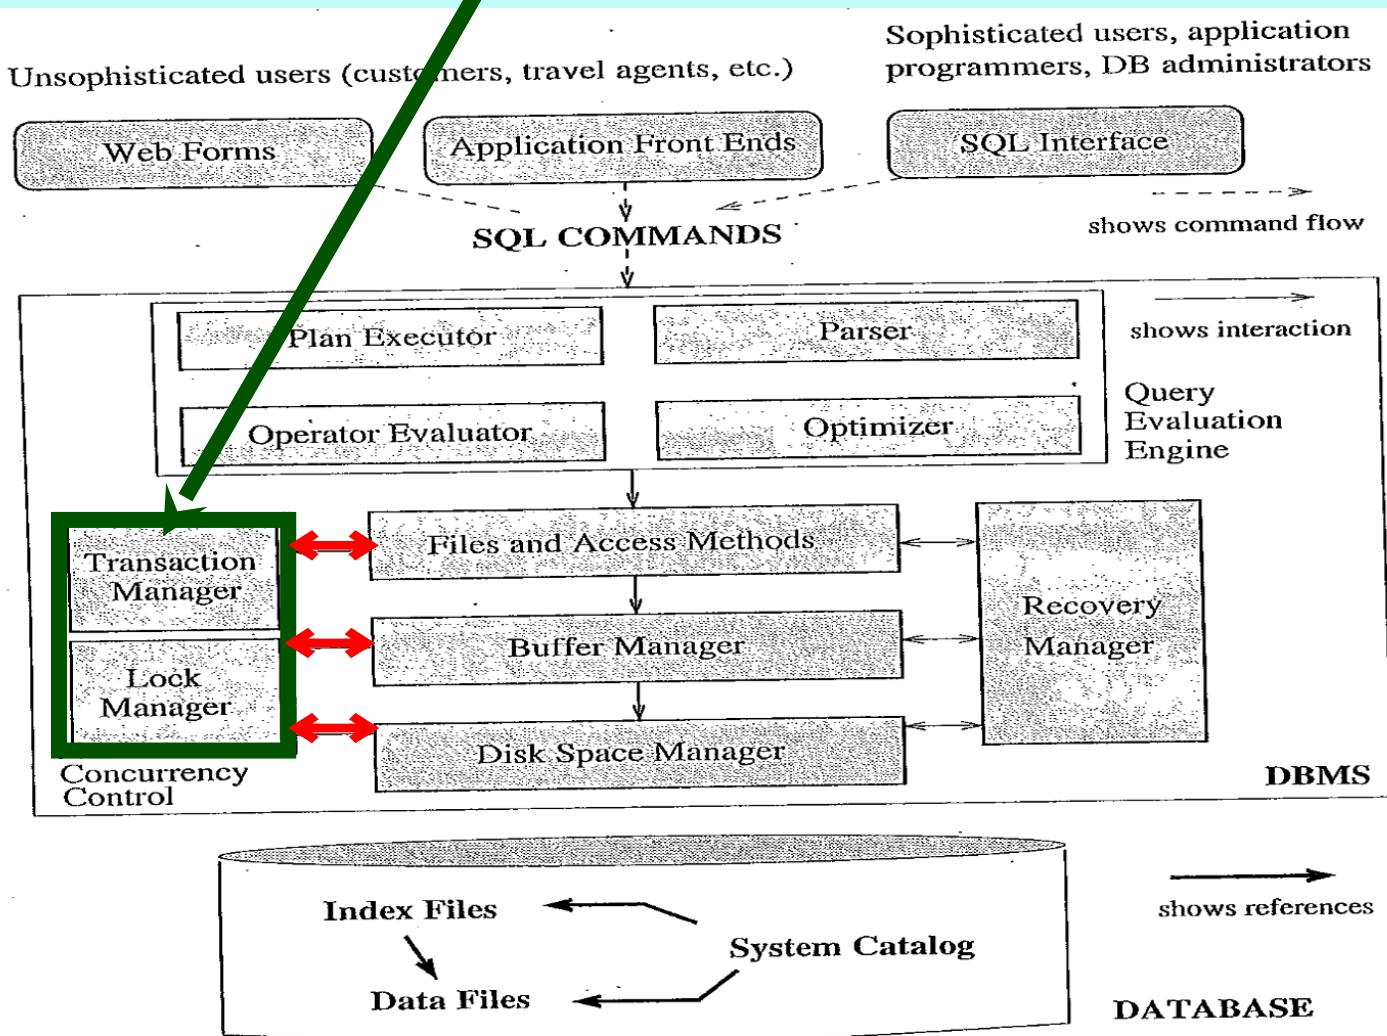
COSC 3380

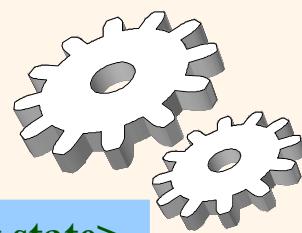
Lecture 17

Concurrency Control



A DBMS - Concurrency Control





A Sample *Transaction*

Initial state

<Consistent state>

```
1: Begin_Transaction
2:     get (K1, K2, CHF) from terminal
3:     Select BALANCE Into S1 From ACCOUNT Where ACCOUNTNR = K1;
4:     S1 := S1 - CHF;
5:     Update ACCOUNT Set BALANCE = S1 Where ACCOUNTNR = K1;
6:     Select BALANCE Into S2 From ACCOUNT Where ACCOUNTNR = K2;
7:     S2 := S2 + CHF;
8:     Update ACCOUNT Set BALANCE = S2 Where ACCOUNTNR = K2;
9:     Insert Into BOOKING(ACCOUNTNR, DATE, AMOUNT, TEXT)
          Values (K1, today, -CHF, 'Transfer');
10:    Insert Into BOOKING(ACCOUNTNR, DATE, AMOUNT, TEXT)
          Values (K2, today, CHF, 'Transfer');
12:    If S1<0 Then Abort_Transaction
11: End_Transaction
```

Final state

<Consistent state>

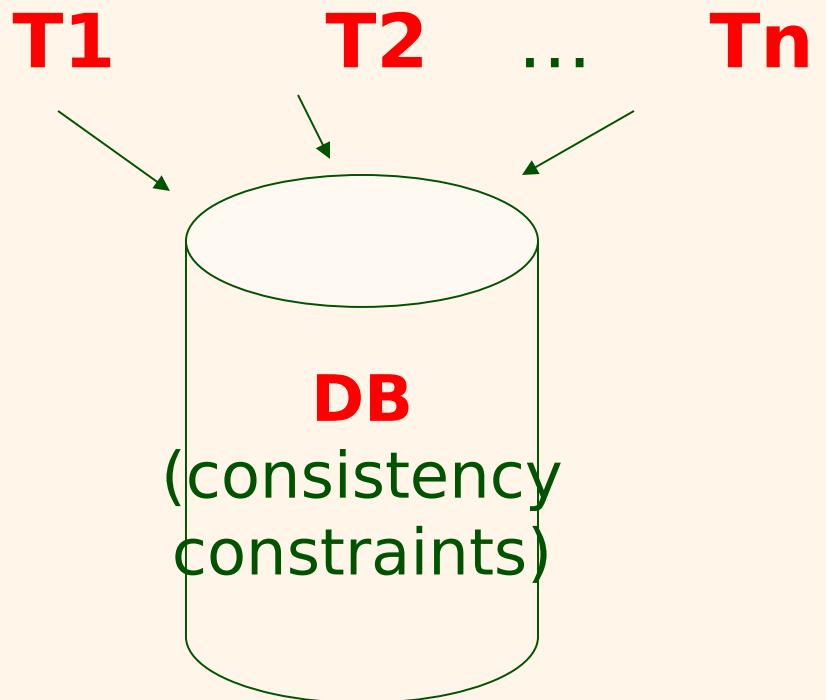
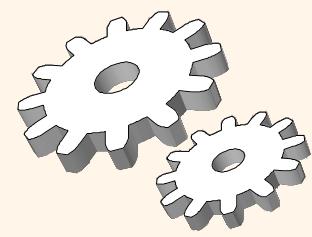
Transaction = Program that takes database from one consistent state
to another consistent state

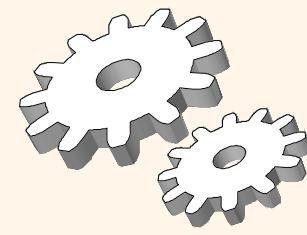
TA, Fernando (A – L).

TA, Alvaro (M – Z).

**Please compare CANVAS vs. TEAMS Attendance.
Print screens of students in CANVAS but not in the TEAMS meeting.
(4.15.2024 Attendance X missing LastName.docx)**

Concurrency Control





The Scheduler

- ❖ Bases its actions on Concurrency Control Algorithms
- ❖ Ensures computer's central processing unit (**CPU**) is used **Efficiently**
- ❖ Facilitates **data Isolation** to ensure that two Transactions do not update same data intent at **same time**

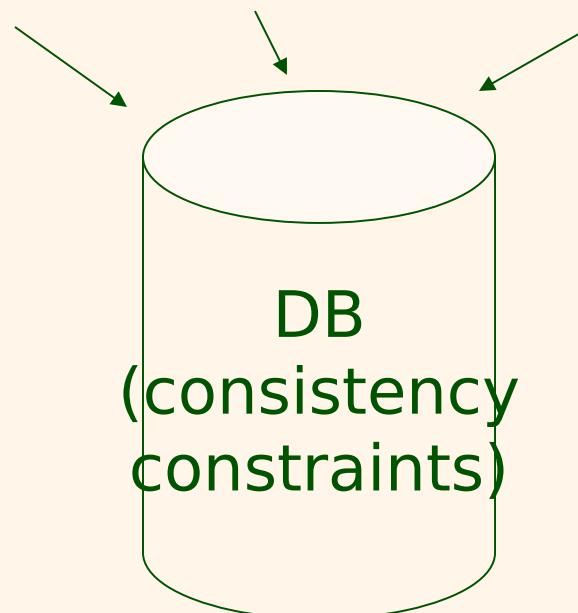
	TRANSACTIONS		RESULT
	T1	T2	
Operations	Read	Read	No conflict
	Read	Write	Conflict
	Write	Read	Conflict
	Write	Write	Conflict

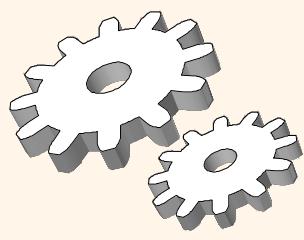


Schedules

Schedule = an interleaving of actions (**read/write** from a set of transactions T_1, T_2, \dots, T_n , where the actions of any single transaction are in the original order)

Complete Schedule T_1, T_2, \dots, T_n add Commit or Abort at the end





Complete Schedule

Transactions

T1: T2:

read(acc1)

read(acc1)
acc1 := acc1 + 20
write(acc1)
commit

read(acc1)

sum := sum + acc1

write(sum)

commit

Schedule

read1(acc1)

read2(acc1)

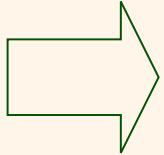
write2(acc1)

commit2

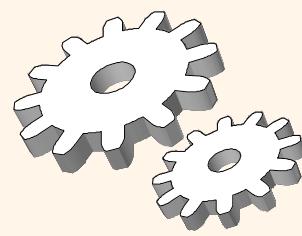
read1(acc1)

write1(sum)

commit1



Initial State of **DB** + **Schedule** → Final State of **DB**

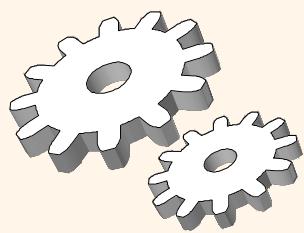


Serial Schedule

- ❖ One **Transaction** at a time, **no interleaving**

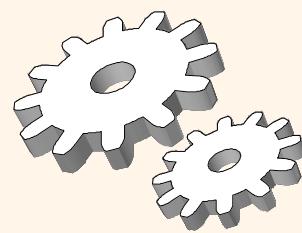
```
T1:                                T2:  
                                         read(acc1)  
                                         acc1 := acc1 + 20  
                                         write(acc1)  
                                         commit  
read(acc1)  
read(acc1)  
sum := sum + acc1  
write(sum)  
commit
```

- ❖ Final state consistent (if **Transactions** are)
- ❖ **Different Serial Schedules** give different final states



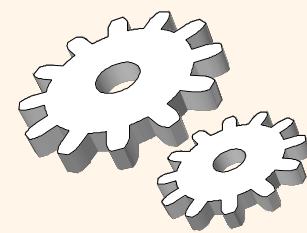
Serializable Schedule

- ❖ **Schedule** with interleaved transactions that produces the same result as **some Serial Schedule**
- ❖ "Good" **Schedules**



Checking **Serializability**

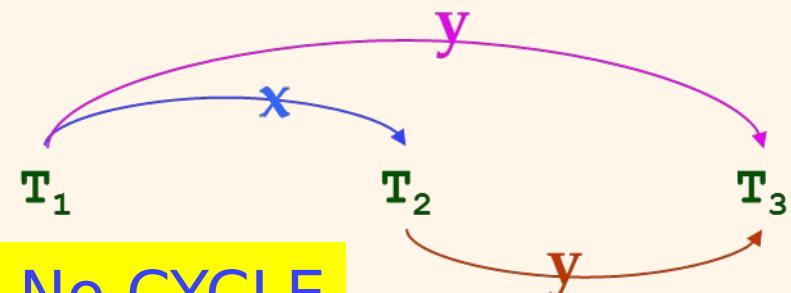
- ❖ Idea: which **actions** can be swapped in a **Schedule**?
- ❖ The following **cannot be swapped** without changing the result (conflict)
 - Actions within the **same Transaction**
 - Actions in **different Transactions** *on the same object* if at least one action is a **write** operation



Serializability Graph

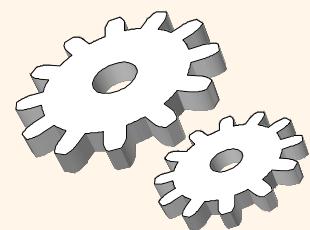
- ❖ Node for each **Transaction T_i**
- ❖ Edge from T_i to T_j if there is an action of T_i that precedes and “conflicts” with an action of T_j
- ❖ **Theorem:**
A **Schedule is Serializable** iff it's **Serializable Graph** is acyclic
(NO CYCLE!).

Acyclic?



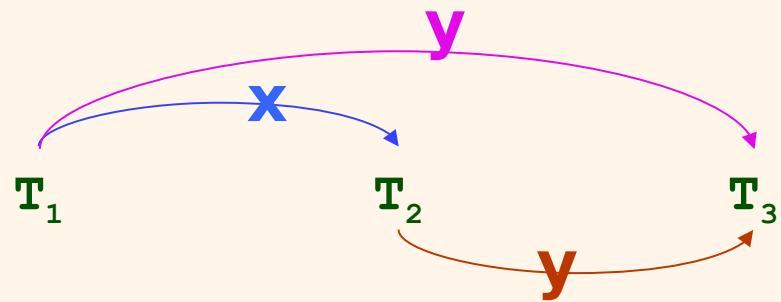
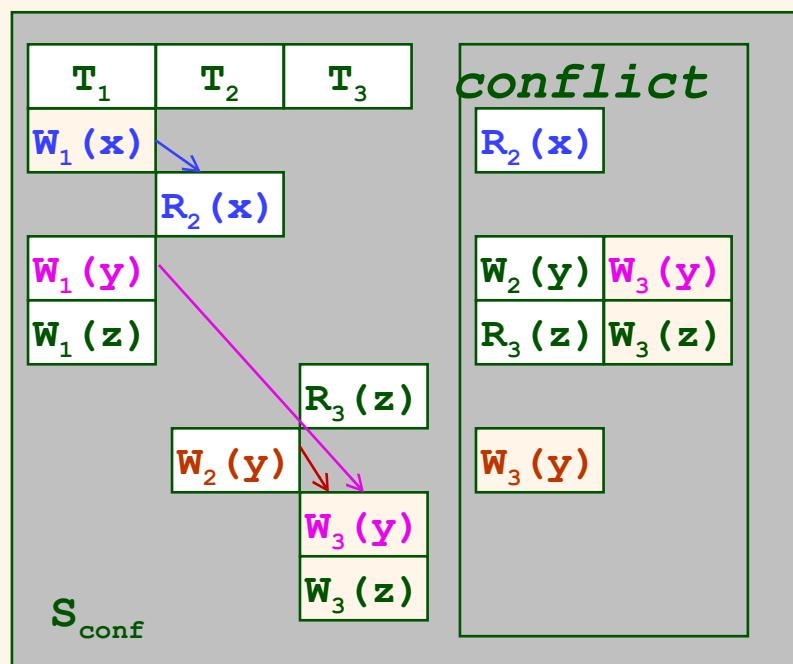
YES! No CYCLE

Example: Serializable Schedule ?



Node for each Transaction T_i

Edge from T_i to T_j if there is an action of T_i that precedes and “conflicts” with an action of T_j



Serializability Graph

ACYCLIC – i.e., NO CYCLE

SERIALIZABLE

1. Build the **Serializability Graph**.

2. Is it **acyclic**?

3. Is this schedule **Serializable**?

A **serial schedule** is a schedule in which transactions are executed one at a time. Serial schedules have no concurrent transactions. Every transaction begins, executes, and commits or rolls back before the next transaction begins. All transactions in a serial schedule are isolated.

Any schedule that is equivalent to a serial schedule is a **serializable schedule**. A serializable schedule can be transformed into a serial schedule by switching the relative order of reads in different transactions. The order of all operations within a single transaction, and reads and writes in different transactions, cannot be changed.

Serializable schedules generate the same result as the equivalent serial schedule. Therefore, concurrent transactions in a serializable schedule are isolated.

PARTICIPATION
ACTIVITY

18.2.3: Serial and serializable schedules.

ACYCLIC – i.e., NO CYCLE

SERIALIZABLE

Serial Schedule

T_1	T_2
read X $Z = X / 3$ write Z commit	
	read Y $X = Y + 2$ write X commit

Serializable Schedule

T_1	T_2
read X $Z = X / 3$ write Z commit	read Y $X = Y + 2$ write X commit

Schedule a (SERIAL)

T1

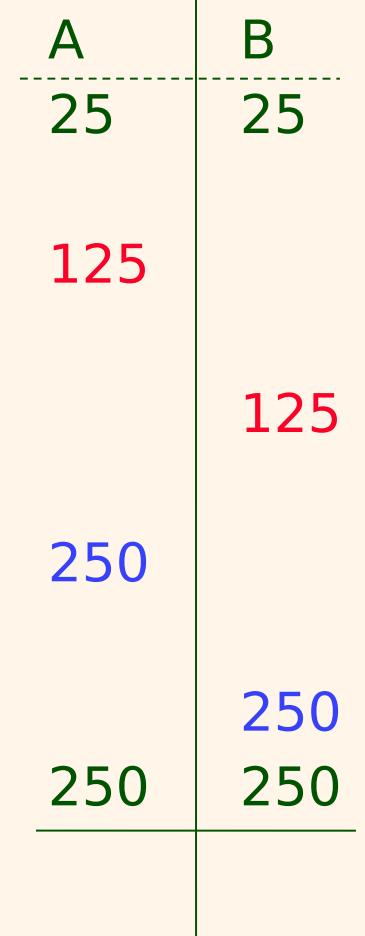
Read(A); A \leftarrow A+100
Write(A);
Read(B); B \leftarrow B+100;
Write(B);

A×2;

B×2;

T2

Read(A);A \leftarrow
Write(A);
Read(B);B \leftarrow
Write(B);



Constraint: A=B

Schedule b (SERIAL)

T1

A×2;

B×2;

Read(A); A \leftarrow A+100
Write(A);
Read(B); B \leftarrow B+100;
Write(B);

T2

Read(A); A \leftarrow

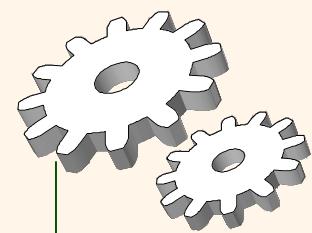
Write(A);

Read(B); B \leftarrow

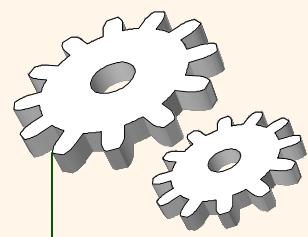
Write(B);

A	B
25	25
50	50
50	50
150	150
150	150
150	150

Constraint: A=B



Schedule c (Good or Bad)



T1

Read(A); A \leftarrow A+100

Write(A);

A \times 2;

Read(B); B \leftarrow B+100;

Write(B);

B \times 2;

T2

Read(A); A \leftarrow

Write(A);

Read(B); B \leftarrow

Write(B);

A B
25 25

125

250

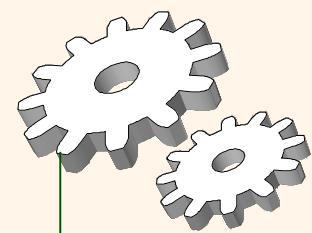
125

250

250 250

Constraint: A=B

Schedule d (Good or Bad)



T1

Read(A); A \leftarrow A+100
Write(A);

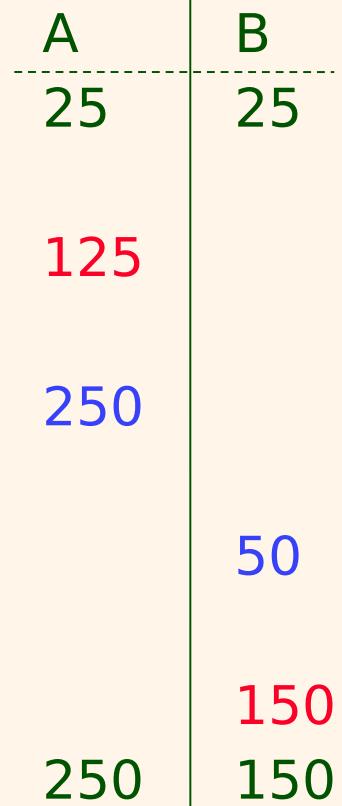
A \times 2;

B \times 2;

Read(B); B \leftarrow B+100;
Write(B);

T2

Read(A); A \leftarrow
Write(A);
Read(B); B \leftarrow
Write(B);



Constraint: A=B

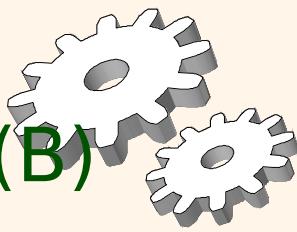
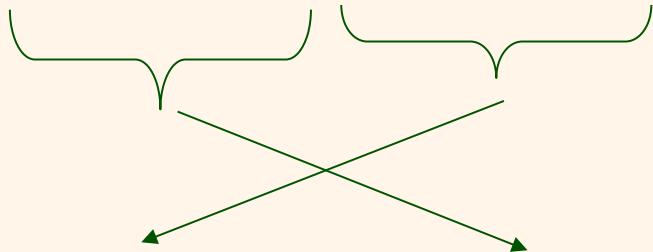
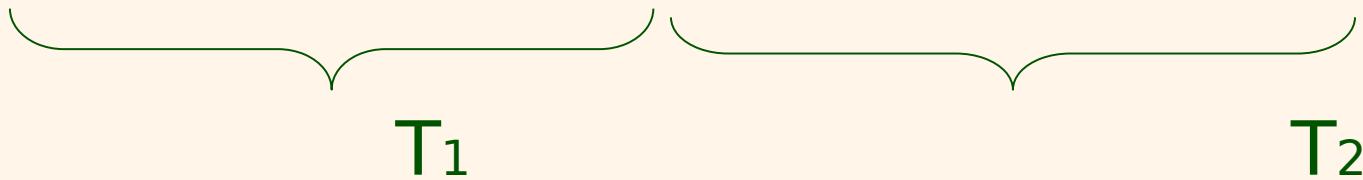


- ❖ Want **Schedules** that are “**good**”, regardless of
 - initial state and
 - Transaction semantics
- ❖ Only look at **order of** read and writes

Example:

Sc=r₁(A)w₁(A)r₂(A)w₂(A)r₁(B)w₁(B)r₂(B)w₂(B)

Any schedule that is equivalent to a serial schedule is a **serializable schedule**. A serializable schedule can be transformed into a serial schedule by switching the relative order of reads in different transactions. The order of all operations within a single transaction, and reads and writes in different transactions, cannot be changed.

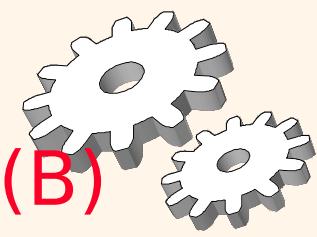

$$Sc = r_1(A)w_1(A)r_2(A)w_2(A)r_1(B)w_1(B)r_2(B)w_2(B)$$

$$Sa = r_1(A)w_1(A) \quad r_1(B)w_1(B)r_2(A)w_2(A)r_2(B)w_2(B)$$


- ❖ The following **cannot be swapped** without changing the result (conflict)
 - Actions within the **same Transaction**
 - Actions in **different Transactions** *on the same object* if at least one action is a **write** operation

Any schedule that is equivalent to a serial schedule is a **serializable schedule**. A serializable schedule can be transformed into a serial schedule by switching the relative order of reads in different transactions. The order of all operations within a single transaction, and reads and writes in different transactions, cannot be changed.

However, for **Sd**:

$$\mathbf{Sd} = r_1(A)w_1(A)r_2(A)w_2(A) \quad r_2(B)w_2(B)r_1(B)w_1(B)$$



Schedule d (Good or Bad?)

T1

T2

Bad

Read(A); $A \leftarrow A + 100$

Write(A);

Read(A); $A \leftarrow A \times 2$;

Write(A);

Read(B); $B \leftarrow B \times 2$;

Write(B);

Read(B); $B \leftarrow B + 100$;

Write(B);

A B
25 25

125

250

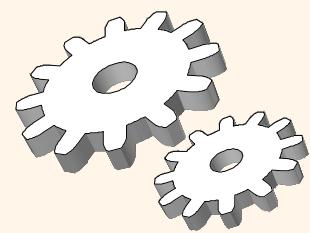
50

150

250 150

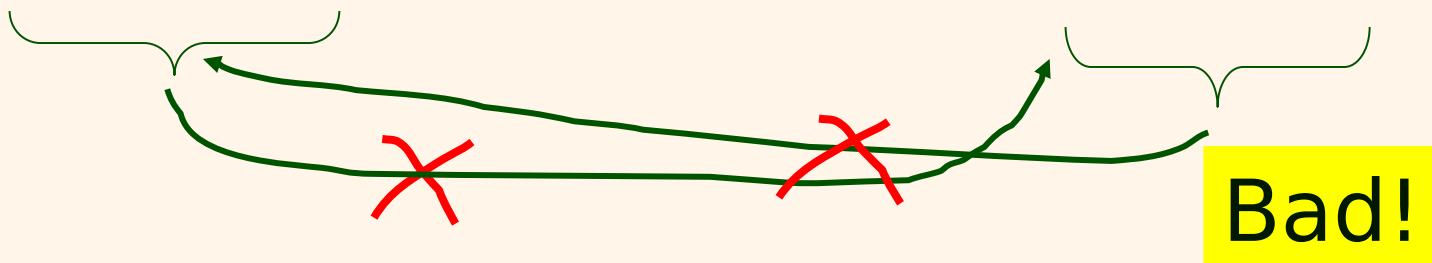
least one action is a write or

Constraint: $A=B$



However, for ~~S_d~~:

$$S_d = r_1(A)w_1(A)r_2(A)w_2(A) \quad r_2(B)w_2(B)r_1(B)w_1(B)$$



The following cannot be swapped without changing the result:

- ✖ Actions within the same Transaction
- ✖ Actions in different Transactions *on the same object if at*

Any schedule that is equivalent to a serial schedule is a **serializable schedule**. A serializable schedule can be transformed into a serial schedule by switching the relative order of reads in different transactions. The order of all operations within a single transaction, and reads and writes in different transactions, cannot be changed.

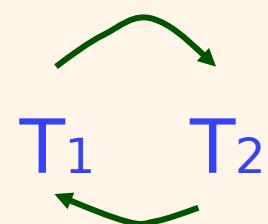
Serializability Graph

❖ Edge from T_i to T_j if there is an action of T_i that precedes and “conflicts” with an action of T_j



$S_d = r_1(A)w_1(A)r_2(A)w_2(A) \quad r_2(B)w_2(B)r_1(B)w_1(B)$

- ❖ $T_1 \rightarrow T_2$
- ❖ Also, $T_2 \rightarrow T_1$



CYCLIC – i.e., A CYCLE
NOT SERIALIZABLE

T1	T2
$r(A)$	
$w(A)$	
	$r(A)$
	$w(A)$
	$r(B)$
	$w(B)$
$r(B)$	
$w(B)$	

- S_d cannot be rearranged into a Serial Schedule
- S_d is not “equivalent” to any Serial Schedule
- S_d is “Bad”

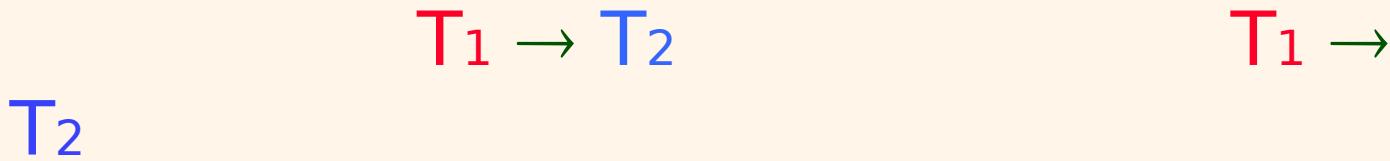
Serializability Graph

❖ Edge from T_i to T_j if there is an action of T_i that precedes and “conflicts” with an action of T_j



Returning to **Sc**

Sc= $r_1(A)w_1(A)r_2(A)w_2(A)r_1(B)w_1(B)r_2(B)w_2(B)$



T1	T2
$r(A)$	
$w(A)$	
	$r(A)$
	$w(A)$
$r(B)$	
$w(B)$	
	$r(B)$
	$w(B)$

⌚ no cycles \Rightarrow **Sc** is “equivalent” to a **Serial Schedule**

(in this case T_1, T_2)

T_1 T_2

ACYCLIC – i.e., NO CYCLE

SERIALIZABLE

Question

Draw the schedules in a vertical timeline manner.

Draw the Serializability Graph for the schedule S below:

w1(x) r1(x) w3(z) r2(x) r2(y) r4(z) r3(x) w3(y)

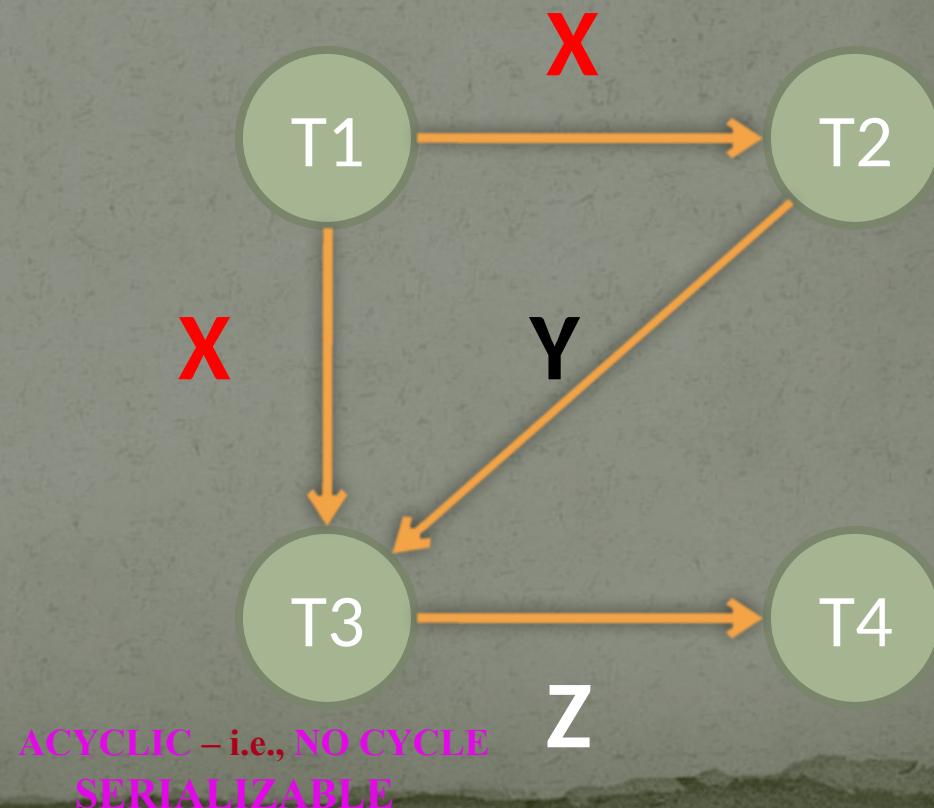
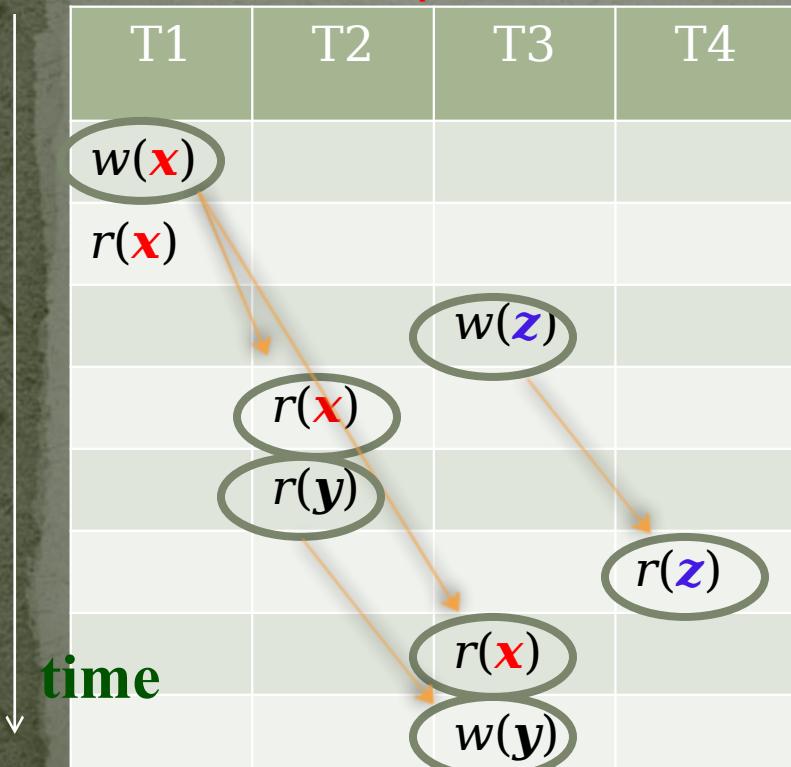
- ❖ Node for each Transaction T_i
- ❖ Edge from T_i to T_j if there is an action of T_i that precedes and "conflicts" with an action of T_j
- ❖ Theorem: A Schedule is Serializable iff its Serializability Graph is acyclic (NO CYCLE!).

If schedule S is serializable give an equivalent serial schedule, if not explain why.

Transactions: T1, T2, T3, T4

Directed edges: (T1, T2), (T2, T3), (T1, T3), (T1, T4)

S is Serializable with equivalent serial schedule: T1, T2, T3, T4



Match the schedule type to the example schedule.

Serial schedule

Serializable schedule

Non-serializable schedule

Serial schedule

T₁ T₂
----- -----

read X

Y = X + 4

write Y

commit

read X

X = X / 8

write X

commit

T₁ ends before T₂ begins. Schedules that contain no concurrent transactions are serial.

Correct

- ❖ Edge from T_i to T_j if there is an action of T_i that precedes and "conflicts" with an action of T_j

Match the schedule type to the example schedule.

Serial schedule

Serializable schedule

Non-serializable schedule

Serializable schedule

T_1	T_2
-----	-----
read X	
$Y = X + 4$	
write Y	
commit	

read X
$X = X / 8$
write X

commit

Correct

'write Y' in T_1 does not conflict with any T_2 operations. Therefore 'write Y' and 'commit' in T_1 can move prior to all T_2 operations, resulting in an equivalent serial schedule.

- ❖ Edge from T_i to T_j if there is an action of T_i that precedes and "conflicts" with an action of T_j

Match the schedule type to the example schedule.

Serial schedule

Serializable schedule

Non-serializable schedule

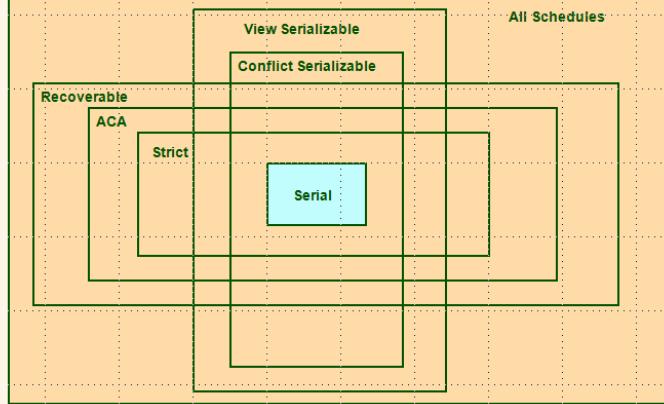
Non-serializable schedule

T_1	T_2
-----	-----
read X	
$X = X + 4$	
	read X
	$X = X / 8$
	write X
write X	
commit	
	commit

'write X' in T_2 conflicts with both 'read X' and 'write X' in T_1 , and therefore cannot change relative order with either. Since a T_2 operation must occur between two T_1 operations, the schedule cannot be serialized.

Correct

Conflicts



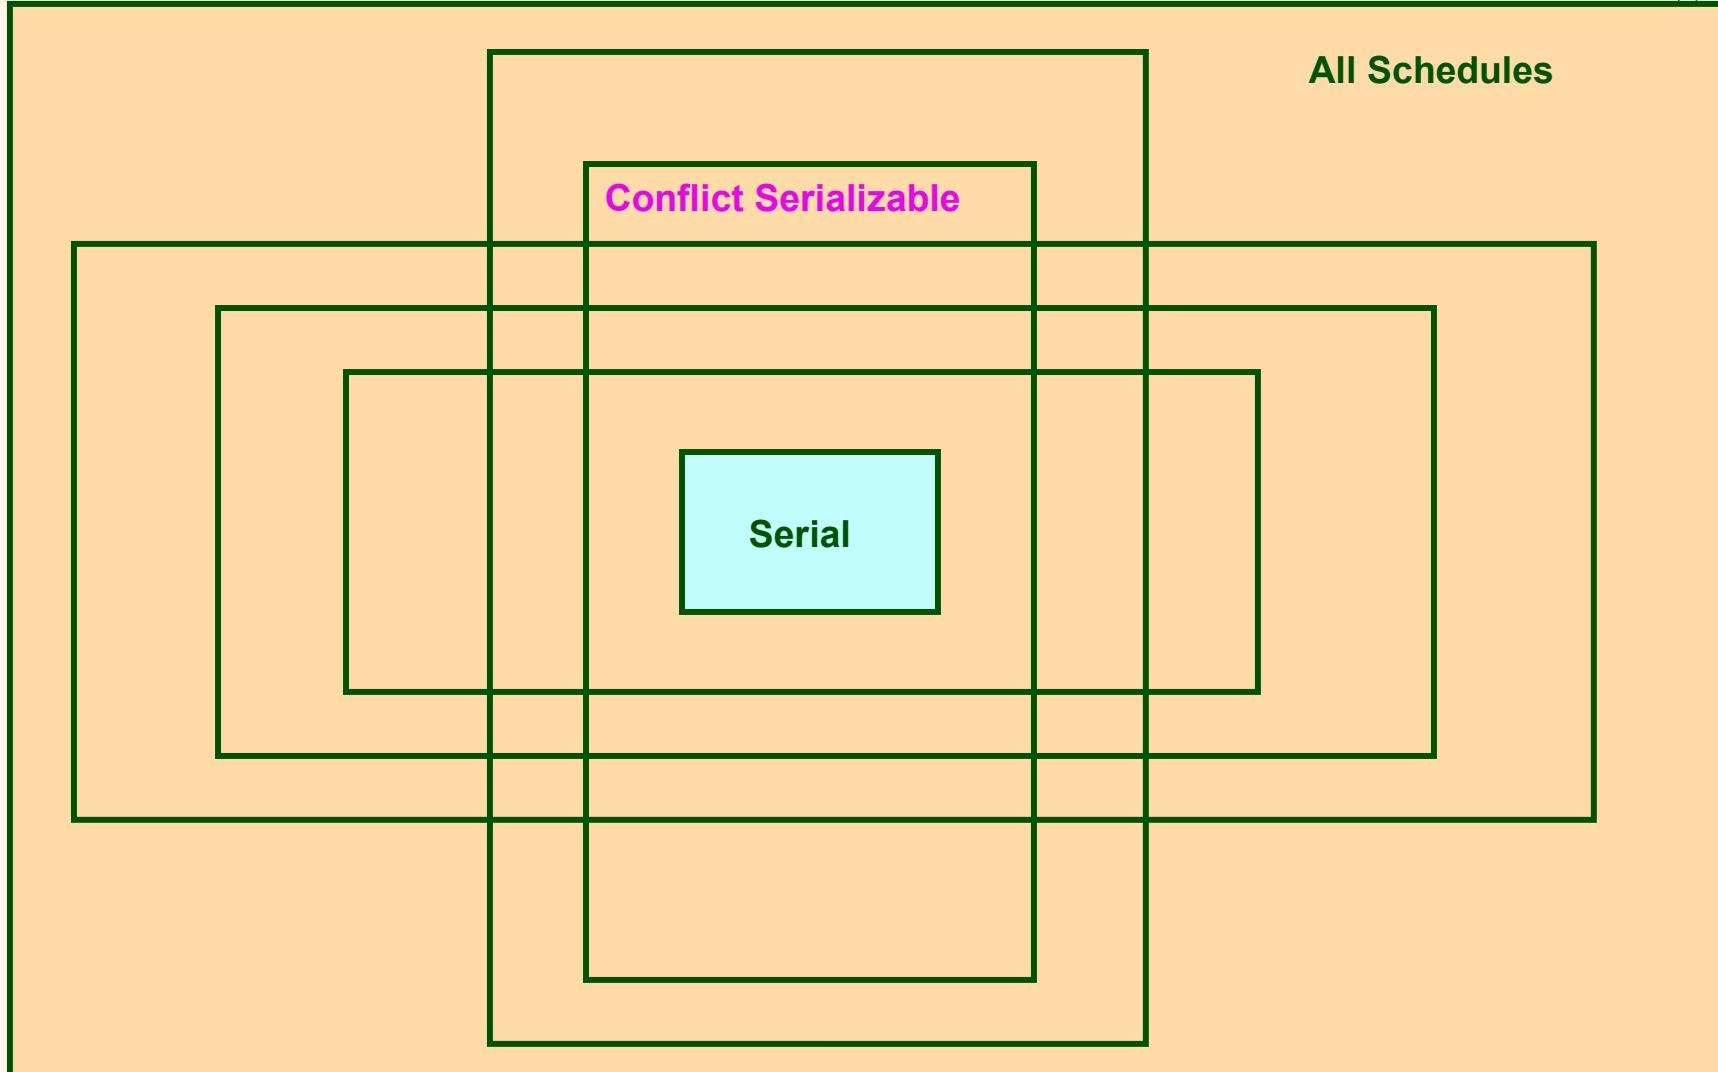
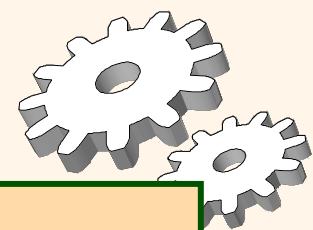
Conflicting **actions**: pairs of **actions** on **same object** from **different Transactions** where at **least one is write**

Two Schedules are ***conflict-equivalent*** if they have the same conflicts

A **Schedule is Conflict Serializable** if it is ***conflict-equivalent*** to a **Serial Schedule**

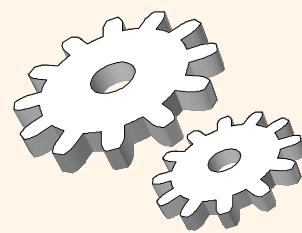
Every **Conflict Serializable Schedule** is

Venn Diagram for *Schedules*



A **Schedule is Conflict Serializable** if it is conflict-equivalent to a **Serial Schedule**

Example - Conflict

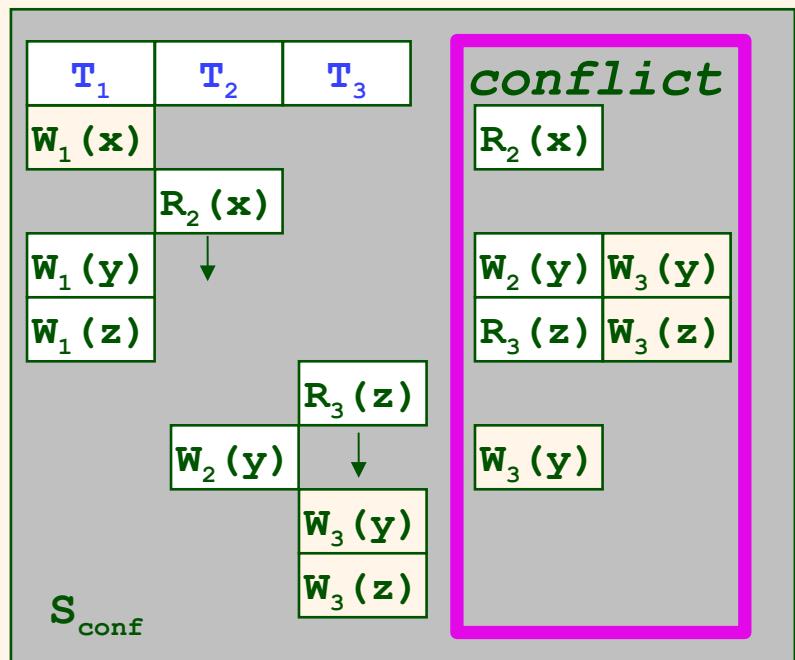


Serializable

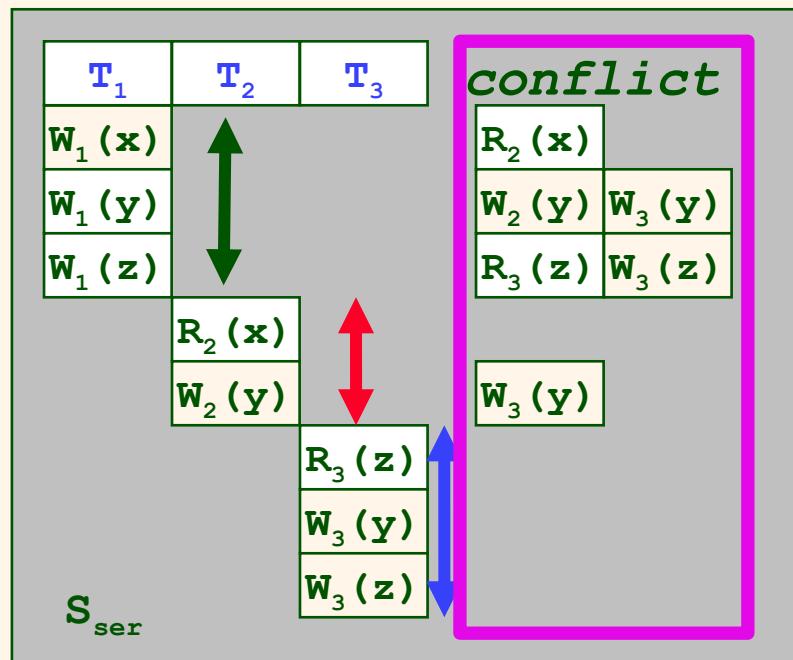
How many objects? 3: x, y, and z

How many transactions? 3: T₁, T₂, and T₃

same conflicts, thus conflict equivalent

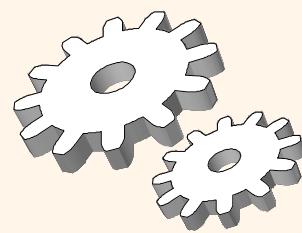


Conflict Serializable



Serial schedule

A schedule is *conflict-serializable* if it is conflict-equivalent to a *serial schedule*



Dependency Graph

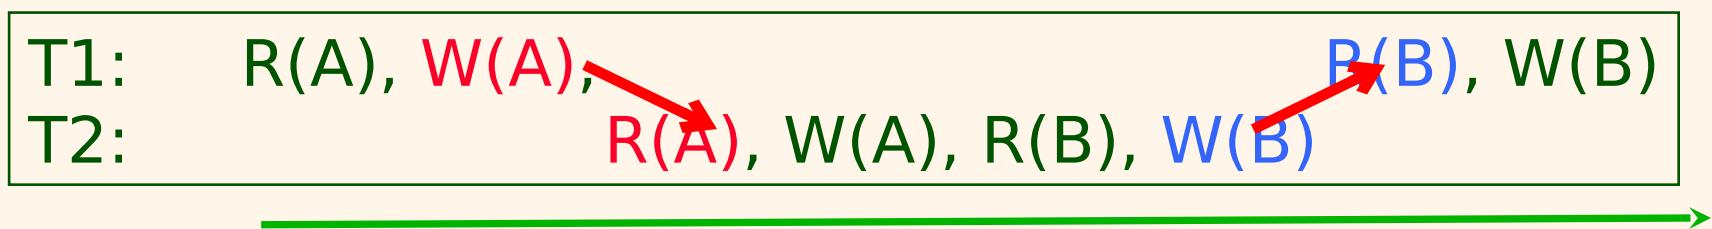
- ❖ **Dependency graph**: One node per **Xact**; edge from ***Ti*** to ***Tj*** if ***Tj* reads/writes** an object last **written** by ***Ti***.
- ❖ **Theorem**: Schedule is **Conflict Serializable** if and only if its **Dependency Graph** is **acyclic**



Example

Dependency Graph: One node per **Xact**; edge from T_i to T_j if T_j **reads/writes** an object last **written** by T_i .

- ❖ A **Schedule** that is **not Conflict Serializable**:



Dependency Graph

CYCLIC – i.e., A CYCLE

NOT Conflict Serializable

- ❖ The **cycle** in the **Graph** reveals the problem.
- ❖ The output of T_1 depends on T_2 , and vice-

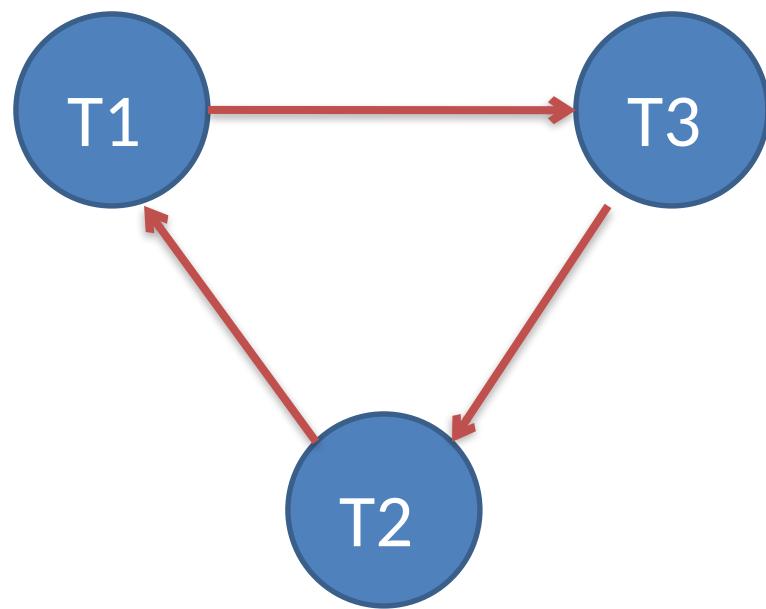
Draw the Precedence/Dependency graph for this schedule.

Node for each transaction T_i

Edge from T_i to T_j if there is an

action of T_i that precedes and “conflicts” with an action of T_j

	T1	T2	T3
0	start		
1	read X		
2		start	
3		read Y	
4	write X		
5			
6		start	
7		read X	
8		write X	
9			commit
10		read X	
11		write Y	
12		write X	
13		commit	
14	read Y		
15	write Y		
	commit		



CYCLIC – i.e., A CYCLE

NOT Conflict Serializable

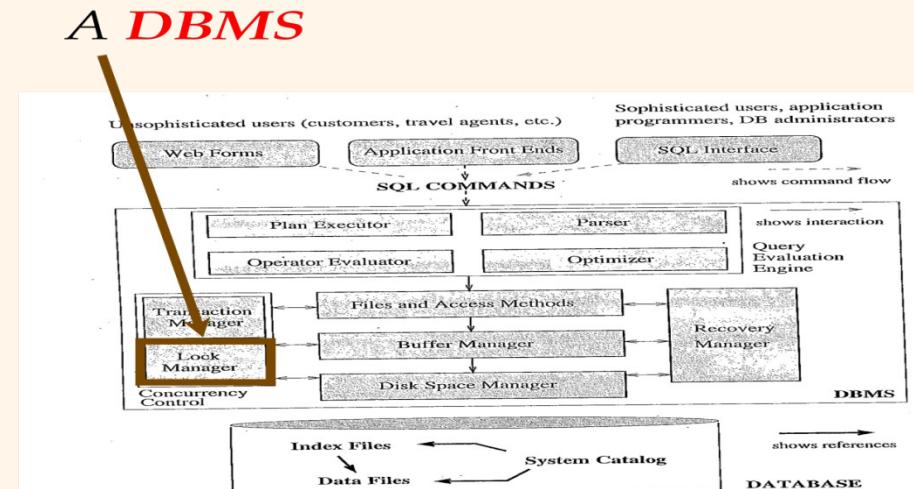
Concurrency Control with Locking Methods

Lock

- Guarantees exclusive use of a **data** item to a current Transaction
- Required to prevent another Transaction from **reading** inconsistent **data**

Lock Manager

- Responsible for assigning and policing the **Locks** used by Transactions



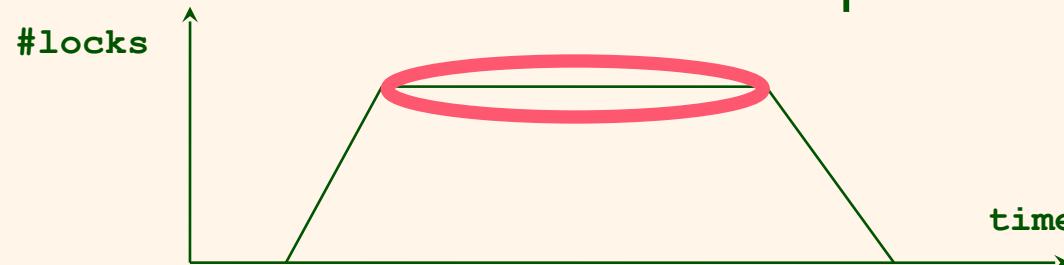
TWO-PHASE LOCKING TO ENSURE

Serializability

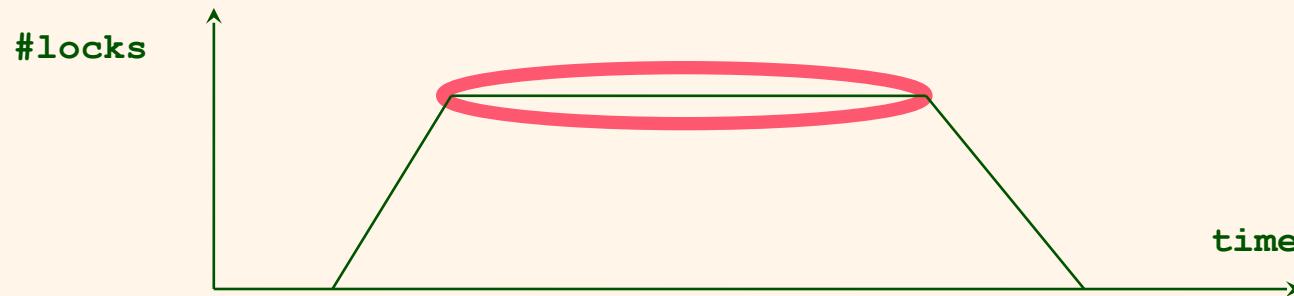
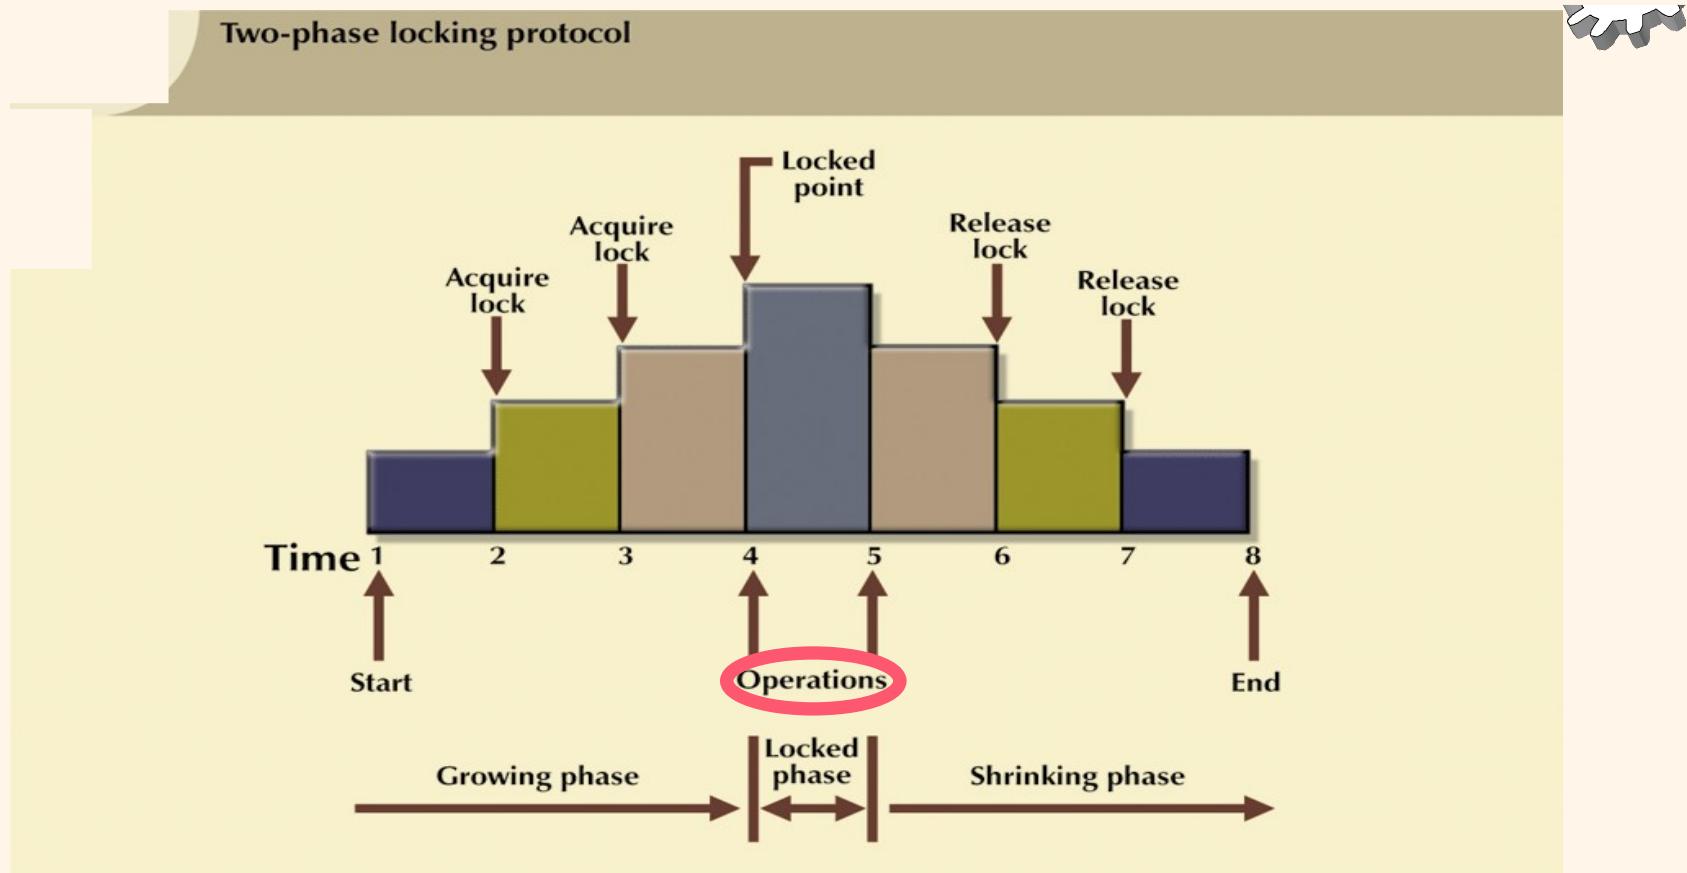


❖ Governed by the following rules:

- Each Xact must obtain a **S (*shared*) Lock** on **object** before **Reading**, and an **X (*exclusive*) Lock** on **object** before **Writing**.
- Two **Transactions** cannot have conflicting **Locks**
- No **unlock** operation can precede a lock operation in the same Transaction
- No **data** are affected until all **Locks** are obtained—that is, until **Transaction** is in its locked point



Serializability

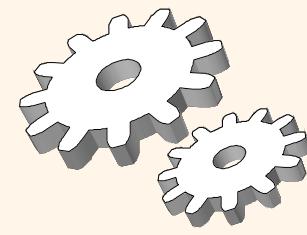




Deadlocks

- ❖ Condition that occurs when two Transactions wait for each other to **unlock data**
- ❖ Possible only if one of the Transactions wants to obtain an **exclusive Lock (X)** on a **data item**
- No **Deadlock** condition can exist among **shared Locks(S)**

Deadlocks



How a Deadlock Condition Is Created

TIME	TRANSACTION	REPLY	LOCK STATUS	
0			Data X	Data Y
1	T1:LOCK(X)	OK	Unlocked	Unlocked
2	T2:LOCK(Y)	OK	Locked	Unlocked
3	T1:LOCK(Y)	WAIT	Locked	Locked
4	T2:LOCK(X)	WAIT	Locked	Locked
5	T1:LOCK(Y)	WAIT	Locked	Locked
6	T2:LOCK(X)	WAIT	Locked	Locked
7	T1:LOCK(Y)	WAIT	Locked	Locked
8	T2:LOCK(X)	WAIT	Locked	Locked
9	T1:LOCK(Y)	WAIT	Locked	Locked
...
...
...
...



LOCK-Based Concurrency



Control

Strict Two-phase Locking (Strict 2PL) Pr

- Each Xact must obtain a **S (shared) Lock** on **object** before **Reading**, and an **X (exclusive) Lock** on **object** before **Writing**.
- All **Locks** held by a **Transaction** are released when the **Transaction** completes
- If an **Xact** holds an **X Lock** on an **object**, no other **Xact** can get a **Lock** (**S** or **X**) on that **object**.
- ❖ Strict 2PL allows only **Serializable Schedules**.



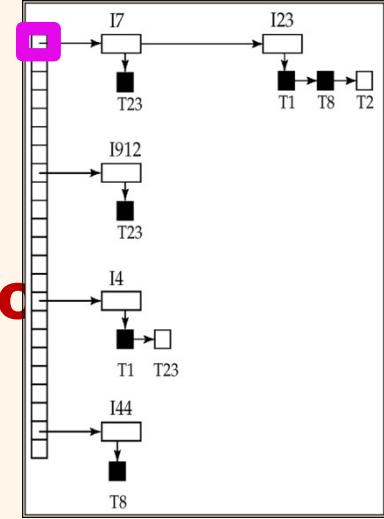
Lock Management



Lock and **Unlock** requests handled by the **Lock Manager**

Lock Table entry:

- Number of transactions currently holding a **Lock**
- Type of **Lock** held (Shared or eXclusive)
- Pointer to queue of **Lock** requests



Locking and **Unlocking** have to be atomic operations

Lock upgrade: transaction that holds shared **lock**

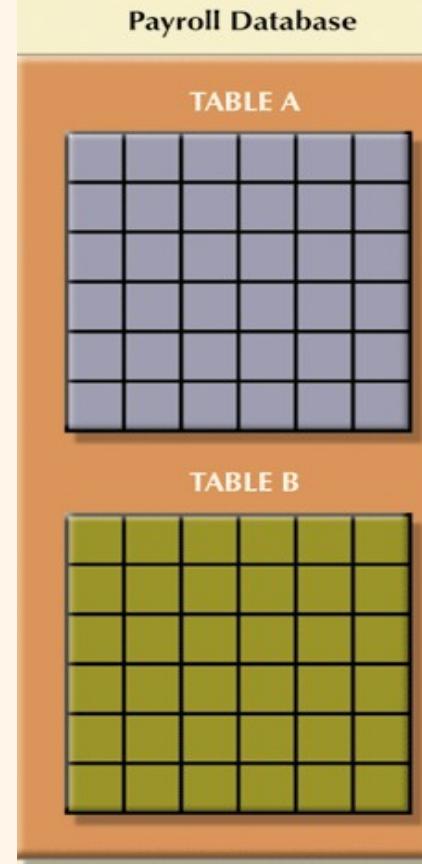
The concurrency system monitors active transactions, determines when locks are needed, and issues requests to grant and release locks.

Requests are sent to the **lock manager**, a component of the concurrency system that tracks, grants, and releases locks.

By requesting shared and exclusive locks as needed, the concurrency system implements the isolation level specified for each transaction.

By minimizing lock scope and duration, the concurrency system reduces the duration of transactions that are waiting for locked data.

Lock Granularity



❖ DATABASE-level Lock

- Entire **DATABASE** is locked

❖ TABLE-level Lock

- Entire **TABLE** is locked

❖ PAGE-level Lock

- Entire disk **PAGE** is locked

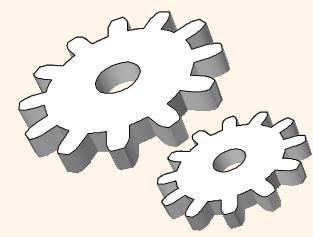
❖ Row-level Lock

- Allows concurrent **Transactions** to access different **Rows** of same **TABLE**, even if Rows are located on same **PAGE**

❖ Field-level Lock

- Allows concurrent **Transactions** to access same **Row**, as long as they require use of different **Fields**

Deadlock Detection



Create a **Waits-For Graph (WFG)**:

- Nodes are Transactions
- There is an edge from T_i to T_j if T_i is waiting for T_j to release a lock

Periodically check for **CYCLES** in the **Waits-For Graph**



How a Deadlock Condition Is Created				
TIME	TRANSACTION	REPLY	LOCK STATUS	
0			Data X	Data Y
1	T1:LOCK(X)	OK	Unlocked	Unlocked
2	T2:LOCK(Y)	OK	Locked	Unlocked
3	T1:LOCK(Y)	WAIT	Locked	Locked
4	T2:LOCK(X)	WAIT	Locked	Locked
5	T1:LOCK(Y)	WAIT	Locked	Locked
6	T2:LOCK(X)	WAIT	Locked	Locked
7	T1:LOCK(Y)	WAIT	Locked	Locked
8	T2:LOCK(X)	WAIT	Locked	Locked
9	T1:LOCK(Y)	WAIT	Locked	Locked
...
...
...
...

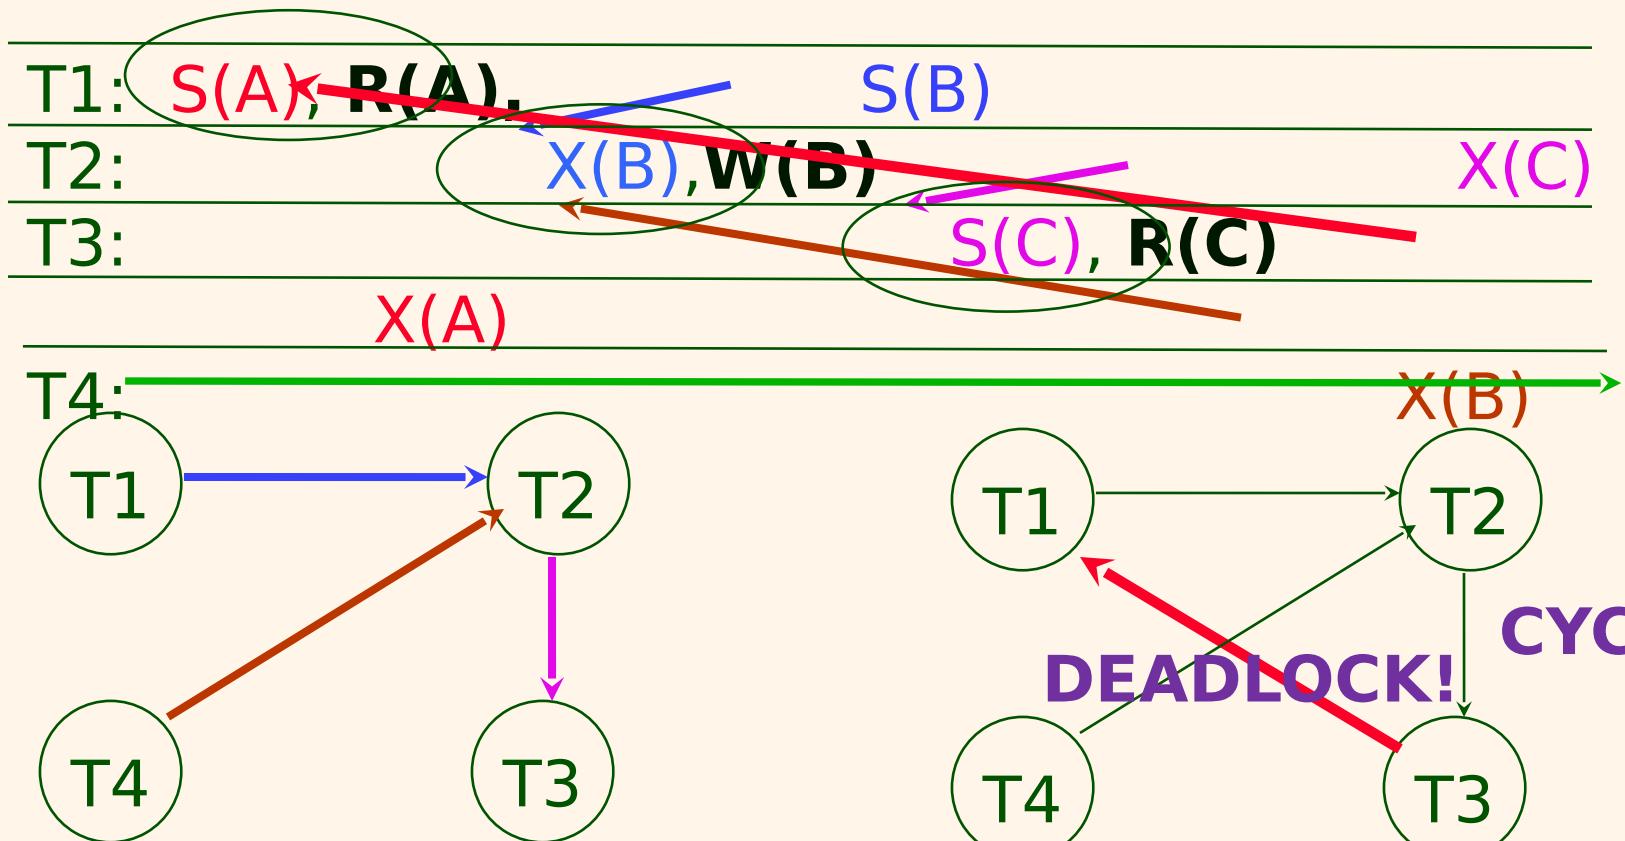


Deadlock Detection

There is an edge from T_i to T_j if
 T_i is waiting for T_j to release
a lock

T_1 is waiting on T_2 for B
 T_2 is waiting on T_3 for C
 T_4 is waiting on T_2 for B
 T_3 is waiting on T_1 for A

Example:



7. Is there in cycle in this Wait-For Graph? Is there a DEADLOCK?

In a **deadlock** detection, transactions are allowed to wait and they are not aborted until a deadlock has been detected. (Compared to other schemas where aborts transaction may have been aborted prematurely.)

Schedule S1: T1:R(X), T2:W(X), T3:W(Y), T1:W(Y), T1:Commit, T2:Commit, T3:Commit

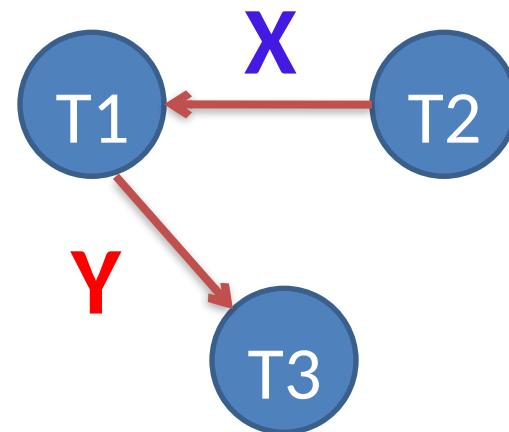
There is an **edge from T_i to T_j if T_i is waiting for T_j to release a lock**

T1	S(x)Rx			X(y)Wy	commit		
T2		X(x)Wx			commit		
T3			X(y)Wy			commit	

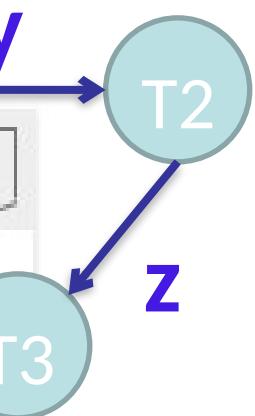
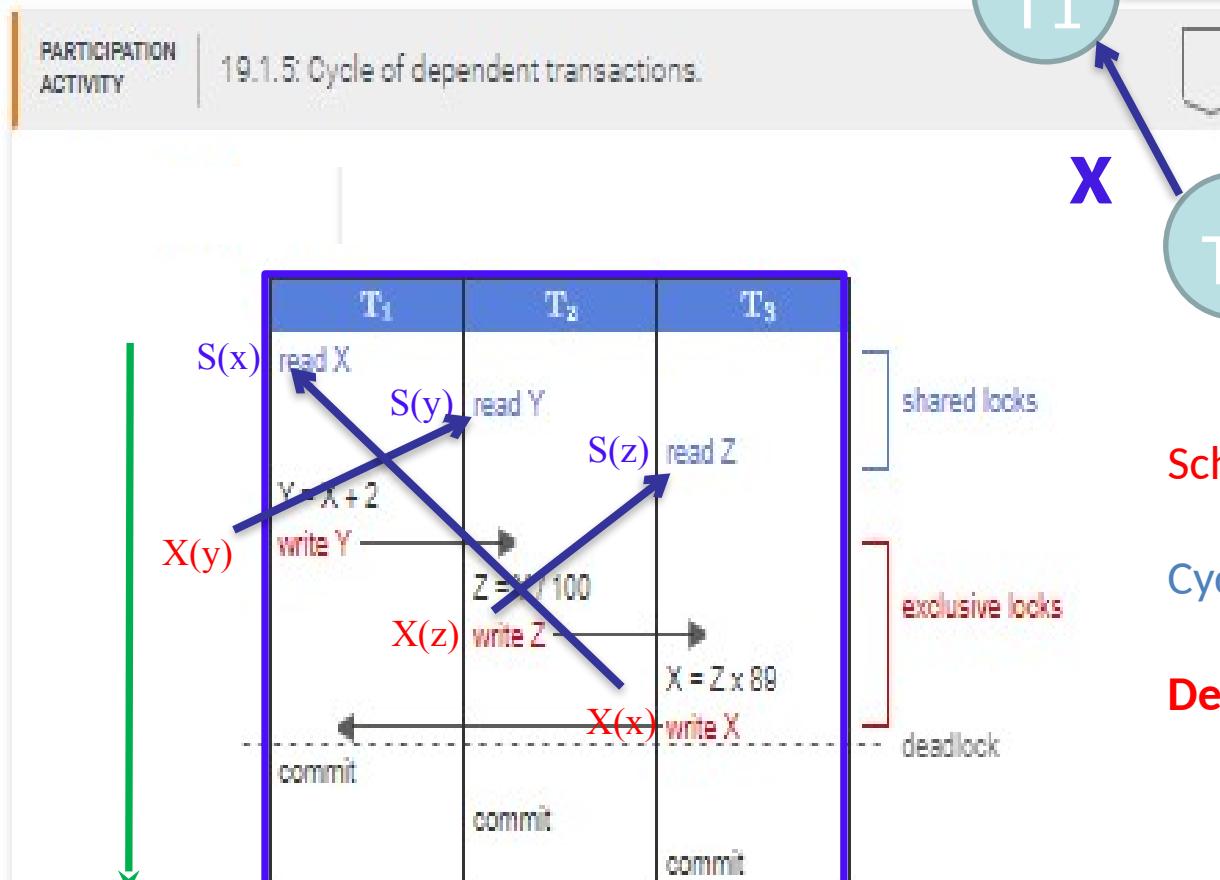
Schedule S1:

Acyclic

Not a **Deadlock**



- **Cycle detection.** The concurrency system periodically checks for cycles of dependent transactions. When a cycle is detected, the concurrency system selects and rolls back the 'cheapest' transaction. The cheapest transaction might, for example, have the fewest rows locked or most recent start time. The rollback breaks the deadlock.
- There is an edge from T_i to T_j if T_i is waiting for T_j to release a lock



Schedule:
Cycle
Deadlock

Start

T ₁	T ₂
<pre>SELECT Continent FROM Country WHERE CountryID = 266;</pre>	<pre>UPDATE Country SET IndepYear = 2003 WHERE CountryID = 266;</pre>

If lock scope is a single row, and T₁ takes an exclusive lock on CountryID 266, what happens to T₂?

Pick

If lock scope is a block, and T₁ takes an exclusive lock on CountryID 266, select the row(s) that are locked.

<input type="checkbox"/>	266	Gabon	Libreville	1960	Africa
<input type="checkbox"/>	430	Liberia	Monrovia	1847	Africa
<input type="checkbox"/>	368	Iraq	Baghdad	1932	Asia
<input type="checkbox"/>	108	Burundi	Gitega	1962	Africa
<input type="checkbox"/>	616	Poland	Warsaw	1918	Europe
<input type="checkbox"/>	752	Sweden	Stockholm	1523	Europe
<input type="checkbox"/>	296	Kiribati	Tarawa	1979	Oceania
<input type="checkbox"/>	496	Mongolia	Ulaanbaatar	1911	Asia
<input type="checkbox"/>	360	Indonesia	Jakarta	1945	Asia

1

2

3

4

TA time (Alvaro)
(CA 19.1.1 Step 1 – Concurrency)

TA time (Alvaro)
(CA 19.1.1 Step 2 – Concurrency)



01:07:42

Take control Pop out Chat 110 People Raise View Rooms Apps More Camera Mic Share Leave



AU



TN

Urtaza, Alvaro A

Nguyen, Tri...

View all

**Participants**

Invite someone or dial a number

Share invite

In this meeting (110)

Mute all

 Hilford, Victoria
Organizer

RA Adhikari, Rohit

MA Ahmed, Mohamed A

AA Akram, Ali

SA Alvarez, Stephanie

OA Anayor-Achu, Ogochukwu E

HA Avci, Hatice Kubra

RA Aysola, Riya

AB Bahl, Anish

SB Banza, Sean Paolo B

HB Bui, Hieu

Burger, Jake

VC Carrillo-Zepeda, Victor E

DC Carter, Deric

☰ zyBooks My library > COSC 3380: Database Systems home > 19.1: Concurrency

zyBooks catalog

Help/FAQ

Alvaro Urtaza

<input checked="" type="checkbox"/>	499	Montenegro	Podgorica	2006	Europe
<input type="checkbox"/>	266	Gabon	Libreville	1960	Africa
<input type="checkbox"/>	44	Bahamas	Nassau	1973	NAmerica
<input type="checkbox"/>	484	Mexico	Mexico City	1810	NAmerica
<input type="checkbox"/>	4	Afghanistan	Kabul	1919	Asia
<input type="checkbox"/>	308	Grenada	St. George's	1974	NAmerica
<input type="checkbox"/>	798	Tuvalu	Funafuti	1978	Oceania

1

2

3

4

Check

Next

✓ Expected:

- (a) The lock allows T₂ to read the row with CountryID 499.
690, 860, 499

Since T₁ and T₂ access different rows, T₂ is not affected by T₁'s exclusive lock on the row with CountryID 690. Thus, the lock allows T₂ to read the row with CountryID 499.

T₁ takes an exclusive lock on the block containing the row with CountryID 690. Thus, all rows in the entire block are locked: 690, 860, and 499.

View solution ▾ (Instructors only)

Feedback?



01:09:53

Take control Pop out Chat 110 People Raise View Rooms Apps More Camera Mic Share Leave

**Participants**

Invite someone or dial a number

Share invite

In this meeting (110)

Mute all

 Hilford, Victoria
Organizer

RA Adhikari, Rohit

MA Ahmed, Mohamed A

AA Akram, Ali

SA Alvarez, Stephanie

OA Anayor-Achu, Ogochukwu E

HA Avci, Hatice Kubra

RA Aysola, Riya

AB Bahl, Anish

SB Banza, Sean Paolo B

HB Bui, Hieu

BC Burger, Jake

VC Carrillo-Zepeda, Victor E

DC Carter, Deric

≡ zyBooks My library > COSC 3380: Database Systems home > 19.1: Concurrency

Consider the following transaction schedule:

T ₁	T ₂
read W	read X
X = 79 - W write X	W = X + 82 write W
commit	commit



Does a deadlock occur when using aggressive locking?

 No
 Yes

cycle detection?

 Yes
 No

1 2 3 4

Check **Next**

Expected: No, Yes

When the database uses aggressive locking:

Each transaction requests all locks when the transaction starts, so the transactions cannot participate in a deadlock.

When the database uses cycle detection:

Transactions first take shared locks when reading, then request exclusive locks when writing.

T₁ reads and creates a shared lock on W. T₂ reads and creates a shared lock on X.T₂ waits for T₁ to release the shared lock on W. T₁ waits for T₂ to release the shared lock on X. Deadlock occurs.The concurrency system detects the cycle between T₁ and T₂, rolls back the 'cheapest' transaction, and breaks the deadlock.View solution (Instructors only)

zyBooks catalog help/Forgot Alvaro Urtaza

View all

Feedback

Concurrency Control Practice Questions

A database administrator is updating one of the tables in the database. Which of the following techniques can prevent other database administrators from doing concurrent transactions to the same table?

- a. Mirroring
- b. Scoping
- c. Logging
- d. Locking

A ____ is a part of the concurrency system that monitors, grants, and releases locks.

a. Lock Administrator

b. Lock System

c. Lock Manager

d. Lock Optimizer

In a ___, all transactions come to a halt and remain at a standstill until one of the transactions is aborted.

a. timeout

b. deadlock

c. two-phase locking

d. locking

Which deadlock management technique automatically rolls back a transaction when a lock is not released in a fixed period of time?

- a. Timeout
- b. Aggressive locking
- c. Data ordering
- d. Cycle detection

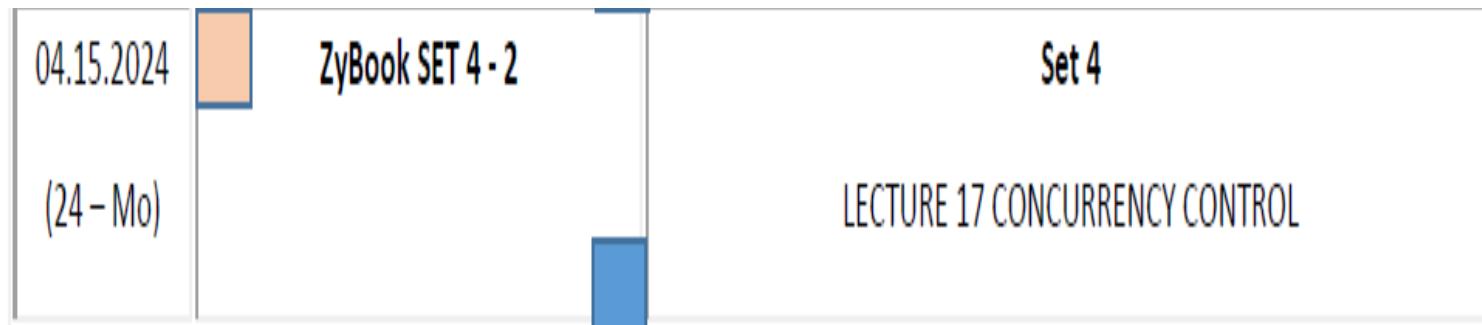
Refer to the table and schedule below. The UPDATE statement in transaction T1 holds an exclusive lock on the Class table. With strict two-phase locking, when does the SELECT statement in transaction T2 execute?

T ₁	T ₂
UPDATE Class SET TeacherID = 32412 X(Class) WHERE ClassID = 80; ROLLBACK;	SELECT CourseTitle FROM Class WHERE ClassID = 80;

- Strict Two-phase Locking (Strict 2PL) Protocol:
- Each Xact must obtain a S (shared) Lock on object before Reading, and an X (exclusive) Lock on object before Writing.
 - All Locks held by a Transaction are released when the Transaction completes
 - If an Xact holds an X Lock on an object, no other Xact can get a Lock (S or X) on that object.
- Strict 2PL allows only Serializable Schedules.
- 
- 63

- Immediately after the UPDATE statement in T1 executes
- After the ROLLBACK statement in T1 executes
- At the same time as the UPDATE statement in T1 executes
- The SELECT statement in T2 will never execute

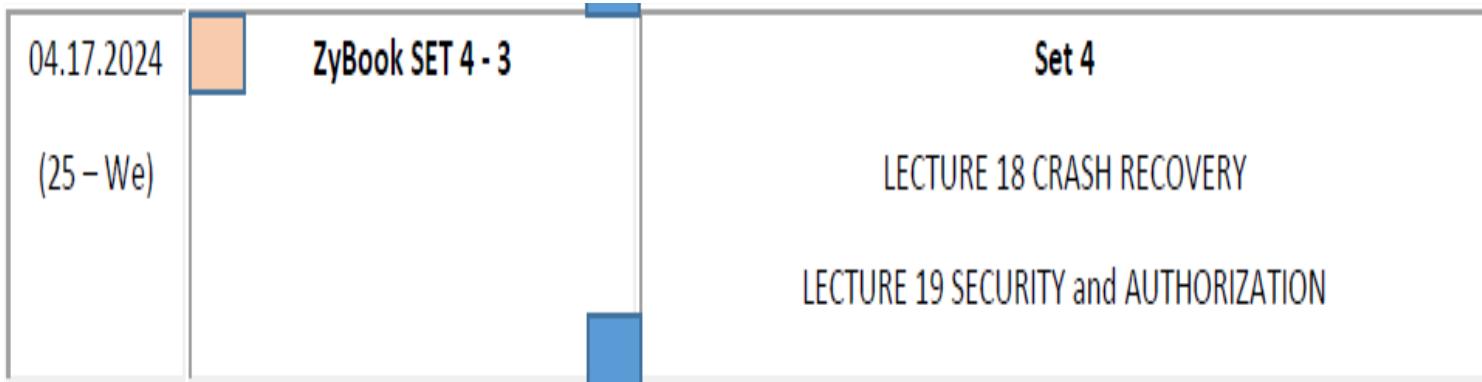
At 5:00 PM .



-
- 17. SET 4 Empty 
 - 19. SET 4 - 2:CONCURRENCY CONTROL Hidden  0%  0% 
-

VH work on
SET 4 – 2: Concurrency

Next



17. SET 4

Empty



20. SET 4 - 3:CRASH RECOVERY

Hidden



0%



0%



0%



21. SET 4 - 4:SECURITY and AUTHORIZATION

Hidden



0%



0%



VH, unHIDE

From 5:05 to 5:15 PM – 5 minutes.

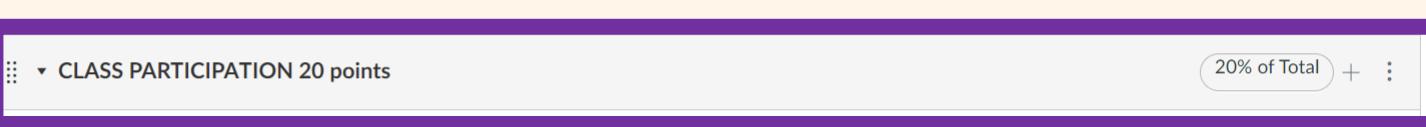


04.15.2024
(24 - Mo)

ZyBook SET 4 - 2

Set 4

LECTURE 17 CONCURRENCY CONTROL



CLASS PARTICIPATION 20 points

20% of Total + :

CONCURRENCY



Class 24 END PARTICIPATION

Not available until Apr 15 at 5:05pm | Due Apr 15 at 5:15pm

VH, publish

This is a synchronous online class.

Attendance is required.

Recording or distribution of class materials is prohibited.

1. At the beginning of selected classes there is an assessment in the first 10 minutes. (beige BOX in the Detailed Syllabus)
2. At the end of selected classes there is an assessment in the last 10 minutes. (blue BOX in the Detailed Syllabus)
3. ZyBook sections will be downloaded and used for 30% of Total Score on the dates specified in the Detailed Syllabus.
4. EXAMS are in CANVAS. No late EXAMS.
5. I have to be present in TEAMS in order to take any graded assignment assigned during that class.

At 5:15 PM.

End Class 24

VH, unHIDE ZyBook SET 4 – 3,4 : RECOVERY & SECURITY.

<input type="checkbox"/> 20. SET 4 - 3:CRASH RECOVERY	Hidden	 0%	 0%	 0%	
<input type="checkbox"/> 21. SET 4 - 4:SECURITY and AUTHORIZATION	Hidden	 0%	 0%		

VH, Download Attendance Report
Rename it:
4.15.2024 Attendance Report FINAL TEAMS

CLASS PARTICIPATION 20 points

20% of Total + :

 Class 25 BEGIN PARTICIPATION Not available until Apr 17 at 4:00pm Due Apr 17 at 4:07pm	VH, publish
 Class 25 END PARTICIPATION Not available until Apr 17 at 5:05pm Due Apr 17 at 5:15pm	VH, publish

VH, upload Class 24 to CANVAS.