



UNIVERSITY of **HOUSTON**

DEPARTMENT OF COMPUTER SCIENCE

**COSC 3380 Spring 2024**

**Database Systems**

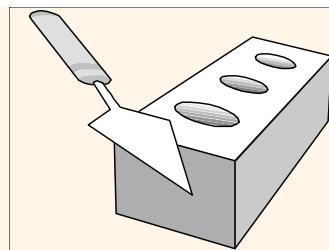
**M & W 4:00 to 5:30 PM**

**Prof. Victoria Hilford**

**PLEASE TURN your webcam ON (must have)**

**NO CHATTING during LECTURE**

**VH, UNhide Section 20, 21**



**COSC 3380**

**4 to 5:30**

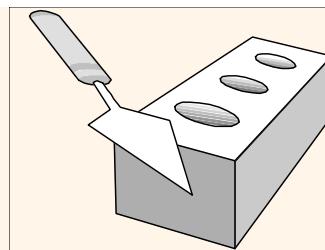
**PLEASE**

**LOG IN**

**CANVAS**

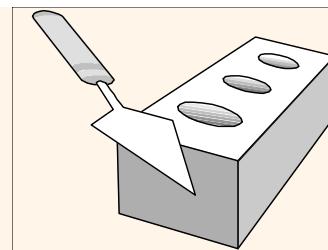
**Please close all other windows.**

# COSC 3380



04.17.2024 (25 – We)	ZyBook SET 4 - 3	Set 4 LECTURE 18 CRASH RECOVERY LECTURE 19 SECURITY and AUTHORIZATION
04.22.2024 (26 – Mo)		EXAM 4 Practice (PART of 20 points)
04.24.2024 (27 – We)	TA Download ZyBook SET 4 Sections (4 PM) (PART of 30 points)	EXAM 4 Review (PART of 20 points)
04.29.2024 (28 – Mo)		EXAM 4 (PART of 50 points)

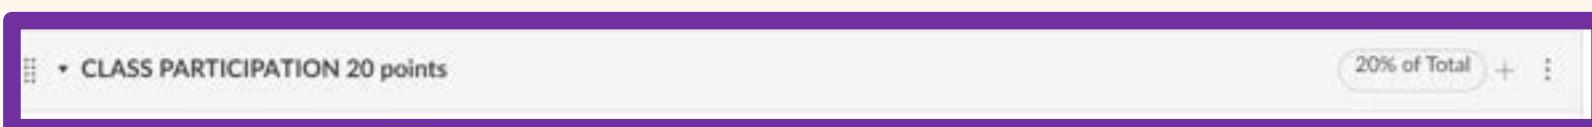
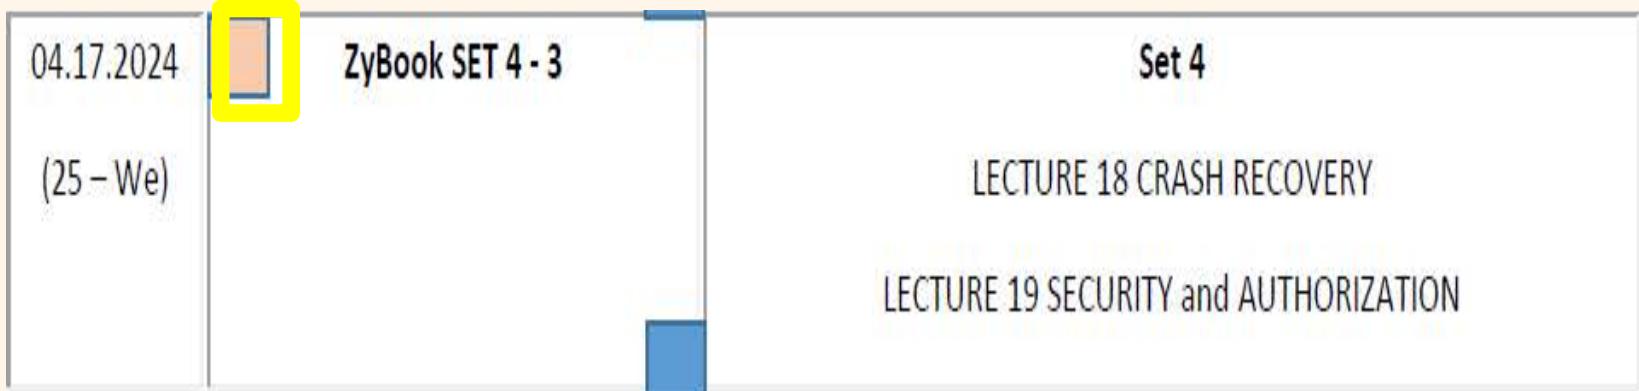
# COSC 3380



## Class 25

04.17.2024  (25 - We)	ZyBook SET 4 - 3	Set 4
		LECTURE 18 CRASH RECOVERY
		LECTURE 19 SECURITY and AUTHORIZATION

From 4:00 to 4:07 PM – 5 minutes.



## RECOVERY



This is a synchronous online class.

Attendance is required.

Recording or distribution of class materials is prohibited.

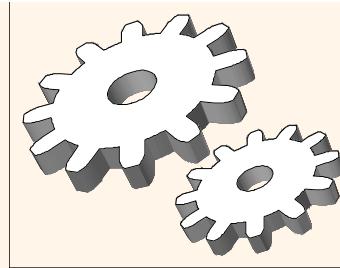
1. At the beginning of selected classes there is an assessment in the first 10 minutes. (beige BOX in the Detailed Syllabus)

2. At the end of selected classes there is an assessment in the last 10 minutes. (blue BOX in the Detailed Syllabus)

3. ZyBook sections will be downloaded and used for 30% of Total Score on the dates specified in the Detailed Syllabus.

4. EXAMS are in CANVAS. No late EXAMS.

5. I have to be present in TEAMS in order to take any graded assignment assigned during that class.

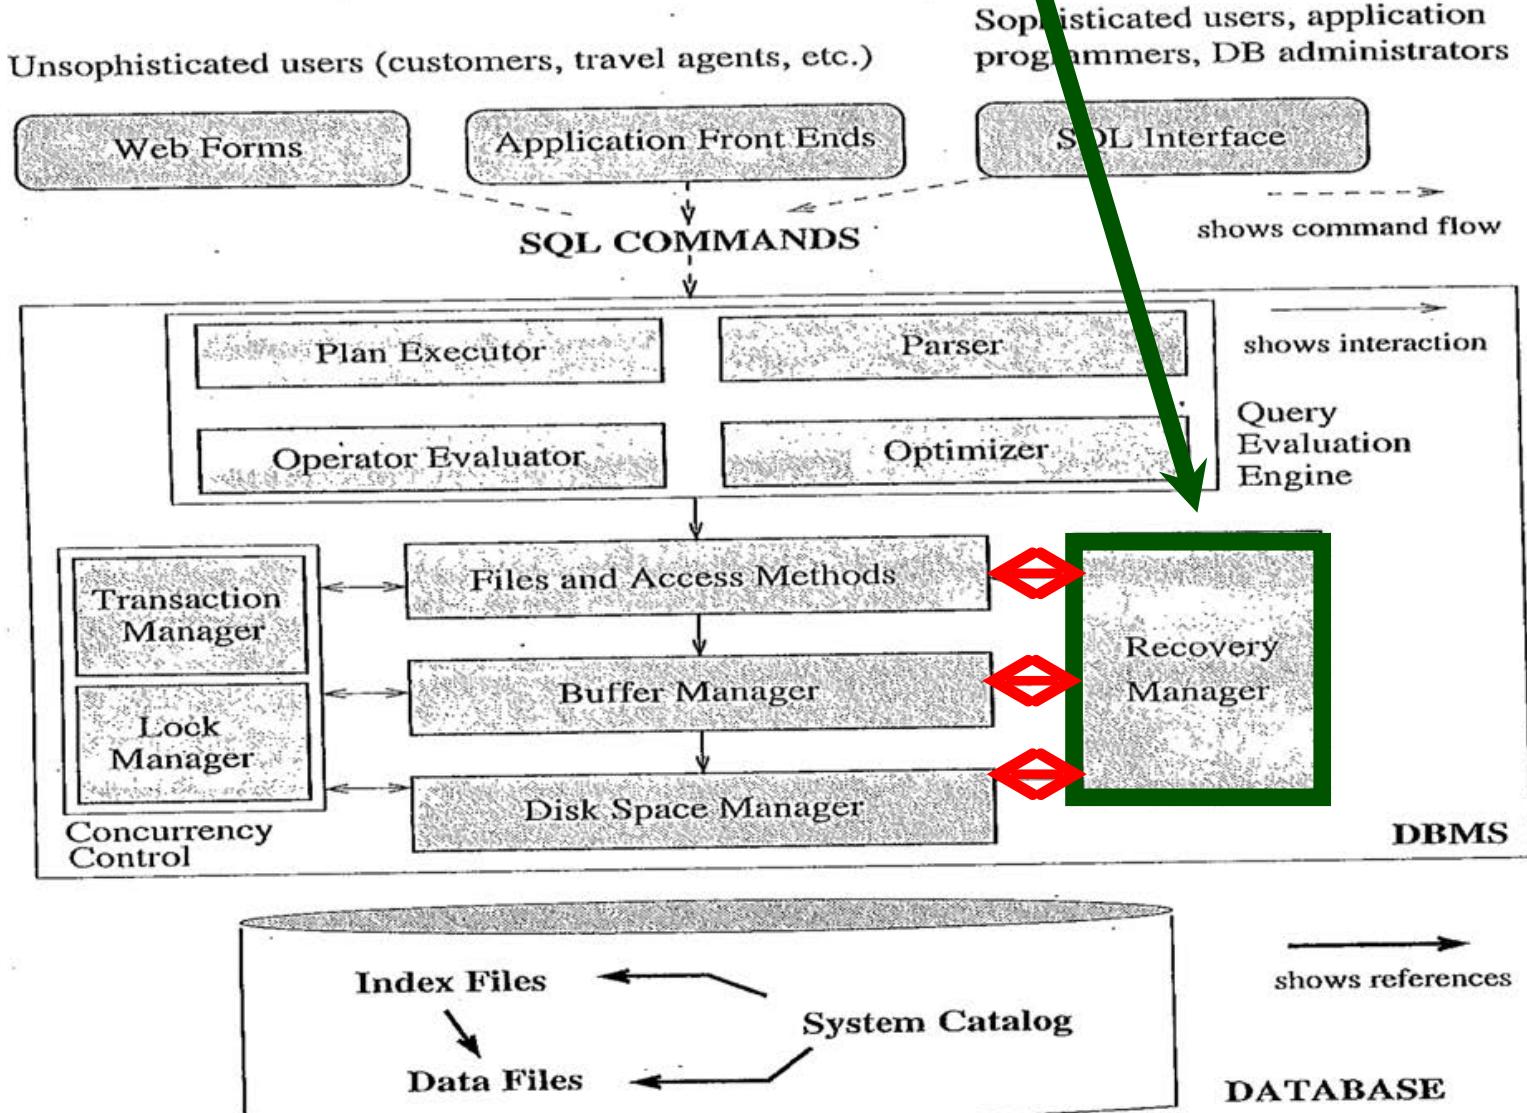
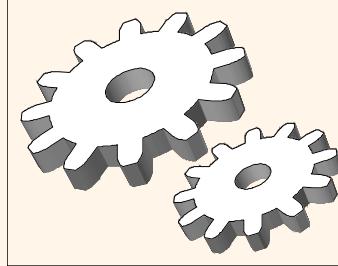


# COSC 3380

## Lecture 18

# *Crash Recovery*

# DBMS - Recovery Manager



## Failure scenarios

The recovery system supports atomic transactions by ensuring partial transaction results are not saved in a database. The recovery system supports durable transactions by ensuring committed transactions are not lost due to hardware or software failures.

The recovery system must manage three failure scenarios:

1. A **transaction failure** results in a rollback. The application program may initiate a rollback due to logical errors. The database system may initiate rollback due to deadlock or insufficient disk space. The operating system may initiate rollback if a hardware or software component fails. Regardless of the cause, the recovery system restores all data changed by the transaction to the original values.

2. A **system failure** includes a variety of events resulting in the loss of main memory. Databases initially write to main memory blocks, which are lost when an application, the operating system, or the database system fails. Blocks are subsequently saved on storage media, which normally survive a system failure. If main memory is lost before blocks are written to storage media, data written by committed transactions might be lost. In this event, the recovery system:

- Recovers data written to main memory, but not storage media, by committed transactions.
- Rolls back data written to storage media by uncommitted transactions.

D  
A

3. A **storage media failure** occurs when the database is corrupted or the database connection is lost.

Many storage systems automatically make redundant copies of data. If one copy of data is corrupted, the storage system automatically switches to a backup copy without intervention by the database or operating system. Nevertheless, storage media do fail occasionally. Alternatively, the connection between the processor and database server might fail. Either way, the database is unavailable to the application.

In principle, recovery from storage media failure is similar to recovery from system failure. In both cases, the recovery system must restore committed transactions and roll back uncommitted transactions. However, storage media failure may cause massive data loss with lengthy recovery times. Consequently, most commercial databases implement special techniques for rapid recovery from storage media failure.

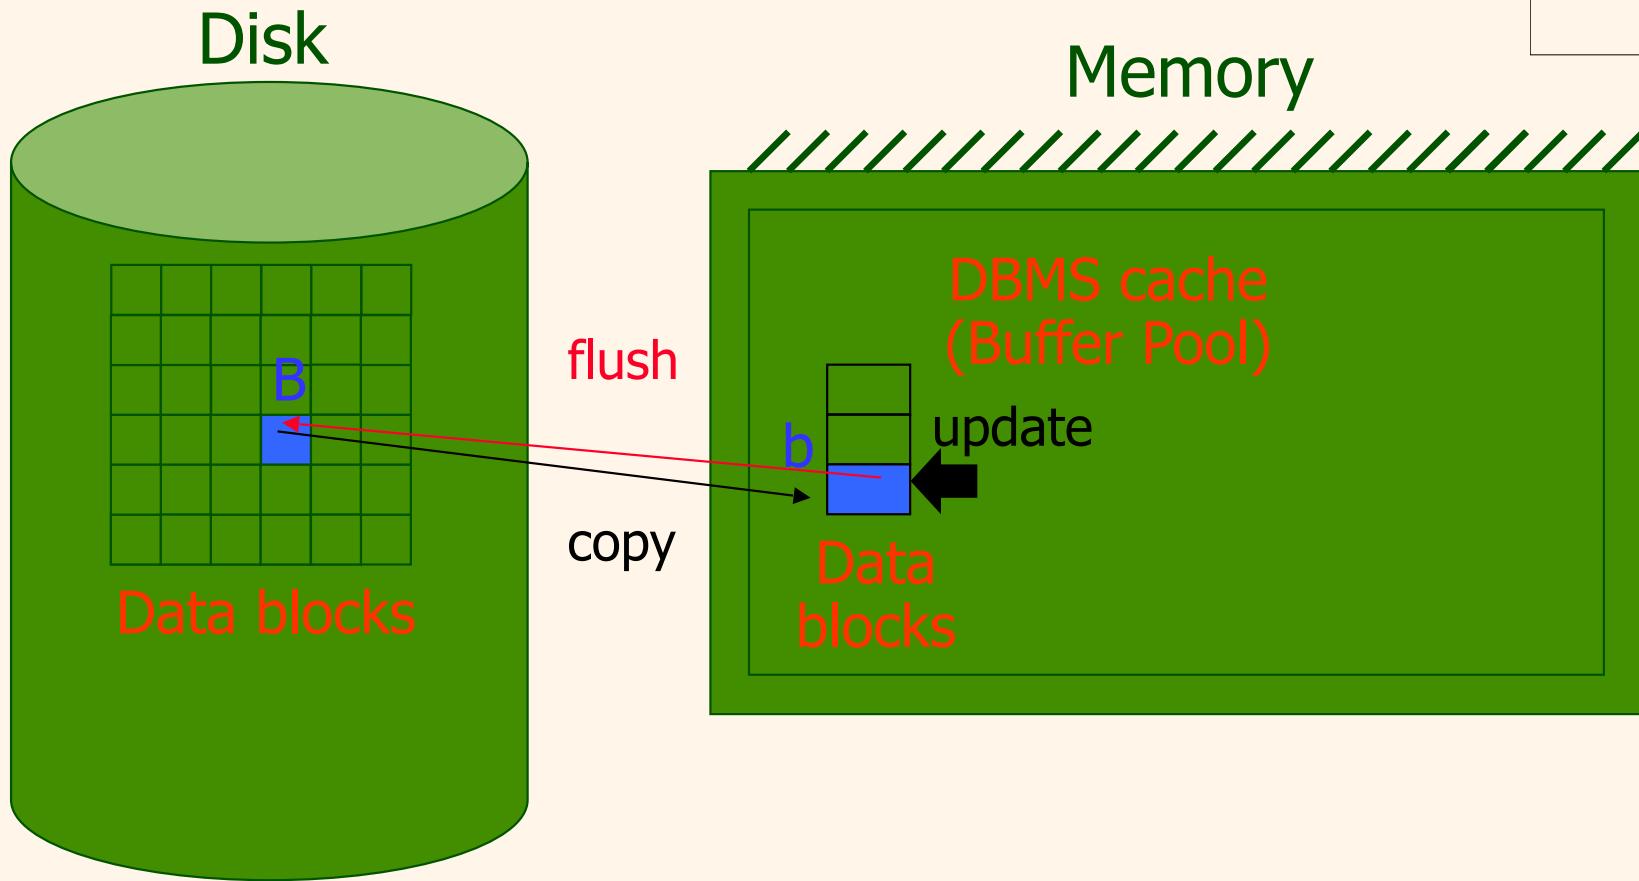
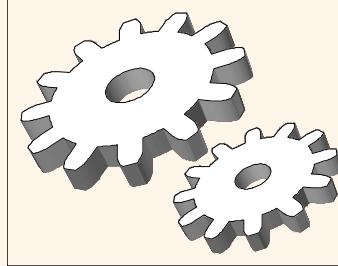
D  
A

**TA, Jordan (A – L).**

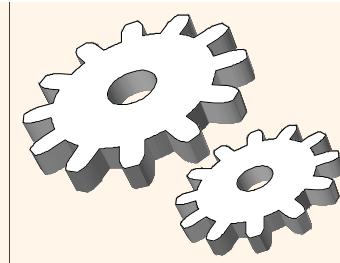
**TA, Alvaro (M – Z).**

**Please compare CANVAS vs. TEAMS Attendance.  
Print screens of students in CANVAS but not in the TEAMS meeting.  
(4.17.2024 Attendance X missing LastName.docx)**

# Physical View



- (1) **copy (data)** from the disk to the cache)
- (2) **update** the cached **data**
- (3) **flush** the **data** (from the cache to the disk)

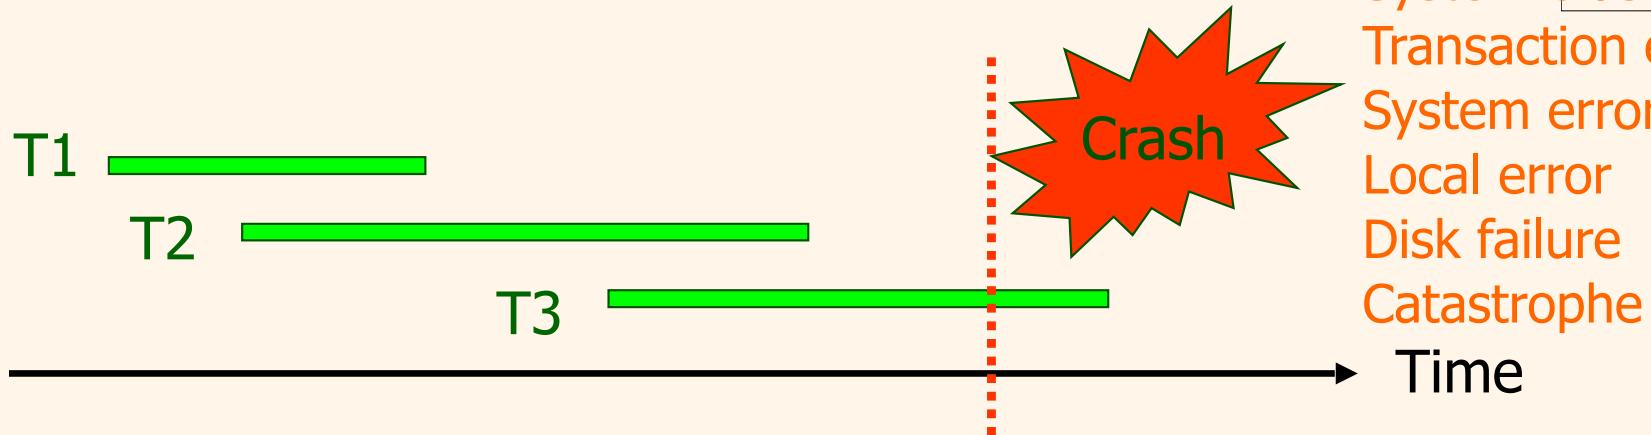


# The *ACID* properties

- ❖ **A**tomicity: **All** actions in the Xact happen, **or none** happen.
- ❖ **C**onsistency: If each Xact is consistent, and the DB starts consistent, it ends up consistent.
- ❖ **I**solation: Execution of one Xact is isolated from that of other Xacts.
- ❖ **D**urability: If a Xact commits, its effects **persist**.
- ❖ The **Recovery Manager** guarantees **Atomicity & Durability**.



## Why “Database Recovery Techniques”?

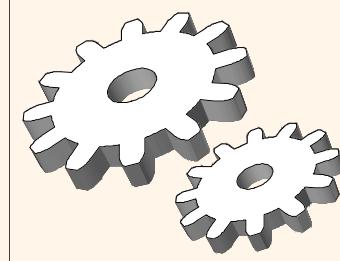


- ## ❖ A CI D properties of Transactions

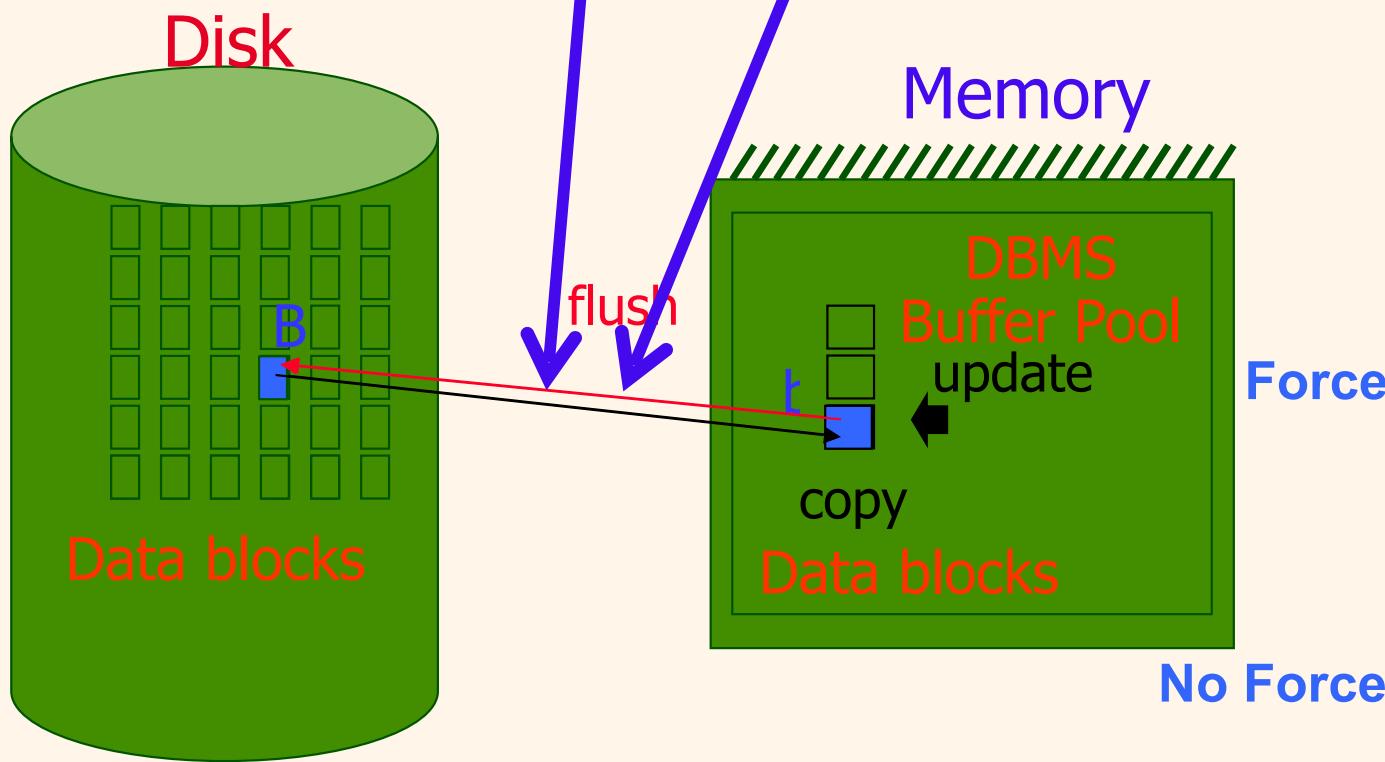
# **Database Management System should guarantee:**

- **A**tomicity : Transactions can be **aborted.** **UNDO**  
~ T3
  - **D**urability : Applied changes by Transactions  
**must not be lost.** **REDO**  
~ T1, T2

# Handling the Buffer Pool (Memory)

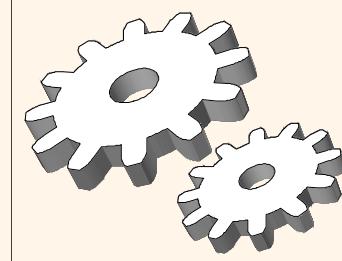


- ❖ Force every write to disk?
  - Poor response time.
  - But provides Durability.
- ❖ Steal buffer-pool frames from uncommitted Transactions?
  - If not, poor throughput.
  - If so, how can we ensure Atomicity?



No Steal	Steal
Trivial	
	Desired

# More on *Steal* and *Force*

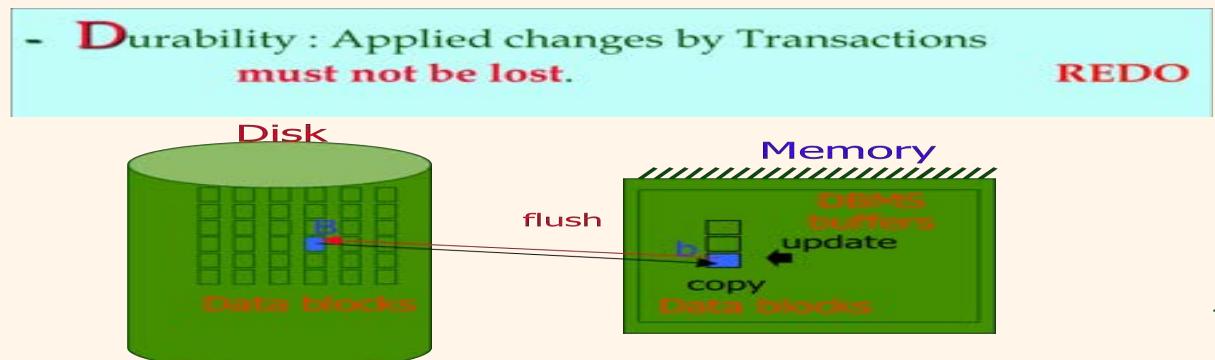


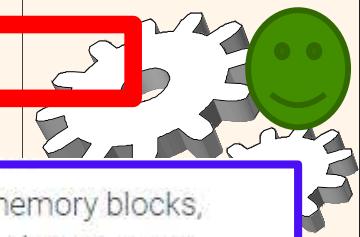
## » STEAL (why enforcing Atomicity is hard)

- *To steal frame F:* Current page in *F* (say *P*) is written to disk; some **Xact** holds lock on *P*.
  - What if the **Xact** with the lock on *P* Aborts?
  - Must remember the **old value of P** at steal time (to support **UNDO**ing the write to page *P*).
    - **Atomicity : Transactions can be aborted.**
    - UNDO**

## » NO FORCE (why enforcing Durability is hard)

- What if system crashes before a modified page is written to disk?
  - Write as little as possible, in a convenient place, at **Commit time**, to support **REDO**ing modifications.





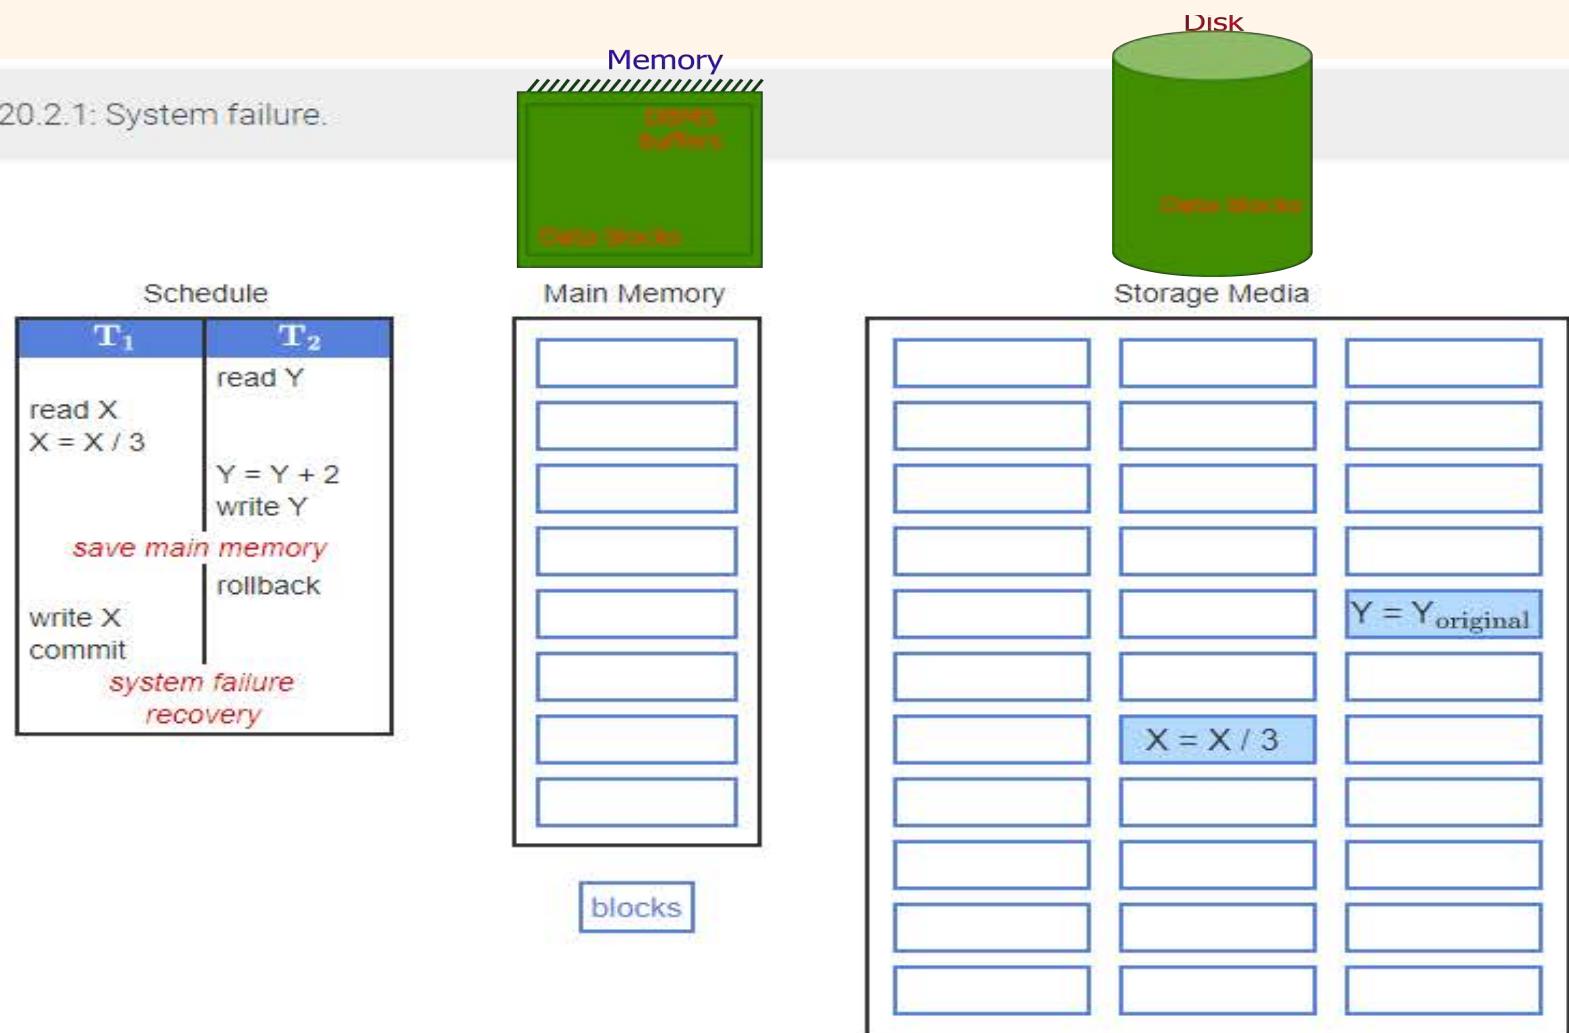
2. A **system failure** includes a variety of events resulting in the loss of main memory. Databases initially write to main memory blocks, which are lost when an application, the operating system, or the database system fails. Blocks are subsequently saved on storage media, which normally survive a system failure. If main memory is lost before blocks are written to storage media, data written by committed transactions might be lost. In this event, the recovery system:

- Recovers data written to main memory, but not storage media, by committed transactions.
- Rolls back data written to storage media by uncommitted transactions.

D  
A

PARTICIPATION ACTIVITY

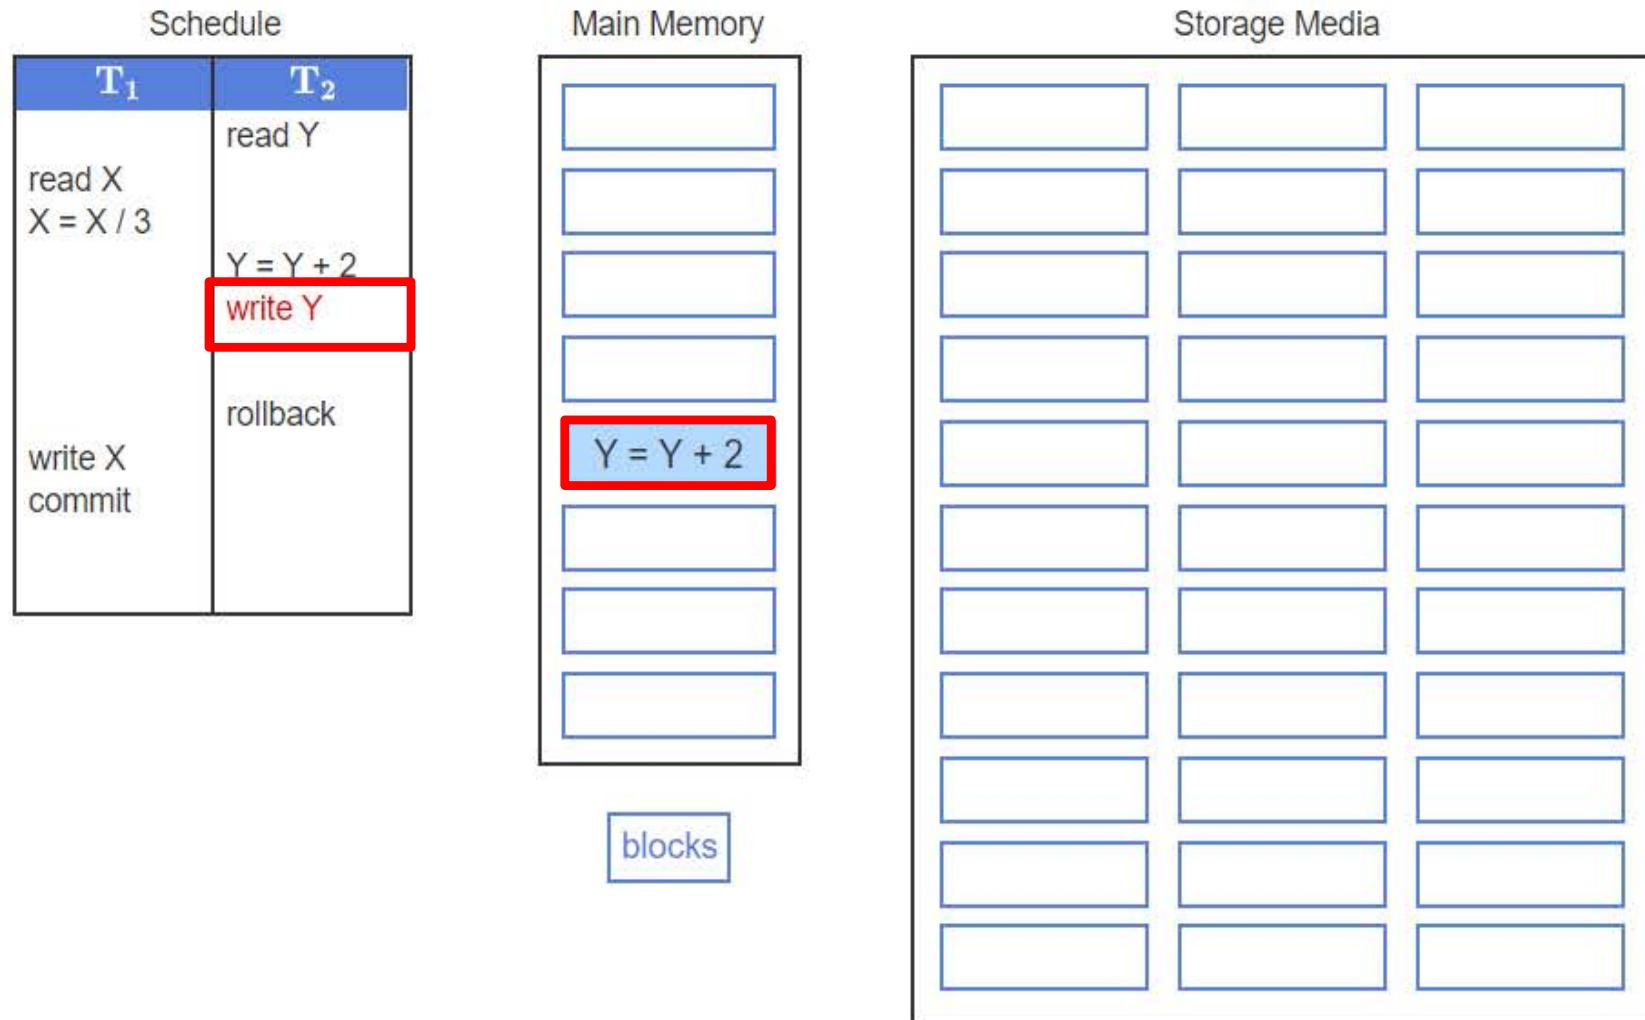
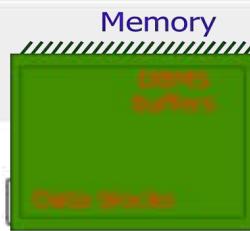
20.2.1: System failure.



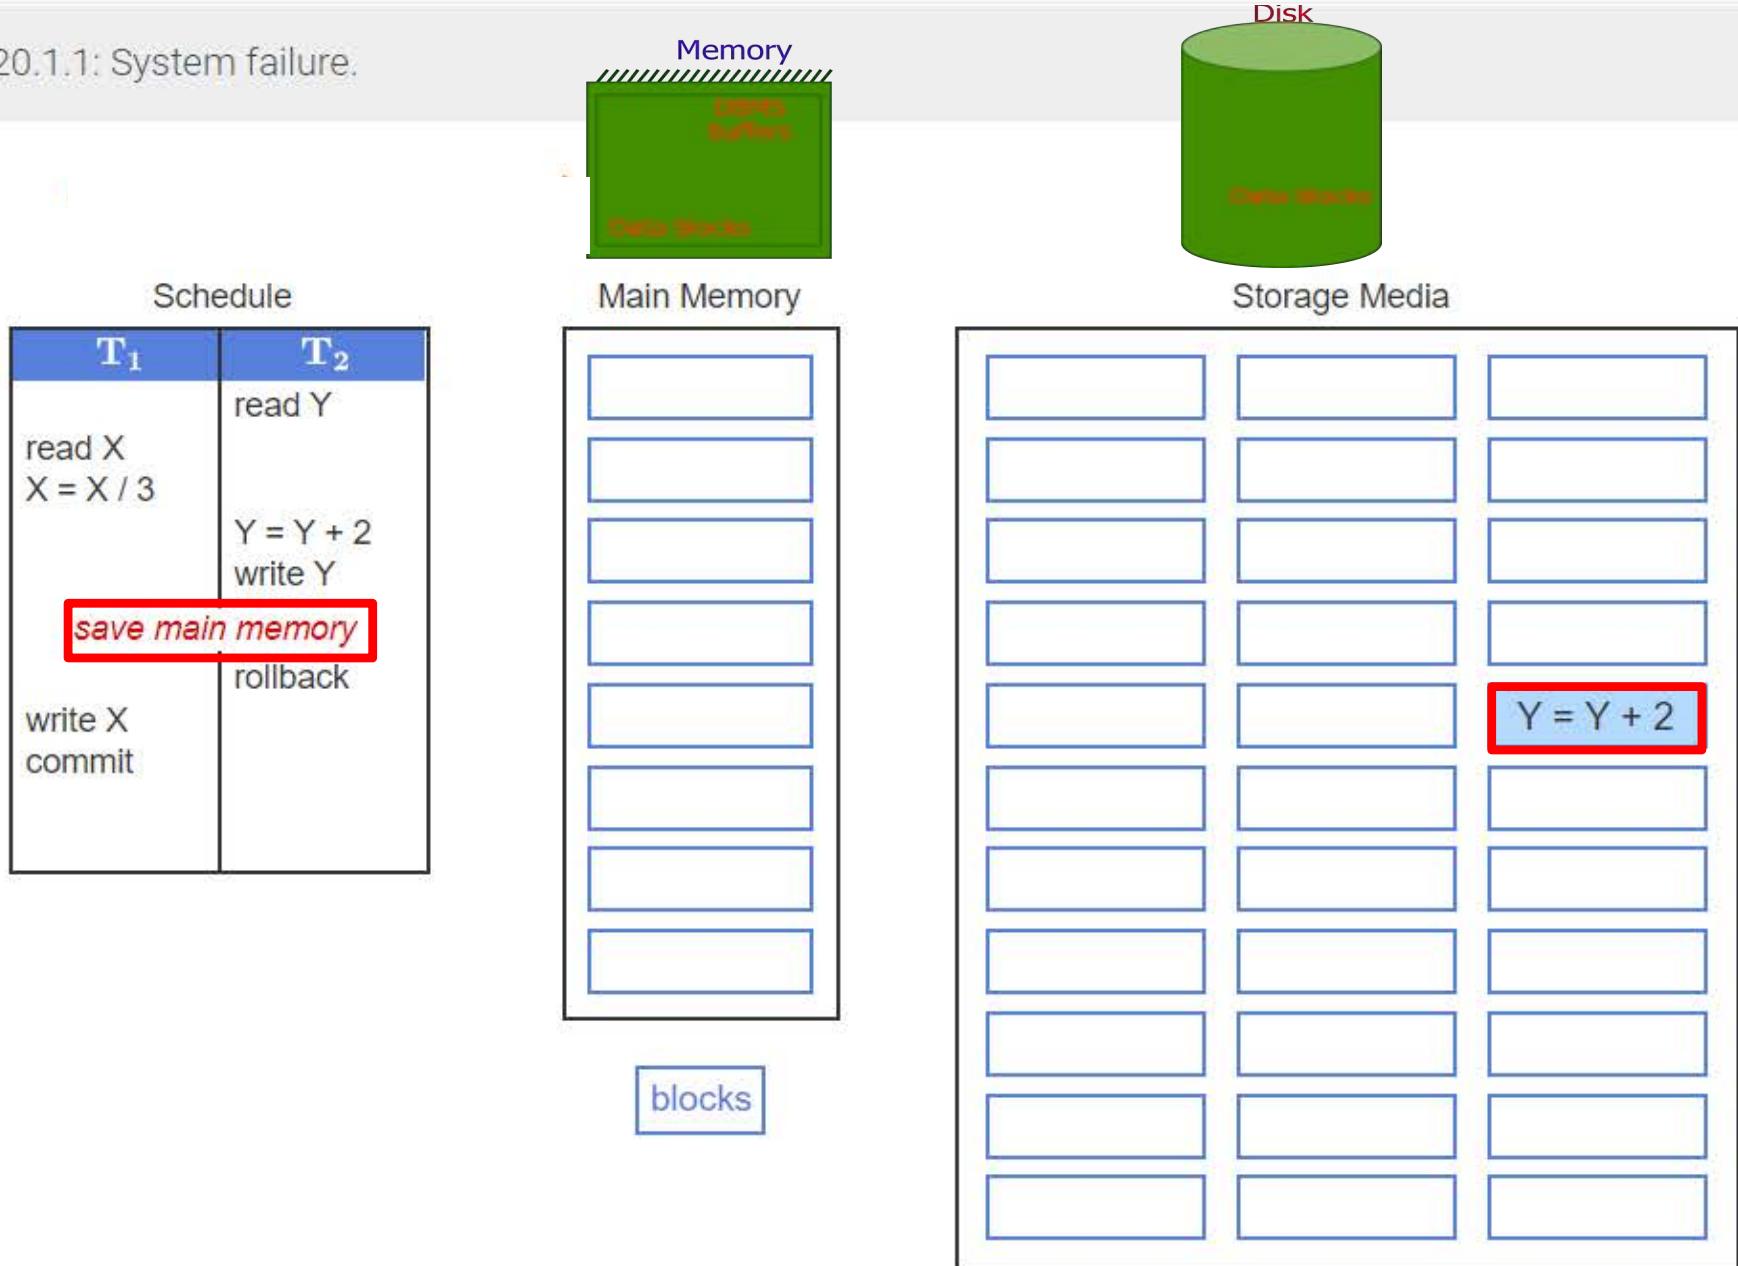
## 20.1.1: System failure.



The database processes two transactions.



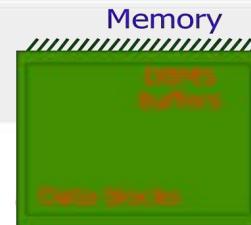
The database writes a new Y value to a block in main memory.



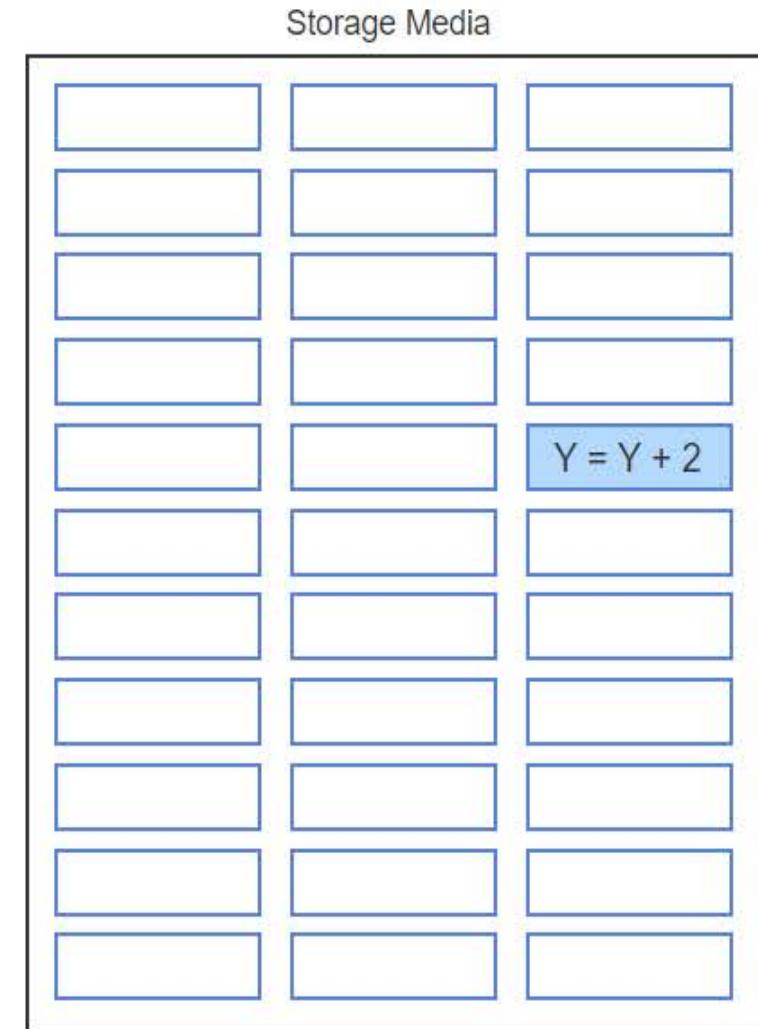
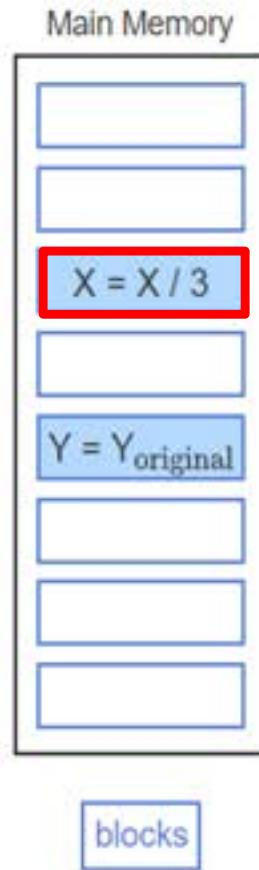
The database saves main memory to storage media.  
All saved blocks in main memory are released for reuse.



The database rolls back  $T_2$ , restoring Y's original value in main memory.

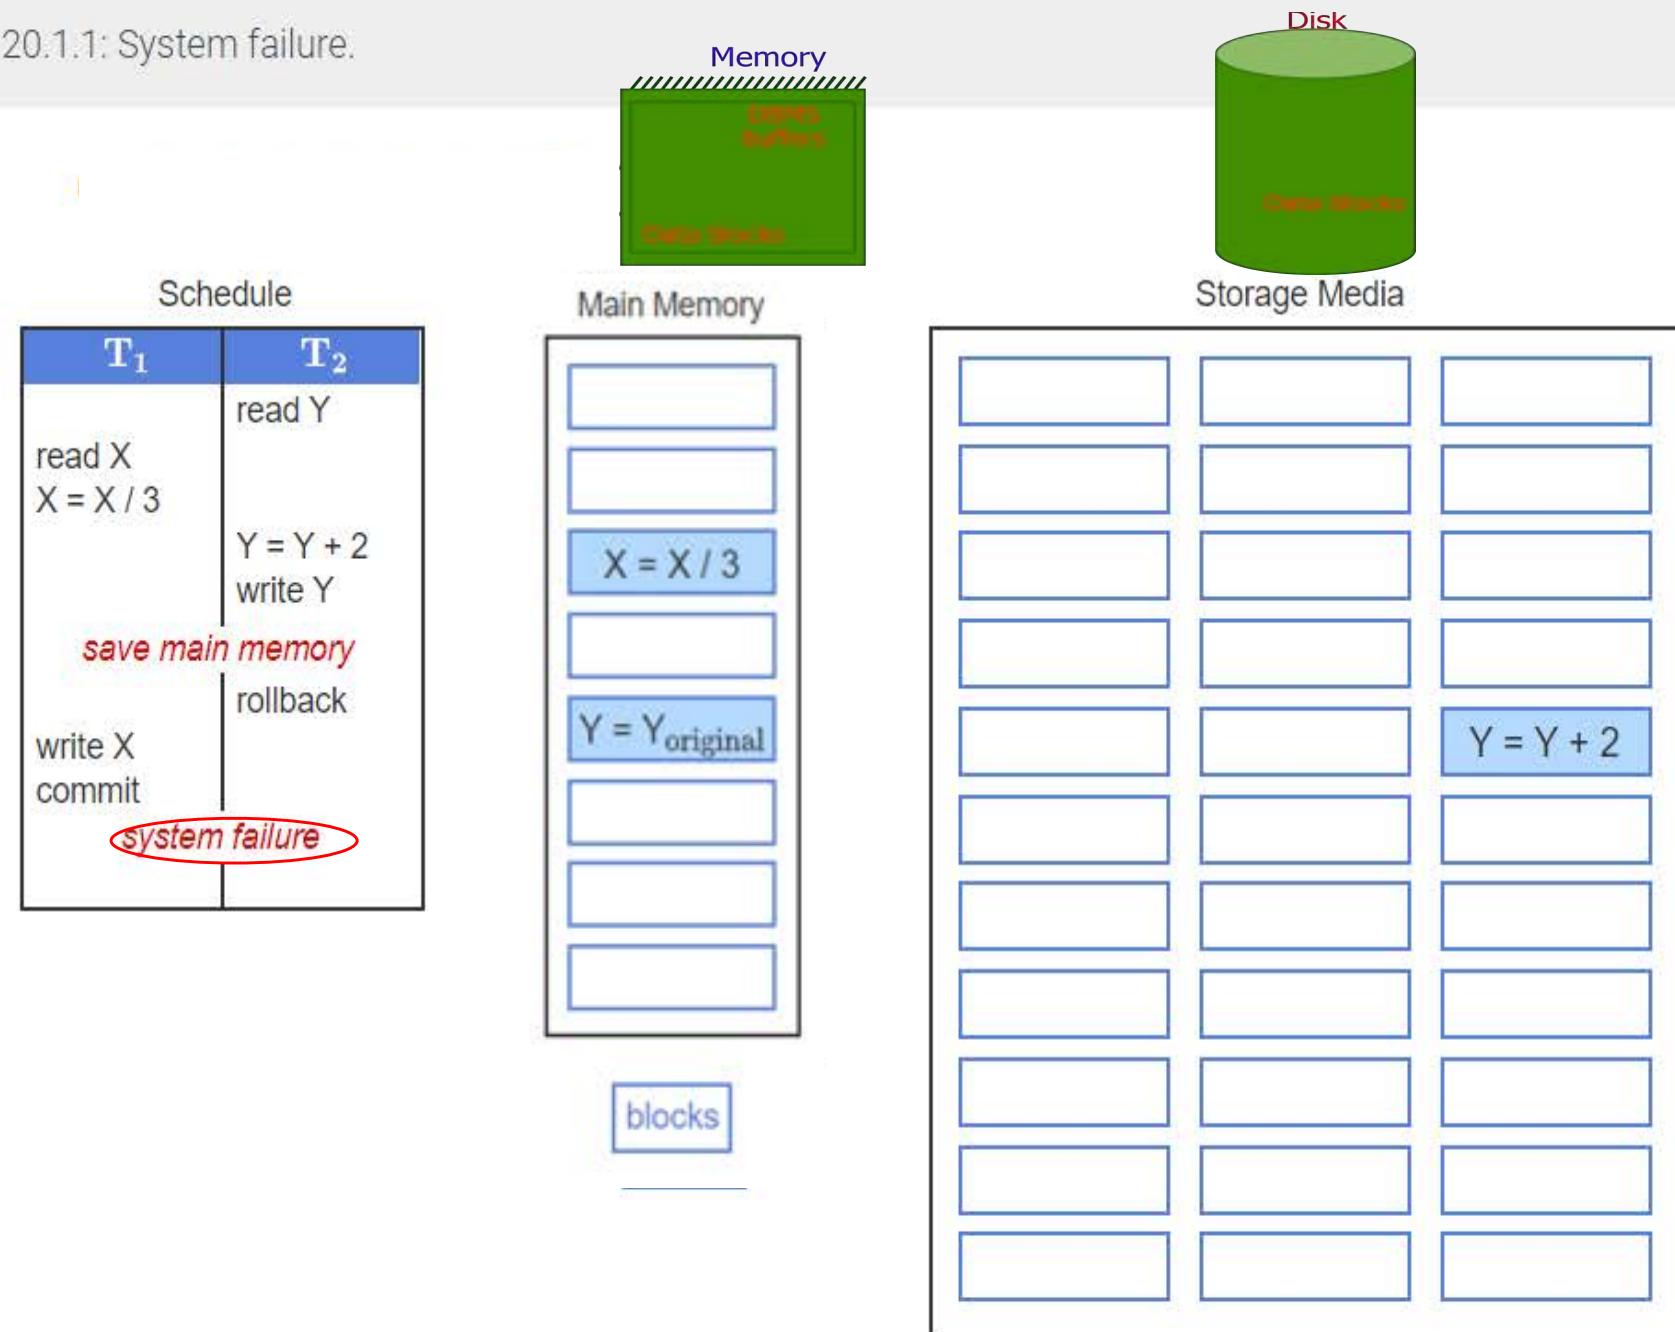


Schedule	
T <sub>1</sub>	T <sub>2</sub>
read X	read Y
X = X / 3	
	Y = Y + 2
	write Y
<i>save main memory</i>	
	rollback
<b>write X</b>	
<b>commit</b>	



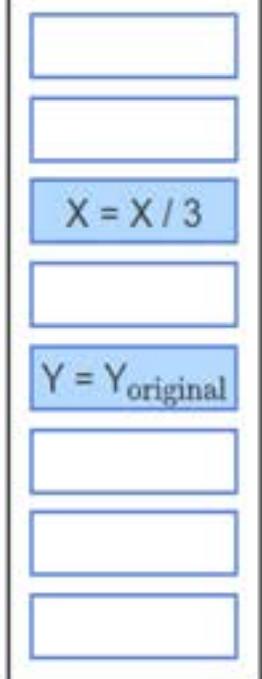
The database writes a new X value in main memory and commits T<sub>1</sub>.

## 20.1.1: System failure.



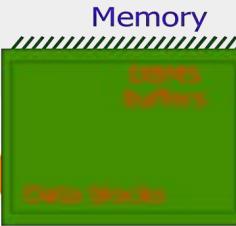
Before main memory is saved to storage media, the operating system fails, and main memory is lost.

Main Memory

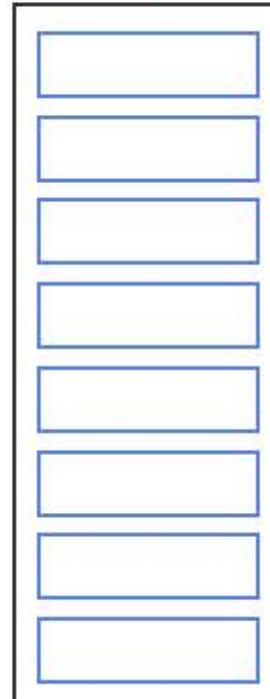


Schedule

$T_1$	$T_2$
read X	read Y
$X = X / 3$	$Y = Y + 2$
	write Y
<i>save main memory</i>	
	rollback
write X	
commit	
<i>system failure recovery</i>	



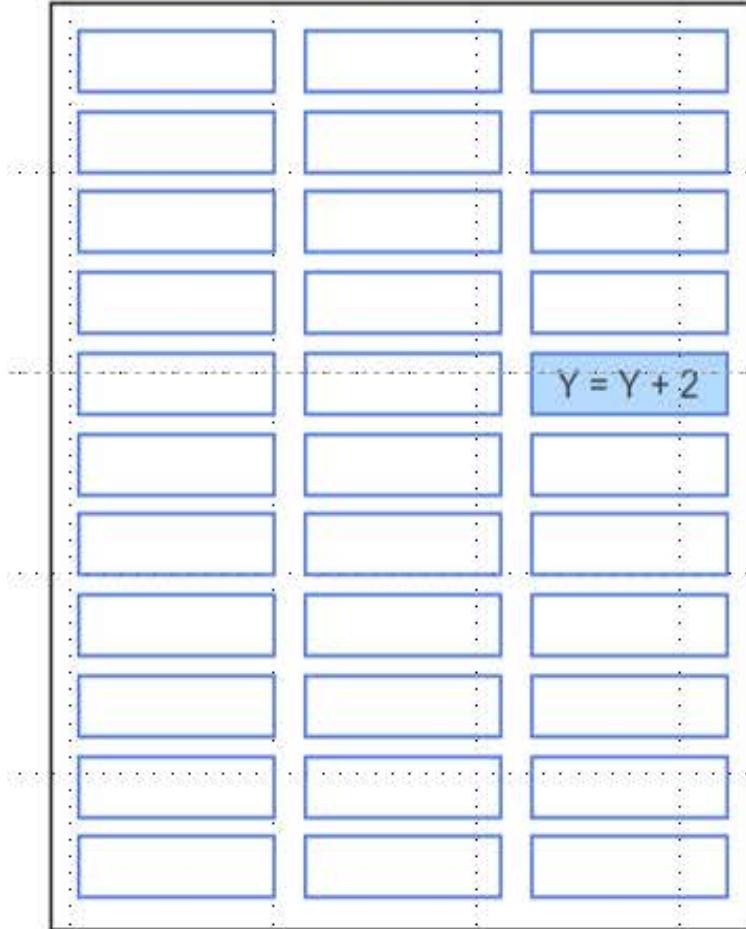
Main Memory



Disk



Storage Media



**REDO**

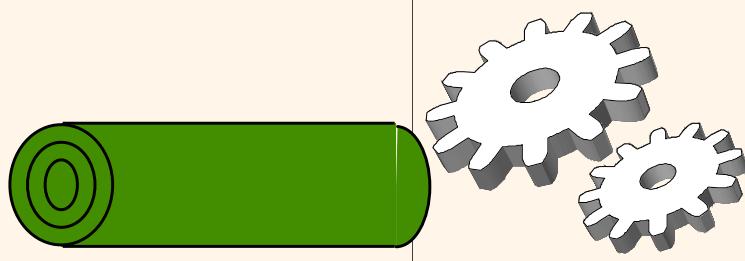
**UNDO**

The recovery system must restore ' $X = X / 3$ ' and roll back ' $Y = Y + 2$ ' in storage media.

D

A

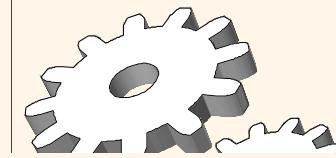
# Basic Idea: Logging



- ❖ Record **REDO** and **UNDO** information, for every UPDATE, in a *Log*.
  - Sequential writes to *Log* (put it on a separate disk).
  - Minimal info (diff) written to *Log*, so multiple UPDATEs fit in a single *Log* page.
- ❖ *Log*: An ordered list of **REDO/UNDO** actions
  - *Log* records contain
    - UPDATE information<XID, pageID, offset, length, old data, new data>
    - and additional **control info**: COMMIT; ABORT; ...

A Transaction Log										
TRL_ID	TRX_NUM	PREV_PTR	NEXT_PTR	OPERATION	TABLE	ROW_ID	ATTRIBUTE	BEFORE_VALUE	AFTER_VALUE	
341	101	Null	352	START	*****Start					
352	101	341	363	UPDATE	PRODUCT	1558-QW1	PROD_QOH	25	23	
363	101	352	365	UPDATE	CUSTOMER	FOOT1	CUST_COST	323.75	615.75	
365	101	363	Null	COMMIT	***** End of					

# The Transaction Log



Transaction Log stores (LSN - Log Sequence #):

- A record for the beginning of Transaction (START)
- For each Transaction component (SQL statement):
  - Type of operation being performed (UPDATE)
  - Names of objects affected by Transaction
  - “Before” and “after” values for updated fields
  - Pointers to previous and next Transaction Log entries (LSN) for the same Transaction
- Ending (COMMIT) of the Transaction
- A record for the end of Transaction (END)

A Transaction Log										
TRX	TRX	PREV	NEXT	OPERATION	TABLE	ROW	ATTRIBUTE	BEFORE	AFTER	
341	101	Null	352	START						
					****Start Transaction					
363	101	352	365	UPDATE	CUSTOMER	QW1 10011	QOH CUST	525.75	615.73	
365	101	363	Null	COMMIT	**** End of Transaction					
↑ TRX_NUM = Transaction number (Note: the transaction number is auto-										



# The Transaction Log

- ❖ The following is a simplified sales example, which *updates a product's quantity on hand (PROD\_QOH) and the customer's 10011 balance when the customer buys 2 units of product 1558-QW1 priced at \$43.99 per unit (for a total of \$87.98) and charges the purchase to his/her account:*

**START;**

**UPDATE PRODUCT**

SET PROD\_QOH = PROD\_QOH - 2  
WHERE PROD\_CODE = '1558-QW1';

**UPDATE CUSTOMER**

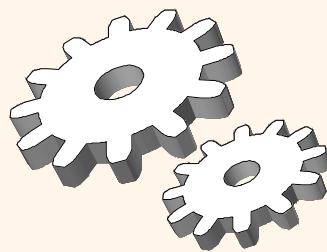
SET CUST\_BALANCE = CUST\_BALANCE + 87.98  
WHERE CUST\_NUMBER = '10011';

**COMMIT;  
END;**

How many Log entries would this Transaction produce?

**5**

# The Transaction Log



LSN – Log Sequence Number

↓

A Transaction Log

TRL_ID	TRX_NUM	PREV_PTR	NEXT_PTR	OPERATION	TABLE	ROW_ID	ATTRIBUTE	BEFORE_VALUE	AFTER_VALUE
341	101	Null	352	START	****Start Transaction				
352	101	341	363	UPDATE	PRODUCT	1558-QW1	PROD_QOH	25	23
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	525.75	615.73
365	101	363	367	COMMIT	**** End of Transaction				
367	101	365	Null				END		

```

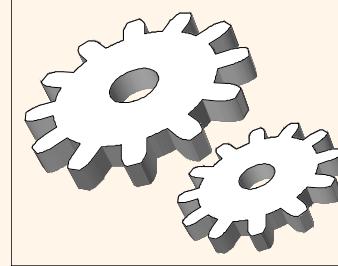
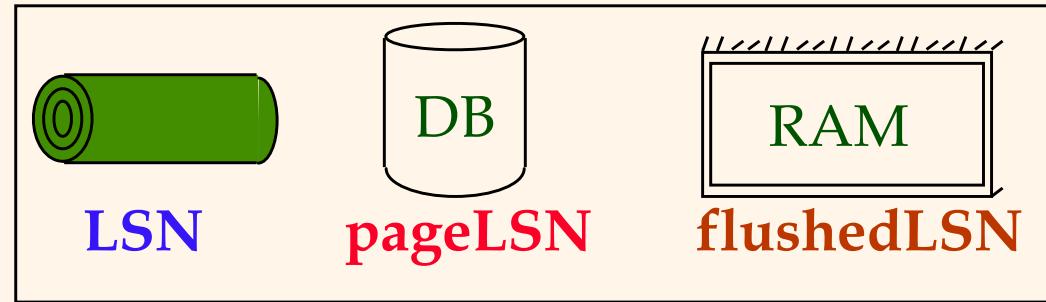
START;
UPDATE PRODUCT
SET PROD_QOH = PROD_QOH - 2
WHERE PROD_CODE = '1558-QW1';

UPDATE CUSTOMER
SET CUST_BALANCE = CUST_BALANCE + 87.98
WHERE CUST_NUMBER = '10011';

COMMIT;
END;

```

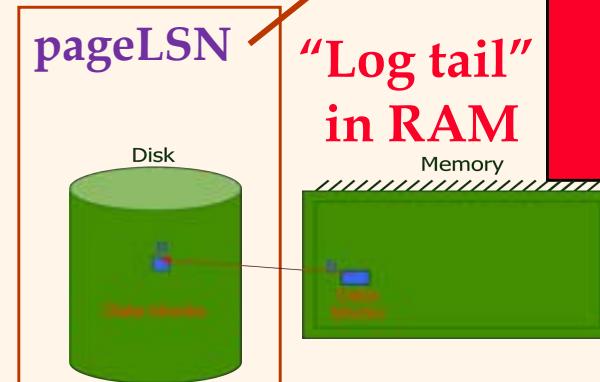
# *WAL & the Log*



- ❖ Each **Log** record has a unique **Log Sequence Number (LSN)**
  - LSNs always increasing.
- ❖ Each **data page** contains a **pageLSN**
  - The LSN of the *first Log record* update to that **page**.
- ❖ System keeps track of **flushedLSN**
  - The max LSN flushed so far.
- ❖ **WAL**: *Before* a **data page** is written, **pageLSN**  $\leq$  **flushedLSN**

Log records  
flushed to disk

for an



System Log Table											
Time	User	SQL	OpType	Statement	Table	RowID	OldValue	NewValue	Index	Condition	Notes
2023-01-01 09:00:00	user1	SELECT	READ	SELECT * FROM CUSTOMER;	CUSTOMER	12345	John Doe, 123 Main St	John Doe, 456 Elm St	INDEX	CustomerName	
2023-01-01 09:00:05	user1	UPDATE	UPDATE	UPDATE CUSTOMER SET ADDRESS = '456 Elm St' WHERE ID = 12345;	CUSTOMER	12345	John Doe, 123 Main St	John Doe, 456 Elm St	INDEX	CustomerName	
2023-01-01 09:00:10	user1	SELECT	READ	SELECT * FROM CUSTOMER;	CUSTOMER	12345	John Doe, 456 Elm St	John Doe, 456 Elm St	INDEX	CustomerName	
2023-01-01 09:00:15	user1	UPDATE	UPDATE	UPDATE CUSTOMER SET ADDRESS = '123 Main St' WHERE ID = 12345;	CUSTOMER	12345	John Doe, 456 Elm St	John Doe, 123 Main St	INDEX	CustomerName	

# Log Records

## Log Record fields:

prevLSN

XID

type

pageID

length

offset

before-image

after-image

UPDATE  
records  
only

Possible Log record types:

- ❖ UPDATE

- ❖ Start

- ❖ Commit

- ❖ Abort

- ❖ End (signifies end of Commit or Abort)

- ❖ Compensation Log Records (CLRs)

- for UNDO actions (LSN of the next UNDO undonextLSN)

A Transaction Log											
TRL_ID	TRX_NUM	PREV_PTR	NEXT_PTR	OPERATION	TABLE	ROW_ID	ATTRIBUTE	BEFORE_VALUE	AFTER_VALUE		
341	101	Null	352	START	****Start Transaction						
352	101	341	363	UPDATE	PRODUCT	1558-0001	PROD_QTY	25	23		
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	525.75	615.73		
365	101	363	Null	COMMIT	**** End of Transaction						

TRL\_ID = Transaction log record ID  
TRX\_NUM = Transaction number  
(Note: The transaction number is automatically assigned by the DBMS.)

PTR = Pointer to a transaction log record ID

## Recovery log



A **recovery log** is a file containing a sequential record of all database operations. The log and database are stored on different storage media so the log survives database failures. The recovery system uses the log to restore the database after a failure.

The log contains transaction and data identifiers. Transaction identifiers are assigned by the database system for internal use by the recovery and concurrency systems. Data identifiers correspond to table name, column name, and primary key value, but are compressed or encoded for efficiency.

The recovery log contains four types of records:

1. An **update record** indicates a transaction has changed data. Update records include the transaction identifier, the data identifier, the original data value, and the new data value. Update records may optionally track insert and delete operations.
2. A **compensation record**, also known as an **undo record**, indicates data has been restored to the original value during a rollback. Compensation records include the transaction identifier, the data identifier, and the restored (original) data value. Compensation records are necessary because the log must capture every database change, including changes executed by a rollback.
3. A **transaction record** indicates a transaction boundary. Three types of transaction records exist: start, commit, and rollback. Transaction records include the 'start', 'commit', or 'rollback' indicator and the transaction identifier.
4. A **checkpoint record** indicates that all data in main memory has been saved on storage media. When the database executes a checkpoint, transaction processing is suspended while all unsaved data and log records are written to storage media. In the event of a system failure, the recovery system reads only log records following the last checkpoint, rather than the entire file. Checkpoint records include a 'checkpoint' indicator along with the identifiers of all transactions that are active (uncommitted) at the time of the checkpoint.

The above four record types provide a complete history of database operations and enable recovery from all three failure scenarios.

### PARTICIPATION ACTIVITY

#### 20.1.3: Recovery log records.



Schedule	
T <sub>1</sub>	T <sub>2</sub>
read X X = X / 3	read Y
	Y = Y + 2 write Y
	save main memory
	rollback
write X commit	

Memory	Recovery Log
	start T <sub>2</sub>
	start T <sub>1</sub>
	update T <sub>2</sub> , Y <sub>ID</sub> , Y <sub>original</sub> , Y <sub>new</sub>
	checkpoint T <sub>1</sub> , T <sub>2</sub>
	undo T <sub>2</sub> , Y <sub>ID</sub> , Y <sub>original</sub>
	rollback T <sub>2</sub>
	update T <sub>1</sub> , X <sub>ID</sub> , X <sub>original</sub> , X <sub>new</sub>
	commit T <sub>1</sub>

### Log Records

Possible Log record types:

- ♦ UPDATE
- ♦ Start
- ♦ Commit
- ♦ Abort
- ♦ End (signifies end of Commit or Abort)
- ♦ Compensation Log Records (CLRs)
  - \* for UNDO actions (0.5% of the next UNDO undone(LSN))

### Log Record fields:

UPDATE records only

- prevLSN
- XID
- type
- pageID
- length
- offset
- before-image
- after-image



## Recovery log

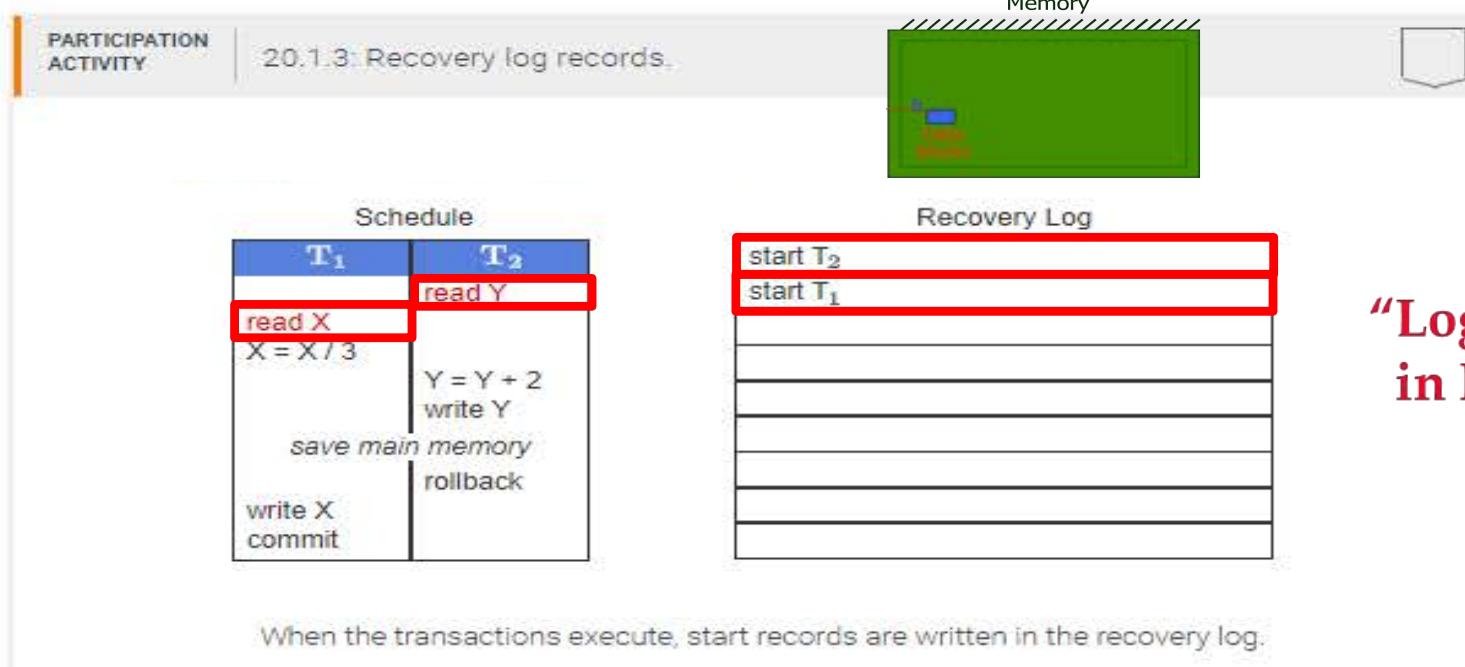
A **recovery log** is a file containing a sequential record of all database operations. The log and database are stored on different storage media so the log survives database failures. The recovery system uses the log to restore the database after a failure.

The log contains transaction and data identifiers. Transaction identifiers are assigned by the database system for internal use by the recovery and concurrency systems. Data identifiers correspond to table name, column name, and primary key value, but are compressed or encoded for efficiency.

The recovery log contains four types of records:

1. An **update record** indicates a transaction has changed data. Update records include the transaction identifier, the data identifier, the original data value, and the new data value. Update records may optionally track insert and delete operations.
2. A **compensation record**, also known as an **undo record**, indicates data has been restored to the original value during a rollback. Compensation records include the transaction identifier, the data identifier, and the restored (original) data value. Compensation records are necessary because the log must capture every database change, including changes executed by a rollback.
3. A **transaction record** indicates a transaction boundary. Three types of transaction records exist: start, commit, and rollback. Transaction records include the 'start', 'commit', or 'rollback' indicator and the transaction identifier.
4. A **checkpoint record** indicates that all data in main memory has been saved on storage media. When the database executes a checkpoint, transaction processing is suspended while all unsaved data and log records are written to storage media. In the event of a system failure, the recovery system reads only log records following the last checkpoint, rather than the entire file. Checkpoint records include a 'checkpoint' indicator along with the identifiers of all transactions that are active (uncommitted) at the time of the checkpoint.

The above four record types provide a complete history of database operations and enable recovery from all three failure scenarios.



## Recovery log

A **recovery log** is a file containing a sequential record of all database operations. The log and database are stored on different storage media so the log survives database failures. The recovery system uses the log to restore the database after a failure.

The log contains transaction and data identifiers. Transaction identifiers are assigned by the database system for internal use by the recovery and concurrency systems. Data identifiers correspond to table name, column name, and primary key value, but are compressed or encoded for efficiency.

The recovery log contains four types of records:

1. An **update record** indicates a transaction has changed data. Update records include the transaction identifier, the data identifier, the original data value, and the new data value. Update records may optionally track insert and delete operations.
2. A **compensation record**, also known as an **undo record**, indicates data has been restored to the original value during a rollback. Compensation records include the transaction identifier, the data identifier, and the restored (original) data value. Compensation records are necessary because the log must capture every database change, including changes executed by a rollback.
3. A **transaction record** indicates a transaction boundary. Three types of transaction records exist: start, commit, and rollback. Transaction records include the 'start', 'commit', or 'rollback' indicator and the transaction identifier.
4. A **checkpoint record** indicates that all data in main memory has been saved on storage media. When the database executes a checkpoint, transaction processing is suspended while all unsaved data and log records are written to storage media. In the event of a system failure, the recovery system reads only log records following the last checkpoint, rather than the entire file. Checkpoint records include a 'checkpoint' indicator along with the identifiers of all transactions that are active (uncommitted) at the time of the checkpoint.

The above four record types provide a complete history of database operations and enable recovery from all three failure scenarios.



Schedule	
T <sub>1</sub>	T <sub>2</sub>
read X X = X / 3	read Y
	Y = Y + 2
	write Y
save main memory	
	rollback
write X	
commit	

Recovery Log
start T <sub>2</sub>
start T <sub>1</sub>
update T <sub>2</sub> , Y <sub>ip</sub> , Y <sub>original</sub> , Y <sub>new</sub>

**"Log tail"  
in RAM**

When T<sub>2</sub> writes Y, an update record is written in the log.

## Recovery log

A **recovery log** is a file containing a sequential record of all database operations. The log and database are stored on different storage media so the log survives database failures. The recovery system uses the log to restore the database after a failure.

The log contains transaction and data identifiers. Transaction identifiers are assigned by the database system for internal use by the recovery and concurrency systems. Data identifiers correspond to table name, column name, and primary key value, but are compressed or encoded for efficiency.

The recovery log contains four types of records:

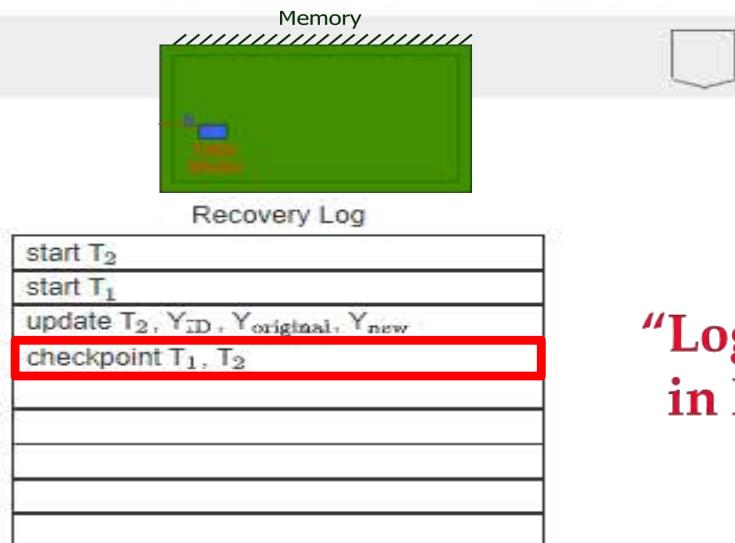
1. An **update record** indicates a transaction has changed data. Update records include the transaction identifier, the data identifier, the original data value, and the new data value. Update records may optionally track insert and delete operations.
2. A **compensation record**, also known as an **undo record**, indicates data has been restored to the original value during a rollback. Compensation records include the transaction identifier, the data identifier, and the restored (original) data value. Compensation records are necessary because the log must capture every database change, including changes executed by a rollback.
3. A **transaction record** indicates a transaction boundary. Three types of transaction records exist: start, commit, and rollback. Transaction records include the 'start', 'commit', or 'rollback' indicator and the transaction identifier.
4. A **checkpoint record** indicates that all data in main memory has been saved on storage media. When the database executes a checkpoint, transaction processing is suspended while all unsaved data and log records are written to storage media. In the event of a system failure, the recovery system reads only log records following the last checkpoint, rather than the entire file. Checkpoint records include a 'checkpoint' indicator along with the identifiers of all transactions that are active (uncommitted) at the time of the checkpoint.

The above four record types provide a complete history of database operations and enable recovery from all three failure scenarios.

PARTICIPATION ACTIVITY

20.1.3: Recovery log records.

Schedule	
T <sub>1</sub>	T <sub>2</sub>
read X X = X / 3	read Y
	Y = Y + 2
	write Y
<i>save main memory</i>	
	rollback
write X commit	



**"Log tail"  
in RAM**

When the system saves all updates from main memory to storage media, a checkpoint record is written in the log.

## Recovery log

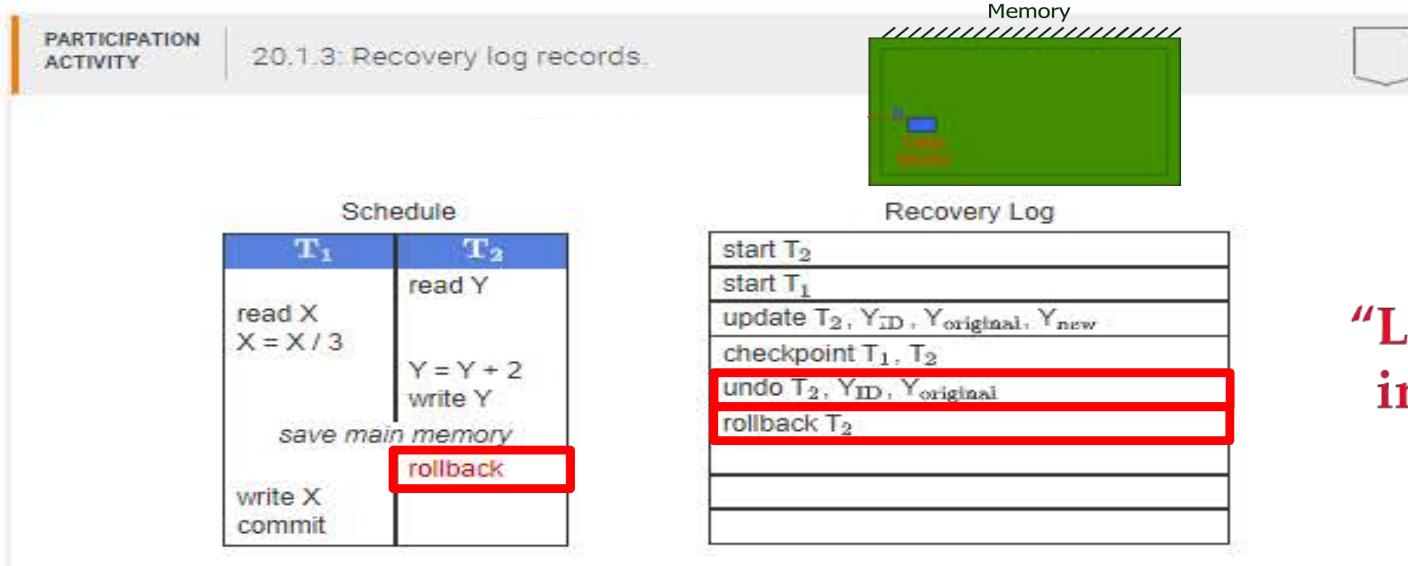
A **recovery log** is a file containing a sequential record of all database operations. The log and database are stored on different storage media so the log survives database failures. The recovery system uses the log to restore the database after a failure.

The log contains transaction and data identifiers. Transaction identifiers are assigned by the database system for internal use by the recovery and concurrency systems. Data identifiers correspond to table name, column name, and primary key value, but are compressed or encoded for efficiency.

The recovery log contains four types of records:

1. An **update record** indicates a transaction has changed data. Update records include the transaction identifier, the data identifier, the original data value, and the new data value. Update records may optionally track insert and delete operations.
2. A **compensation record**, also known as an **undo record**, indicates data has been restored to the original value during a rollback. Compensation records include the transaction identifier, the data identifier, and the restored (original) data value. Compensation records are necessary because the log must capture every database change, including changes executed by a rollback.
3. A **transaction record** indicates a transaction boundary. Three types of transaction records exist: start, commit, and rollback. Transaction records include the 'start', 'commit', or 'rollback' indicator and the transaction identifier.
4. A **checkpoint record** indicates that all data in main memory has been saved on storage media. When the database executes a checkpoint, transaction processing is suspended while all unsaved data and log records are written to storage media. In the event of a system failure, the recovery system reads only log records following the last checkpoint, rather than the entire file. Checkpoint records include a 'checkpoint' indicator along with the identifiers of all transactions that are active (uncommitted) at the time of the checkpoint.

The above four record types provide a complete history of database operations and enable recovery from all three failure scenarios.



## Recovery log

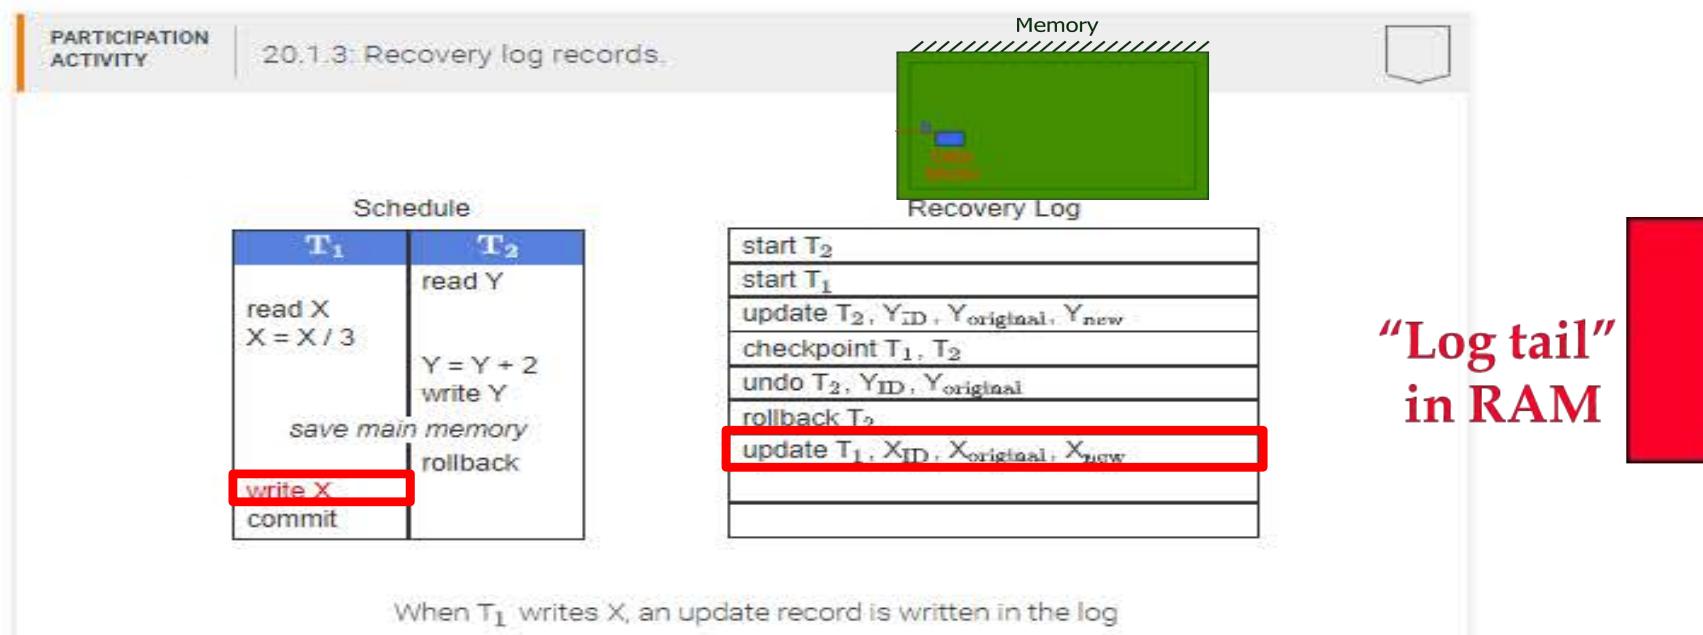
A **recovery log** is a file containing a sequential record of all database operations. The log and database are stored on different storage media so the log survives database failures. The recovery system uses the log to restore the database after a failure.

The log contains transaction and data identifiers. Transaction identifiers are assigned by the database system for internal use by the recovery and concurrency systems. Data identifiers correspond to table name, column name, and primary key value, but are compressed or encoded for efficiency.

The recovery log contains four types of records:

1. An **update record** indicates a transaction has changed data. Update records include the transaction identifier, the data identifier, the original data value, and the new data value. Update records may optionally track insert and delete operations.
2. A **compensation record**, also known as an **undo record**, indicates data has been restored to the original value during a rollback. Compensation records include the transaction identifier, the data identifier, and the restored (original) data value. Compensation records are necessary because the log must capture every database change, including changes executed by a rollback.
3. A **transaction record** indicates a transaction boundary. Three types of transaction records exist: start, commit, and rollback. Transaction records include the 'start', 'commit', or 'rollback' indicator and the transaction identifier.
4. A **checkpoint record** indicates that all data in main memory has been saved on storage media. When the database executes a checkpoint, transaction processing is suspended while all unsaved data and log records are written to storage media. In the event of a system failure, the recovery system reads only log records following the last checkpoint, rather than the entire file. Checkpoint records include a 'checkpoint' indicator along with the identifiers of all transactions that are active (uncommitted) at the time of the checkpoint.

The above four record types provide a complete history of database operations and enable recovery from all three failure scenarios.



## Recovery log

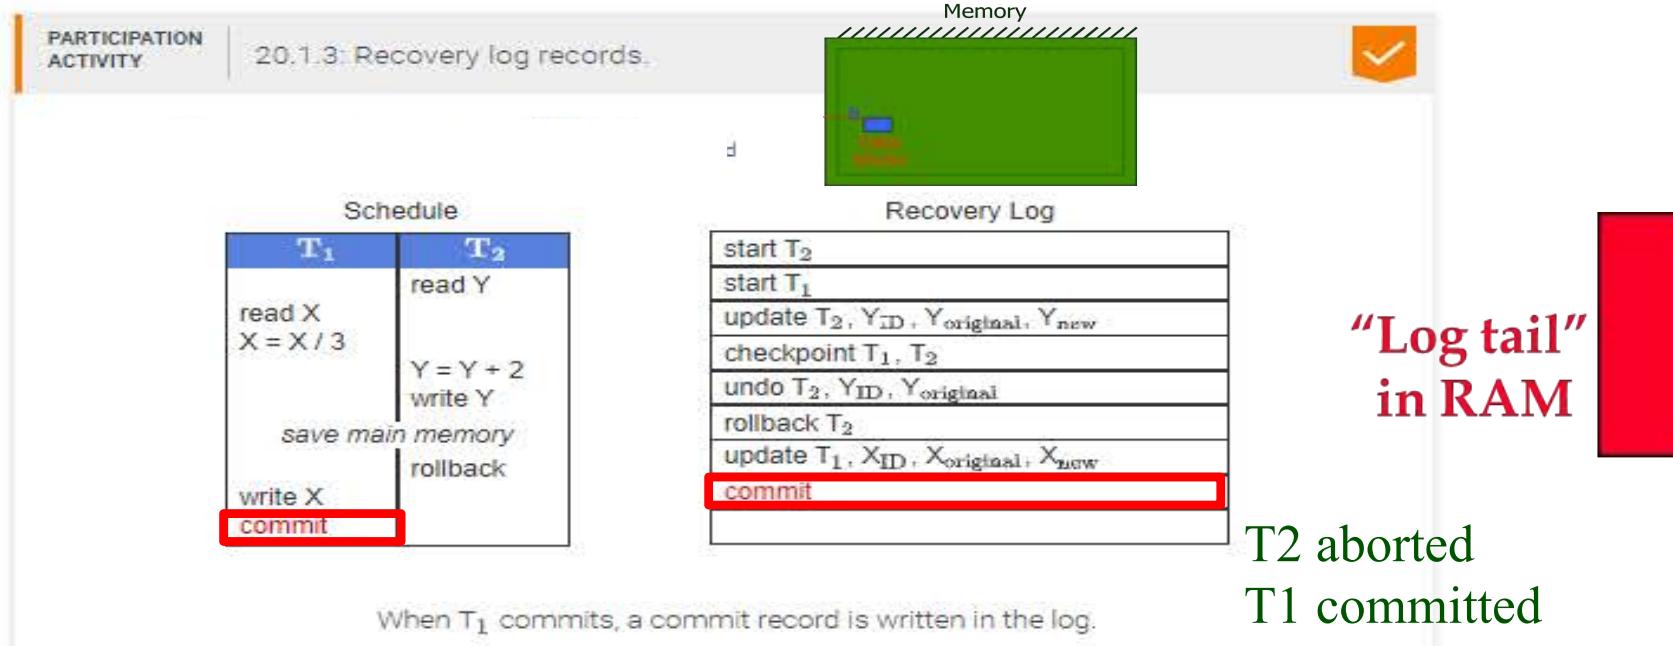
A **recovery log** is a file containing a sequential record of all database operations. The log and database are stored on different storage media so the log survives database failures. The recovery system uses the log to restore the database after a failure.

The log contains transaction and data identifiers. Transaction identifiers are assigned by the database system for internal use by the recovery and concurrency systems. Data identifiers correspond to table name, column name, and primary key value, but are compressed or encoded for efficiency.

The recovery log contains four types of records:

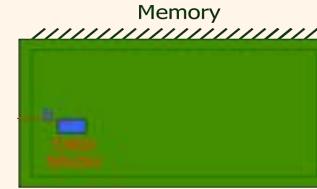
1. An **update record** indicates a transaction has changed data. Update records include the transaction identifier, the data identifier, the original data value, and the new data value. Update records may optionally track insert and delete operations.
2. A **compensation record**, also known as an **undo record**, indicates data has been restored to the original value during a rollback. Compensation records include the transaction identifier, the data identifier, and the restored (original) data value. Compensation records are necessary because the log must capture every database change, including changes executed by a rollback.
3. A **transaction record** indicates a transaction boundary. Three types of transaction records exist: start, commit, and rollback. Transaction records include the 'start', 'commit', or 'rollback' indicator and the transaction identifier.
4. A **checkpoint record** indicates that all data in main memory has been saved on storage media. When the database executes a checkpoint, transaction processing is suspended while all unsaved data and log records are written to storage media. In the event of a system failure, the recovery system reads only log records following the last checkpoint, rather than the entire file. Checkpoint records include a 'checkpoint' indicator along with the identifiers of all transactions that are active (uncommitted) at the time of the checkpoint.

The above four record types provide a complete history of database operations and enable recovery from all three failure scenarios.



# Other *Log* - Related Data Structures

Dirty Page Table  
pageID, recLSN




## Dirty Page Table:

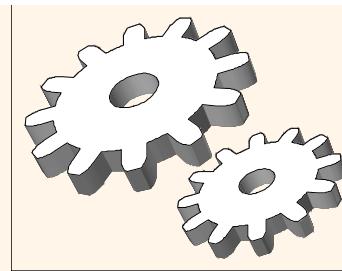
- One entry per Dirty Page in **buffer pool**.
- Contains **recLSN** or **firstLSN** -- the **LSN** of the **Log** record which first caused the **Data Page** to be Dirty.

Transaction Table		
transID	status	lastLSN

## Transaction Table:

- One entry per active Xact.
- Contains **XID**, **status** (Running/Committed/Aborted), and the **lastLSN** of the **Log** record for that **Xact lastLSN**.

# *Normal Execution of an Xact*



- ❖ Series of reads & writes, followed by Commit or Abort.
    - We will assume that write is atomic on disk.
      - In practice, additional details to deal with non-atomic writes.

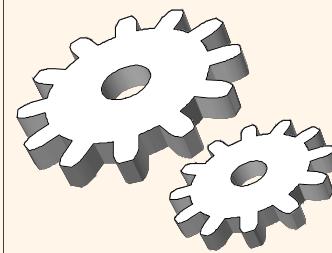
Strict 2PL allows only **Serializable Schedules**.



16

- ❖ STEAL, NO-FORCE buffer management, with **WAL**.

	No Steal	Steal
Force	Trivial	
Force		Desired



# Transaction Recovery

A Transaction Log for Transaction Recovery Examples

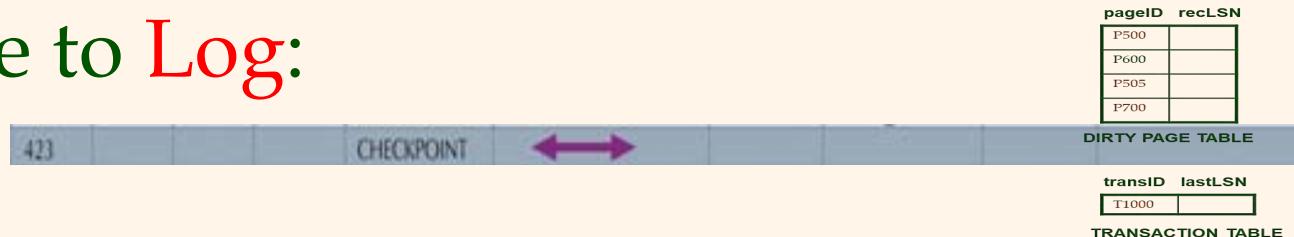
TRL ID	TRX NUM	PREV PTR	NEXT PTR	OPERATION	TABLE	ROW ID	ATTRIBUTE	BEFORE VALUE	AFTER VALUE
341	101	Null	352	START	**** Start Transaction				
352	101	341	363	UPDATE	PRODUCT	54778-2T	PROD_QOH	45	43
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	615.72	675.62
365	101	363	Null	COMMIT	**** End of Transaction			REDO D	
397	106	Null	405	START	**** Start Transaction				
405	106	397	415	INSERT	INVOICE	1009			1009,10016, ...
415	106	405	419	INSERT	LINE	1009,1			1009,1, 89-WRE-Q,1, ...
419	106	415	427	UPDATE	PRODUCT	89-WRE-Q	PROD_QOH	12	11
423				CHECKPOINT					
427	106	419	431	UPDATE	CUSTOMER	10016	CUST_BALANCE	0.00	277.55
431	106	427	457	INSERT	ACCT_TRANSACTION	10007			1007,18-JAN-2004, ...
457	106	431	Null	COMMIT	**** End of Transaction			REDO D	
521	155	Null	525	START	**** Start Transaction				
525	155	521	528	UPDATE	PRODUCT	2232/QWE	PROD_QOH	6	26
528	155	525	Null	COMMIT	**** End of Transaction			REDO D	
***** C * R * A * S * H *****									

Status of transactions before the CRASH?

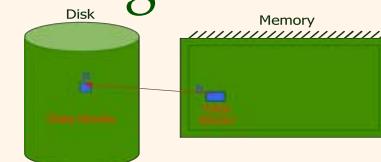


# Checkpointing

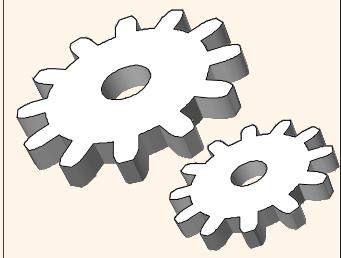
Periodically, the **DBMS** creates a checkpoint, in order to minimize the time taken to recover in the event of a system CRASH. Write to Log:



- **begin\_checkpoint** record: Indicates when **checkpoint** began.
- **end\_checkpoint** record: Contains current *Tact Table* and *Dirty Page Table*. This is a `fuzzy checkpoint':
  - Other *Xacts* continue to run; so these tables accurate only as of the time of the **begin\_checkpoint** record.
  - No attempt to force **dirty pages** to **disk**; effectiveness of **checkpoint** limited by **oldest unwritten change to a dirty page**. (So it's a good idea to periodically **flush dirty pages** to **disk**!)
- Store **LSN** of **checkpoint** record in a **safe place** (*Master Record*).



# The Big Picture: What's Stored Where



## LogRecords

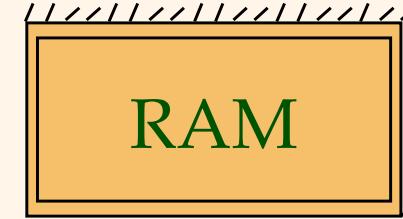
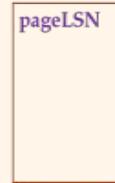
prevLSN  
XID  
type  
pageID  
length  
offset  
before-image  
after-image

**UPDATE**  
records  
only



Data pages  
each  
with a  
pageLSN

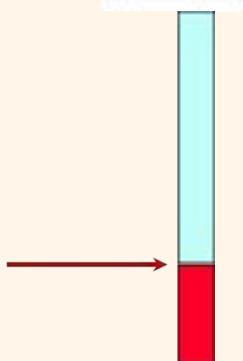
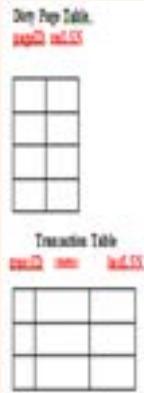
**Master record**  
(checkpoint)



**Dirty Page Table**  
firstLSN

**Xact Table**  
lastLSN  
status

**flushedLSN**



# *Simple Transaction Abort*

For now, consider an explicit ABORT of a Xact.

- No CRASH involved.

LSN	LOG
00	update: T1 writes P2
10	update: T1 writes P1
20	update: T2 writes P5
30	update: T3 writes P3
40	T3 commit
50	update: T2 writes P5
60	update: T2 writes P3
70	T2 abort

We want to “play back” the log in reverse order, **UNDO**ing UPDATES.

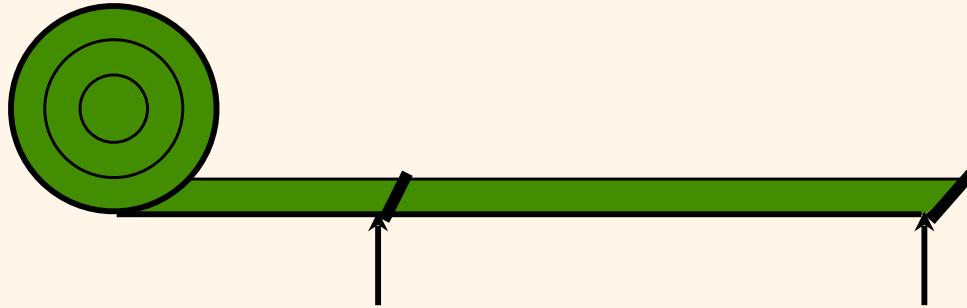
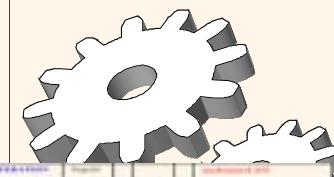
- Get lastLSN of Xact from Tact Table.

Transaction Table		
transID	status	lastLSN
T2	Abort	70
T1	Active	10

- Can follow chain of Log records backward via the prevLSN field.
  - Before starting UNDO, write an *Abort* log record (CLR).
    - For recovering from CRASH during UNDO!

LSN_N M	TRN_NU M	PREV_LSN	NEXT_PTR	OPERA TION	PagID				undonextLSN
00	T1		10	Updte	P2				
10	T1	00		Updte	P1				
20	T2		30	Updte	P3				
30	T3		40	Updte	P3				
40	T3	30		Updte	P3				
50	T2	30	60	Updte	P3				
60	T2	50	70	Updte	P3				
70	T2	60		Updte	P3				
<b>80</b>	<b>T2</b>	<b>70</b>	<b>90</b>	<b>CLR</b>	<b>P3</b>				<b>50</b>
<b>90</b>	<b>T2</b>	<b>80</b>	<b>100</b>	<b>CLR</b>	<b>P5</b>				<b>20</b>
<b>100</b>	<b>T2</b>	<b>90</b>	<b>110</b>	<b>CLR</b>	<b>P5</b>				
<b>110</b>	<b>T2</b>	<b>100</b>		<b>END</b>					

# Abort



LSN	Time	Op	Value	Op	Value	Op	Value
100	T1	Set	100	Update	100	Update	100
110	T1	Set	110	Update	110	Update	110
120	T2	Set	120	Update	120	Update	120
130	T2	Set	130	Update	130	Update	130
140	T2	Set	140	Update	140	Update	140
150	T2	Set	150	Update	150	Update	150
160	T2	Set	160	Update	160	Update	160
170	T2	Set	170	Update	170	Update	170
180	T2	Set	180	Update	180	Update	180
190	T2	Set	190	Update	190	Update	190
200	T2	Set	200	Update	200	Update	200
210	T2	Set	210	Update	210	Update	210
220	T2	Set	220	Update	220	Update	220
230	T2	Set	230	Update	230	Update	230
240	T2	Set	240	Update	240	Update	240
250	T2	Set	250	Update	250	Update	250
260	T2	Set	260	Update	260	Update	260
270	T2	Set	270	Update	270	Update	270
280	T2	Set	280	Update	280	Update	280
290	T2	Set	290	Update	290	Update	290
300	T2	Set	300	Update	300	Update	300
310	T2	Set	310	Update	310	Update	310
320	T2	Set	320	Update	320	Update	320
330	T2	Set	330	Update	330	Update	330
340	T2	Set	340	Update	340	Update	340
350	T2	Set	350	Update	350	Update	350
360	T2	Set	360	Update	360	Update	360
370	T2	Set	370	Update	370	Update	370
380	T2	Set	380	Update	380	Update	380
390	T2	Set	390	Update	390	Update	390
400	T2	Set	400	Update	400	Update	400
410	T2	Set	410	Update	410	Update	410
420	T2	Set	420	Update	420	Update	420
430	T2	Set	430	Update	430	Update	430
440	T2	Set	440	Update	440	Update	440
450	T2	Set	450	Update	450	Update	450
460	T2	Set	460	Update	460	Update	460
470	T2	Set	470	Update	470	Update	470
480	T2	Set	480	Update	480	Update	480
490	T2	Set	490	Update	490	Update	490
500	T2	Set	500	Update	500	Update	500
510	T2	Set	510	Update	510	Update	510
520	T2	Set	520	Update	520	Update	520
530	T2	Set	530	Update	530	Update	530
540	T2	Set	540	Update	540	Update	540
550	T2	Set	550	Update	550	Update	550
560	T2	Set	560	Update	560	Update	560
570	T2	Set	570	Update	570	Update	570
580	T2	Set	580	Update	580	Update	580
590	T2	Set	590	Update	590	Update	590
600	T2	Set	600	Update	600	Update	600
610	T2	Set	610	Update	610	Update	610
620	T2	Set	620	Update	620	Update	620
630	T2	Set	630	Update	630	Update	630
640	T2	Set	640	Update	640	Update	640
650	T2	Set	650	Update	650	Update	650
660	T2	Set	660	Update	660	Update	660
670	T2	Set	670	Update	670	Update	670
680	T2	Set	680	Update	680	Update	680
690	T2	Set	690	Update	690	Update	690
700	T2	Set	700	Update	700	Update	700
710	T2	Set	710	Update	710	Update	710
720	T2	Set	720	Update	720	Update	720
730	T2	Set	730	Update	730	Update	730
740	T2	Set	740	Update	740	Update	740
750	T2	Set	750	Update	750	Update	750
760	T2	Set	760	Update	760	Update	760
770	T2	Set	770	Update	770	Update	770
780	T2	Set	780	Update	780	Update	780
790	T2	Set	790	Update	790	Update	790
800	T2	Set	800	Update	800	Update	800
810	T2	Set	810	Update	810	Update	810
820	T2	Set	820	Update	820	Update	820
830	T2	Set	830	Update	830	Update	830
840	T2	Set	840	Update	840	Update	840
850	T2	Set	850	Update	850	Update	850
860	T2	Set	860	Update	860	Update	860
870	T2	Set	870	Update	870	Update	870
880	T2	Set	880	Update	880	Update	880
890	T2	Set	890	Update	890	Update	890
900	T2	Set	900	Update	900	Update	900
910	T2	Set	910	Update	910	Update	910
920	T2	Set	920	Update	920	Update	920
930	T2	Set	930	Update	930	Update	930
940	T2	Set	940	Update	940	Update	940
950	T2	Set	950	Update	950	Update	950
960	T2	Set	960	Update	960	Update	960
970	T2	Set	970	Update	970	Update	970
980	T2	Set	980	Update	980	Update	980
990	T2	Set	990	Update	990	Update	990
1000	T2	Set	1000	Update	1000	Update	1000

- ❖ Before restoring old value of a **page**, write a **CLR** (**Compensation Log Record**):
  - You continue logging while you **UNDO!!**
  - CLR has one extra field: **undonextLSN**
    - Points to the next LSN to **UNDO** (i.e. the prevLSN of the record we're currently undoing).
  - CLRs *never UNDONE* (but they might be **REDOOne** when repeating history: guarantees **Atomicity!**)
- ❖ At end of **UNDO**, write an “**END**” log record.



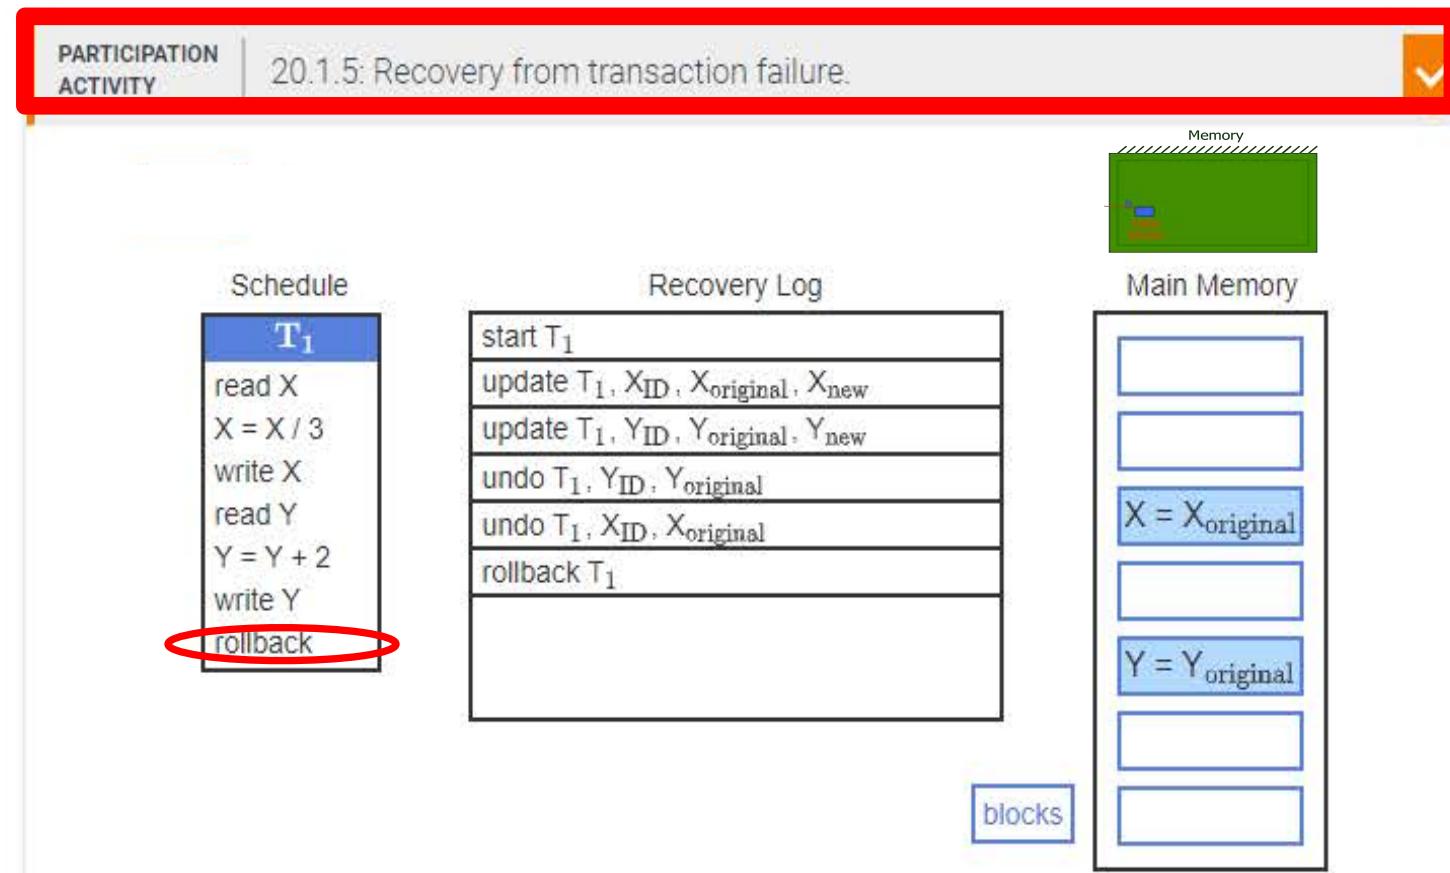
## Recovery from transaction failure

Recovery from a transaction failure is relatively simple. A database component, such as the concurrency system, instructs the recovery system to roll back transaction T. The recovery system reads the recovery log backwards, searching for update records for T. For each update record, the recovery system:

- Restores data D to original value V in the update record.
- Writes a compensation record for T, D, and V to the end of the log.

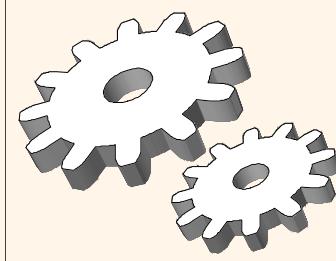
Eventually, the recovery system reads the 'start T' transaction record and writes a 'rollback T' record to the end of the log. At this point, the rollback is complete.

The recovery system reads the log backwards because data must be restored in reverse order of transaction operations.



### A Transaction Log for Transaction Recovery Examples

TRX_ID	TRX_NUM	PREV_PTR	NEXT_PTR	OPERATION	TABLE	ROW_ID	ATTRIBUTE	BEFORE_VALUE	AFTER_VALUE
341	101	Null	352	START	**** Start Transaction				
352	101	341	363	UPDATE	PRODUCT	54778-2T	PROD_QOH	45	43
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	615.73	675.62
365	101	363	Null	COMMIT	**** End of Transaction				
397	106	Null	405	START	**** Start Transaction				
405	106	397	415	INSERT	INVOICE	1009			1009,10016, -
415	106	405	419	INSERT	LINE	1009,1			1009,1, 89-WRE-Q,1, -
419	106	415	427	UPDATE	PRODUCT	89-WRE-Q	PROD_QOH	12	11
423				CHECKPOINT					
427	106	419	431	UPDATE	CUSTOMER	10016	CUST_BALANCE	0.00	277.55
431	106	427	457	INSERT	ACCT_TRANSACTION	10007			1007,18-JAN-2004, -
457	106	431	Null	COMMIT	**** End of Transaction				
521	155	Null	525	START	**** Start Transaction				
525	155	521	528	UPDATE	PRODUCT	2232/QWE	PROD_QOH	6	26
528	155	525	Null	COMMIT	**** End of Transaction				
***** C *R*A* S* H *****									



#### PARTICIPATION ACTIVITY

#### 20.1.5: Recovery from transaction failure.

Schedule

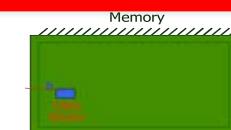
```

T1
read X
X = X / 3
write X
read Y
Y = Y + 2
write Y
rollback

```

Recovery Log

start T <sub>1</sub>



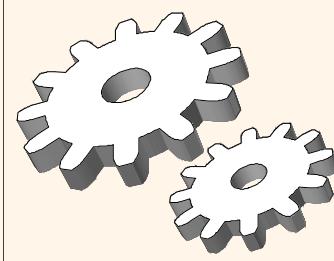
Main Memory



As the transaction executes, a start record is written to the recovery log.

### A Transaction Log for Transaction Recovery Examples

TRX_ID	TRX_NUM	PREV_PTR	NEXT_PTR	OPERATION	TABLE	ROW_ID	ATTRIBUTE	BEFORE_VALUE	AFTER_VALUE
341	101	Null	352	START	**** Start Transaction				
352	101	341	363	UPDATE	PRODUCT	54778-2T	PROD_QOH	45	43
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	615.73	675.62
365	101	363	Null	COMMIT	**** End of Transaction				
397	106	Null	405	START	**** Start Transaction				
405	106	397	415	INSERT	INVOICE	1009			1009,10016, -
415	106	405	419	INSERT	LINE	1009,1			1009,1, 89-WRE-Q,1, -
419	106	415	427	UPDATE	PRODUCT	89-WRE-Q	PROD_QOH	12	11
423				CHECKPOINT					
427	106	419	431	UPDATE	CUSTOMER	10016	CUST_BALANCE	0.00	277.55
431	106	427	457	INSERT	ACCT_TRANSACTION	10007			1007,18-JAN-2004, -
457	106	431	Null	COMMIT	**** End of Transaction				
521	155	Null	525	START	**** Start Transaction				
525	155	521	528	UPDATE	PRODUCT	2232/QWE	PROD_QOH	6	26
528	155	525	Null	COMMIT	**** End of Transaction				
***** C *R*A* S* H *****									



#### PARTICIPATION ACTIVITY

##### 20.1.5: Recovery from transaction failure.

###### Schedule

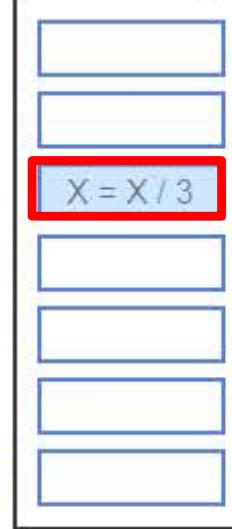
```

T1
read X
X = X / 3
write X
read Y
Y = Y + 2
write Y
rollback
    
```

###### Recovery Log

start T<sub>1</sub>  
**update T<sub>1</sub>, X<sub>ID</sub>, X<sub>original</sub>, X<sub>new</sub>**

###### Main Memory

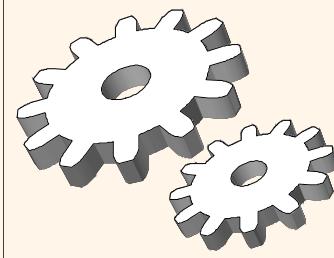


blocks

Write statements generate update records in the log and save data in main memory.

### A Transaction Log for Transaction Recovery Examples

TRX_ID	TRX_NUM	PREV_PTR	NEXT_PTR	OPERATION	TABLE	ROW_ID	ATTRIBUTE	BEFORE_VALUE	AFTER_VALUE
341	101	Null	352	START	**** Start Transaction				
352	101	341	363	UPDATE	PRODUCT	54778-2T	PROD_QOH	45	43
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	615.73	675.62
365	101	363	Null	COMMIT	**** End of Transaction				
397	106	Null	405	START	**** Start Transaction				
405	106	397	415	INSERT	INVOICE	1009			1009,10016, -
415	106	405	419	INSERT	LINE	1009,1			1009,1, 89-WRE-Q,1, -
419	106	415	427	UPDATE	PRODUCT	89-WRE-Q	PROD_QOH	12	11
423				CHECKPOINT					
427	106	419	431	UPDATE	CUSTOMER	10016	CUST_BALANCE	0.00	277.55
431	106	427	457	INSERT	ACCT_TRANSACTION	10007			1007,18-JAN-2004, -
457	106	431	Null	COMMIT	**** End of Transaction				
521	155	Null	525	START	**** Start Transaction				
525	155	521	528	UPDATE	PRODUCT	2232/QWE	PROD_QOH	6	26
528	155	525	Null	COMMIT	**** End of Transaction				
***** C *R*A* S* H *****									



#### PARTICIPATION ACTIVITY

##### 20.1.5: Recovery from transaction failure.

###### Schedule

```

T1
read X
X = X / 3
write X
read Y
Y = Y + 2
write Y
rollback
    
```

###### Recovery Log

start T <sub>1</sub>
update T <sub>1</sub> , X <sub>ID</sub> , X <sub>original</sub> , X <sub>new</sub>
<b>update T<sub>1</sub>, Y<sub>ID</sub>, Y<sub>original</sub>, Y<sub>new</sub></b>

###### Main Memory

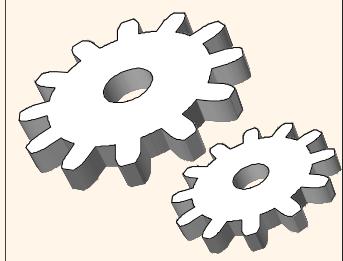
X = X / 3
<b>Y = Y + 2</b>

blocks

Write statements generate update records in the log and save data in main memory.

### A Transaction Log for Transaction Recovery Examples

TRX_ID	TRX_NUM	PREV_PTR	NEXT_PTR	OPERATION	TABLE	ROW_ID	ATTRIBUTE	BEFORE_VALUE	AFTER_VALUE
341	101	Null	352	START	**** Start Transaction				
352	101	341	363	UPDATE	PRODUCT	54778-2T	PROD_QOH	45	43
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	615.73	675.62
365	101	363	Null	COMMIT	**** End of Transaction				
397	106	Null	405	START	**** Start Transaction				
405	106	397	415	INSERT	INVOICE	1009			1009,10016, -
415	106	405	419	INSERT	LINE	1009,1			1009,1, 89-WRE-Q,1, -
419	106	415	427	UPDATE	PRODUCT	89-WRE-Q	PROD_QOH	12	11
423				CHECKPOINT					
427	106	419	431	UPDATE	CUSTOMER	10016	CUST_BALANCE	0.00	277.55
431	106	427	457	INSERT	ACCT_TRANSACTION	10007			1007,18-JAN-2004, -
457	106	431	Null	COMMIT	**** End of Transaction				
521	155	Null	525	START	**** Start Transaction				
525	155	521	528	UPDATE	PRODUCT	2232/QWE	PROD_QOH	6	26
528	155	525	Null	COMMIT	**** End of Transaction				
***** C *R*A* S* H *****									



#### PARTICIPATION ACTIVITY

20.1.5: Recovery from transaction failure.

Schedule

```
T1
read X
X = X / 3
write X
read Y
Y = Y + 2
write Y
rollback
```

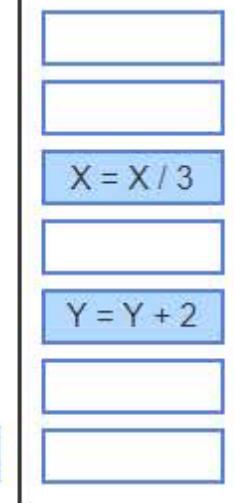
Recovery Log

start T <sub>1</sub>
update T <sub>1</sub> , X <sub>ID</sub> , X <sub>original</sub> , X <sub>new</sub>
update T <sub>1</sub> , Y <sub>ID</sub> , Y <sub>original</sub> , Y <sub>new</sub>

Memory



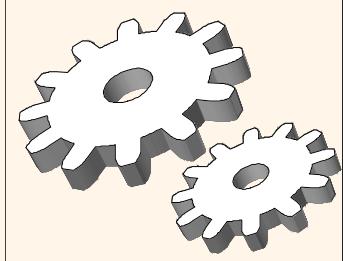
Main Memory



When the system or programmer initiates rollback, the recovery system reads the log in reverse.

### A Transaction Log for Transaction Recovery Examples

TRX_ID	TRX_NUM	PREV_PTR	NEXT_PTR	OPERATION	TABLE	ROW_ID	ATTRIBUTE	BEFORE_VALUE	AFTER_VALUE
341	101	Null	352	START	**** Start Transaction				
352	101	341	363	UPDATE	PRODUCT	54778-2T	PROD_QOH	45	43
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	615.73	675.62
365	101	363	Null	COMMIT	**** End of Transaction				
397	106	Null	405	START	**** Start Transaction				
405	106	397	415	INSERT	INVOICE	1009			1009,10016, -
415	106	405	419	INSERT	LINE	1009,1			1009,1, 89-WRE-Q,1, -
419	106	415	427	UPDATE	PRODUCT	89-WRE-Q	PROD_QOH	12	11
423				CHECKPOINT					
427	106	419	431	UPDATE	CUSTOMER	10016	CUST_BALANCE	0.00	277.55
431	106	427	457	INSERT	ACCT_TRANSACTION	10007			1007,18-JAN-2004, -
457	106	431	Null	COMMIT	**** End of Transaction				
521	155	Null	525	START	**** Start Transaction				
525	155	521	528	UPDATE	PRODUCT	2232/QWE	PROD_QOH	6	26
528	155	525	Null	COMMIT	**** End of Transaction				
***** C *R*A* S* H *****									



#### PARTICIPATION ACTIVITY

20.1.5: Recovery from transaction failure.

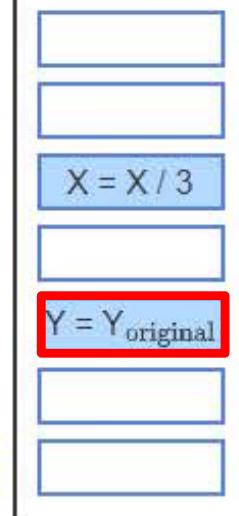
Schedule

T <sub>1</sub>
read X
X = X / 3
write X
read Y
Y = Y + 2
write Y
rollback

Recovery Log

start T <sub>1</sub>
update T <sub>1</sub> , X <sub>ID</sub> , X <sub>original</sub> , X <sub>new</sub>
update T <sub>1</sub> , Y <sub>ID</sub> , Y <sub>original</sub> , Y <sub>new</sub>
undo T <sub>1</sub> , Y <sub>ID</sub> , Y <sub>original</sub>

Main Memory

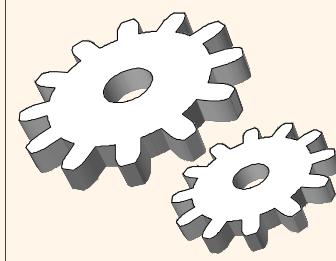


blocks

Y is restored to original value in main memory, and a compensation record is written to the log.

### A Transaction Log for Transaction Recovery Examples

TRX_ID	TRX_NUM	PREV_PTR	NEXT_PTR	OPERATION	TABLE	ROW_ID	ATTRIBUTE	BEFORE_VALUE	AFTER_VALUE
341	101	Null	352	START	**** Start Transaction				
352	101	341	363	UPDATE	PRODUCT	54778-2T	PROD_QOH	45	43
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	615.73	675.62
365	101	363	Null	COMMIT	**** End of Transaction				
397	106	Null	405	START	**** Start Transaction				
405	106	397	415	INSERT	INVOICE	1009			1009,10016, -
415	106	405	419	INSERT	LINE	1009,1			1009,1, 89-WRE-Q,1, -
419	106	415	427	UPDATE	PRODUCT	89-WRE-Q	PROD_QOH	12	11
423				CHECKPOINT					
427	106	419	431	UPDATE	CUSTOMER	10016	CUST_BALANCE	0.00	277.55
431	106	427	457	INSERT	ACCT_TRANSACTION	10007			1007,18-JAN-2004, -
457	106	431	Null	COMMIT	**** End of Transaction				
521	155	Null	525	START	**** Start Transaction				
525	155	521	528	UPDATE	PRODUCT	2232/QWE	PROD_QOH	6	26
528	155	525	Null	COMMIT	**** End of Transaction				
***** C *R*A* S* H *****									



#### PARTICIPATION ACTIVITY

20.1.5: Recovery from transaction failure.

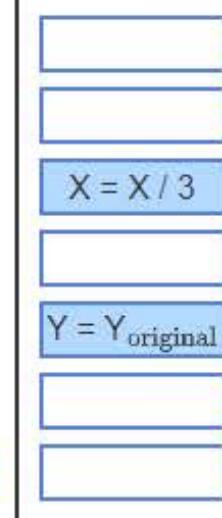
Schedule

T <sub>1</sub>
read X
X = X / 3
write X
read Y
Y = Y + 2
write Y
rollback

Recovery Log

start T <sub>1</sub>
update T <sub>1</sub> , X <sub>ID</sub> , X <sub>original</sub> , X <sub>new</sub>
update T <sub>1</sub> , Y <sub>ID</sub> , Y <sub>original</sub> , Y <sub>new</sub>
undo T <sub>1</sub> , Y <sub>ID</sub> , Y <sub>original</sub>

Main Memory

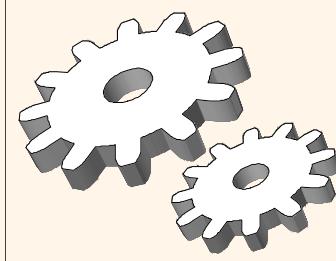


blocks

Y is restored to original value in main memory, and a compensation record is written to the log.

### A Transaction Log for Transaction Recovery Examples

TRX_ID	TRX_NUM	PREV_PTR	NEXT_PTR	OPERATION	TABLE	ROW_ID	ATTRIBUTE	BEFORE_VALUE	AFTER_VALUE
341	101	Null	352	START	**** Start Transaction				
352	101	341	363	UPDATE	PRODUCT	54778-2T	PROD_QOH	45	43
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	615.73	675.62
365	101	363	Null	COMMIT	**** End of Transaction				
397	106	Null	405	START	**** Start Transaction				
405	106	397	415	INSERT	INVOICE	1009			1009,10016, -
415	106	405	419	INSERT	LINE	1009,1			1009,1, 89-WRE-Q,1, -
419	106	415	427	UPDATE	PRODUCT	89-WRE-Q	PROD_QOH	12	11
423				CHECKPOINT					
427	106	419	431	UPDATE	CUSTOMER	10016	CUST_BALANCE	0.00	277.55
431	106	427	457	INSERT	ACCT_TRANSACTION	10007			1007,18-JAN-2004, -
457	106	431	Null	COMMIT	**** End of Transaction				
521	155	Null	525	START	**** Start Transaction				
525	155	521	528	UPDATE	PRODUCT	2232/QWE	PROD_QOH	6	26
528	155	525	Null	COMMIT	**** End of Transaction				
***** C *R*A* S* H *****									



#### PARTICIPATION ACTIVITY

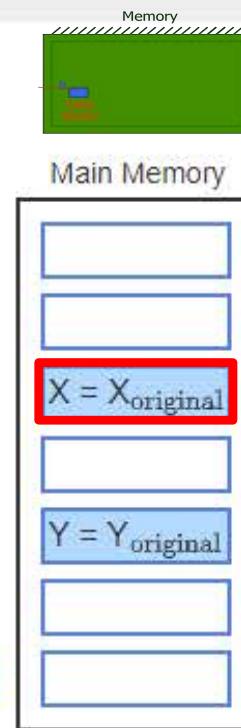
##### 20.1.5: Recovery from transaction failure.

Schedule

<b>T<sub>1</sub></b>
read X
X = X / 3
write X
read Y
Y = Y + 2
write Y
<b>rollback</b>

Recovery Log

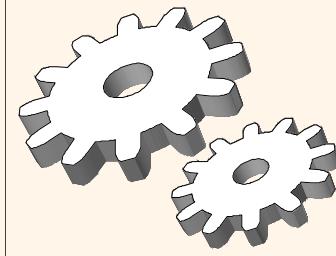
start T <sub>1</sub>
update T <sub>1</sub> , X <sub>ID</sub> , X <sub>original</sub> , X <sub>new</sub>
update T <sub>1</sub> , Y <sub>ID</sub> , Y <sub>original</sub> , Y <sub>new</sub>
undo T <sub>1</sub> , Y <sub>ID</sub> , Y <sub>original</sub>
<b>undo T<sub>1</sub>, X<sub>ID</sub>, X<sub>original</sub></b>



X is restored to original value in main memory, and a compensation record is written to the log.

### A Transaction Log for Transaction Recovery Examples

TRX_ID	TRX_NUM	PREV_PTR	NEXT_PTR	OPERATION	TABLE	ROW_ID	ATTRIBUTE	BEFORE_VALUE	AFTER_VALUE
341	101	Null	352	START	**** Start Transaction				
352	101	341	363	UPDATE	PRODUCT	54778-2T	PROD_QOH	45	43
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	615.73	675.62
365	101	363	Null	COMMIT	**** End of Transaction				
397	106	Null	405	START	**** Start Transaction				
405	106	397	415	INSERT	INVOICE	1009			1009,10016, -
415	106	405	419	INSERT	LINE	1009,1			1009,1, 89-WRE-Q,1, -
419	106	415	427	UPDATE	PRODUCT	89-WRE-Q	PROD_QOH	12	11
423				CHECKPOINT					
427	106	419	431	UPDATE	CUSTOMER	10016	CUST_BALANCE	0.00	277.55
431	106	427	457	INSERT	ACCT_TRANSACTION	10007			1007,18-JAN-2004, -
457	106	431	Null	COMMIT	**** End of Transaction				
521	155	Null	525	START	**** Start Transaction				
525	155	521	528	UPDATE	PRODUCT	2232/QWE	PROD_QOH	6	26
528	155	525	Null	COMMIT	**** End of Transaction				
***** C *R*A* S* H *****									



#### PARTICIPATION ACTIVITY

20.1.5: Recovery from transaction failure.



Schedule

T <sub>1</sub>
read X
X = X / 3
write X
read Y
Y = Y + 2
write Y
rollback

Recovery Log

start T <sub>1</sub>
update T <sub>1</sub> , X <sub>ID</sub> , X <sub>original</sub> , X <sub>new</sub>
update T <sub>1</sub> , Y <sub>ID</sub> , Y <sub>original</sub> , Y <sub>new</sub>
undo T <sub>1</sub> , Y <sub>ID</sub> , Y <sub>original</sub>
undo T <sub>1</sub> , X <sub>ID</sub> , X <sub>original</sub>
rollback T <sub>1</sub>

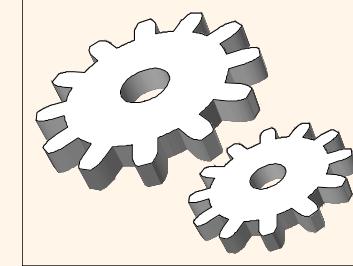
Main Memory

X = X <sub>original</sub>
Y = Y <sub>original</sub>

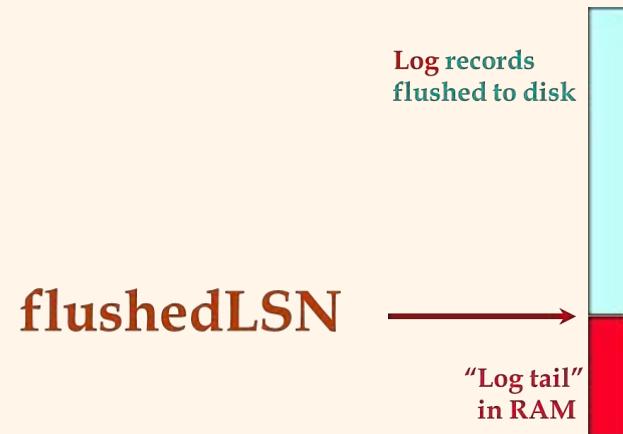
blocks

Finally, the rollback record is written to the log, and rollback is complete

# *Transaction Commit*

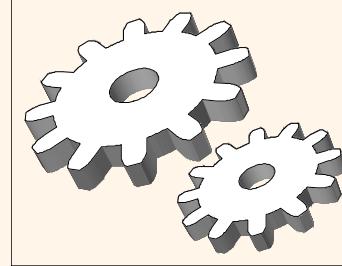


- ❖ Write **COMMIT** record to **Log**.
  - ❖ All **Log** records up to **Xact's lastLSN** are **flushed**.
    - Guarantees that **flushedLSN  $\geq$  lastLSN**.
    - Note that **Log** flushes are sequential, synchronous writes to disk.
    - Many **Log** records per **Log Page**.
  - ❖ **COMMIT()** returns.  
flushedLSN \_\_\_\_\_  
“Log ta  
in RA
  - ❖ Write **END** record to **Log**.



TRX_ID	TRX_NUM	PREV_PTR	NEXT_PTR	OPERATION	TABLE	ROW_ID	ATTRIBUTE	BEFORE_VALUE	AFTER_VALUE
341	101	Null	352	START	*****Start Transaction				
352	101	341	363	UPDATE	PRODUCT	1558-QW1	PROD_QOH	25	23
363	101	352	365	UPDATE	CUSTOMER	10011	CUST_BALANCE	525.75	615.73
365	101	363	Null	COMMIT	***** End of Transaction				

# Crash Recovery: Big Picture

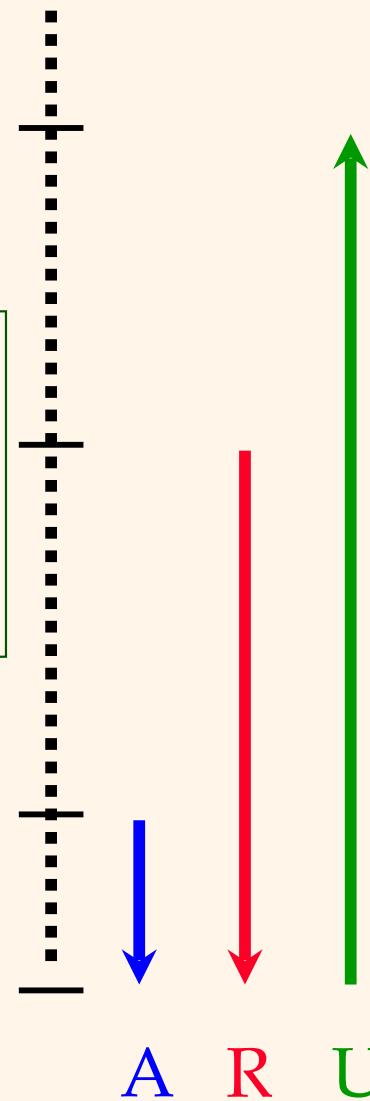


LAST log rec.  
lastLSN of  
Tact Table  
active at crash

FIRST log rec.  
recLSN in  
Dirty Page  
Table after  
Analysis

Checkpoint

CRASH

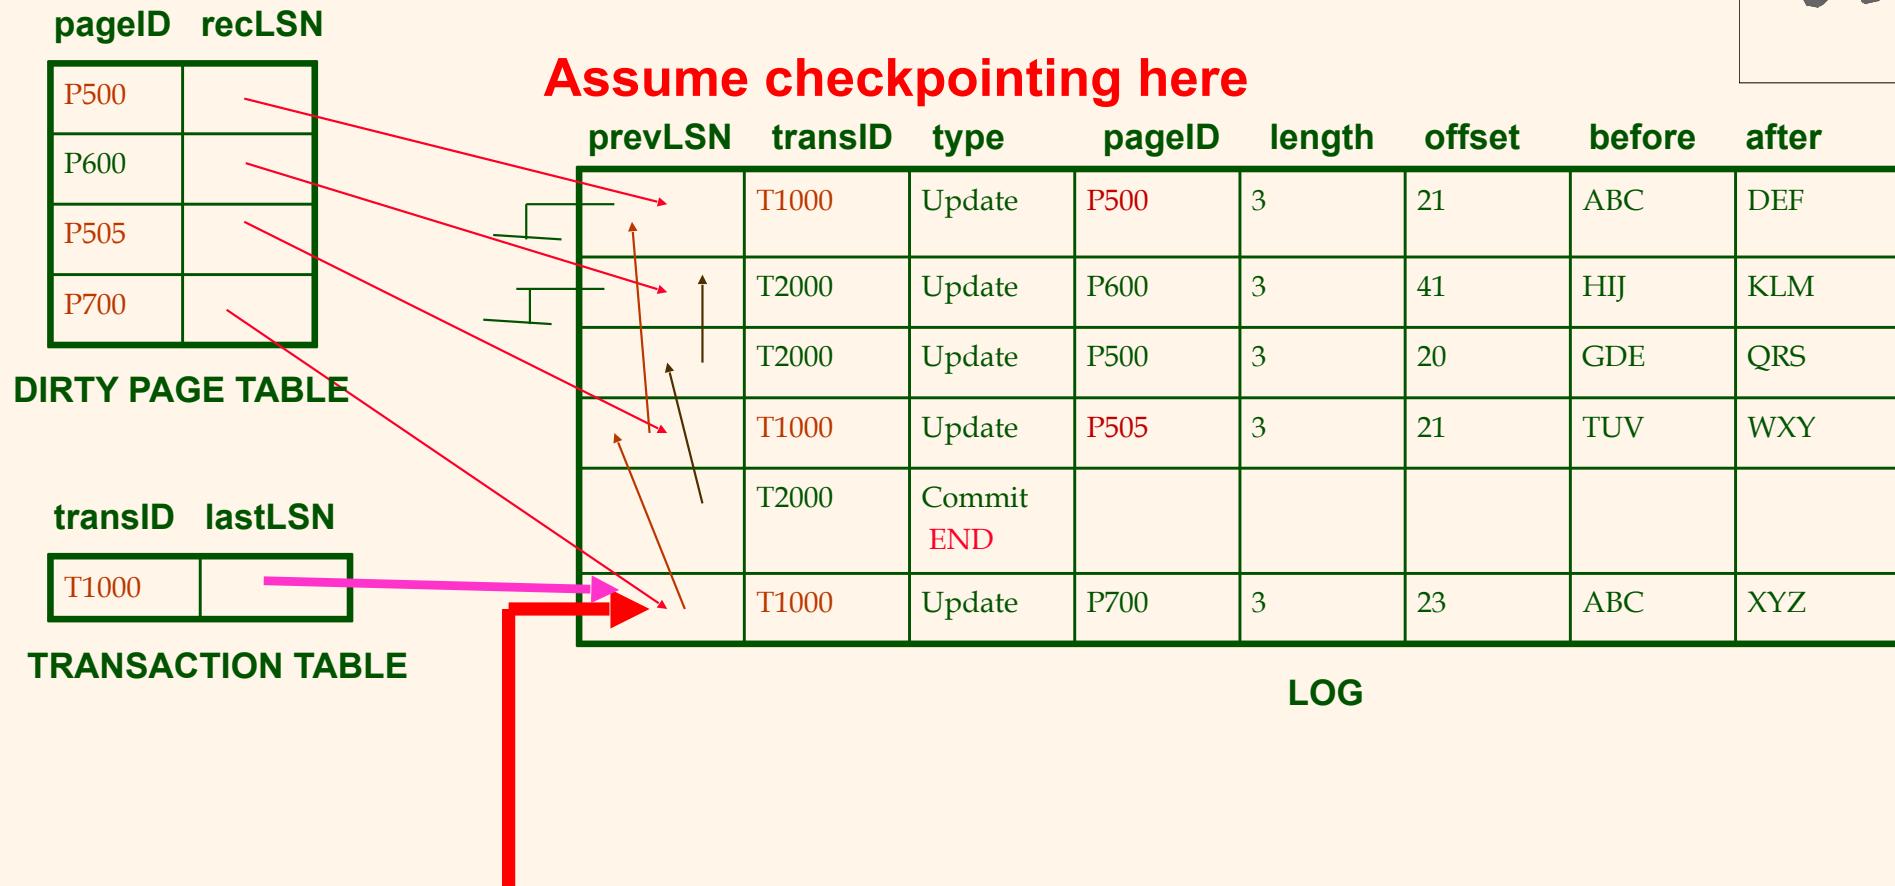
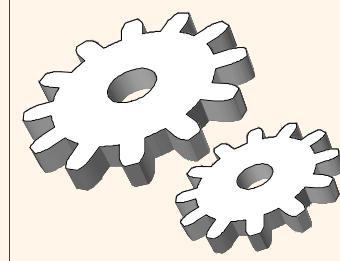


Start from a **Checkpoint** (found via **Master Record**).

Three Phases. Need to:

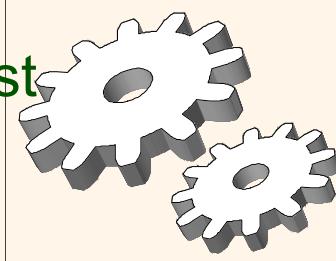
- Figure out which Xacts **COMMIT**ted since **Checkpoint**, which **FAILED** (**Analysis**).
- **REDO all** actions.  
    ?(repeat history) (**R**).
- **UNDO** effects of **FAILED Xacts** (**U**).

# Recovery: The Analysis Phase



**System crashes before this log entry is written to stable storage!**

# Recovery – DIRTY PAGE & TRANSACTION TABLES are lost



pageID recLSN

--	--

DIRTY PAGE TABLE

transID lastLSN

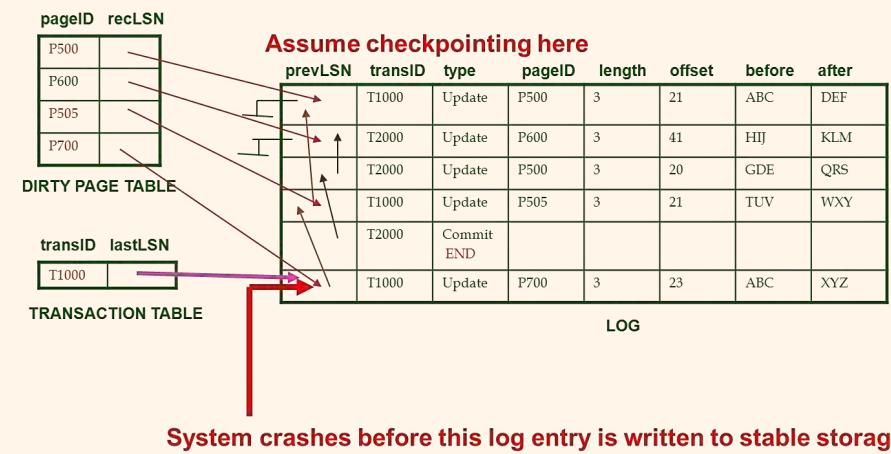
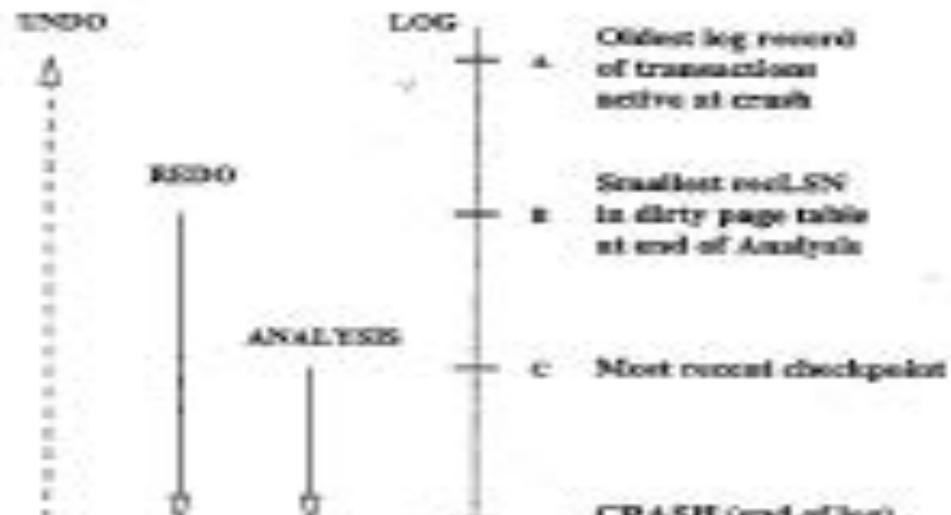
--	--

TRANSACTION TABLE

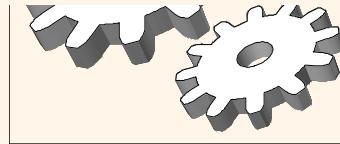
Assume checkpointing here

prevLSN	transID	type	pageID	length	offset	before	after
	T1000	Update	P500	3	21	ABC	DEF
	T2000	Update	P600	3	41	HIJ	KLM
	T2000	Update	P500	3	20	GDE	QRS
	T1000	Update	P505	3	21	TUV	WXY
	T2000	Commit					

LOG



# Recovery - Analysis — DIRTY PAGE & TRANSACTION TABLES are lost



pageID recLSN

P500	
P600	
P505	

DIRTY PAGE TABLE

Assume checkpointing here

prevLSN	transID	type	pageID	length	offset	before	after
	T1000	Update	P500	3	21	ABC	DEF
	T2000	Update	P600	3	41	HIJ	KLM
	T2000	Update	P500	3	20	GDE	QRS
	T1000	Update	P505	3	21	TUV	WXY
	T2000	Commit END					

transID lastLSN

T1000	
T2000	

TRANSACTION TABLE

LOG

P500	
P600	
P505	
P700	

DIRTY PAGE TABLE

TRANSACTION TABLE

Assume checkpointing here

prevLSN	transID	type	pageID	length	offset	before	after
	T1000	Update	P500	3	21	ABC	DEF
	T2000	Update	P600	3	41	HIJ	KLM
	T2000	Update	P500	3	20	GDE	QRS
	T1000	Update	P505	3	21	TUV	WXY
	T2000	Commit END					
	T1000	Update	P700	3	23	ABC	XYZ

LOG

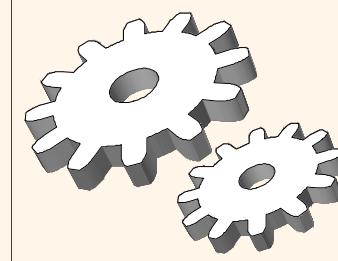
System crashes before this log entry is written to stable storage!

This log record will cause T2000 to be removed from the TRANSACTION TABLE!

Analysis phase recognizes that the only active transaction at the time of the crash is T1000, with lastLSN equal to the LSN of the fourth record in this figure.

NOTE that page P600 is NOT dirty, yet it is included in the DIRTY PAGE TABLE!

# Recovery - Redo – D



pageID recLSN

P500	
P600	
P505	

DIRTY PAGE TABLE

Assume checkpointing here

prevLSN	transID	type	pageID	length	offset	before	after
	T1000	Update	P500	3	21	ABC	DEF
T2000	Update	P600	3	41	HIJ	KLM	
T2000	Update	P500	3	20	GDE	QRS	
	T1000	Update	P505	3	21	TUV	WXY
	T2000	Commit					

transID lastLSN

T1000	

TRANSACTION TABLE

LOG

Smallest recLSN of all pages in the DIRTY PAGE TABLE is the **first record**.

It identifies the oldest update that may not have been written to disk prior to the **CRASH**.

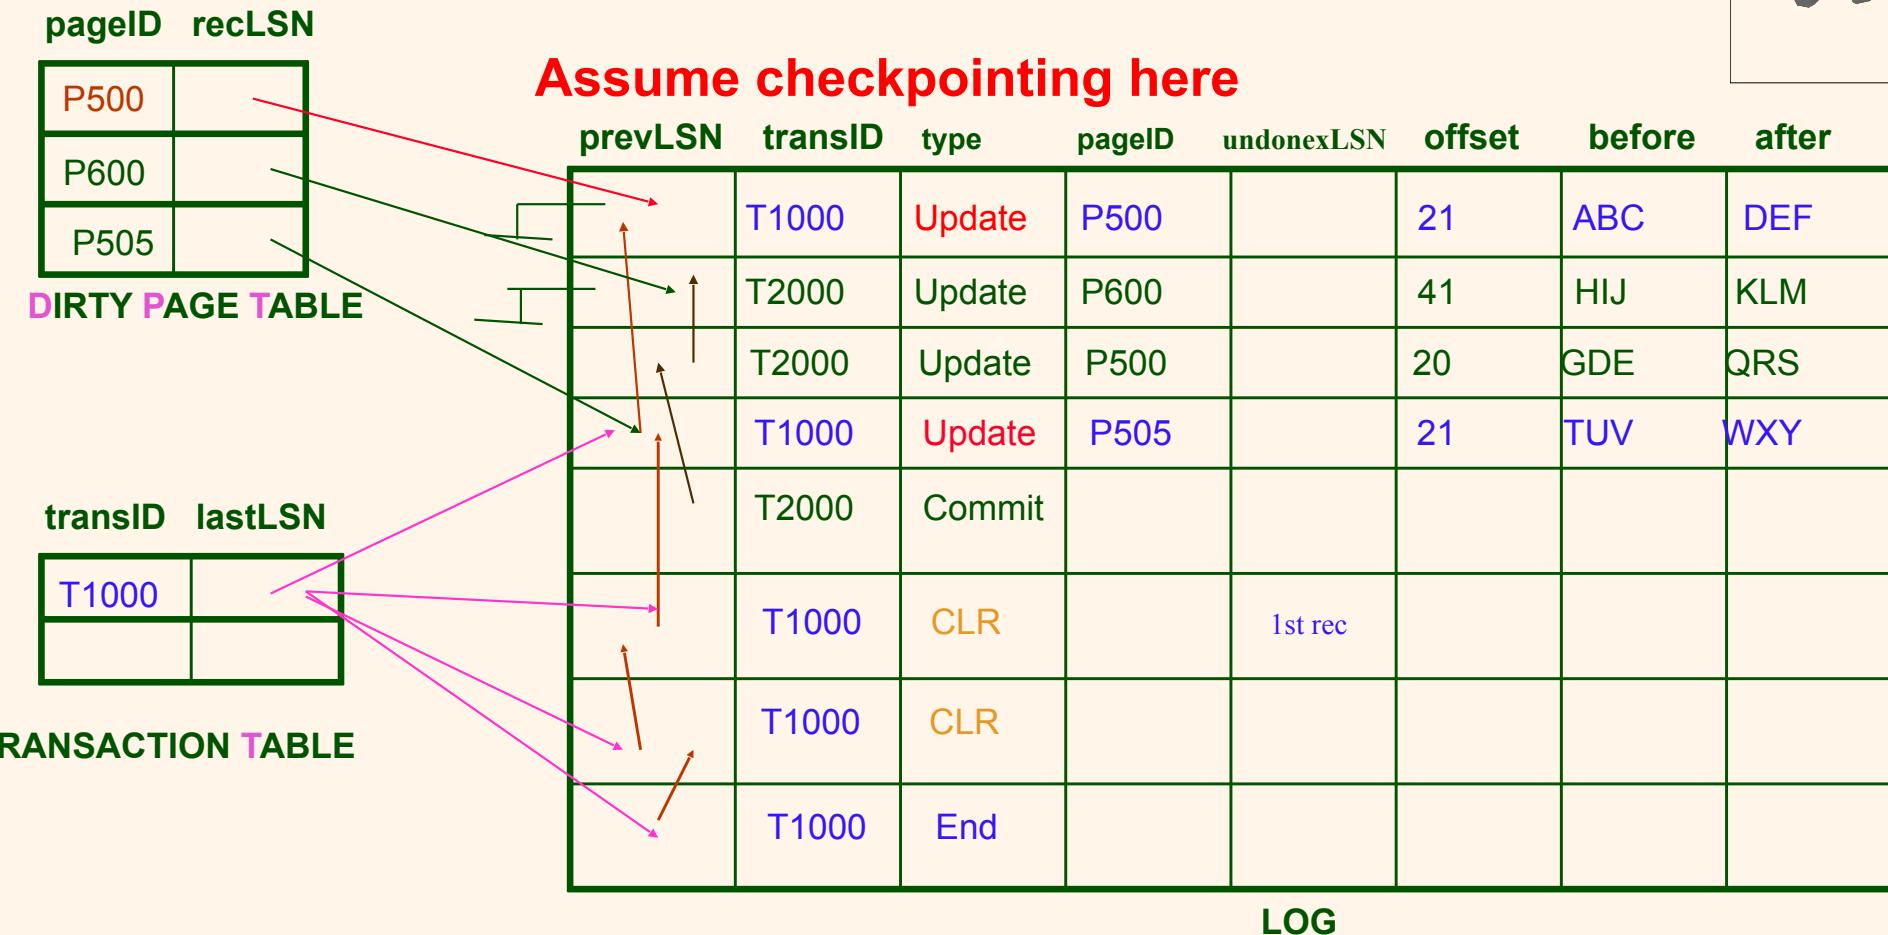
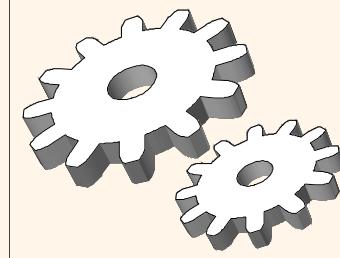
**REDO** fetches page P500 and compares the LSN of this log record with the **page LSN** on the page and , because we assumed that this page was not written to disk before the crash, finds that the **pageLSN** is less. The update is therefore reapplied; bytes 21 though 23 are changed to ‘DEF’, and the **pageLSN** is set to the LSN of this **UPDATE** log record.

Starting from there, **REDO** scans forward until the end of the **Log**. **REDO** will check to see if this logged action needs to be redone.

**The second log record**. The affected page, P600, is fetched and the **pageLSN** is compared to the LSN of the update log record. In this case, because we assumed that **P600 was written to disk before the crash**, they are equal, and the update does not have to be redone.

The remaining log records are processed similarly, bringing the system back to the exact state it was at the time of the **CRASH**.

# Recovery - Undo – A



The **UNDO** phase scans backward from the end of the log. The goal of this phase is to undo the actions of all transactions active at the time of the crash, that is, to effectively **ABORT** them.

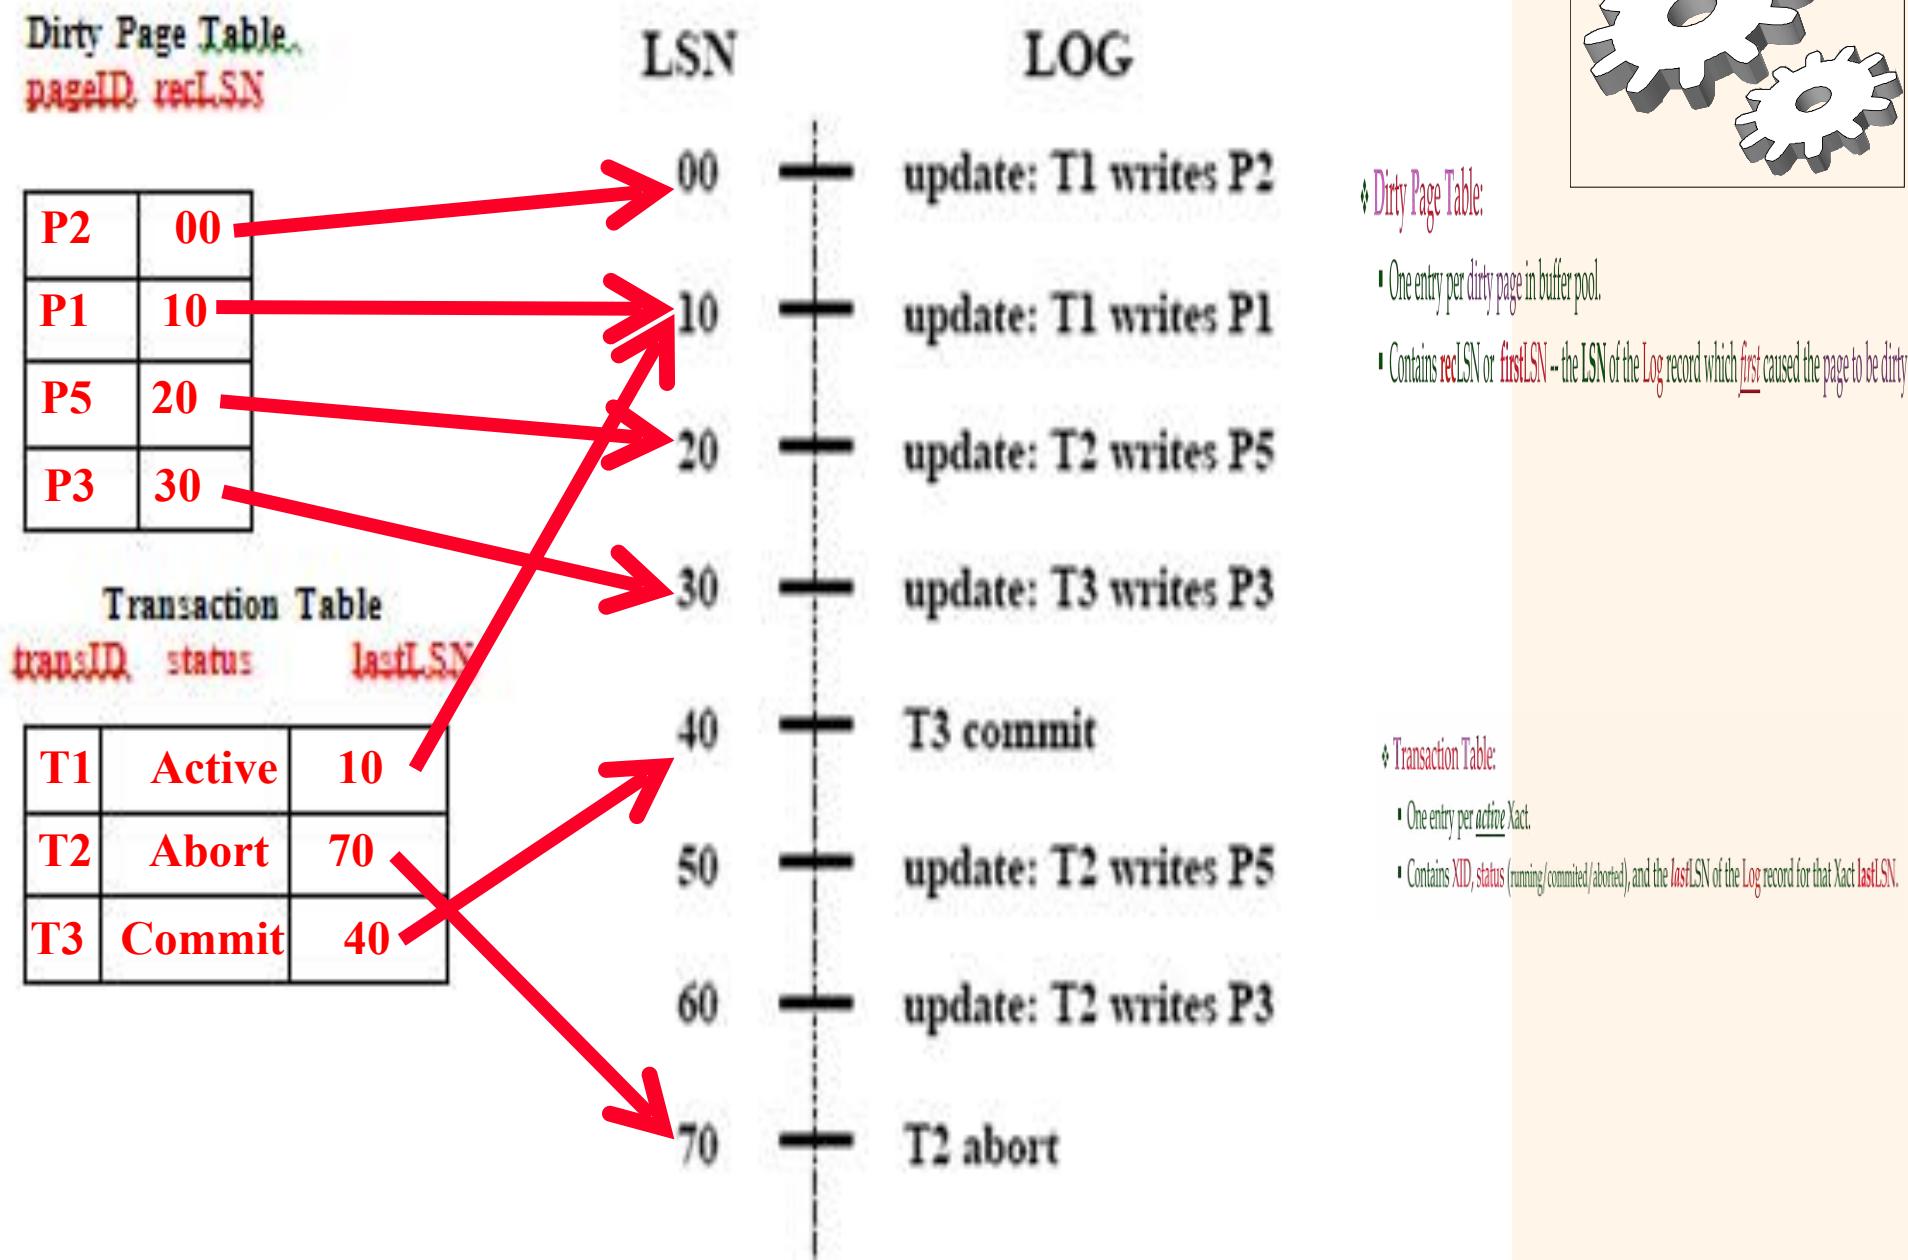
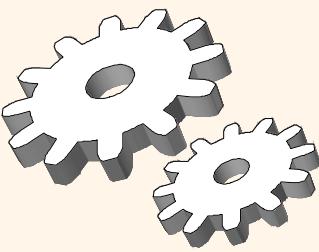
In our case, T1000 is the only transaction that needs to be aborted. From the **TRANSACTION TABLE** we get the LSN of its **LAST** log record, which is the fourth update log record. The update is undone, and a **CLR record** is written with **undoNextLSN** equal to the **LSN** of the first log record. The next record to be undone for transaction T1000 is the **first log record**. After this is undone, a **CLR log** record and an **END** log record for T1000 are written, and the **UNDO** phase is complete.

3.

**Consider the execution shown below:**

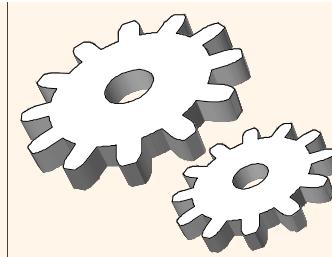
LSN	LOG
00	update: T1 writes P2
10	update: T1 writes P1
20	update: T2 writes P5
30	update: T3 writes P3
40	T3 commit
50	update: T2 writes P5
60	update: T2 writes P3
70	T2 abort

1. Extend the figure to show prevLSN and undonextLSN values.
  2. Describe the actions taken to rollback transaction  $T_2$ .
  3. Show the log after  $T_2$  is rolled back, including all prevLSN and undonextLSN values in log records.



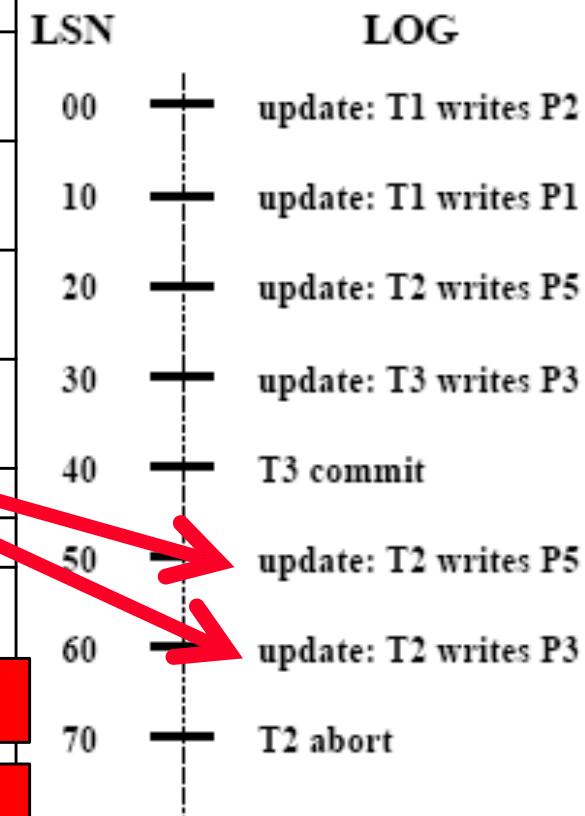
Fill in these DSs!

# The Transaction Log

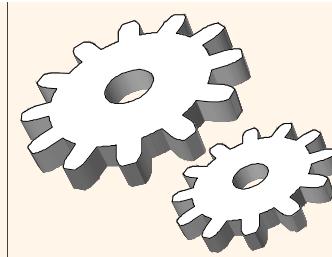


1. Extend the figure to show undonextLSN values for T2.

LSN	TRX_NUM	PREV_LSN	NEXT_PTR	OPERATION	PageID			undonextLSN
00	T1		10	Update	P2			
10	T1	00		Update	P1			
20	T2		50	Update	P5			
30	T3		40	Update	P3			
40	T3	30		Commit				
50	T2	20	60	Update	P5		20	
60	T2	50	70	Update	P3		50	
70	T2	60		Abort				
80		70	90	CER	10			
90		80	100	CER	10			
100		90	110	CER	10			
110		100	END					

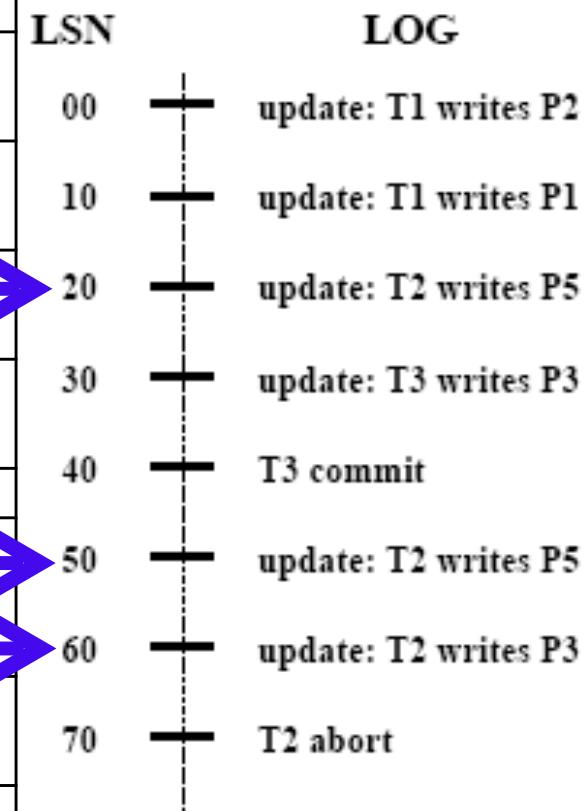


# The Transaction Log



2. Describe the actions taken to **ROLLBACK T2**.

LSN	TRX_NUM	PREV_LSN	NEXT_PTR	OPERATION	PageID			undonextLSN
00	T1		10	Update	P2			
10	T1	00		Update	P1			
20	T2		50	Update	P5			
30	T3		40	Update	P3	UNDO		
40	T3	30		Commit				
50	T2	20	60	Update	P5	UNDO 20		
60	T2	50	70	Update	P3	UNDO 50		
70	T2	60		Abort				
80		70	90	CER	10			
90		80	100	CER	10			
100		90	110	CER	10			
110		100		LND				

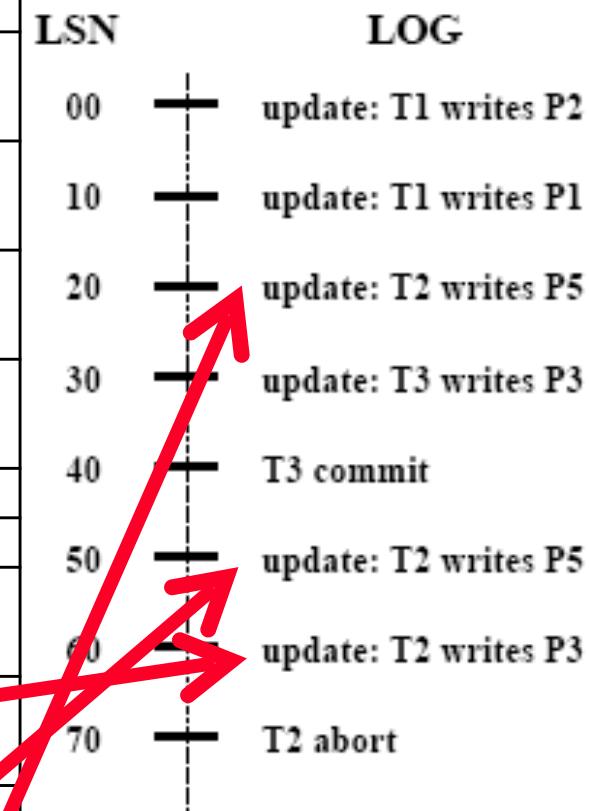




# The Transaction Log

3. Show the Log after T2 is rolled back, including prevLSN and undonextLSN values in Log records. For OPERATION enter CLR. After that add END Log record.

LSN	TRX_NUM	PREV_LSN	NEXT_PTR	OPERATION	PageID			undonextLSN
00	T1		10	Update	P2			
10	T1	00		Update	P1			
20	T2		50	Update	P5			
30	T3		40	Update	P3			
40	T3	30		Commit				
50	T2	20	60	Update	P5			20
60	T2	50	70	Update	P3			50
70	T2	60	80	Abort				
80								
90								
100				CLR	10			
110				CLR	10			
110				END				



Consider the execution shown below.

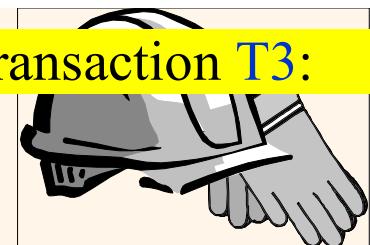
a. Rewrite with more details the log entries till LSN 120:

LSN	LOG
00	begin_checkpoint
10	end_checkpoint
20	update: T1 writes P1
30	update: T3 writes P4
40	update: T2 writes P2
50	update: T3 writes P3
60	T2 commit
70	update: T3 writes P2
80	T2 end
90	update: T3 writes P1
100	update: T1 writes P5
110	update: T3 writes P2
120	T3 abort
<del>CRASH, RESTART</del>	

LSN	Type	prevLSN	undoneLSN
00	begin_checkpoint		
10	end_checkpoint		
20	update		
30	update		
40	update		
50	update	30	30
60	commit	40	
70	update	50	50
80	end	60	
90	update	70	70
100	update	20	20
110	update	90	90
120	abort	110	

Consider the execution shown below.

b. Describe the actions taken to rollback transaction T3:

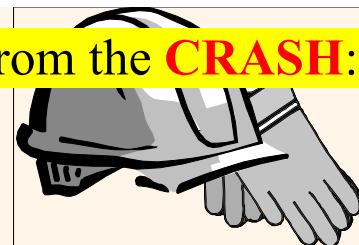


LSN	LOG
00	begin_checkpoint
10	end_checkpoint
20	update: T1 writes P1
30	update: T3 writes P4
40	update: T2 writes P2
50	update: T3 writes P3
60	T2 commit
70	update: T3 writes P2
80	T2 end
90	update: T3 writes P1
100	update: T1 writes P5
110	update: T3 writes P2
120	T3 abort
	<del>CRASH, RESTART</del>

1. Restore P2 to the before-image stored in LSN 110
2. Restore P1 to the before-image stored in LSN 90
3. Restore P2 to the before-image stored in LSN 70
4. Restore P3 to the before-image stored in LSN 50
5. Restore P4 to the before-image stored in LSN 30

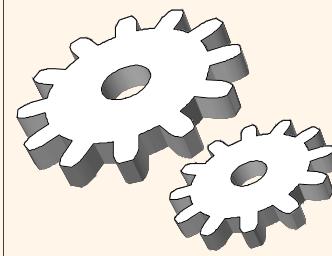
Consider the execution shown below.

c. Show the tail end of the log after the recovery from the **CRASH**:

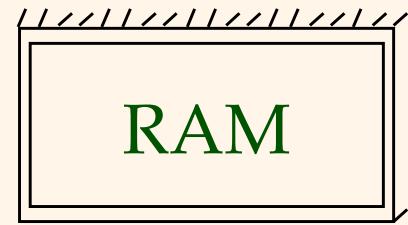


LSN	LOG
00	begin_checkpoint
10	end_checkpoint
20	update: T1 writes P1
30	update: T3 writes P4
40	update: T2 writes P2
50	update: T3 writes P3
60	T2 commit
70	update: T3 writes P2
80	T2 end
90	update: T3 writes P1
100	update: T1 writes P5
110	update: T3 writes P2
120	T3 abort
	<b>CRASH, RESTART</b>

LSN	nextLSN	prevLSN	tansID	Type	pageID	undonextLSN
130	140	120	T3	CLR	P2	90
140	150	130	T3	CLR	P1	70
150	160	140	T3	CLR	P2	50
160	170	150	T3	CLR	P3	30
170	180	160	T3	CLR	P4	
180		170	T3	END		
190	200	100	T1	CLR	P5	20
200	210	190	T1	CLR	P1	
210		200	T1	END		



# Example of Recovery



Tact Table

lastLSN

status

Dirty Page Table

recLSN

flushedLSN

ToUndo

LSN      LOG

00 begin\_checkpoint

05 end\_checkpoint

10 update: T1 writes P5

20 update T2 writes P3

30 T1 Abort

40 CLR: UNDO T1 LSN 10

45 T1 END

50 update: T3 writes P1

60 update: T2 writes P5

CRASH, RESTART

prevLSNs

How many **UPDATES** need to be **UNDOne** by **T2**?

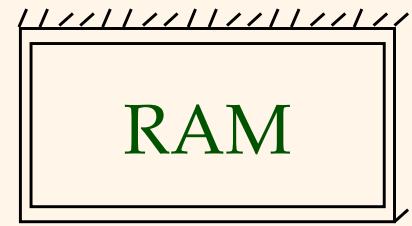
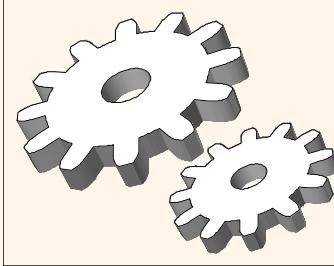
2!

How many **UPDATES** need to be **UNDOne** by **T3**?

1!

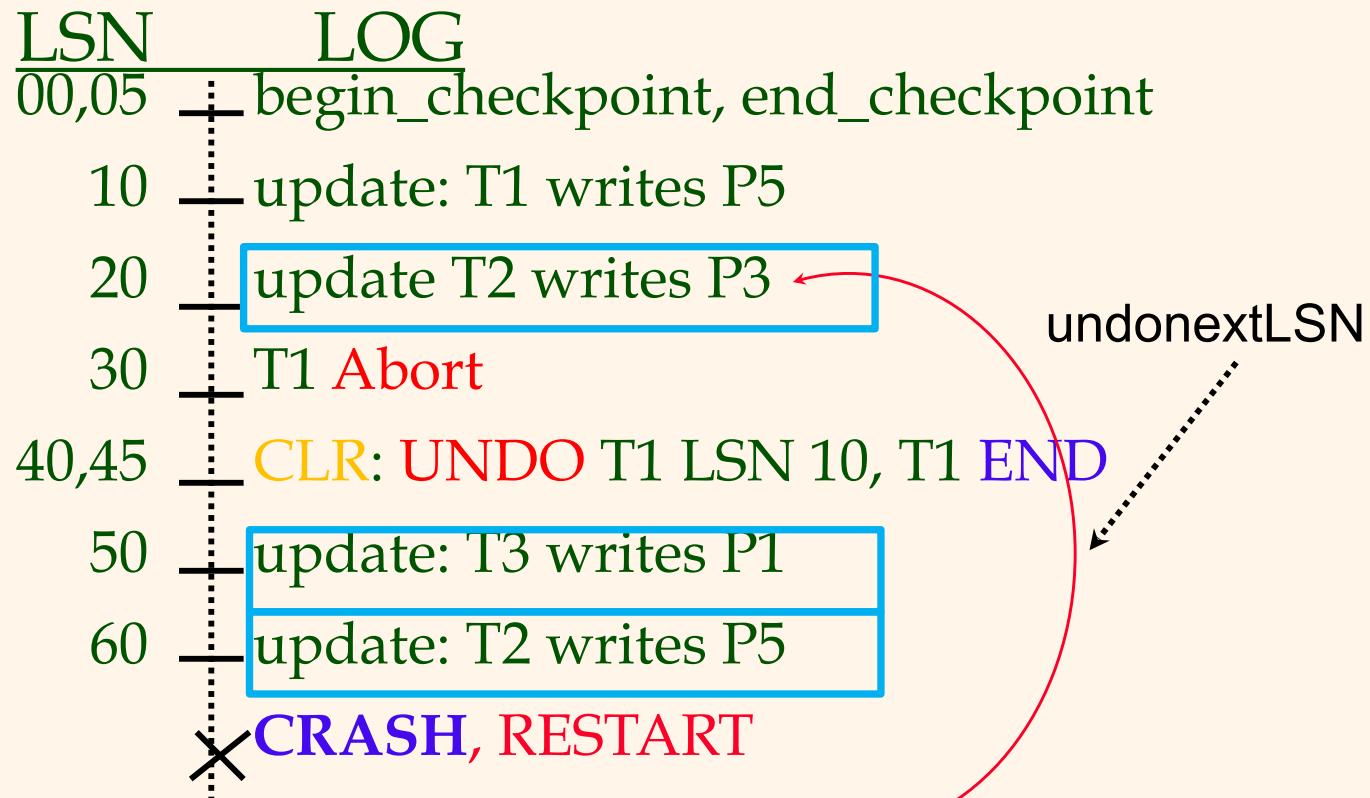
How many **Log Entries**? 5!

# Example: CRASH During Restart!



Xact Table  
lastLSN  
status  
Dirty Page Table  
recLSN  
flushedLSN

ToUndo

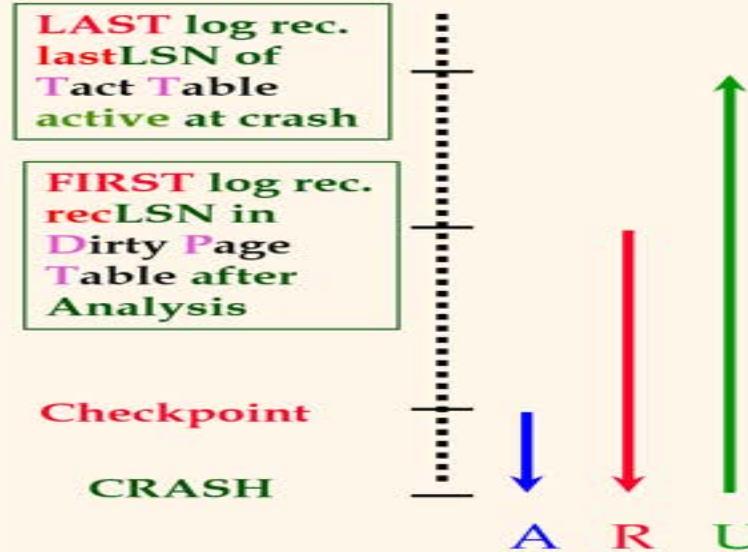




# Crash Recovery: Big Picture



Start from a **Checkpoint** (found via **Master Record**).



Three Phases. Need to:

- Figure out which Xacts **COMMIT**ted since **Checkpoint**, which **FAILED** (**Analysis**).
- **REDO all actions.**
  - ◆ (repeat history) (**R**).
- **UNDO effects of FAILED Xacts** (**U**).

51



pageID	recLSN

DIRTY PAGE TABLE

transID	lastLSN

TRANSACTION TABLE

LSN	prevLSN	undoneLSN	transID	type	pageID	offset	before	after

LOG

Given the schedule below, complete the recovery log to match the schedule.

Schedule	
T <sub>1</sub>	T <sub>2</sub>
	read Y
read X	
X = X * 6	
	Y = Y - 3
write X	
save main memory	
rollback	
	write Y
	commit

Recovery Log
(A)
start T <sub>1</sub>
update T <sub>1</sub> , X <sub>ID</sub> , X <sub>original</sub> , X <sub>new</sub>
checkpoint T <sub>1</sub> , T <sub>2</sub>
undo T <sub>1</sub> , X <sub>ID</sub> , X <sub>original</sub>
rollback T <sub>1</sub>
(B)
(C)

- (A) Pick
- (B) Pick
- (C) Pick

1	2
---	---

## TA time (Jordan) (CA 20.2.1 Step 1 – Recovery)

## TA time (Jordan) (CA 20.2.1 Step 2 – Recovery)



JN

Nizam, Jaer



Yu, Jordan T

**Participants**

Invite someone or dial a number

**Share invite**

in this meeting (106)

Mute all

 Hilford, Victoria  
Organizer

RA Adhikari, Rohit

AA Akram, Ali

SA Altaf, Sameer

SA Alvarez, Stephanie

OA Anayor-Achu, Ogochukwu E

HA Avci, Hatice Kubra

RA Aysola, Riya

AB Bahi, Anish

SB Banza, Sean Paolo B

HB Bui, Hieu

Burger, Jake

VC Carrillo-Zepeda, Victor E

CW Christopher W... (Unverified)

**zyBooks** My library > COSC 3380: Database Systems Formal > 20.2: Recovery

**QUESTION**

undo Y	start T <sub>1</sub>
load X	update T <sub>2</sub> , Y <sub>0</sub> ; X <sub>original</sub> , Y <sub>new</sub>
X + X / 4	checkpoint T <sub>1</sub> , T <sub>2</sub>
V = V + 2	(A)
write Y	rollback T <sub>3</sub>
save main memory	(B)
rollback	(C)
write X	
commit	

**ANSWER**

1    2    3    4

**Check**    **Next**

✓ Expected

(A) undo T<sub>3</sub>, Y<sub>0</sub>; X<sub>original</sub>  
 (B) update T<sub>1</sub>, X<sub>0</sub>; X<sub>original</sub>, X<sub>new</sub>  
 (C) commit T<sub>1</sub>

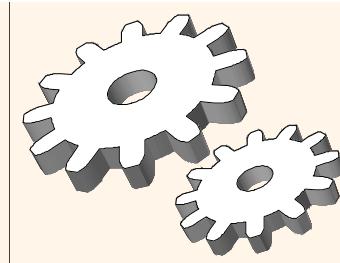
(A) When T<sub>3</sub> rolls back, the write Y operation must be reversed. An undo T<sub>3</sub>, Y<sub>0</sub>; Y<sub>original</sub> record is written in the log before the rollback record. The undo record includes the transaction identifier T<sub>3</sub>, the data identifier Y<sub>0</sub>, and the restored data value Y<sub>original</sub>.

(B) When T<sub>1</sub> writes X an update T<sub>1</sub>, X<sub>0</sub>; X<sub>original</sub>, X<sub>new</sub> record is written in the log. The update record includes the transaction identifier T<sub>1</sub>, the data identifier X<sub>0</sub>, the original data value X<sub>original</sub>, and the new data value X<sub>new</sub>.

(C) When T<sub>1</sub> commits, a commit T<sub>1</sub> record is written in the log.

**View solution**    **Feedback!**



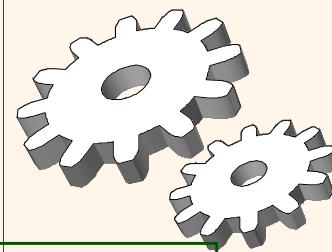


# COSC 3380

## Lecture 19

### *Security and Authorization*

<input type="checkbox"/> 21. SET 4 - 4:SECURITY and AUTHORIZATION	0%	0%	
<input type="checkbox"/> 21.1 MySQL	0%		
<input type="checkbox"/> 21.2 MySQL architecture	0%		
<input type="checkbox"/> 21.3 MySQL Workbench: Import and export		No activities	
<input type="checkbox"/> 21.4 MySQL Workbench: Stored procedures and functions		No activities	
<input type="checkbox"/> 21.5 LAB - MySQL Workbench review (Sakila)	0%		



# SQL GRANT Command

GRANT privileges ON object TO users [WITH GRANT OPTION]

- ❖ The following **privileges** can be specified:
  - ❖ **SELECT**: Can read all columns (including those added later via ALTER TABLE command).
  - ❖ **INSERT(col-name)**: Can insert tuples with non-null or non-default values in this column.
    - ❖ **INSERT** means same right with respect to all columns.
  - ❖ **DELETE**: Can delete tuples.
  - ❖ **REFERENCES (col-name)**: Can define foreign keys (in other tables) that refer to this column.
- ❖ If a **user** has a **privilege WITH GRANT OPTION**, they can pass **privilege** on to other **users** (with or without passing on the **GRANT OPTION**).
- ❖ Only **owner** can execute CREATE, ALTER, and DROP.

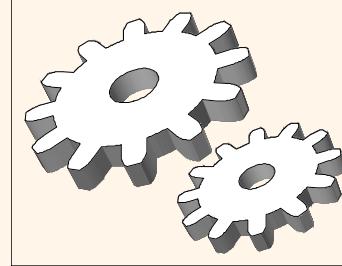
# *SQL GRANT and REVOKE of Privileges*



- ❖ **GRANT** INSERT, SELECT ON Sailors TO Horatio
  - Horatio can INSERT tuples or query Sailors.
- ❖ **GRANT** DELETE ON Sailors TO Yuppy WITH GRANT OPTION
  - Yuppy can DELETE tuples, and also can authorize others to do so.
- ❖ **GRANT** UPDATE (*rating*) ON Sailors TO Dustin
  - Dustin can UPDATE (only) the *rating* field of Sailors tuples.
- ❖ **GRANT** SELECT ON ActiveSailors TO Guppy, Yuppy
  - This does NOT allow the 'uppies to query Sailors directly!
- ❖ **REVOKE:** When a privilege is revoked from X, it is also **revoked** from all users who got it *solely* from X.

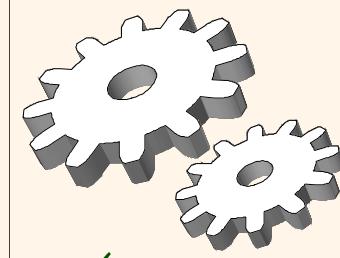
GRANT privileges ON object TO users [WITH GRANT OPTION]

# *GRANT/REVOKE* on Views



- ❖ If the **creator** of a **View** loses the **SELECT** privilege on an **underlying table**, the **View** is dropped!
- ❖ If the **creator** of a **View** loses a privilege held **WITH GRANT OPTION** on an **underlying table**, (s)he loses the privilege on the **View** as well; so do **users** who were **GRANTED** that privilege on the **View**!

# *Views and Security*

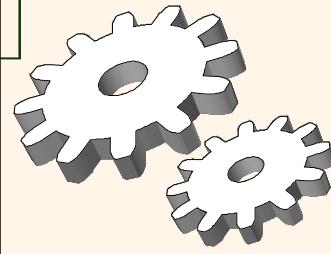


Views can be used to present necessary information (or a summary), while hiding details in **underlying Relation(s)**.

- Given **ActiveSailors**, but not **Sailors** or **Reserves**, we can find sailors who have a reservation, but not the *bid*'s of boats that have been reserved.

**Creator of View** has a privilege on the **View** if (s)he has the privilege on all **underlying Tables/Relations**.

Together with **GRANT/REVOKE** commands, **Views** are a very powerful **access control Tool**.



# Views and Security

4. Consider the Employee(ENAME, DEPT, SALARY) below:

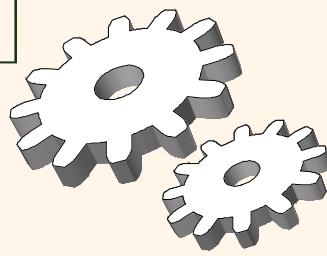
ENAME	DEPT	SALARY
CHUCK NORRIS	TOYS	100000
JOHN SMITH	TOYS	300000
HANK GREEN	TOYS	200000
BOB WHITE	CS	100000
CHRIS BROWN	CS	50000

For authorization purposes you also define Views **EmployeeNames** (with ENAME as the only attribute) and **DeptInfo** with fields DEPT and AVGSALARY. The latter lists the average salary for each department.

- a. Show the View definition statements for **EmployeeName** and **DeptInfo**.

```
CREATE VIEW EmployeeNames (ename)
AS SELECT E.ename
FROM Employees E
```

```
CREATE VIEW DeptInfo (dept, avgsalary)
AS SELECT DISTINCT E.dept, AVG (E.salary) AS avgsalary
FROM Employees E
GROUP BY E.dept
```



# Views and Security

4. Consider the Employee(ENAME, DEPT, SALARY) below:

ENAME	DEPT	SALARY
CHUCK NORRIS	TOYS	100000
JOHN SMITH	TOYS	300000
HANK GREEN	TOYS	200000
BOB WHITE	CS	100000
CHRIS BROWN	CS	50000

```
CREATE VIEW EmployeeNames (ename)
AS SELECT E.ename
FROM Employees E
```

EmployeeName

ENAME
CHUCK NORRIS
JOHN SMITH
HANK GREEN
BOB WHITE
CHRIS BROWN

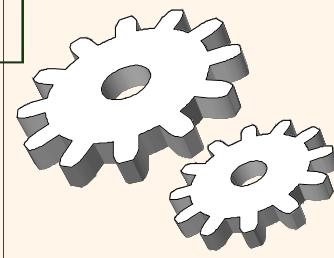
```
CREATE VIEW DeptInfo (dept, avgsalary)
AS SELECT DISTINCT E.dept, AVG (E.salary) AS avgsalary
FROM Employees E
GROUP BY E.dept
```

DeptInfo

DEPT	AVGSALARY
CS	75000
TOYS	200000

b.

Show contents of Views EmployeeName and DeptInfo.



# Views and Security

## Employee

ENAME	DEPT	SALARY
CHUCK NORRIS	TOYS	100000
JOHN SMITH	TOYS	300000
HANK GREEN	TOYS	200000
BOB WHITE	CS	100000
CHRIS BROWN	CS	50000

## DeptInfo

DEPT	AVGSALARY
TOYS	200000
CS	75000

## EmployeeName

ENAME
CHUCK NORRIS
JOHN SMITH
HANK GREEN
BOB WHITE
CHRIS BROWN

CANNOT create accounts ;

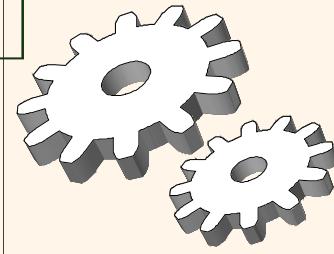
- c. What **privileges** should be granted to a user **who needs to know only** average department salaries for the **Toy** and **CS** departments?

GRANT SELECT on DeptInfo to public;

```
SQL> GRANT SELECT on DeptInfo to public;
Grant succeeded.

SQL>
```

Note that it is impossible to allow the user to access only the average salaries of Toy and CS departments. Only another View will allow it.



# Views and Security

## Employee

ENAME	DEPT	SALARY
CHUCK NORRIS	TOYS	100000
JOHN SMITH	TOYS	300000
HANK GREEN	TOYS	200000
BOB WHITE	CS	100000
CHRIS BROWN	CS	50000

DeptInfo

DEPT	AVGSALARY
TOYS	200000
CS	75000

EmployeeName

ENAME
CHUCK NORRIS
JOHN SMITH
HANK GREEN
BOB WHITE
CHRIS BROWN

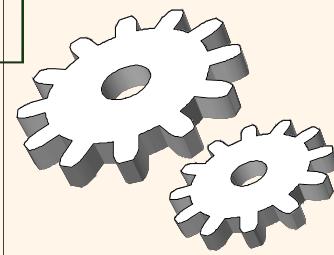
- d. You are to **authorize** your secretary **to fire people** (you will probably tell him whom to fire, but you want to be able to delegate this task), **CANNOT** create accounts ;

GRANT DELETE on Employee to public;

```
SQL> GRANT DELETE on Employee to public;
```

```
Grant succeeded.
```

```
SQL> |
```



## DeptInfo

DEPT	AVGSALARY
TOYS	200000
CS	75000

- e. Assume existence of user Assistant Manager **assistant**. What privilege should you grant to **assistant** who needs to know only average department salaries? **Express it in proper SQL.**

GRANT **SELECT** ON **DeptInfo** TO **assistant**      CANNOT create accounts ;

## *Recovery Practice Questions*

Refer to the sequence below. What is the isolation level of transaction E?

```
session begins  
SET GLOBAL TRANSACTION  
ISOLATION LEVEL SERIALIZABLE;  
session ends
```

```
session begins  
SET SESSION TRANSACTION  
ISOLATION LEVEL REPEATABLE READ;  
transaction A  
transaction B  
SET TRANSACTION  
ISOLATION LEVEL READ UNCOMMITTED;  
transaction C  
SET TRANSACTION  
ISOLATION LEVEL READ COMMITTED;  
transaction D  
transaction E  
session ends
```

### Isolation levels

Relational databases allow database administrators and application programmers to specify strict or relaxed levels of isolation for each transaction. The SQL standard defines four isolation levels:

1. **SERIALIZABLE** transactions run in a serializable schedule with concurrent transactions. Isolation is guaranteed.
2. **REPEATABLE READ** transactions read only committed data. After the transaction reads data, other transactions cannot update the data. REPEATABLE READ prevents most types of isolation violations but allows phantom reads.
3. **READ COMMITTED** transactions read only committed data. After the transaction reads data, other transactions can update the data. READ COMMITTED allows nonrepeatable and phantom reads.
4. **READ UNCOMMITTED** transactions read uncommitted data. READ UNCOMMITTED processes concurrent transactions efficiently but allows a broad range of isolation violations, including dirty, nonrepeatable, and phantom reads.

- a. Serializable
- b. Repeatable read
- c. Read committed
- d. Read uncommitted

A transaction \_\_\_\_ indicates the starting and ending statement of a database transaction.

a. limit

b. boundary

c. coverage

d. scope

\_\_\_\_\_ allows all executed database instructions to be rolled back and to be restored in a prior transaction state.

- a. SET TRANSACTION
- b. COMMIT
- c. SAVEPOINT
- d. RELEASE SAVEPOINT

A document that sequentially writes all the database operations is known as a/an \_\_\_\_\_.

a. recovery log

b. system log

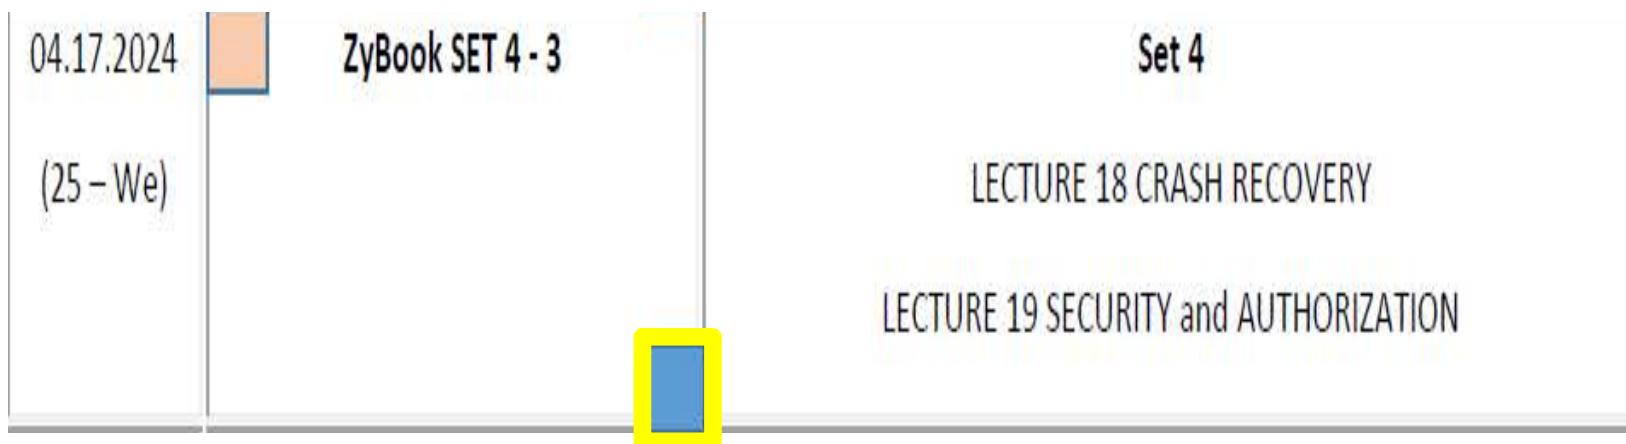
c. application log

d. task log

A recovery system should manage which three failure scenarios?

- a. Transaction failure, Memory failure, and Storage media failure
- b. Transaction failure, System failure, and Storage media failure
- c. Transaction failure, System failure, and Application failure
- d. Transaction failure, Program failure, and Storage media failure

**At 5:00 PM .**



<input type="checkbox"/> 17. SET 4	Empty	▼
<input type="checkbox"/> 20. SET 4 - 3:CRASH RECOVERY	Hidden  0%  0%  0%	▼
<input type="checkbox"/> 21. SET 4 - 4:SECURITY and AUTHORIZATION	Hidden  0%  0%	▼

VH work on  
SET 4 – 3, 4: Recovery & Security

# Next

04.22.2024

(26 - Mo)

## EXAM 4 Practice

(PART of 20 points)

17. SET 4

Empty

v

18. SET 4 - 1:TRANSACTIONS

Hidden



09

0%

1

□ 19. SET 4 - 2:CONCURRENCY CONTROL

Hidden



109

0%

7

20. SET 4 - 3:CRASH RECOVERY

Hidden



0%

0°

0%

7

#### 21. SET 4 - 4:SECURITY and AUTHORIZATION

Hidden



0%

09

1

# From 5:05 to 5:15 PM – 5 minutes.

04.17.2024

ZyBook SET 4 - 3

Set 4

(25 – We)

LECTURE 18 CRASH RECOVERY

LECTURE 19 SECURITY and AUTHORIZATION

CLASS PARTICIPATION 20 points

20% of Total

Not available until Apr 17 at 5:05pm | Due Apr 17 at 5:15pm

## SECURITY

Class 25 END PARTICIPATION

Not available until Apr 17 at 5:05pm | Due Apr 17 at 5:15pm

VH, publish

This is a synchronous online class.  
Attendance is required.  
Recording or distribution of class materials is prohibited.

- At the beginning of selected classes there is an assessment in the first 10 minutes. (beige BOX in the Detailed Syllabus)
- At the end of selected classes there is an assessment in the last 10 minutes. (blue BOX in the Detailed Syllabus)
- ZyBook sections will be downloaded and used for 30% of Total Score on the dates specified in the Detailed Syllabus.
- EXAMS are in CANVAS. No late EXAMS.
- I have to be present in TEAMS in order to take any graded assignment assigned during that class.

**At 5:15 PM.**

## **End Class 25**

**VH, Download Attendance Report  
Rename it:  
4.17.2024 Attendance Report FINAL TEAMS**

The screenshot shows a Canvas assignment page. At the top left, there is a section titled "CLASS PARTICIPATION 20 points". To the right of this, a button says "20% of Total +". Below this, there is a section titled "EXAM 4 Practice - Requires Respondus LockDown Browser + Webcam". Underneath this section, the text reads: "CLASS PARTICIPATION [20 points] Module | Not available until Apr 22 at 4:05pm | Due Apr 22 at 5:05pm | 100 pts". On the far right of the page, the text "VH, publish" is overlaid in red and blue.

**VH, upload Class 25 to CANVAS.**