

COSC 4368

Fundamentals of Artificial Intelligence

Constraint Satisfaction Problems
September 18th, 2023

Constraint Satisfaction Problems

- Regular search problems:
 - Search in a state space
 - Each state is indivisible --- a black box with no internal structure
- Constraint satisfaction problem
 - Each state has a factored representation: a set of variables, each has a value
 - Goal condition consists of a set of sub-conditions: each variable should have a value that satisfies all the constraints on the variable


What is a Constraint Satisfaction Problem (CSP)

- A CSP consists of three components:
 - A finite set of variables:
 - A finite set of domains:
 - One for each variable; finite or infinite domain
 - A finite set of constraints:
 - Each constraint involves some subset of the variables
 - Specifies allowable combinations of values for that subset
 - E.g., $<(,), >$

What is a Constraint Satisfaction Problem (CSP)

- **State space:** each state is defined by an *assignment* of values to some or all of the variables, e.g.,
- **Consistent assignment:** an assignment that does not violate any constraints
- **Complete (Partial) assignment:** an assignment where every variable is (a subset of variables are) assigned
- **Goal:** find a complete and consistent assignment
 - Find an assignment of the variables from their domains such that none of the constraints are violated

Example: Map Coloring



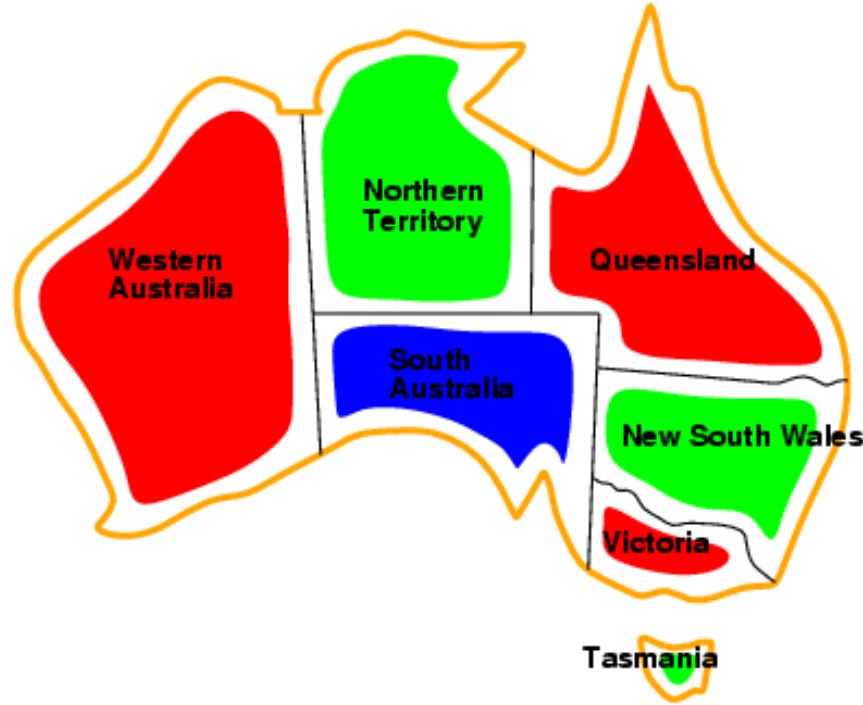
Constraint graph:

Node: variable

Edge: two variables in the same constraint


- **Task:** color each region either red, green, or blue such that no neighboring regions have the same color
- **Variables:**
- **Domains:**
- **Constraints:** adjacent regions must have different colors
 - or

Example: Map Coloring



- **Solution:** complete and consistent assignments that satisfy all constraints
 - Many possible solutions
 - e.g.,

Example: N-queens Problem



- Assume each queen is in one column, which row does each one go in?
- Variables: position of each queen in the column
- Domains: same (row value)
- Constraints:
 - Two queens cannot be in the same row:
 - Two queens cannot be in the same diagonal:

Example: Sudoku

	1	2	3	4	5	6	7	8	9
A			3		2		6		
B	9			3		5			1
C			1	8		6	4		
D			8	1		2	9		
E	7								8
F			6	7		8	2		
G			2	6		9	5		
H	8			2		3			9
I			5		1		3		

	1	2	3	4	5	6	7	8	9
A	4	8	3	9	2	1	6	5	7
B	9	6	7	3	4	5	8	2	1
C	2	5	1	8	7	6	4	9	3
D	5	4	8	1	3	2	9	7	6
E	7	2	9	5	6	4	1	3	8
F	1	3	6	7	9	8	2	4	5
G	3	7	2	6	8	9	5	1	4
H	8	1	4	2	5	3	7	6	9
I	6	9	5	4	1	7	3	8	2

- **Task:** fill in all the remaining squares with digits from 1 to 9 such that no digit appears twice in any row, column, or 3*3 box
- Variables: 81 variables, one for each square
 - A1, A2, A3,..., I7, I8, I9
 - Letter index rows, top to bottom; Digits index columns, left to right
- Domains: nine positive digits, e.g., A1
- Constraints: 27 *Alldiff* constraints, e.g., *Alldiff*{A1-A9}, *Alldiff*{A1-I1},...

Variations of CSP --- Type of Variables

- **Discrete variables**

- Finite domains

- Size complete assignments
 - E.g., map coloring, N-queens problem

- Infinite domains

- E.g., the set of integers
 - E.g., job-scheduling problem when we don't have a deadline
 - Need a constraint language without enumerating the available assignments, e.g,
 - Linear constraints: solvable
 - Nonlinear constraints: no general algorithm

- **Continuous variables**

- E.g., circuit layout
 - Linear programming: constraints must be linear equalities or inequalities; can be solved in polynomial time

Variations of CSP --- Type of Constraints

- Unary constraint involves a single variable
 - e.g.,
- Binary constraint involves two variables
 - e.g.,
- High-order constraints involve 3 or more variables
 - Professors A, B, and C cannot be on a committee together
 - Can always be represented by multiple binary constraints
- Preference (soft) constraints indicate which solutions are preferred
 - Constraints can be violated, but with a higher cost on individual assignments
 - E.g., class-scheduling problem, Prof. A prefers teaching in morning; assigning an afternoon slot may be allowable, but can lead to a higher cost for the overall objective
 - Combination of optimization with CSPs

CSPs Only Need Binary Constraints

- Unary constraint: just change the domain by deleting the values
- Higher order: reduce to binary constraints
- Simple example:
 - Three variables: X, Y, Z
 - Domains:
 - Constraints:
- Create an auxiliary variable, W, taking values as 3-tuples
- Domain:
 - Exactly the tuples satisfying the constraints
- Create three new constraints

How to Solve CSPs

- Example: integer CSP
- Variables A, B, C, D which take values in $\{1, \dots, 10\}$
- Constraints:
 - C1:
 - C2:
 - C3:
- Solution:

Brute Force Search

- For $A=1, \dots, 10$
 - For $B=1, \dots, 10$
 - For $C=1, \dots, 10$
 - For $D=1, \dots, 10$
 - If C1 and C2 and C3
 - WriteSolution(A, B, C, D)

How to reduce the complexity of this search?

Look into the Constraints

- Constraints:
 - C1:
 - C2:
 - C3:
- A more efficient loop:
 - For $A=2, \dots, 10$
 - For $B=1, \dots, A-1$
 - For $C=1, \dots, A-1$
 - For $D=1, \dots, A-1$
 - If C1 and C2 and C3
 - WriteSolution(A, B, C, D)

We Can Do More...

- Constraints:
 - C1:
 - C2:
 - C3:
- **Variable elimination: Remove variable by replacing it by**
- New constraints with 3 variables:
 - C1':
 - C2':
- Find valid assignments for (B, C, D), and set $A=B+C$

CSP as A Standard Search Problem


- A CSP can easily be expressed as a standard search problem
- States defined by the values assigned so far (partial assignments)
 - Initial state: the empty assignment
 - Successor function: assign a value to an unassigned variable provided that it does not violate a constraint
 - Goal test: the current assignment is complete and satisfies all the constraints
 - Path cost: not really relevant

CSP as A Standard Search Problem

- For a CSP with variables of domain size
- Solution can be found at depth

Assign a value to one variable
Branching factor:

Assign a value to 2nd variable
Branching factor:




In total:

- End up with leaves at the n -th level
- But there are only possible complete assignments

What is missing



- Commutativity:
 - A problem is commutative if the order of any given set of actions has no effect on the outcome
 - CSPs are commutative: can reach the same partial assignment regardless of order
 - Example: map coloring for Australia
 - [WA=red then NT=green] same as [NT=green then WA=red]
- In CSP, we need only consider a single variable at each node in the search tree
 - Might choose among SA=red, SA=green, SA=blue
 - Never choose between SA=red and WA=blue
- There are leaves
- (still need to figure out which variable to choose at each node)



Backtracking

- Basic uninformed search algorithm to solve CSPs
- Similar to depth-first search, exploring one node at a time
- Chooses values for one variable at a time
- Backtracks when a variable has no legal values left to assign

Example



Backtracking Search


```
function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({ }, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove {var = value} from assignment
  return failure
```

Improving Backtracking



- Previous improvements on uninformed search
 - Use heuristics (domain knowledge): informed search
- General-purpose ideas can achieve huge gains in speed for backtracking
- Ordering:
 - Which variable should be assigned next?
 - In what order should its values be tried?
- Filtering: can we detect inevitable failure early?

Variable Ordering: Minimum Remaining Values (MRV)





- Aka “most constrained variable” or “fail-first”:
 - Choose the variable with the fewest legal left values in its domain
 - If some variable X has no legal values left, select X and failure will be detected immediately --- avoiding pointless searches through other variables

Degree Heuristic





- MRV doesn't help in choosing the first region to color
- Degree heuristic: select variable that is involved in the largest number of constraints on other unassigned variables
- MRV is usually a more powerful guide, but degree heuristic can be useful as a tie breaker

Value Ordering: Least Constraining Value





- Given a variable, choose the least constraining value
 - Prefer the value that rules out the fewest choices for the neighboring variables in the constraint graph
 - Leave the maximum flexibility for subsequent variable assignments
- Why least rather than most?
 - Only need one solution, look for the most likely values first

Filtering: Forward Checking




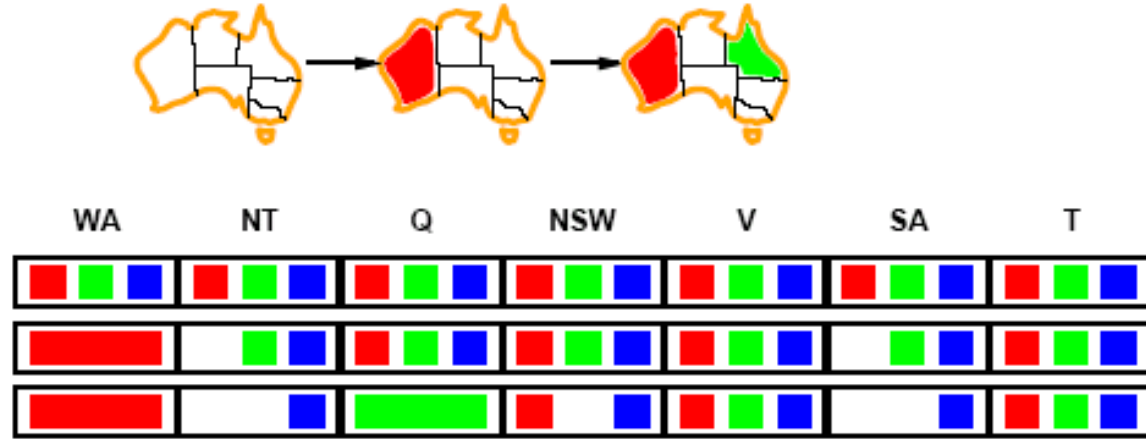
- Keep track of remaining legal values for unassigned variables and cross off bad options
- Terminate search when any variable has no legal values

Filtering: Forward Checking




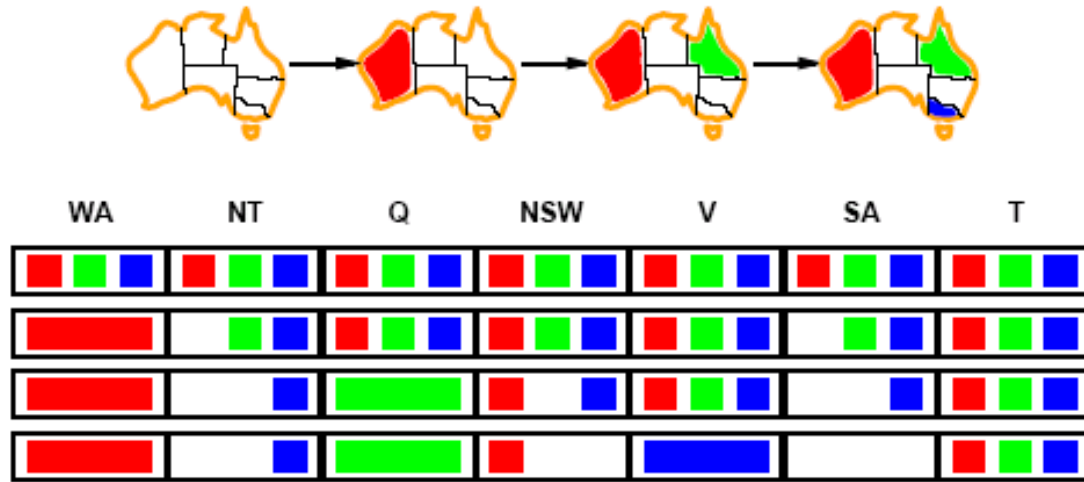
- Assign {WA=red}
- Effects on other variables connected by constraints with WA
 - NT can no longer be red
 - SA can no longer be red

Filtering: Forward Checking



- Assign {Q=green}
- Effects on other variables connected by constraints with Q
 - NT can no longer be green
 - NSW can no longer be green
 - SA can no longer be green
- MRV can also be used for variable ordering


Filtering: Forward Checking





- If assign $\{V=\text{blue}\}$
- Effects on other variables connected by constraints with V
 - NSW can no longer be blue
 - SA is empty
- Forward checking detected that the partial assignment is inconsistent with the constraints and backtracking can occur

Example: 4-queens



	1	2	3	4
1				
2				
3				
4				





Example: 4-queens



Example: 4-queens




Example: 4-queens





Example: 4-queens



	1	2	3	4
1	★	●	●	●
2		●	●	
3		★	●	●
4			●	●





Example: 4-queens



Example: 4-queens




Example: 4-queens




Example: 4-queens

	1	2	3	4
1	★	●	●	●
2		●	★	●
3			●	
4		★	●	●


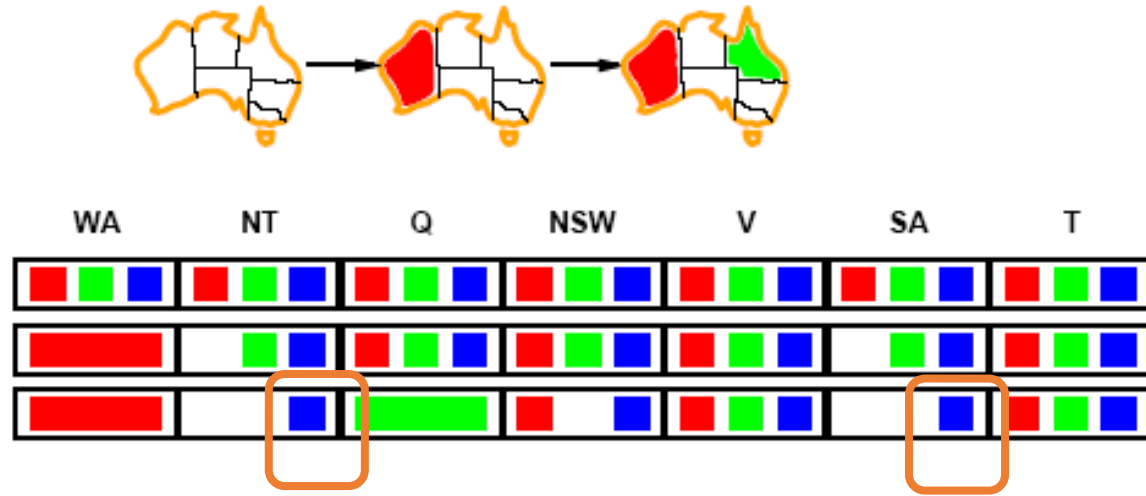


Example: 4-queens

	1	2	3	4
1	★	●	●	●
2		●	★	●
3			●	●
4		★	●	●



Filtering: Constraint Propagation





- Forward checking propagates information from assigned to unassigned variables, but doesn't detect all failures
- NT and SA cannot both be blue
- Why didn't we detect this yet?

Constraint Propagation



- Constraint propagation: a specific type of inference
 - Using constraints to reduce the domain for a variable, which in turn can reduce the domain for another variable, and so on...
 - Reason from constraint to constraint
- Techniques like constraint propagation and forward checking are in effect eliminating parts of the search space
 - Somewhat complementary to search
- Constraint propagation goes further than FC by repeatedly enforcing constraints locally
- Arc-consistency (AC) is a systematic procedure for constraint propagation⁴¹

Arc Consistency





- An arc is consistent if for every value x of X there is some value y of Y consistent with x (without violating the constraint)
- Consider state of search after WA and Q are assigned
- SA NSW is consistent if
 - SA=blue and NSW=red

Arc Consistency





- NSW SA is consistent if
 - NSW=red and SA=blue
 - NSW=blue and SA=???
- Enforcing arc consistency removes blue from NSW

Arc Consistency



- Continue to propagate constraints
 - Check V NSW
 - Not consistent for V=red
 - Remove red from V

Arc Consistency



- Continue to propagate constraints
 - Check SA NT
 - Not consistent for any value of SA
 - Detect failure
- Arc consistency detects failure earlier than forward checking

Arc Consistency Checking

- Can be run as a preprocessor or after each assignment
- AC must be run repeatedly until no inconsistency remains
- Trade-off
 - Requires some overhead to do, but generally more effective than direct search
 - Eliminates large inconsistent parts of the state space more effectively than search can
- Need a systematic method for arc-checking
 - If X loses a value, neighbors of X need to be rechecked
 - Incoming arcs can become inconsistent again (outgoing arcs stay the same)

Arc Consistency Checking

- Can be run as a preprocessor or after each assignment
- AC must be run repeatedly until no inconsistency remains
- Trade-off
 - Requires some overhead to do, but generally more effective than direct search
 - Eliminates large inconsistent parts of the state space more effectively than search can
- Need a systematic method for arc-checking
 - If X loses a value, neighbors of X need to be rechecked
 - Incoming arcs can become inconsistent again (outgoing arcs stay the same)

AC-3 Algorithm

function AC-3(csp) **returns** false if an inconsistency is found and true otherwise

inputs: csp , a binary CSP with components (X, D, C)

local variables: $queue$, a queue of arcs, initially all the arcs in csp

while $queue$ is not empty **do**

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(queue)$

if REVISE(csp, X_i, X_j) **then**

if size of $D_i = 0$ **then return** false

for each X_k **in** $X_i.\text{NEIGHBORS} - \{X_j\}$ **do**

 add (X_k, X_i) to $queue$

return true

function REVISE(csp, X_i, X_j) **returns** true iff we revise the domain of X_i

$revised \leftarrow false$

for each x **in** D_i **do**

if no value y in D_j allows (x, y) to satisfy the constraint between X_i and X_j **then**

 delete x from D_i

$revised \leftarrow true$

return $revised$

Complexity of AC-3

- A binary CSP with n variables has at most n^2 arcs
- Each arc can be inserted in the queue d times (worst case)
 - Only d values to delete
- Consistency of an arc can be checked in $O(d^2)$ time
 - d^2 pairs of values to check
- Complexity is $O(n^3d)$
- Although more expensive than forward checking, AC is usually worthwhile

Local Search for CSPs

- Use complete-state representation
 - Initial state: assign a value to every variable
 - Successor function: change the value of one variable at a time
- Allow states with unsatisfied constraints (unlike backtracking)
- The point of local search is to eliminate the violated constraints
- Variable selection: randomly select any conflicted variable
- Value selection: min-conflicts heuristic
 - Select new value that results in the minimum number of conflicts with other variables

Local Search for CSPs

function MIN-CONFLICTS(csp, max_steps) **returns** a solution or failure

inputs: csp , a constraint satisfaction problem

max_steps , the number of steps allowed before giving up

$current \leftarrow$ an initial complete assignment for csp

for $i = 1$ to max_steps **do**

if $current$ is a solution for csp **then return** $current$


$var \leftarrow$ a randomly chosen conflicted variable from $csp.VARIABLES$

$value \leftarrow$ the value v for var that minimizes CONFLICTS($var, v, current, csp$)

 set $var = value$ in $current$


return $failure$

Example 1



- Use min-conflicts heuristic
 - Randomly select the 4th variable (position of the 4th queen)
 - X=1, two constraints violated
 - X=3, no constraints violated
 - X=4, two constraints violated

Example 2



- A two-step solution for the 8-queens problem
- At each stage, a queen is chosen for reassignment in its column
- The algorithm moves the queen to the min-conflicts square, breaking ties randomly

Advantages of Local Search

- The runtime of min-conflicts is roughly independent of problem size
 - Can solve the millions-queen problem in roughly 50 steps
 - Why?
 - N-queens is easy for local search because of the relatively high density of solutions in state-space
- Local search can be particularly useful in an online setting
 - Airline schedule problem
 - E.g., mechanical problems require that 1 plane is taken out of service
 - Can locally search for another “close” solution in state-space
 - Much better and faster in practice than finding an entirely new schedule

Summary

- CSPs:
 - Special kind of problem: states defined by values of a fixed set of variables, goal test defined by constraints on variable values
- Backtracking = depth-first search with one variable assigned per node
- Heuristics:
 - Variable ordering and value selection help significantly
- Constraint propagation does additional work to constrain values and detect inconsistencies
 - Works effectively when combined with heuristics
- Local search like iterative min-conflicts is often effective in practice