


# Support Vector Machines.

W. Wang<sup>1</sup>

<sup>1</sup>Department of Mathematics  
University of Houston

MATH 4323

# Data Example: What boundary is appropriate?



# Classification with Non-linear Decision Boundaries.

The support vector classifier is a natural classification approach in the

- two-class setting,
- when the **boundary** between the two classes **is linear**.

In practice, we are sometimes faced with **non-linear class boundaries**.


**Example.** Next slide shows:

- **(Left)** a two-dimensional data example,
- **(Right)** separating hyperplane & support vectors result of applying a support vector classifier.

It is clear that a **support vector classifier** or **any linear classifier** will **perform poorly here**.

**Question:** How to deal with **non-linear boundaries**?

# Classification with Non-linear Decision Boundaries.



**FIGURE 9.8.** Left: The observations fall into two classes, with a non-linear boundary between them. Right: The support vector classifier seeks a linear boundary, and consequently performs very poorly.

# Linear Regression: Enlarging Predictor Space.

**Example.** In linear regression with

- predictors  $X_1, X_2$ , and
- response  $Y$

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \epsilon$$

**Question:** How do we model potentially **non-linear** relationship between predictors  $X_1, X_2$  and response  $Y$ ?

**Answer:** **polynomial regression**, where we **enlarge** the predictor space to include **quadratic, cubic & other terms**. E.g. we could fit

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_1^2 + \beta_3 X_2 + \beta_4 X_2^2 + \epsilon$$

extending the original predictor space as follows

$$\{X_1, X_2\} \Rightarrow \{X_1, X_1^2, X_2, X_2^2\}$$

# Support Vector Classifier: Enlarging Predictor Space.

For **support vector classifier**, the problem of **possible non-linear boundaries** is addressed the **same way**:

We **enlarge** the feature space using

- quadratic,
- cubic or higher-order polynomial,
- general non-linear (e.g.  $\log$  or  $\exp$ )  
functions of the predictors.

**Example.** In  $p$ -dimensional case,

- instead of fitting a support vector classifier using  $p$  features

$$\{X_1, X_2, \dots, X_p\},$$

- we could fit a support vector classifier using  $2p$  features

$$\{X_1, X_1^2, X_2, X_2^2, \dots, X_p, X_p^2\}$$

# Support Vector Classifier: Enlarging Predictor Space.

**Example (cont'd).** Logic of obtaining separating hyperplane equation remains the same:

- while in the  $p$ -feature space of  $\{X_1, X_2, \dots, X_p\}$  we had

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p = 0 \quad (1)$$

- in the extended  $2p$ -feature space we have

$$\beta_0 + \beta_{11} X_1 + \beta_{12} X_1^2 + \beta_{21} X_2 + \beta_{22} X_2^2 + \dots + \beta_{p1} X_p + \beta_{p2} X_p^2 = 0 \quad (2)$$

Why does (2) lead to a non-linear decision boundary?

- While in the **enlarged feature space** ( $\{X_1, X_1^2, X_2, X_2^2, \dots\}$ ), the resulting decision boundary is, in fact, **linear**,
- But in the **original feature space**, the decision boundary is of the form  $q(x) = 0$ , where  $q$  is a **quadratic polynomial**  $\implies$  its solutions are generally **non-linear**.

## Enlarging Predictor Space: Example.

**Example.** Below is a similar case of non-linearly separable scenario:




FIGURE 2.3: A not linearly separable dataset.



## Enlarging Predictor Space: Example.

**Example (cont'd).** That's what happens when we enlarge the predictor space by adding a third predictor  $Z = X_1^2 + X_2^2$ :




FIGURE 2.4: The dataset with the new  $z$  dimension.




FIGURE 2.5: A hyperplane dividing the dataset.

There is a separating hyperplane

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 Z = 0$$

in the extended 3D predictor space.

## Enlarging Predictor Space: Example.

**Example (cont'd).** In the meantime, if we were to project that hyperplane **back onto the original 2D space of  $\{X_1, X_2\}$** :




FIGURE 2.3: A not linearly separable dataset.

We get a **circle shape** as a result of equation that is **not linear in  $X_1, X_2$** :


$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 (X_1^2 + X_2^2) = 0$$

# Support Vector Classifier: Enlarging Predictor Space.

**Example.** Let's look at  $p = 2$ , with extended space  $\{X_1, X_1^2, X_2, X_2^2\}$  and the following fitted hyperplane equation:

$$2 - X_1 + 0 \times X_1^2 + 3X_2 + 2X_2^2 = 0$$


This corresponds to the following separating  $2D$  hyperplane:



# Support Vector Classifier: Enlarging Predictor Space.

**Example.** Let's look at  $p = 2$ , with extended space  $\{X_1, X_1^2, X_2, X_2^2\}$  and the following fitted hyperplane equation:

$$1 - 6 \times X_1 + X_1^2 + 2 \times X_2 + X_2^2 = 0$$



Looks familiar, and after some rearrangement:

$$(X_1 - 3)^2 + (X_2 + 1)^2 = 9$$

# Support Vector Classifier: Enlarging Predictor Space.

When enlarging original space of  $p$  features with quadratic terms, which leads to separating hyperplane equation of the form

$$\beta_0 + \beta_{11}X_1 + \beta_{12}X_1^2 + \beta_{21}X_2 + \beta_{22}X_2^2 + \cdots + \beta_{p1}X_p + \beta_{p2}X_p^2 = 0$$

the initial optimization task (6) – (9) (from "Maximal Margin & Support Vector Classifier" slides) becomes

$$\begin{array}{ll} \text{maximize} & M \\ \beta_0, \beta_{11}, \beta_{12}, \dots, \beta_{p1}, \beta_{p2}, \epsilon_1, \dots, \epsilon_n & \end{array} \quad (3)$$

$$\text{subject to } \sum_{j=1}^p \sum_{k=1}^2 \beta_{jk}^2 = 1 \quad (4)$$

$$y_i(\beta_0 + \sum_{j=1}^p \beta_{j1}x_{ij} + \sum_{j=1}^p \beta_{j2}x_{ij}^2) \geq M(1 - \epsilon_i), \quad (5)$$

$$\epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq C, \quad (6)$$

# Support Vector Classifier: Enlarging Predictor Space.

One could also enlarge the feature space with

- higher-order polynomial terms ( $X_1^3, X_1^4, \dots$ )

**Example.** For 2D case of  $(X_1, X_2)$ , one could extend to

$$\{X_1, X_1^2, X_1^3, X_1^4, X_2, X_2^2, X_2^3, X_2^4\}$$

- interaction terms ( $X_i X_j, i \neq j$ ),

**Example.** For 3D case of  $(X_1, X_2, X_3)$ , one could extend to

$$\{X_1, X_1^2, X_2, X_2^2, X_3, X_3^2, X_1 X_2, X_1 X_3, X_2 X_3\}$$

- non-polynomial functions of predictors ( $\log(X_1), \exp(X_1), \dots$ )

**Example.** For 2D case of  $(X_1, X_2)$ , one could extend to

$$\{X_1, \log(X_1), X_2, \log(X_2)\}$$

# Support Vector Machines.

There are infinitely many ways to enlarge the feature space, and one **doesn't want to end up with a huge number of features**, as it renders **computations unmanageable**.

The **support vector machine**, which we present next, allows us a

- **general framework of enlarging the feature space** for a **support vector classifier**,
- in a way that leads to **efficient computations**.

# Support Vector Machines: Optimization Task.

**REMINDER:** For a given set of features  $\{X_1, \dots, X_p\}$ , **support vector classifier** is obtained by finding the coefficients  $\beta_0, \beta_1, \dots, \beta_p$  of **hyperplane equation**

$$\beta_0 + \beta_1 x_1 + \dots + \beta_p x_p = 0$$

that solve the following **optimization task**

$$\underset{\beta_0, \beta_1, \dots, \beta_p, \epsilon_1, \dots, \epsilon_n}{\text{maximize}} \quad M \quad (7)$$

subject to

$$\sum_{j=1}^p \beta_j^2 = 1 \quad (8)$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i), \quad (9)$$

$$\epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C \quad (10)$$



# Support Vector Machines: Inner Products.

It turns out that the solution to the support vector classifier optimization task (7) – (10) involves **only the inner products** of the observations  $\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{ip})$ .

The **inner product** of two  $r$ -dimensional vectors  $\mathbf{a} = (a_1, \dots, a_r) \in \mathbb{R}^r$  and  $\mathbf{b} = (b_1, \dots, b_r) \in \mathbb{R}^r$  is defined as

$$\langle \mathbf{a}, \mathbf{b} \rangle = \sum_{k=1}^r a_k b_k = \mathbf{a}^T \mathbf{b}$$

Thus the **inner product** of two observations

- $\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{ip}) \in \mathbb{R}^p$ , and
- $\mathbf{x}_j = (x_{j1}, x_{j2}, \dots, x_{jp}) \in \mathbb{R}^p$

is given by

$$\langle \mathbf{x}_i, \mathbf{x}_j \rangle = \sum_{k=1}^p x_{ik} x_{jk}$$

# Support Vector Machines: Inner Products.

It can be shown that (derivation is beyond the scope of class)

1. The solution to support vector classifier optimization task (7) – (10) can be represented as

$$\begin{aligned} f(\mathbf{x}) &= \beta_0 + \beta_1 \mathbf{x}_1 + \cdots + \beta_p \mathbf{x}_p \equiv \beta_0 + \sum_{k=1}^p \beta_k \mathbf{x}_k = \\ &= \beta_0 + \sum_{i=1}^n \alpha_i \langle \mathbf{x}, \mathbf{x}_i \rangle \end{aligned} \quad (11)$$

where we have

- ▶  $\mathbf{x} = (x_1, \dots, x_p)^T$ ,  $\mathbf{x}_i = (x_{i1}, \dots, x_{ip})^T$ ,  $i = 1, \dots, n$ ,
- ▶  $n$  parameters  $\alpha_i$ ,  $i = 1, \dots, n$ , one per training observation  $\mathbf{x}_i$ .

**Note:** keep in mind,  $\{X_1, \dots, X_p\}$  is the **original**, not transformed, feature space.

# Support Vector Machines: Inner Products.

2. To estimate the parameters  $\alpha_1, \dots, \alpha_n$  and  $\beta_0$  in eq. (11), all we need are the  $\binom{n}{2}$  inner products  $\langle \mathbf{x}_i, \mathbf{x}_j \rangle$  between all pairs of training observations, meaning

$$\hat{\beta}_0 = g_0(\langle \mathbf{x}_1, \mathbf{x}_2 \rangle, \langle \mathbf{x}_1, \mathbf{x}_3 \rangle, \dots, \langle \mathbf{x}_{n-1}, \mathbf{x}_n \rangle)$$

$$\hat{\alpha}_i = g_i(\langle \mathbf{x}_1, \mathbf{x}_2 \rangle, \langle \mathbf{x}_1, \mathbf{x}_3 \rangle, \dots, \langle \mathbf{x}_{n-1}, \mathbf{x}_n \rangle), \quad i = 1, \dots, n$$

where  $g_0, g_1, \dots, g_n$  are functions resulting from solving the optimization task (7) – (10).

In our scenario,

$$\binom{n}{2} \equiv \frac{n(n-1)}{2} \equiv$$

$\equiv \{\text{the number of pairs among a set of } n \text{ training observations}\}$

# Support Vector Machines: Inner Products.

3. However, it turns out that

- ▶  $\alpha_i \neq 0 \Leftrightarrow$  training observation  $\mathbf{x}_i$  is a **support vector**,
- ▶  $\alpha_i = 0 \Leftrightarrow$  training observation  $\mathbf{x}_i$  is **not a support vector**,

So if  $S$  is the collection of indices for training observations that are **support vectors**, then

$$f(\mathbf{x}) = \beta_0 + \beta_1 \mathbf{x}_1 + \cdots + \beta_p \mathbf{x}_p = \beta_0 + \sum_{i \in S} \alpha_i \langle \mathbf{x}, \mathbf{x}_i \rangle \quad (12)$$

So, in the end of the day,  $\hat{\beta}_0$  and  $\hat{\alpha}_i$  solving the (7) – (10) optimization task via the representation in eq. (11) are:

$$\begin{cases} \hat{\beta}_0 = g_0(\langle \mathbf{x}_1, \mathbf{x}_2 \rangle, \langle \mathbf{x}_1, \mathbf{x}_3 \rangle, \dots, \langle \mathbf{x}_{n-1}, \mathbf{x}_n \rangle) \\ \hat{\alpha}_i = g_i(\langle \mathbf{x}_1, \mathbf{x}_2 \rangle, \langle \mathbf{x}_1, \mathbf{x}_3 \rangle, \dots, \langle \mathbf{x}_{n-1}, \mathbf{x}_n \rangle), \quad i \in S \\ \hat{\alpha}_i = 0, \quad i \notin S \end{cases}$$

# Support Vector Machines: Kernels.

Now, let's replace any appearance of inner product  $\langle \mathbf{x}_i, \mathbf{x}_j \rangle$  in (12) with a **generalization of the inner product**

$$K(\mathbf{x}_i, \mathbf{x}_j),$$

where  $K$  is a function referred to as **kernel**.

That leads to the following solution form:

$$f(\mathbf{x}) = \beta_0 + \sum_{i \in S} \alpha_i K(\mathbf{x}, \mathbf{x}_i) \quad (13)$$

with  $\hat{\beta}_0$  and  $\hat{\alpha}_i$  being **functions of  $K(\mathbf{x}_i, \mathbf{x}_j)$** , instead of **just  $\langle \mathbf{x}_i, \mathbf{x}_j \rangle$** :

$$\begin{cases} \hat{\beta}_0 = g_0(K(\mathbf{x}_1, \mathbf{x}_2), K(\mathbf{x}_1, \mathbf{x}_3), \dots, K(\mathbf{x}_{n-1}, \mathbf{x}_n)) \\ \hat{\alpha}_i = g_i(K(\mathbf{x}_1, \mathbf{x}_2), K(\mathbf{x}_1, \mathbf{x}_3), \dots, K(\mathbf{x}_{n-1}, \mathbf{x}_n)), \quad i \in S \\ \hat{\alpha}_i = 0, \quad i \notin S \end{cases}$$

# Support Vector Machines: Linear Kernel.

A kernel is a function that **quantifies the similarity** of two observations.

**Example.** For two observations  $\mathbf{x}_i, \mathbf{x}_j \in \mathbb{R}^p$ :

$$K(\mathbf{x}_i, \mathbf{x}_j) = \sum_{k=1}^p x_{ik} x_{jk} - \text{linear kernel.}$$

A few questions popping up on that:

1. Why is it called **linear**?


If we plug it into (13):

$$\begin{aligned} f(\mathbf{x}) &= \beta_0 + \sum_{i \in S} \alpha_i K(\mathbf{x}, \mathbf{x}_i) = \beta_0 + \sum_{i \in S} \alpha_i \left( \sum_{k=1}^p x_k x_{ik} \right) = \\ &= \beta_0 + \sum_{k=1}^p \left( \sum_{i \in S} \alpha_i x_{ik} \right) x_k \implies \text{linear function in features } x_1, \dots, x_p \end{aligned}$$

# Support Vector Machines: Linear Kernel.

2. Why does it quantify similarity of  $\mathbf{x}_i \in \mathbb{R}^p$  and  $\mathbf{x}_j \in \mathbb{R}^p$ ?

## Vector Inner Product



$$u = \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} \quad v = \begin{bmatrix} v_1 \\ v_2 \end{bmatrix}$$

$$\begin{aligned} \|\mathbf{u}\|_2 &= \text{length}(\mathbf{u}) \in \mathbb{R} \\ &= \sqrt{u_1^2 + u_2^2} \end{aligned}$$

$$\begin{aligned} \mathbf{u}^\top \mathbf{v} &= \mathbf{v}^\top \mathbf{u} \\ &= u_1 v_1 + u_2 v_2 \\ &= \|\mathbf{u}\|_2 \|\mathbf{v}\|_2 \cos \theta \end{aligned}$$

Quantifies **directional similarity** well - whether vectors point in the

- ▶ **same** direction ( $\cos(\theta) > 0$ )
- ▶ **opposite** directions ( $\cos(\theta) < 0$ )

# Support Vector Machines: Polynomial Kernel.

To introduce **non-linearity** into the boundaries produced by solving optimization task (7) – (10), we can use the following kernels:

- $K(\mathbf{x}_i, \mathbf{x}_j) = (1 + \langle \mathbf{x}_i, \mathbf{x}_j \rangle)^d = (1 + \sum_{k=1}^p x_{ik} x_{jk})^d$  - **polynomial** kernel of **degree  $d$**  ( $d$  - positive integer). It is a **non-linear** kernel.

**Question:** What happened to the idea of manually adding extra features, e.g.  $X_1^2, X_2^2, \dots$ ?

**Answer:** Polynomial kernel **implicitly** adds those for us, **depending on value  $d$  specified**. E.g. for  $d = 2$ :


$$\begin{aligned} f(\mathbf{x}) &= \beta_0 + \sum_{i \in S} \alpha_i K(\mathbf{x}, \mathbf{x}_i) = \beta_0 + \sum_{i \in S} \alpha_i \left(1 + \sum_{k=1}^p x_k x_{ik}\right)^2 = \\ &= \beta_0 + \sum_{i \in S} \alpha_i \left(1 + 2 \sum_{k=1}^p x_{ik} x_k + \left(\sum_{k=1}^p x_{ik} x_k\right)^2\right) \end{aligned}$$

$\Rightarrow$  **Task:** show explicitly that it's a linear function of  $x_1, x_1^2, x_2, x_2^2, \dots$



# Support Vector Machines: Polynomial Kernel.

**Example (cont'd).** Back to our example from the very first slide: here's an SVM fit with **polynomial kernel** of **degree 3**.



It produces much more **flexible boundary shapes**, which appear appropriate for this data.

# Support Vector Machines: Radial Kernel.

Another example of a kernel yielding flexible non-linear boundary is:

- $K(\mathbf{x}_i, \mathbf{x}_j) = \exp(-\gamma \sum_{k=1}^p (x_{ik} - x_{jk})^2)$ ,  $\gamma > 0$  - **radial** kernel.

## Intuition:

If obs.  $\mathbf{x}^* = (x_1^*, \dots, x_p^*)^T$  is far from obs.  $\mathbf{x}_i = (x_{i1}, \dots, x_{ip})^T \implies$

Euclidean distance  $\sum_{k=1}^p (x_k^* - x_{ik})^2$  is **large**  $\implies$

$K(\mathbf{x}^*, \mathbf{x}_i) = \exp(-\gamma \sum_{k=1}^p (x_{ik}^* - x_{ik})^2)$  is **very tiny**.

This goes on to show that

- ▶ Radial kernel is a **similarity measure** between  $\mathbf{x}^*$  &  $\mathbf{x}_i$ .
- ▶ A "**far away**" support vector  $\mathbf{x}_i$  plays **virtually no role** in  $f(\mathbf{x}^*)$


$$f(\mathbf{x}^*) = \beta_0 + \sum_{i \in S} \alpha_i K(\mathbf{x}^*, \mathbf{x}_i)$$

due to **tiny value of**  $K(\mathbf{x}^*, \mathbf{x}_i)$ .

# Support Vector Machines: Radial Kernel.

Hence, the radial kernel has very **local behavior**: only **nearby** training observations determine the **class label** of a **test observation**.

**Example (cont'd)**. Below is the SVM fit for **radial kernel** with  $\gamma = 0.5$ .



# Support Vector Machines: Radial Kernel.

- $\gamma$  is a parameter of the radial based kernel. It determines the spread of the kernel and therefore the decision region.
- When  $\gamma$  is small, the curvature of the decision boundary is very low and thus the decision region is broad.
- When  $\gamma$  is large, the curvature of the decision boundary is high, which may lead to islands of decision-boundaries around data points.

# Kernel vs Enlarging Feature Space: Advantages?

Biggest advantage of **kernels** as opposed to simply enlarging feature space is **computational**:

- With kernels, we only compute  $K(\mathbf{x}_i, \mathbf{x}_j)$  for all  $\binom{n}{2}$  distinct pairs  $\mathbf{x}_i, \mathbf{x}_j$  of training observations, in the **original  $p$ -dimensional space**.
- Kernels are computed **without explicitly working in the enlarged feature space**, which in many SVM applications can be **so large** that **computations are intractable**.

**Example.** Imagine working in 10-dimensional space:

$$\{X_1, X_2, \dots, X_{10}\}$$

and extending it to include quadratic, cubic terms + interactions:

$$\{X_1, X_1^2, X_1^3, \dots, X_1 X_2, X_1 X_3, \dots, X_1 X_2 X_3\}$$

You end up with  $10 \times 3 + \frac{10(10-1)}{2} + 1 = 76$  features.

# SVM: Heart data example.

**Example.** In *Heart* data we have

- 13 predictors (such as age, sex, cholesterol level, etc), and
- predict whether individual has heart disease or not (*AHD* = *Yes/No*, binary response)

```
> Heart <- read.csv("~/Downloads/Heart.csv")
> head(Heart)
```

	X	Age	Sex	ChestPain	RestBP	Chol	...	AHD
1	1	63	1	typical	145	233	...	No
2	2	67	1	asymptomatic	160	286	...	Yes
3	3	67	1	asymptomatic	120	229	...	Yes
4	4	37	1	nonanginal	130	250	...	No
5	5	41	0	nontypical	130	204	...	No
6	6	56	1	nontypical	120	236	...	No

We will use [support vector machines \(SVM\)](#) to build a classifier diagnosing the heart disease status of a patient given other characteristics.

## Heart data example: EDA.

**Example (cont'd).** But first, some **exploratory data analysis**:

```
> dim(Heart)
303  15
> summary(Heart)
```

X	Age	..	Ca	Thal
Min. : 1.0	Min. :29.00	..	Min. :0.0000	fixed : 18
1st Qu.: 76.5	1st Qu.:48.00	..	1st Qu.:0.0000	normal:166
Median :152.0	Median :56.00	..	Median :0.0000	revers:117
Mean :152.0	Mean :54.44	..	Mean :0.6722	NA's : 2
3rd Qu.:227.5	3rd Qu.:61.00	..	3rd Qu.:1.0000	
Max. :303.0	Max. :77.00	..	Max. :3.0000	
		..	NA's :4	

We may witness a **few missing values** among the *Ca* and *Thal* variables. Typically one may want to look into it in more detail, but given that it is just 6 observations out of 303, we may simply discard them:

```
Heart <- na.omit(Heart)
dim(Heart)
297  15
summary(Heart) # Check - does it have NAs now?
```

## SVM in R: *svm()* function of e1071 library.

SVM in R can be performed via *svm()* function of e1071 library:

```
library(e1071)
?svm()
```

**Example (cont'd).** To build a support vector machine with linear kernel ( $\Leftrightarrow$  support vector classifier) for *Heart* data:

```
svm.obj <- svm(AHD ~.,
               data=Heart,
               kernel='linear',
               cost=5)
```

In this *svm()* call:

- *AHD ~ .* - classical *formula* argument,
- *data = Heart* - classical *data* argument,
- *kernel = 'linear'* - SVM kernel (other opt: 'polynomial', 'radial', ...)
- *cost = 5* - the cost of a margin violation ( $\equiv \frac{1}{\text{Budget}} = \frac{1}{C}$ ).
- by **default** has *scale = TRUE* - it *scales* the features



## SVM in R: *predict()*

**Example (cont'd).** To evaluate quality of SVM classifications, we first obtain them via *predict()* function:

```
> predict(svm.obj)
 1    2    3    4    5    6    7    8    9   10   11   12   13   14   15 ...
No Yes Yes  No  No  No Yes  No Yes Yes  No  No  No  No  No ...

> head(data.frame(True=Heart$AHD, Predicted=predict(svm.obj)))
  True Predicted
1   No         No
2  Yes         Yes
3  Yes         Yes
4   No         No
5   No         No
6   No         No

> mean(predict(svm.obj) != Heart$AHD)
[1] 0.1279461
```

This SVM with linear kernel yields 12.8% training error.

## SVM in R: Radial Kernel.

**Example (cont'd).** Now, let's run SVM for radial basis kernel

$$K(\mathbf{x}_i, \mathbf{x}_j) = \exp(-\gamma \sum_{k=1}^p (x_{ik} - x_{jk})^2), \gamma > 0$$

with  $\gamma = 10^{-3}, 10^{-2}$  and  $10^{-1}$ . Value of  $\gamma$  is specified via *gamma* argument to *svm()* function:

```
for (j in 3:1){  
  svm.obj <- svm(AHD ~.,  
                 data=Heart,  
                 kernel='radial',  
                 gamma=10^{-j},  
                 cost=5)  
  
  print(mean(predict(svm.obj) != Heart$AHD)) }  
[1] 0.1481481  
[1] 0.1178451  
[1] 0.006734007
```

## SVM in $R$ : Radial Kernel.

**Example (cont'd).** As we could see above,

training error is decreasing (from 0.148 down to 0.007)

as

$\gamma$  is increasing (from  $10^{-3}$  up to  $10^{-1}$ )

**BUT** it might a sign of **overfitting**  $\implies$  we need a **test error estimate**.

Let's **randomly** divide data 70/30 into

- 70% training data (207 observations), and
- 30% test set (90 observations).

Then let's compare the performance on test set for

- SVM with **linear** kernel,
- SVMs with **radial** kernel and  $\gamma$  values of  $10^{-3}$ ,  $10^{-2}$ ,  $10^{-1}$ .

**NOTE:** these test error estimates are based **just on one train/test data subdivision**. **Cross-validation** should be done for **reliable model comparisons** (see `tune()` function used in Labs #4 and #5).

## SVM in R: Test errors.

```
set.seed(1)
n <- nrow(Heart);  train <- sample(1:n, 0.7*n)

## Support Vector Classifier (Linear Kernel)
svm.obj <- svm(AHD ~., data=Heart, cost=5, kernel='linear',
               subset=train)
mean(predict(svm.obj, newdata=Heart[-train,])
      != Heart$AHD[-train])
[1] 0.2

## Support Vector Machines (Radial Kernel)
for (j in 3:1){
  svm.obj <- svm(AHD ~., data=Heart, cost=5,
                 kernel='radial', gamma=10^{-j},
                 subset=train)
  print(mean(predict(svm.obj, newdata=Heart[-train,])
             != Heart$AHD[-train])))
[1] 0.1666667
[1] 0.2
[1] 0.1555556
```

## False Postives/Negatives.

While misclassification error is an important metric of classification model performance, it **doesn't give us a full picture** of performance quality.

**Example.** In *Default* data set, describing whether a customer defaults on credit card payment or not, we had only  $\approx 3\%$  of customers that defaulted. Hence, if we simply **predicted *Default* = No for each customer**, we'd obtain a **train/test error of  $\leq 3\text{-}4\%$**  just like that.

In general, 1/0 (or *TRUE*/*FALSE*) binary classifiers can make **two types of misclassification errors**:

1. Misclassify a 1 for a 0 (*TRUE* for *FALSE*)  $\implies$  **false negative**.
2. Misclassify a 0 for a 1 (*FALSE* for *TRUE*)  $\implies$  **false positive**.

It is often of interest to determine exactly which of these two types of errors are being made.

# False Postives/Negatives, Confusion Matrix.

**Example (cont'd).** For the *Default* example, our algorithm can

1. **incorrectly assign** an individual who **actually defaults** (*Default = Yes*) to the ***Default = No* category**  $\implies$  **false negative**,
2. **incorrectly assign** an individual who **does not default** (*Default = No*) to the ***Default = Yes* category**  $\implies$  **false positive**.

A **confusion matrix** is a convenient way to display that information:

		<i>True default status</i>		
		No	Yes	Total
<i>Predicted default status</i>	No	9,644	252	9,896
	Yes	23	81	104
Total		9,667	333	10,000

**False positive rate (FPR):** for customers that didn't default, model predicts them **incorrectly** (*Default = Yes*) in just  $\frac{23}{9667} = 0.02\%$  cases  $\implies$   **$FPR = 0.02$** .

**Question:** What's the **false negative rate (FNR)** though?

## False Postives/Negatives, Confusion Matrix.

**Example.** Back to the *Heart* data of predicting whether a patient has a heart disease or not. SVM with radial kernel and  $\gamma = 10^{-1}$  yielded best test set performance (15.5% error). Let's check the **confusion matrix** of its predictions:

```
> table(pred=predict(svm.obj, newdata=Heart[-train,]),
        true=Heart$AHD[-train])
```

	true	
pred	No	Yes
No	42	8
Yes	6	34

**Question #1:** What is the false positive rate? False negative rate?

**Question #2:** Given that we are trying to detect a disease, which metric should be of more importance to us - false positive or false negative rate?

# SVMs with More than Two Classes: One-vs-One.

So far, we've discussed SVMs for **binary** classification, the **two-class setting**. How to extend SVMs onto a **general case of  $K \geq 2$  classes**?

## One-versus-One Classification:

1. Construct  $\binom{K}{2} \left( \equiv \frac{K(K-1)}{2} \right)$  SVMs, for all possible **pairs of classes**.
2. Use those SVMs to get  $\binom{K}{2}$  **class predictions** for a test observation.
3. Assign this test observation **the most frequent class** among the  $\binom{K}{2}$  **predictions**.



# SVMs with More than Two Classes: One-vs-All.

## One-versus-All Classification:

1. For class  $k$ ,  $k = 1, \dots, K$ , fit SVM such that it compares
  - ▶ observations from class  $k$  (class #1,  $y = +1$ ), with
  - ▶ those not belonging to class  $k$  (class #2,  $y = -1$ )

Let  $\beta_{0k}, \beta_{1k}, \dots, \beta_{pk}$  denote hyperplane parameters from such "One-vs-All" SVM for class  $k$ ,  $k = 1, \dots, K$ . Given that  $y_i = +1$  for arbitrary observation  $\mathbf{x} \in$  class  $k$ :

$$\mathbf{x} \in \{\text{Class } k\} \Leftrightarrow \beta_{0k} + \beta_{1k}x_1^* + \beta_{2k}x_2^* + \dots + \beta_{pk}x_p^* > 0$$

2. Given test observation  $\mathbf{x}^* = (x_1^*, x_2^*, \dots, x_p^*)$ , we assign it to class  $k$  such that

$$k = \max_{k=1, \dots, K} \beta_{0k} + \beta_{1k}x_1^* + \beta_{2k}x_2^* + \dots + \beta_{pk}x_p^*$$

as this amounts to a high level of confidence that the test observation belongs to the  $k^{\text{th}}$  class rather than to any of the other classes.

# Which one to choose?

- **One-versus-One.** Fit all  $\binom{K}{2}$  pairwise SVM classifiers. Classify test observation  $x^*$  to the class to which it was most frequently assigned in these  $\binom{K}{2}$  pairwise classifications.
- **One-versus-All.** Fit  $K$  different 2-class SVM classifiers  $\hat{f}_k(x)$ ,  $k = 1, \dots, K$ ; each class versus the rest. Classify  $x^*$  to the class for which  $\hat{f}_k(x)$  is largest.
- Which one to choose? If  $K$  is not too large, use One-versus-One.

## One-vs-One SVM example in *R*.

If the response is a factor containing **more than two levels**, then the `svm()` function will perform **multi-class classification** using the **one-versus-one** approach.

**Example.** We create a simulated data example with  $K = 3$  classes of observations in  $2D$  predictor space. We first generate some data on  $K = 2$  classes with a non-linear class boundary, as follows:

```
set.seed(1)
x1 <- rnorm(200)
x2 <- rnorm(200)

x1[1:100] <- x1[1:100] + 2; x2[1:100] <- x2[1:100] + 2;
x1[101:150] <- x1[101:150] - 2; x2[101:150] <- x2[101:150] - 2

y=c(rep(1,150),
    rep(2,50))

x <- cbind(x1,x2)
```

# One-vs-One SVM example in R.

**Example (cont'd).** Then we add a *3rd class*:

```
set.seed(1)
x=rbind(x,
        matrix(rnorm(50*2), ncol=2))
y=c(y, rep(0,50))
x[y==0,2]=x[y==0,2]+2
dat=data.frame(x=x,
               y=as.factor(y))
```

We now fit an SVM to the data:

```
svmfit=svm(y~., data=dat,
           kernel="radial",
           cost=10,
           gamma=1)
```

# One-vs-One SVM example in *R*.

**Example (cont'd).** Below are the plots of original data (left) and the SVM fit (right).

```
par(mfrow=c(1,2))  
plot(x,col=(y+1))  
plot(svmfit, dat)
```

